



Dependability through Redundancy

Prof. George Canea

School of Computer & Communication Sciences

How to achieve dependability?

- Use modularity ...
- ... and REDUNDANCY for ...
 - *fault tolerance*
 - *high reliability*
 - *high availability*

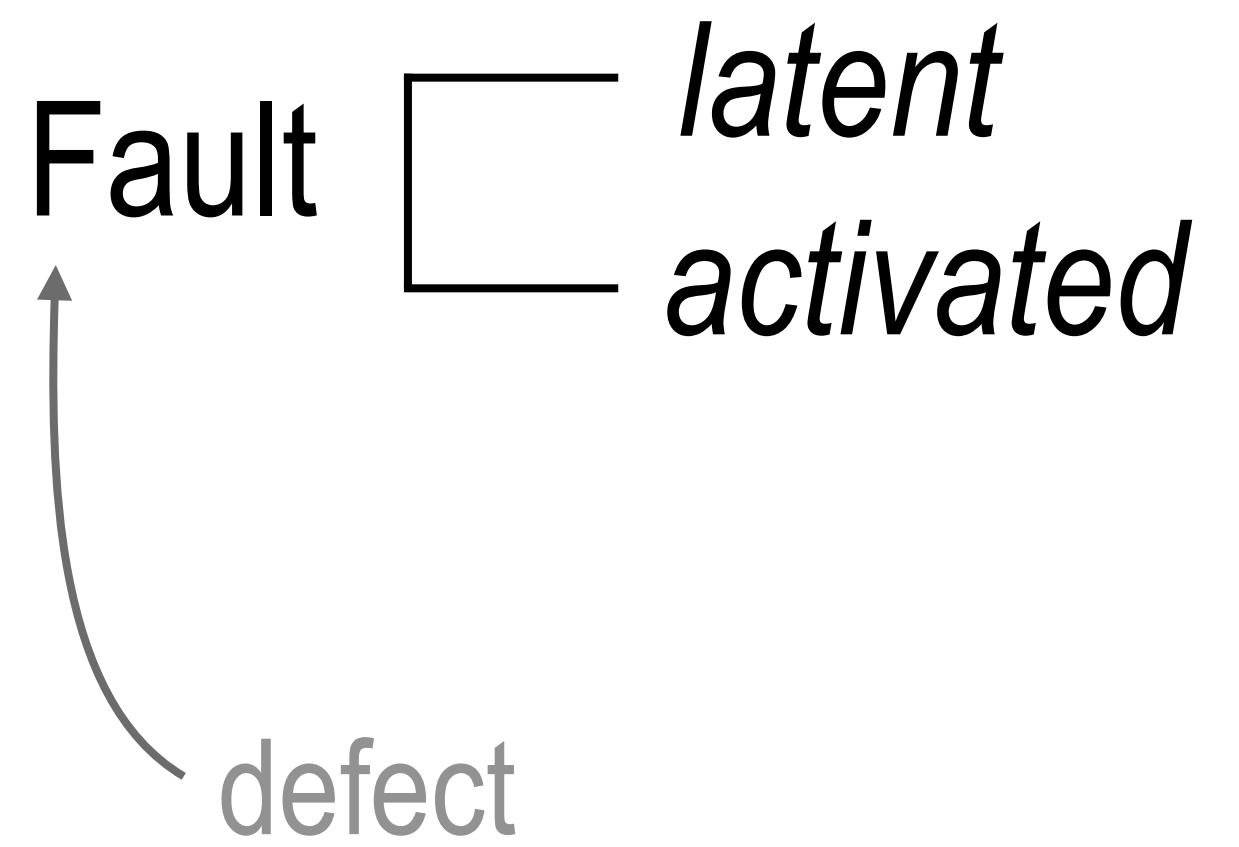
How to achieve dependability?

- Use modularity ...
- ... and REDUNDANCY for ...
 - *fault tolerance*
 - *high reliability*
 - *high availability*

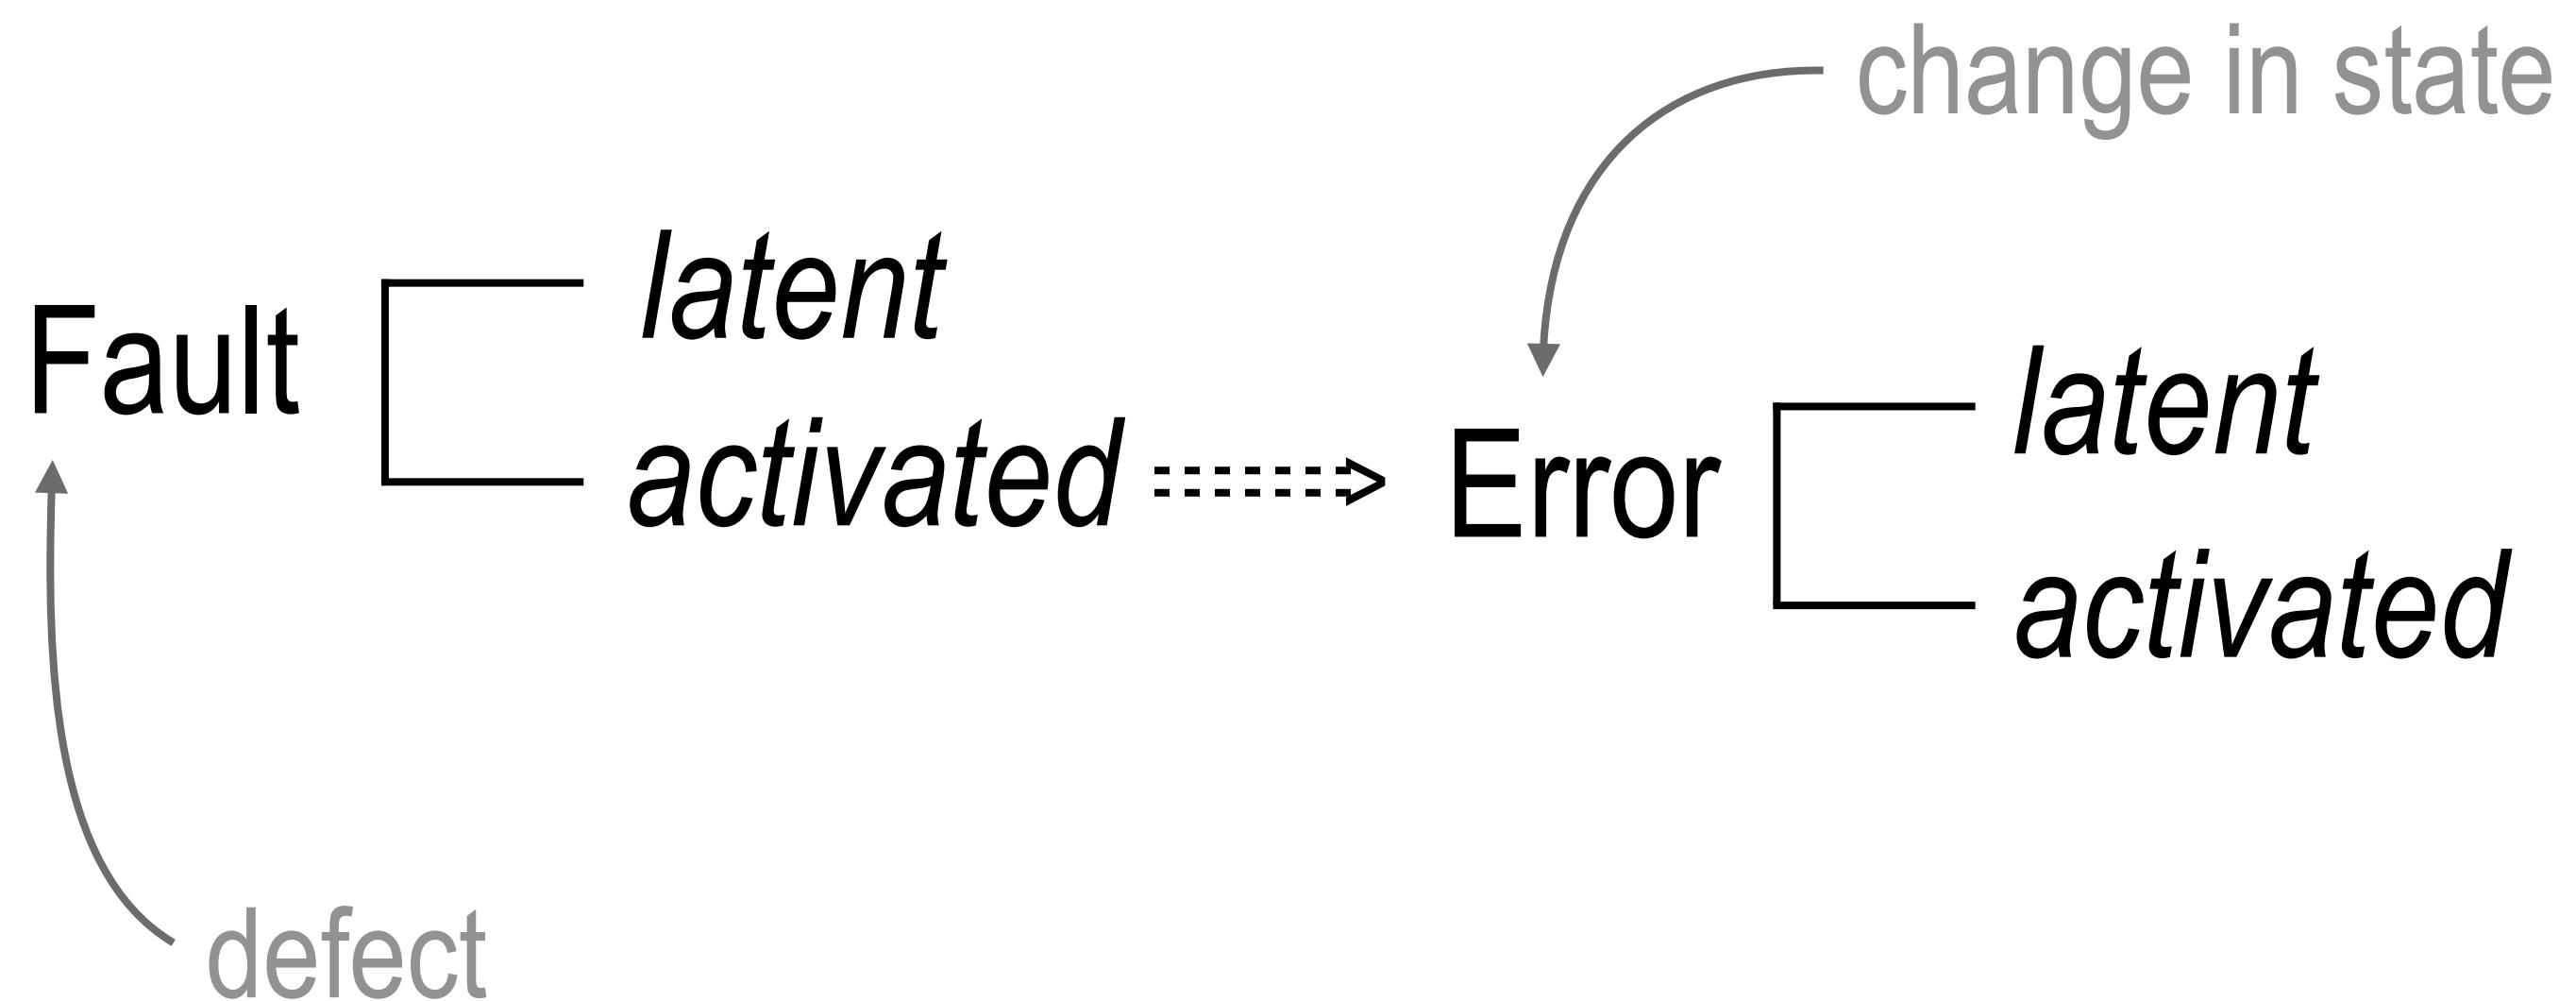
Redundancy = duplication with the purpose of increasing dependability

Fault tolerance

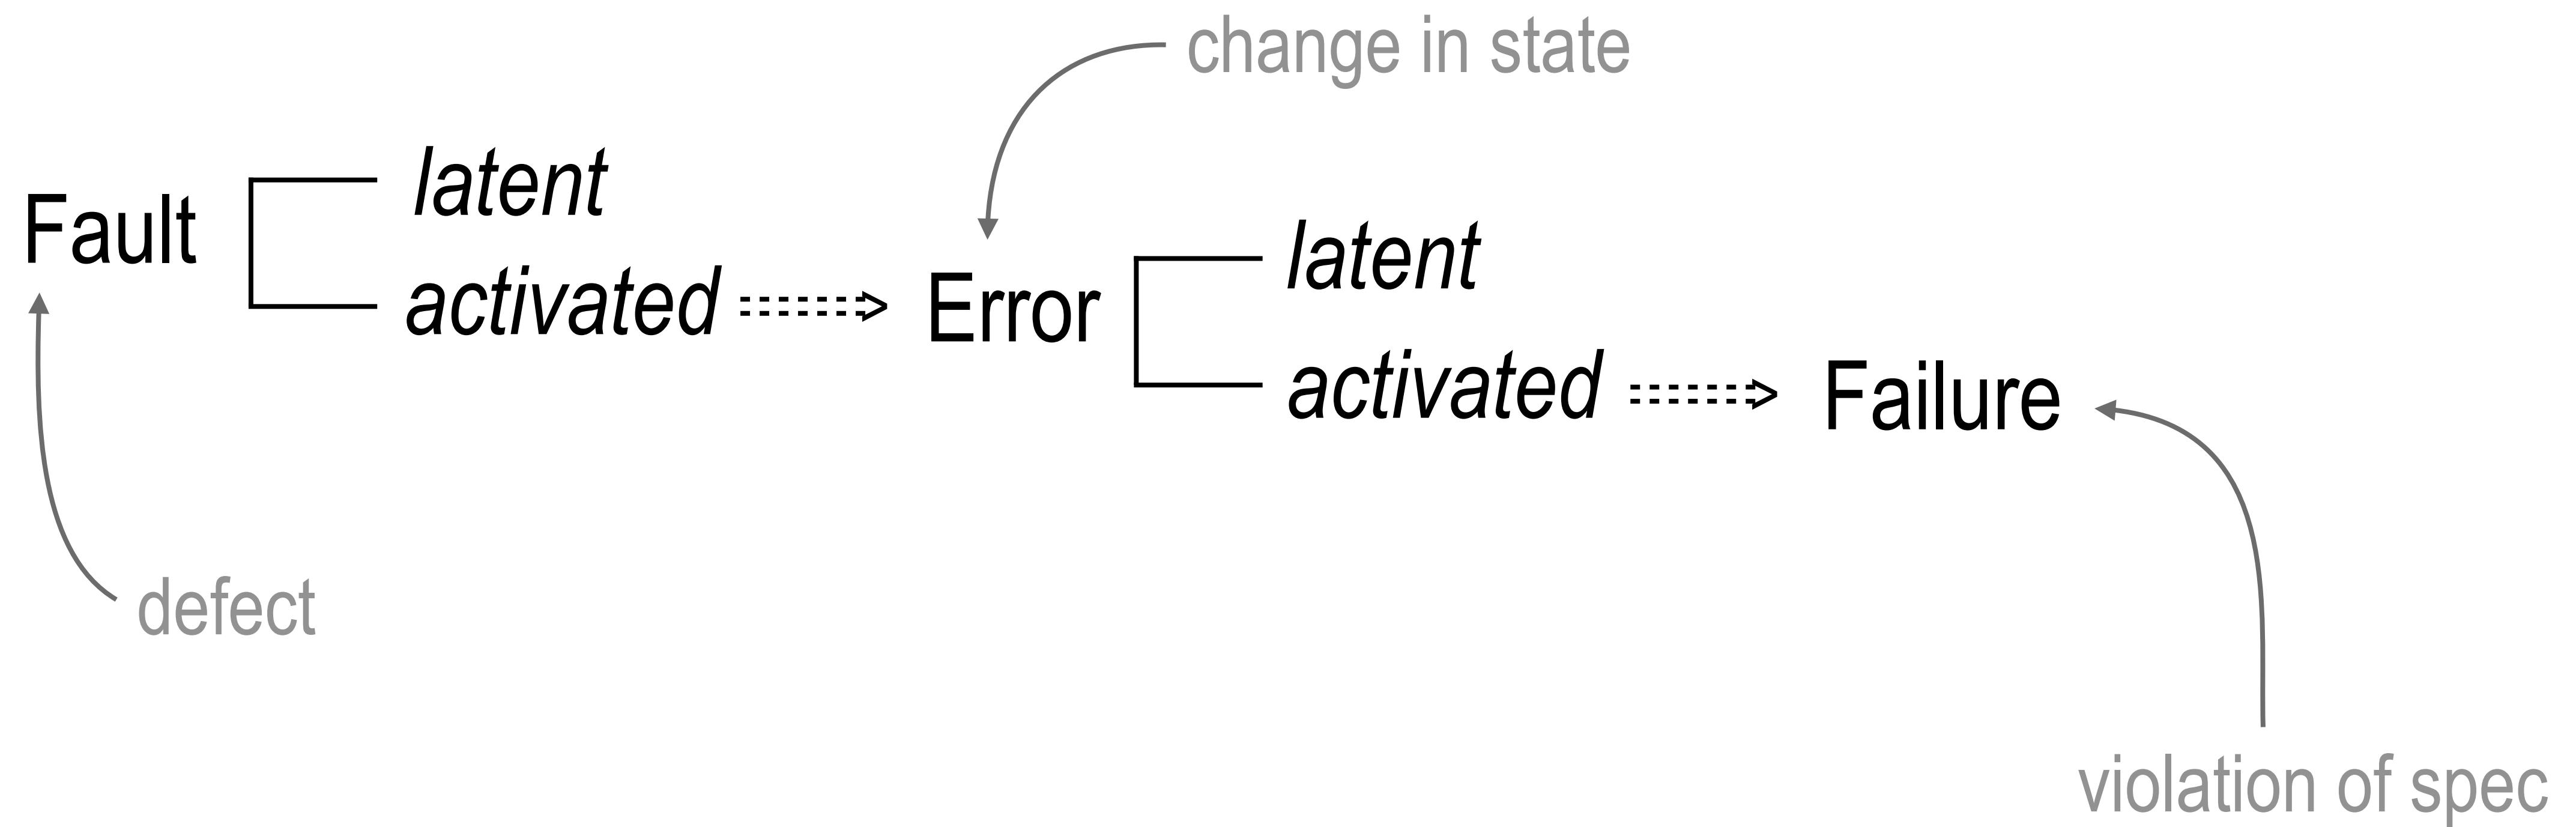
Fault tolerance



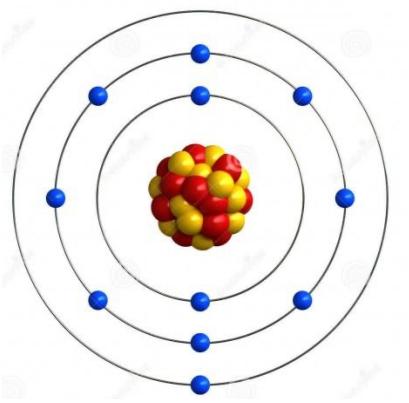
Fault tolerance



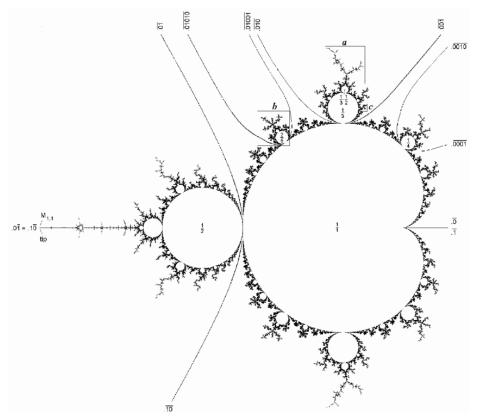
Fault tolerance



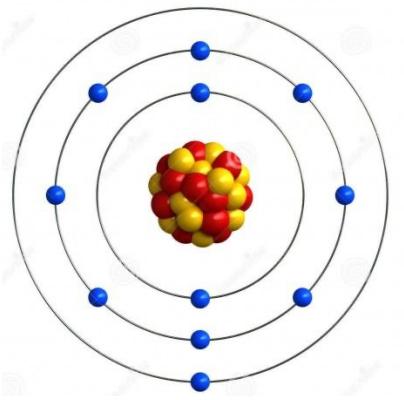
Types of software faults / defects



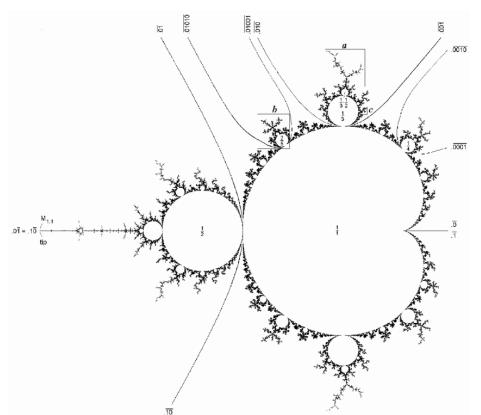
$$\Delta\chi\Delta\rho \geq \frac{\hbar}{2}$$



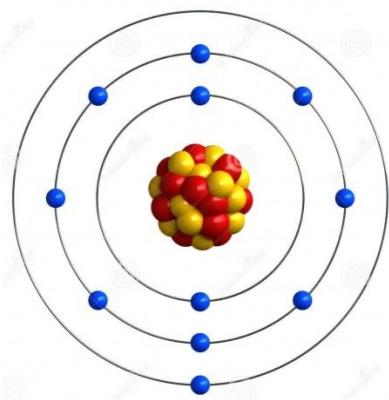
Types of software faults / defects



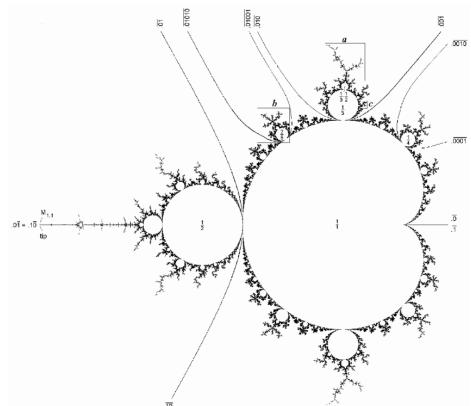
- Bohrbug
 - *clear + easy to reproduce => easy to fix*
- Heisenbug
 - *disappears when you attach with debugger*



Types of software faults / defects



- Bohrbug
 - *clear + easy to reproduce => easy to fix*
- Heisenbug
 - *disappears when you attach with debugger*
- Schrödingbug
 - *starts causing failure once you realize it should*
- Mandelbug
 - *complex, obscure, chaotic, seemingly non-deterministic*



Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

Data/information redundancy

Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

Data/information redundancy

Geographic redundancy

Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

Data/information redundancy

Geographic redundancy

Processing redundancy

Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

Data/information redundancy

Geographic redundancy

Space
Processing redundancy

Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

Data/information redundancy

Geographic redundancy

Processing redundancy

Space

Time

Using redundancy to tolerate faults

- "tolerate" faults = cope with errors or the resulting failures
 - *the actual goal is to tolerate the consequences of faults*
- Using redundancy to cope with errors
 - *forward error correction*
 - *redundant copies/replicas (=coarse-grained ECC)*
 - ...
- Using redundancy to cope with failures
 - *server/service failover*
 - *Internet routing*
 - ...

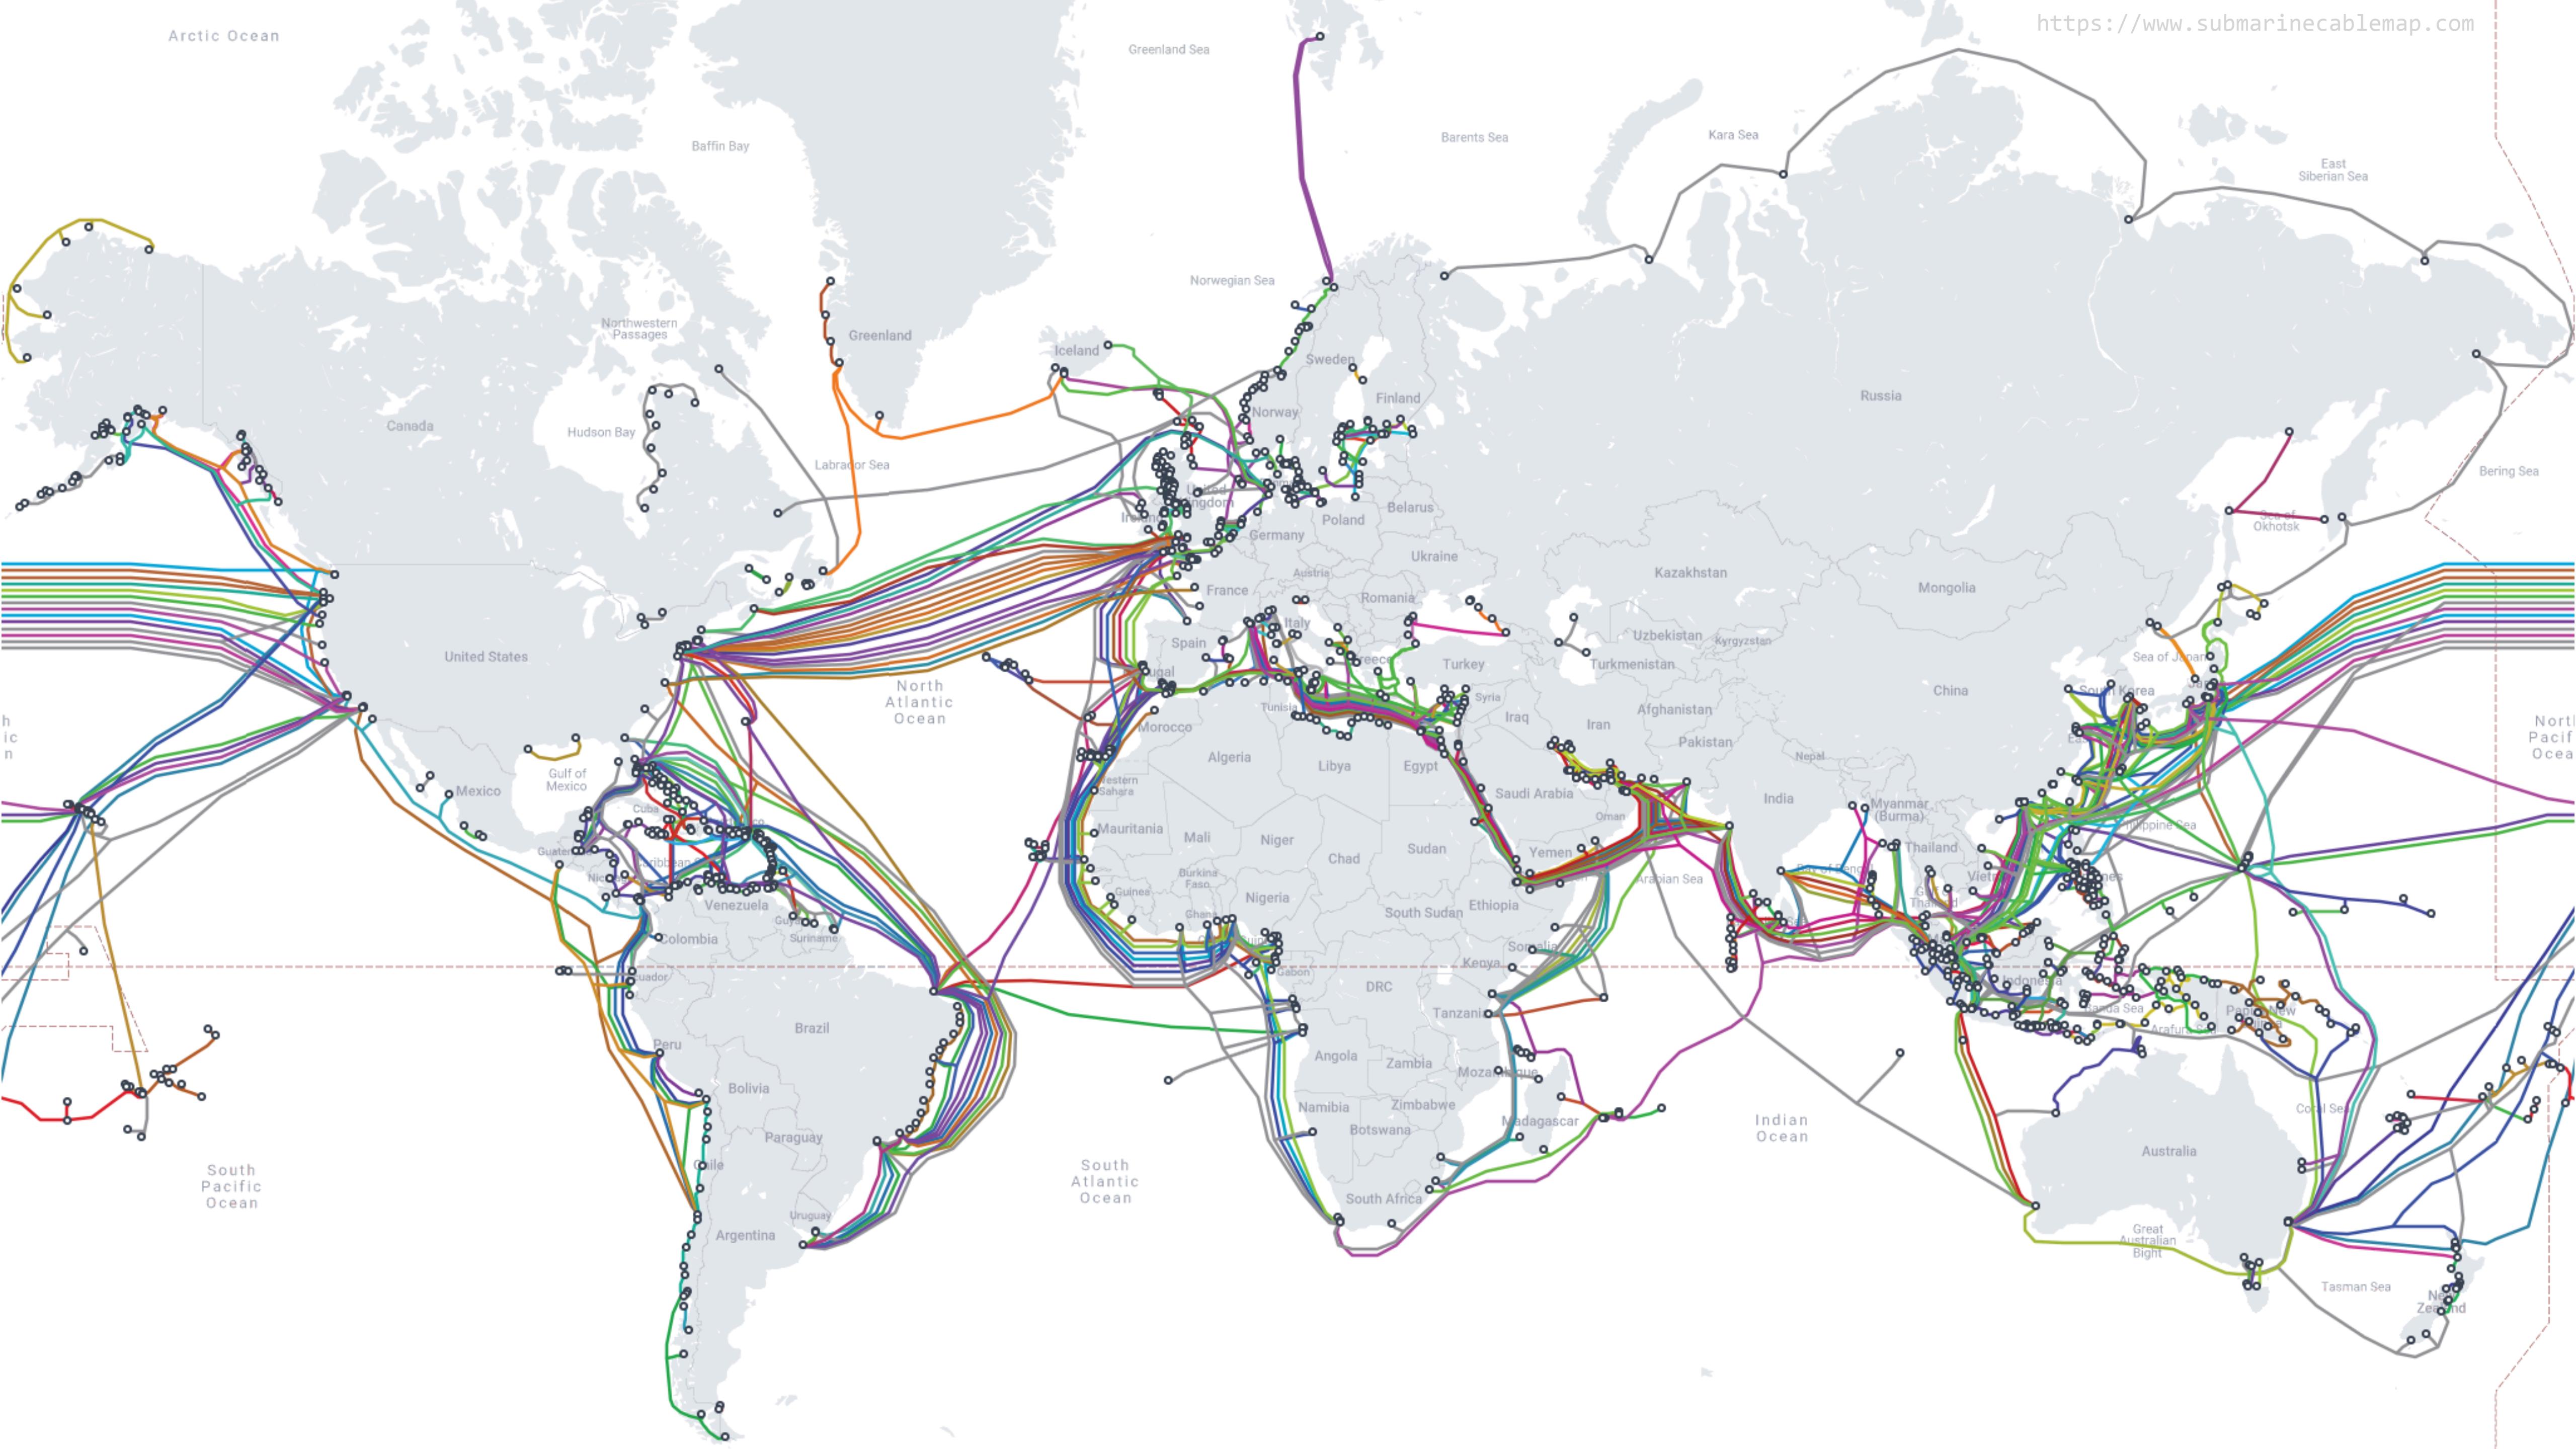
Data/information redundancy

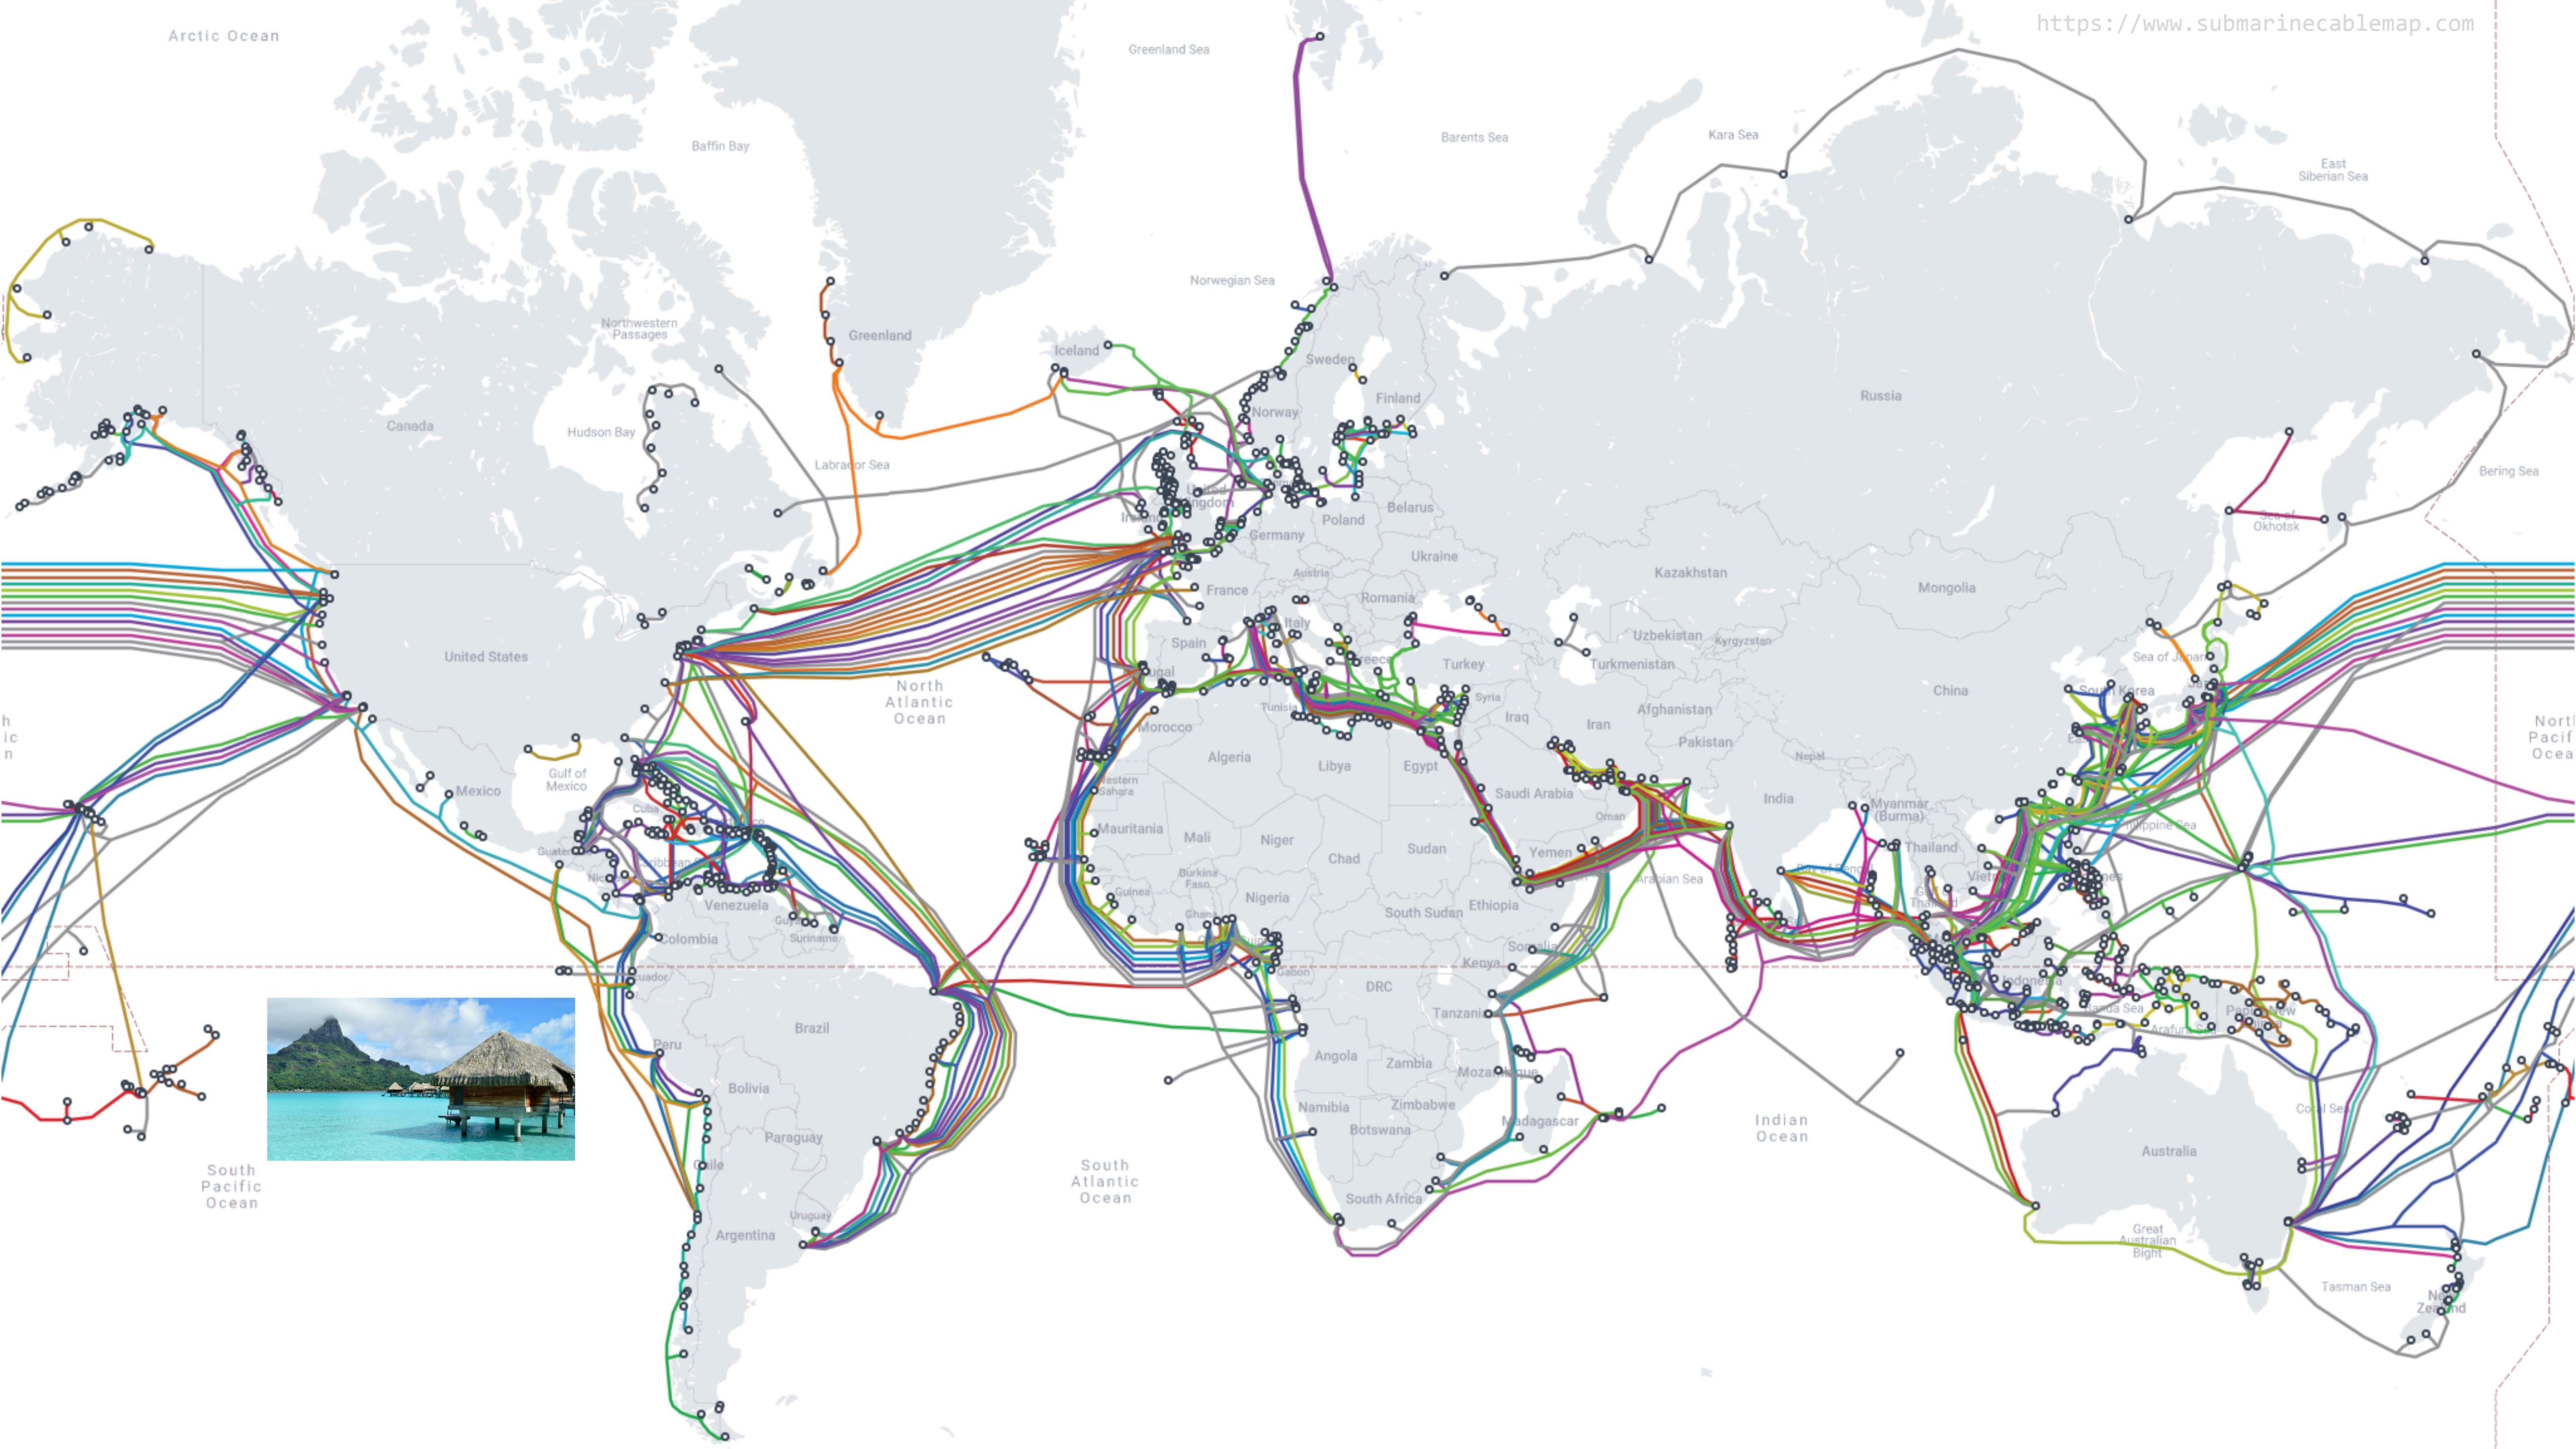
Geographic redundancy

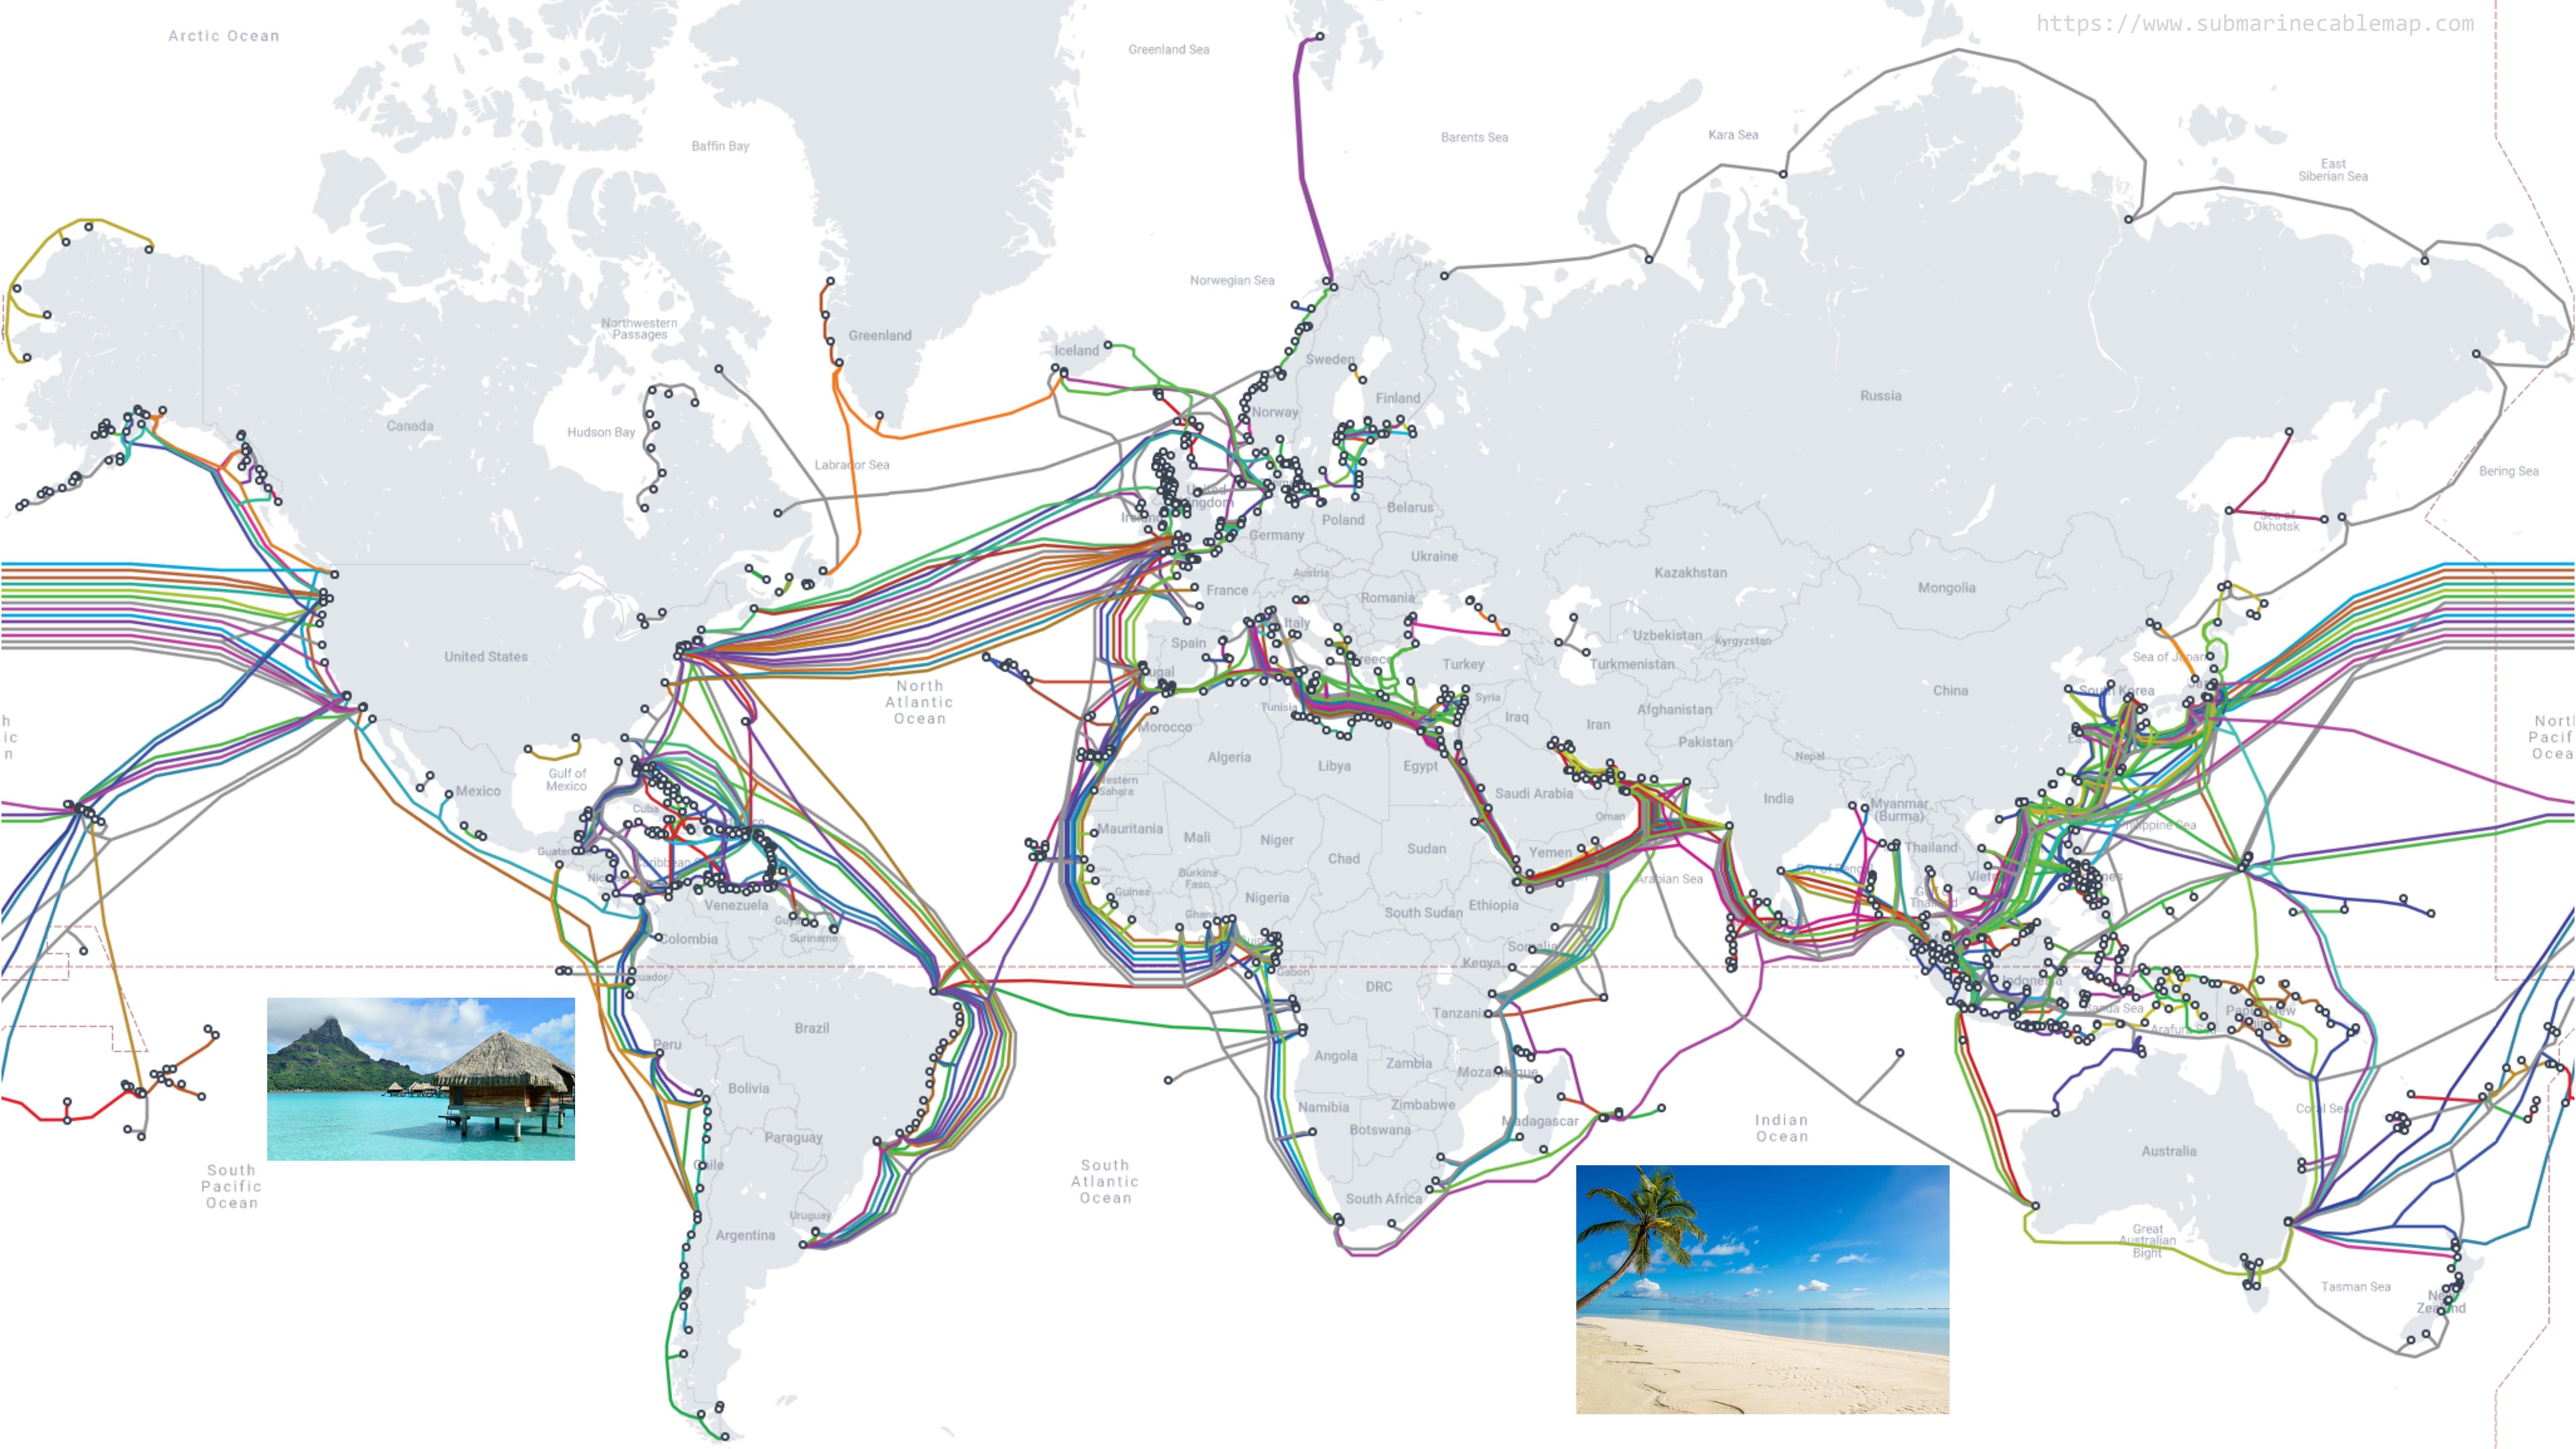
Processing redundancy

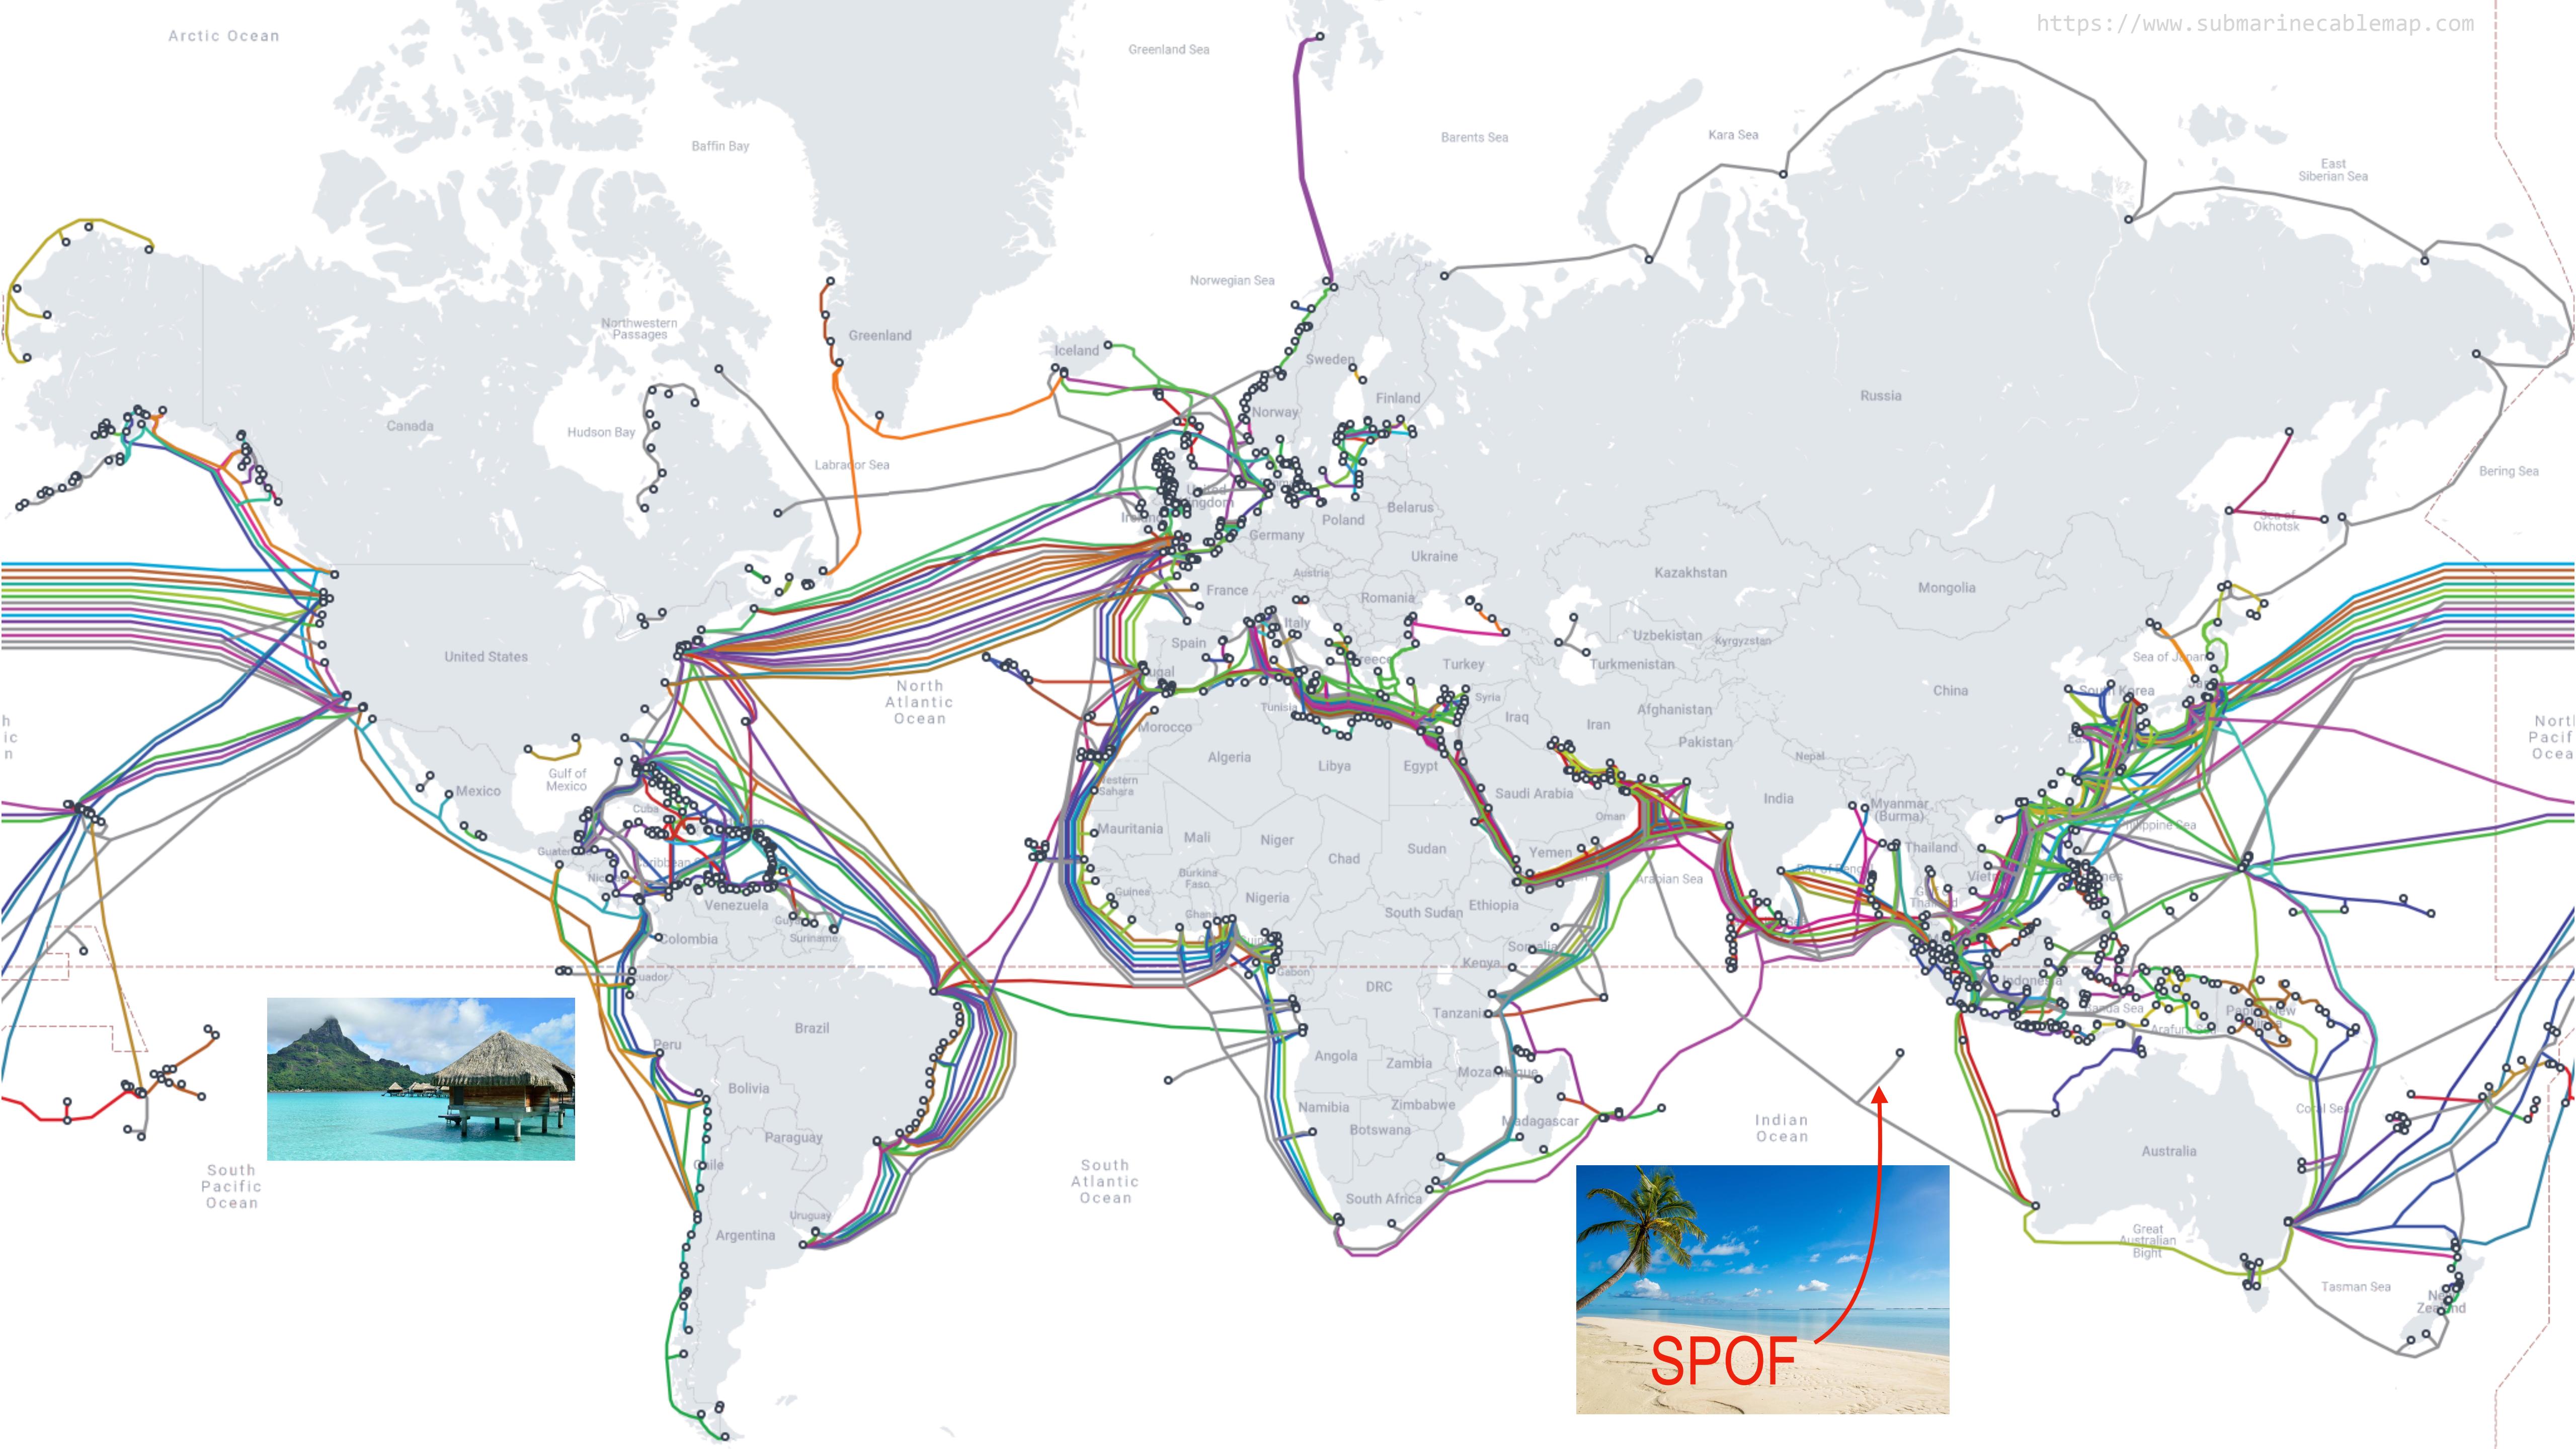
Functional redundancy

Space
Time









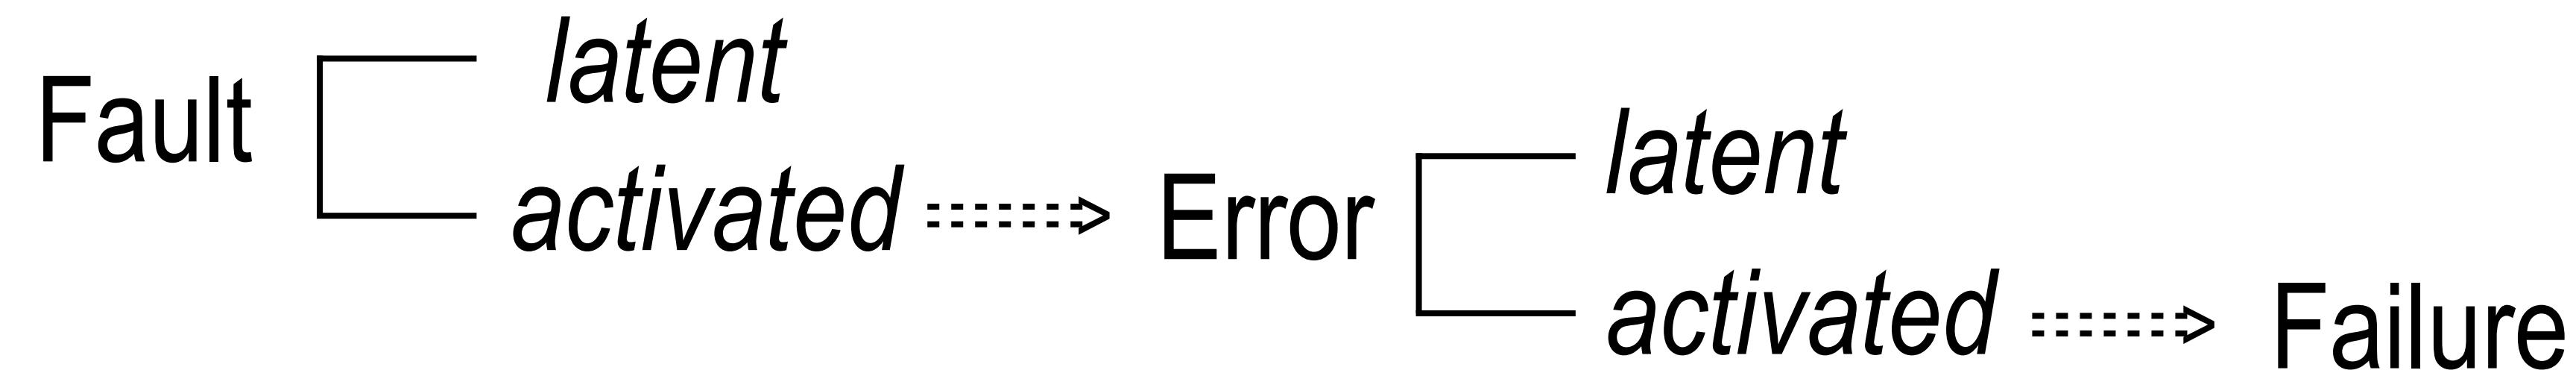
Fault model

- Specification of what could go wrong and what cannot go wrong
 - *Used to predict consequences of failures*
 - *Should also specify what can / cannot happen during recovery*
 - *Remember the single points of failure (SPOFs)*

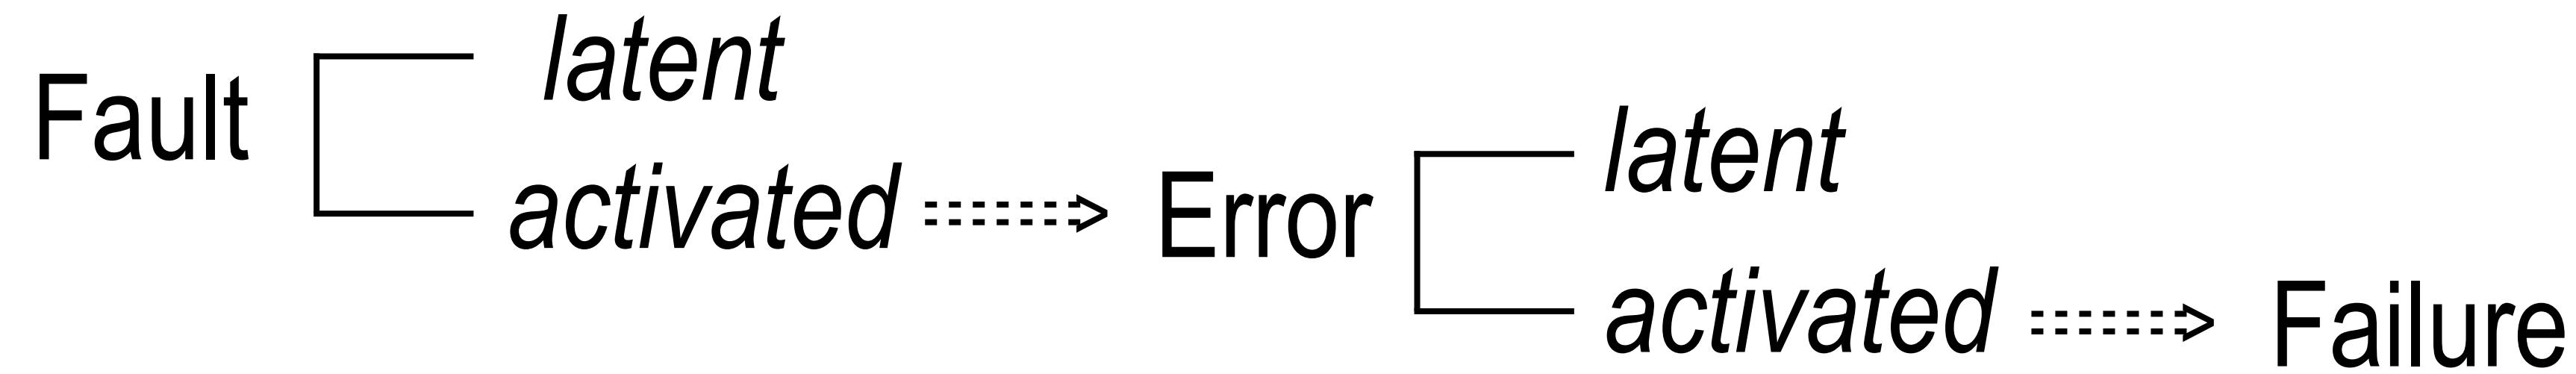
Fault model

- Specification of what could go wrong and what cannot go wrong
 - *Used to predict consequences of failures*
 - *Should also specify what can / cannot happen during recovery*
 - *Remember the single points of failure (SPOFs)*
- Example: N-version programming
 - *use redundancy to tolerate software faults*

Recap: Fault tolerance

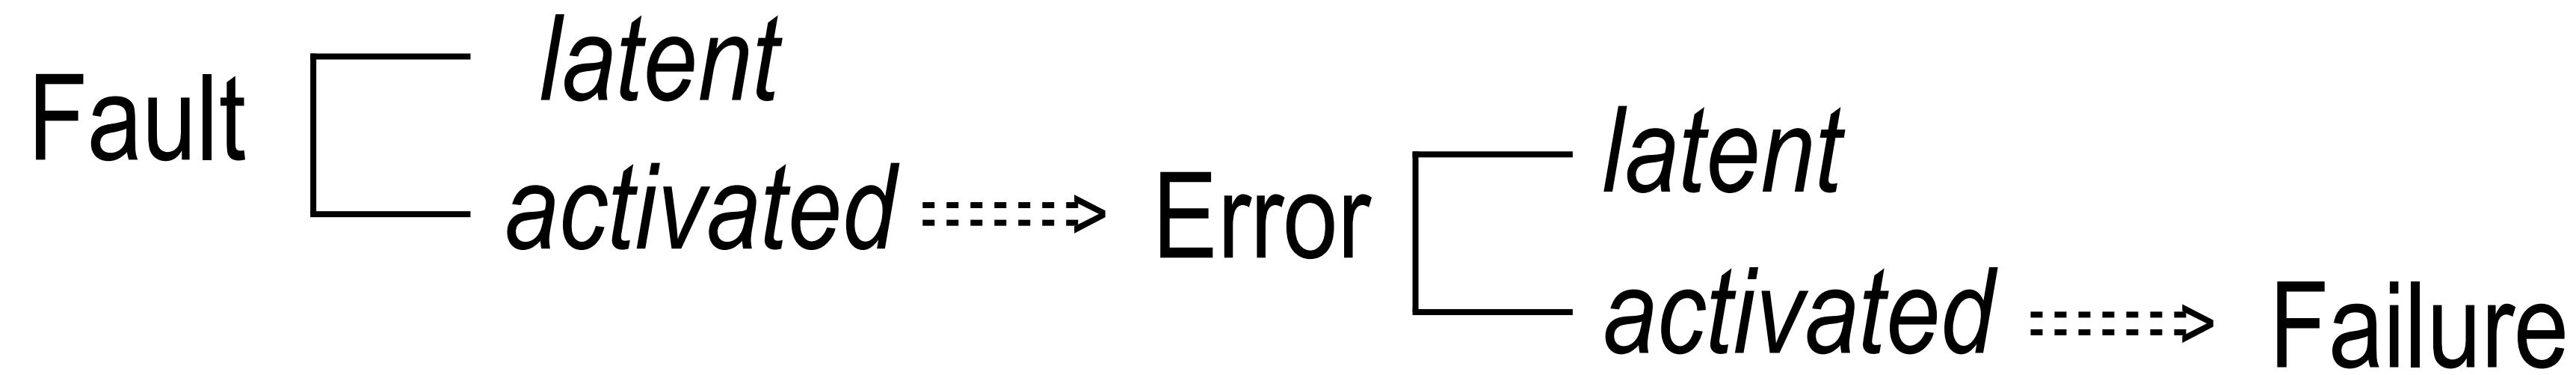


Recap: Fault tolerance



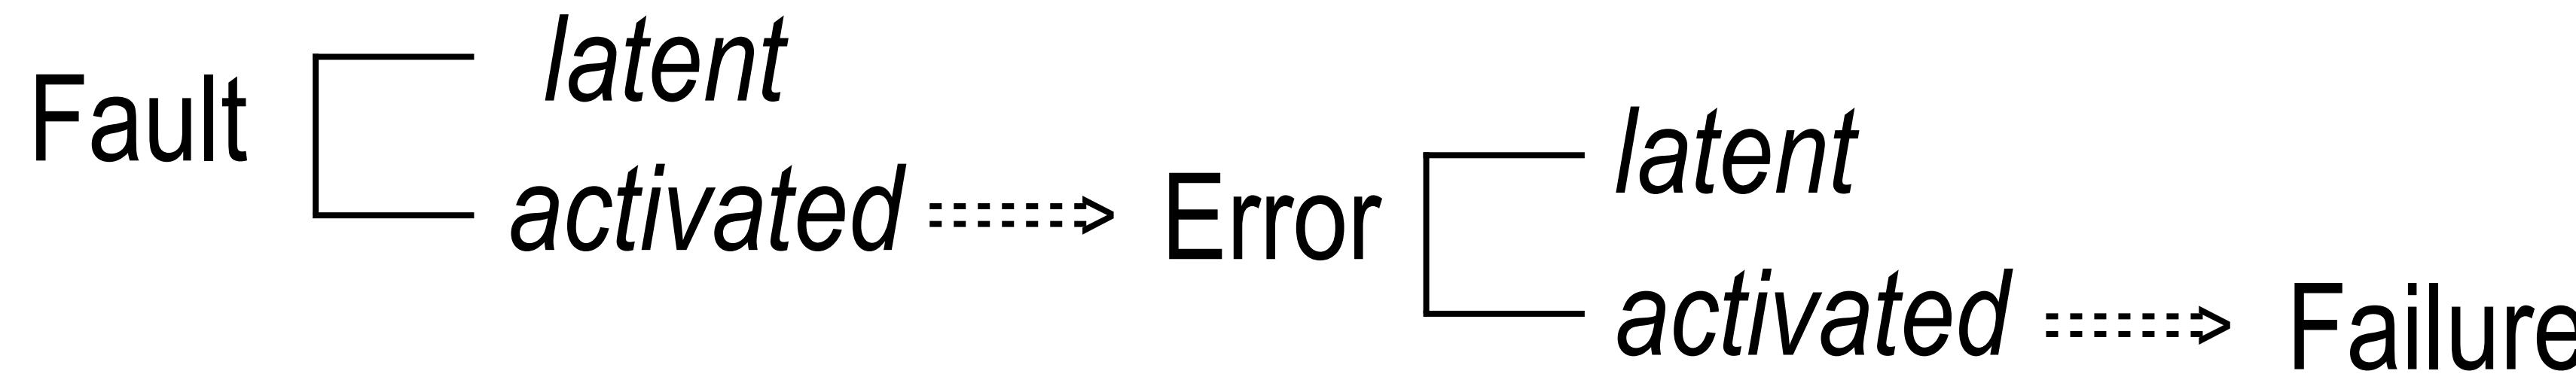
- Different types of software defects
 - *Bohrbug, Heisenbug, ...*

Recap: Fault tolerance



- Different types of software defects
 - *Bohrbug, Heisenbug, ...*
- Redundancy helps tolerate errors and failures
 - *Data redundancy, processing redundancy, ...*

Recap: Fault tolerance



- Different types of software defects
 - *Bohrbug, Heisenbug, ...*
- Redundancy helps tolerate errors and failures
 - *Data redundancy, processing redundancy, ...*
- Fault model = assumptions about what can vs. cannot go wrong

Safety-critical systems

- Safety critical = system whose failure may result in "bad" outcomes
 - *SCADA, aviation, space, automotive, healthcare, ...*

Safety-critical systems

- Safety critical = system whose failure may result in "bad" outcomes
 - *SCADA, aviation, space, automotive, healthcare, ...*
- Fail-safe = failure does not have "bad" consequences
 - *safety-critical \Rightarrow fail-safe*

Dependable systems

Dependable systems

- Availability = readiness for correct service

Dependable systems

- Availability = readiness for correct service
- Reliability = continuity of correct service

Dependable systems

- Availability = readiness for correct service
- Reliability = continuity of correct service
- Safety = absence of catastrophic consequences

Dependable systems

- Availability = readiness for correct service
- Reliability = continuity of correct service
- Safety = absence of catastrophic consequences
- Confidentiality = absence of unauthorized disclosure of information

Dependable systems

- Availability = readiness for correct service
- Reliability = continuity of correct service
- Safety = absence of catastrophic consequences
- Confidentiality = absence of unauthorized disclosure of information
- Integrity = absence of improper system state alterations

Dependable systems

- Availability = readiness for correct service
- Reliability = continuity of correct service
- Safety = absence of catastrophic consequences
- Confidentiality = absence of unauthorized disclosure of information
- Integrity = absence of improper system state alterations
- Maintainability = ability to undergo repairs and modifications

Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$

Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

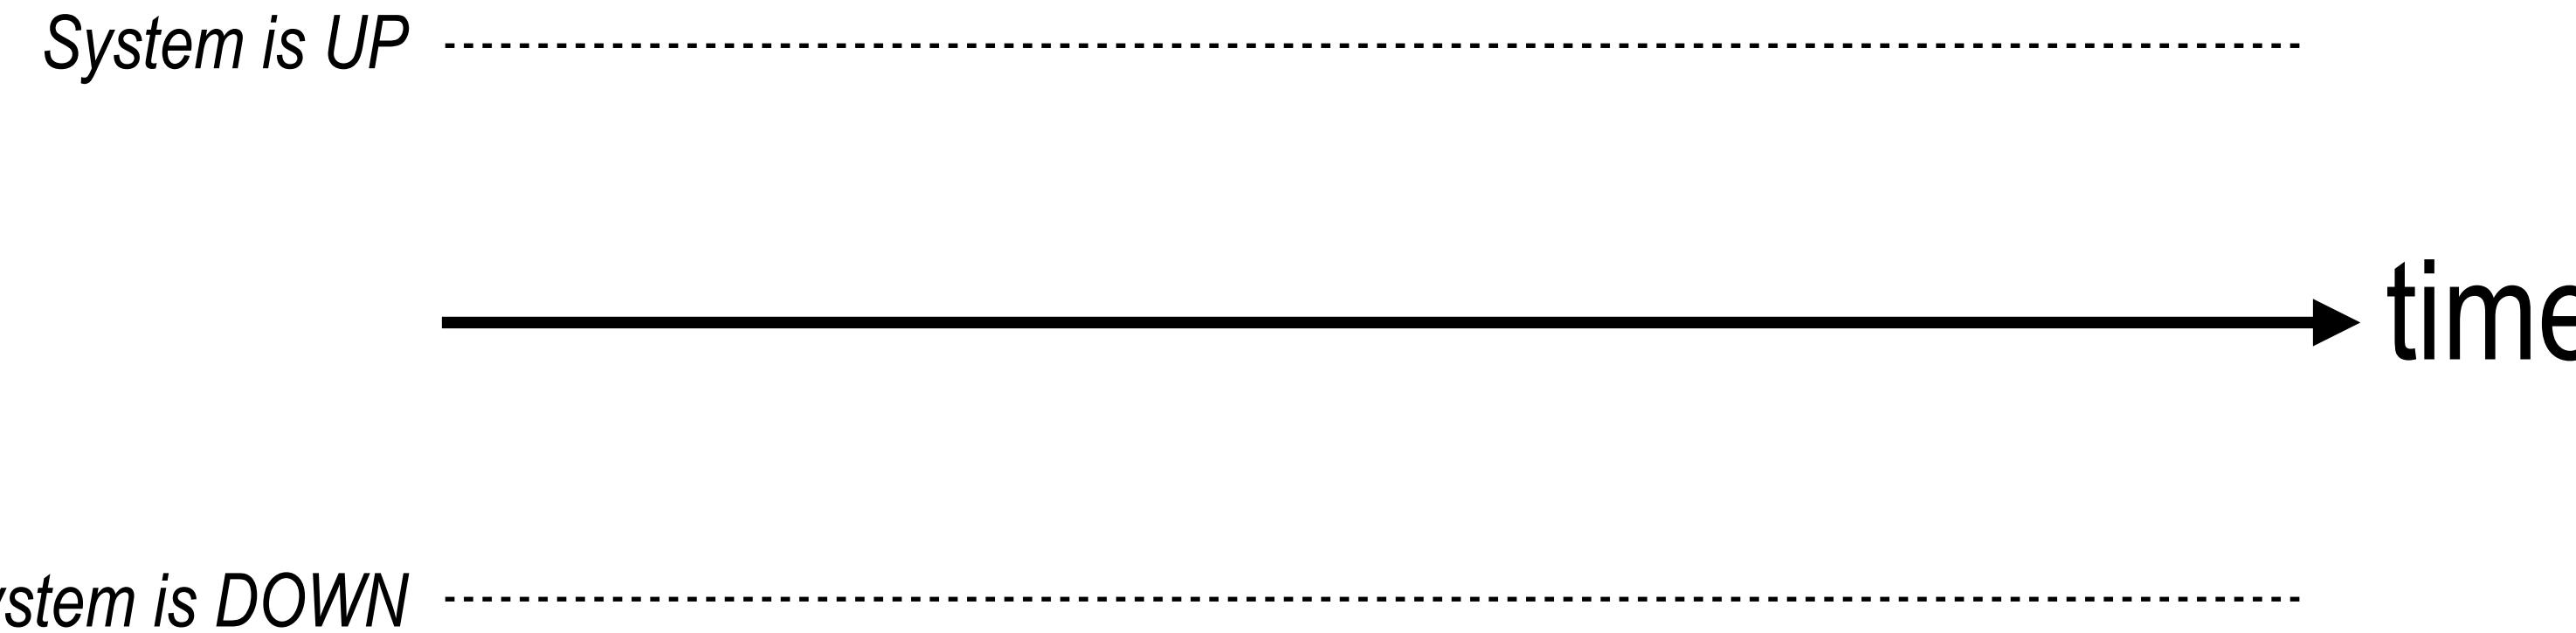
$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$



Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

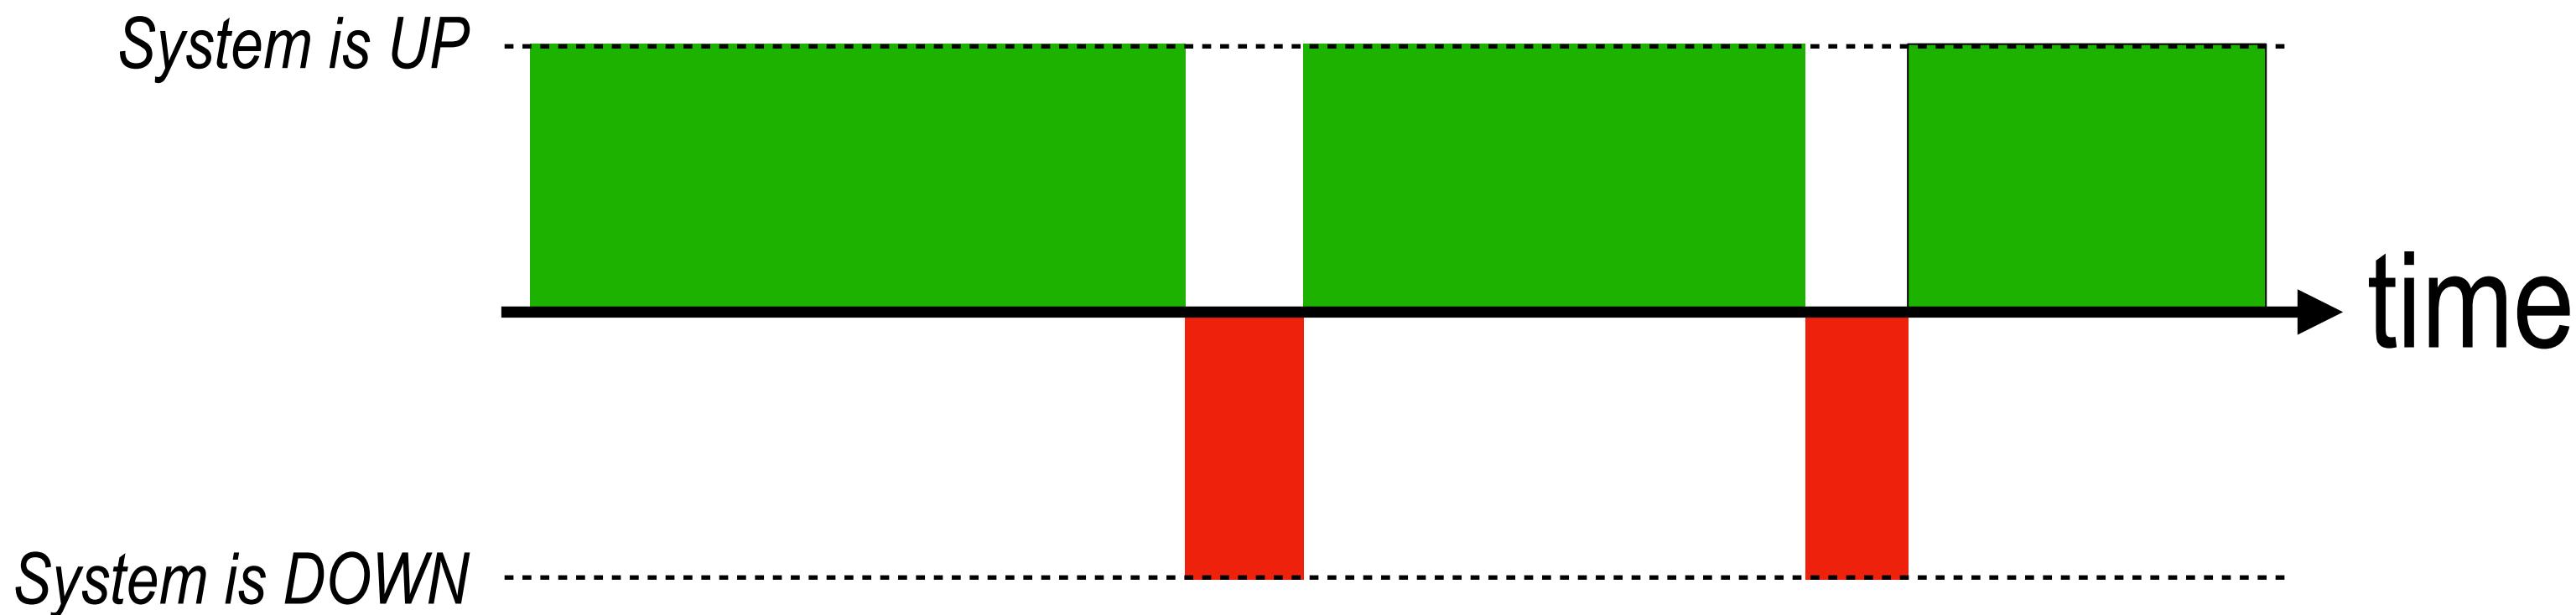
$$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$$



Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

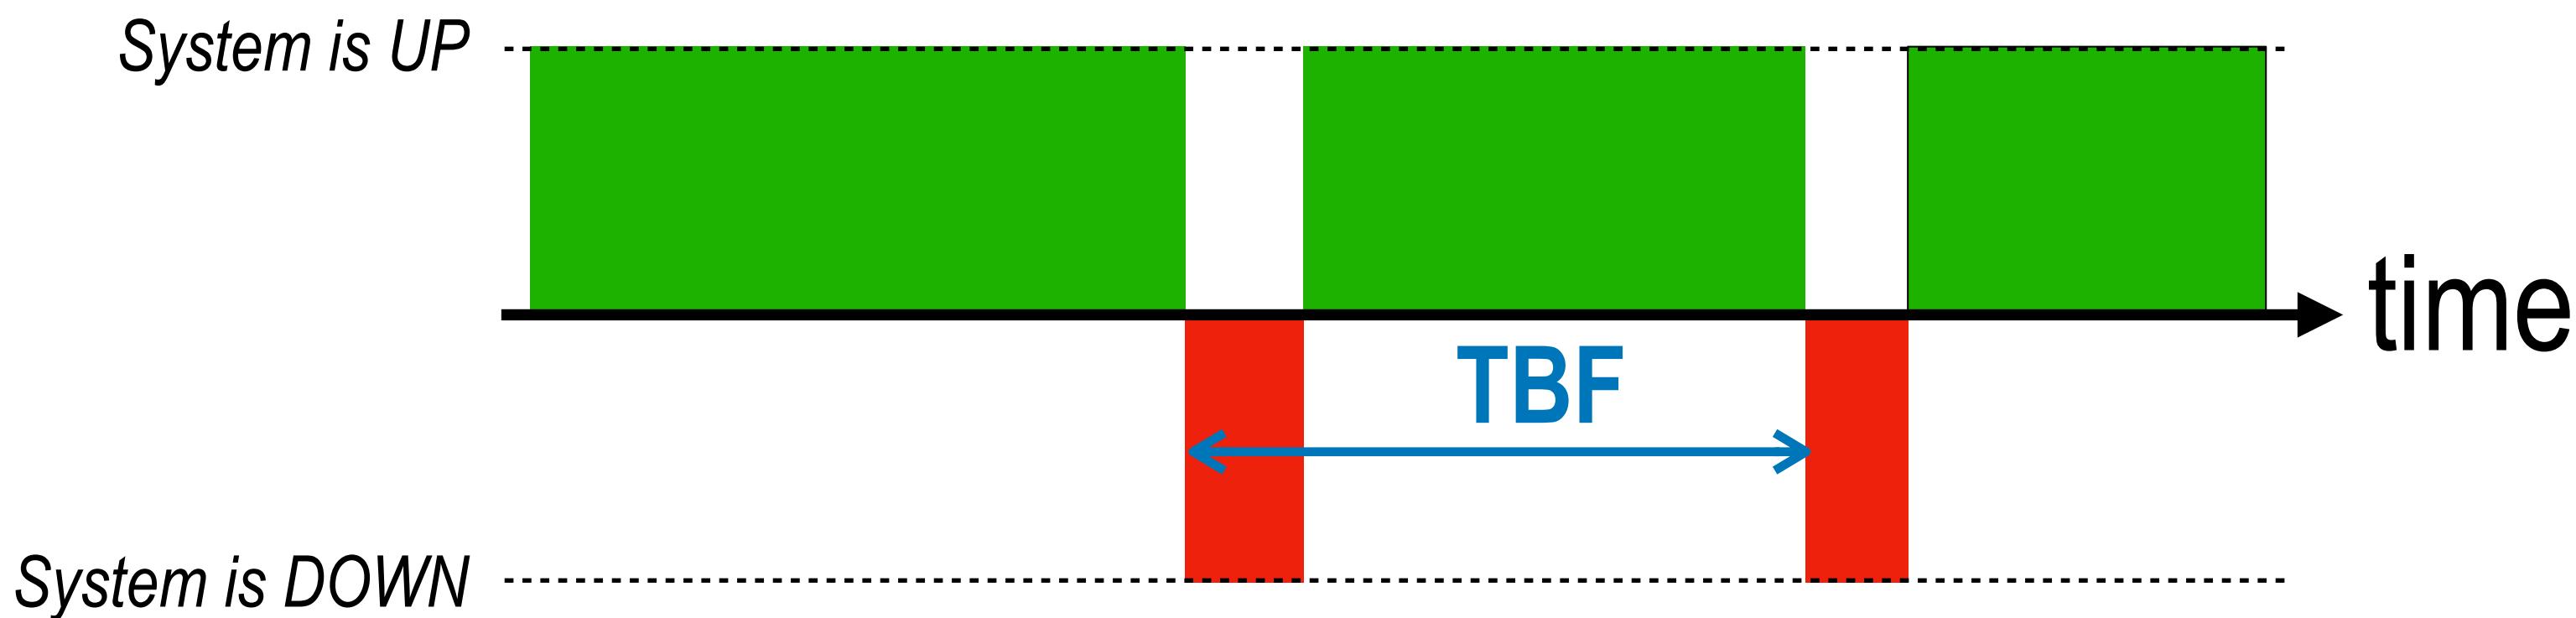
$$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$$



Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

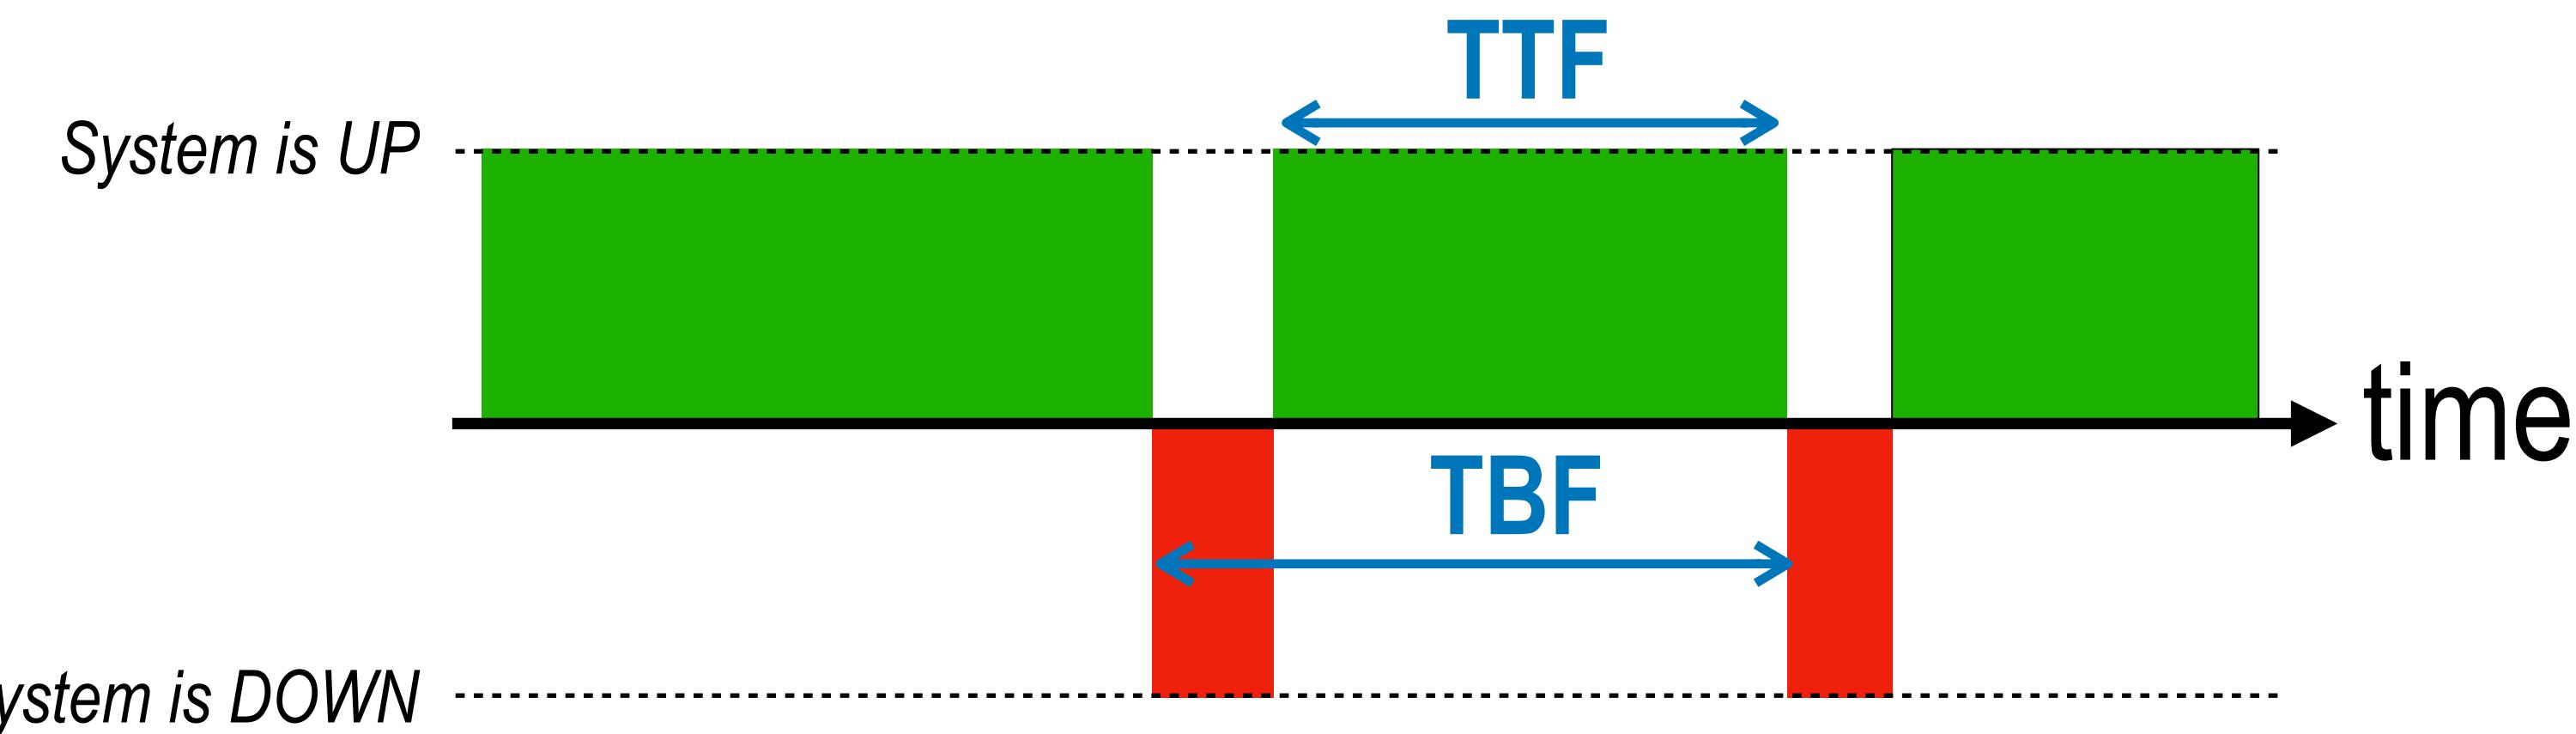
$$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$$



Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

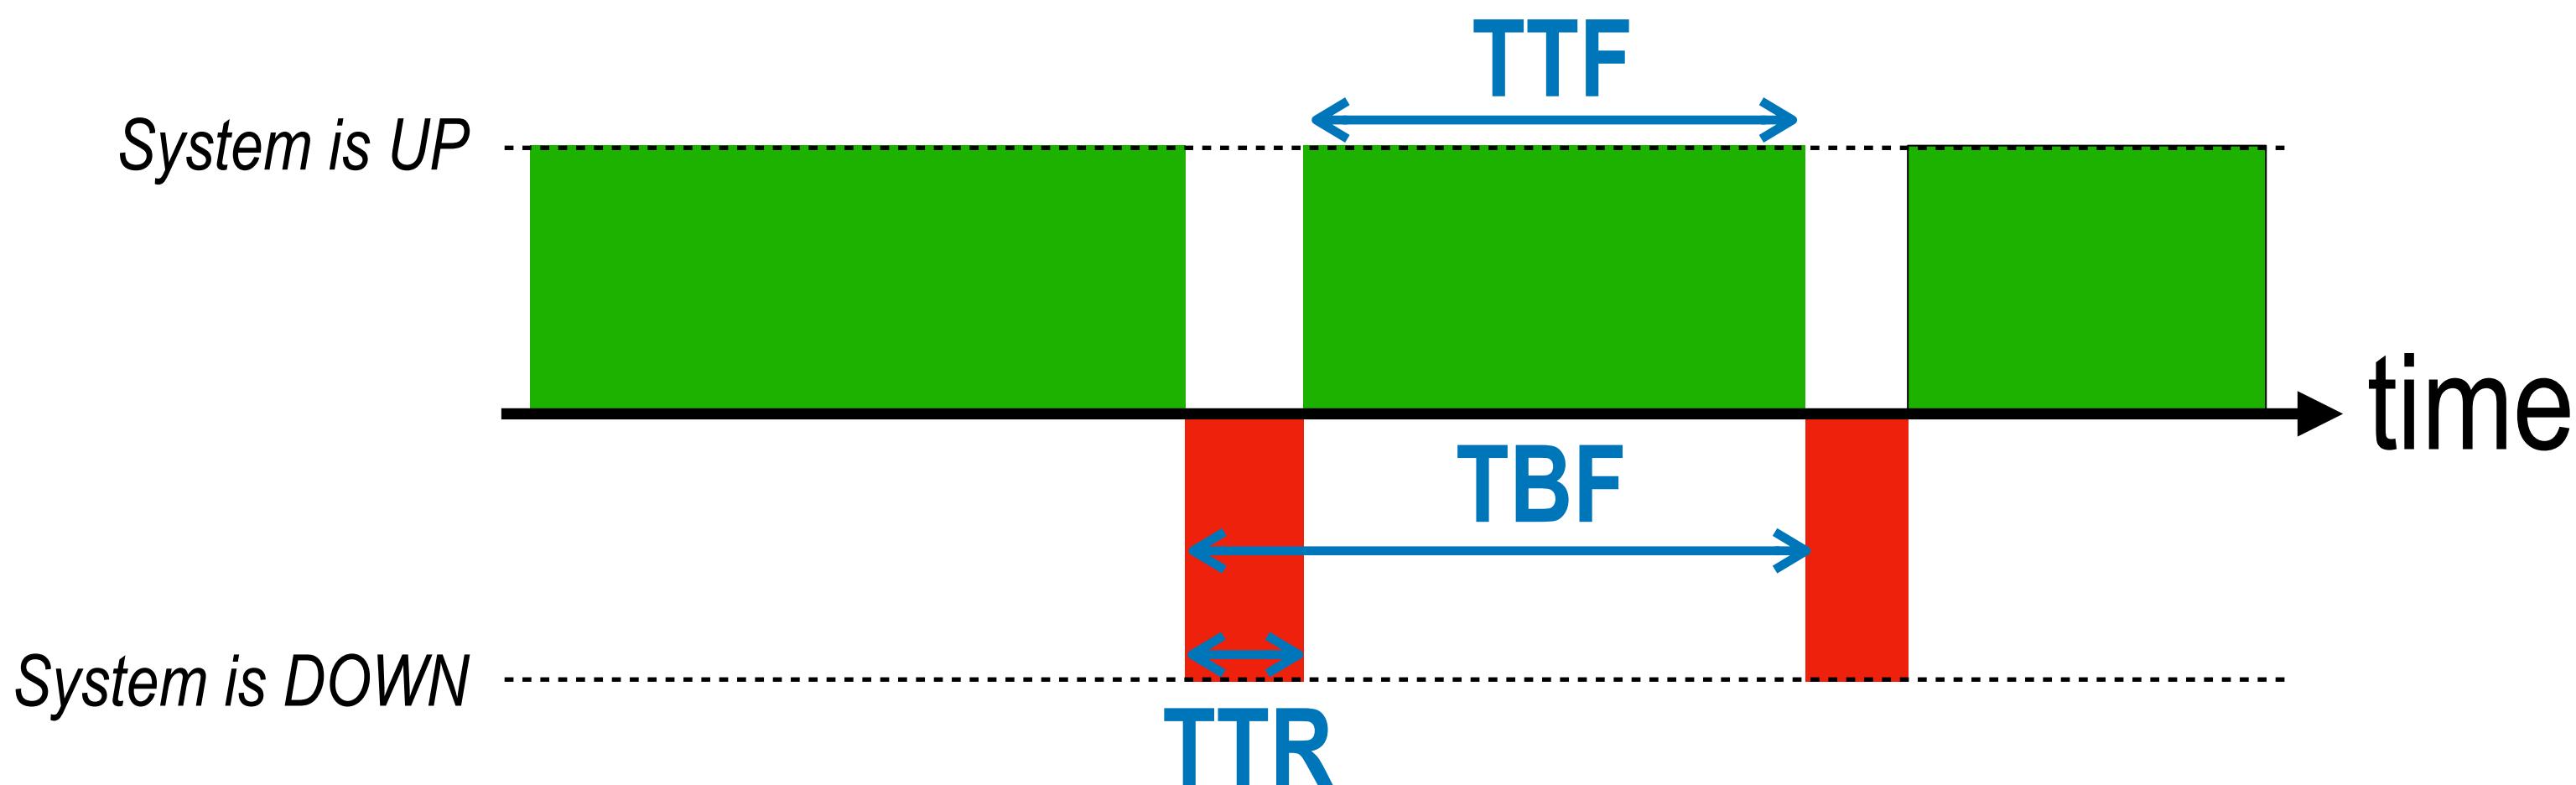
$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$



Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

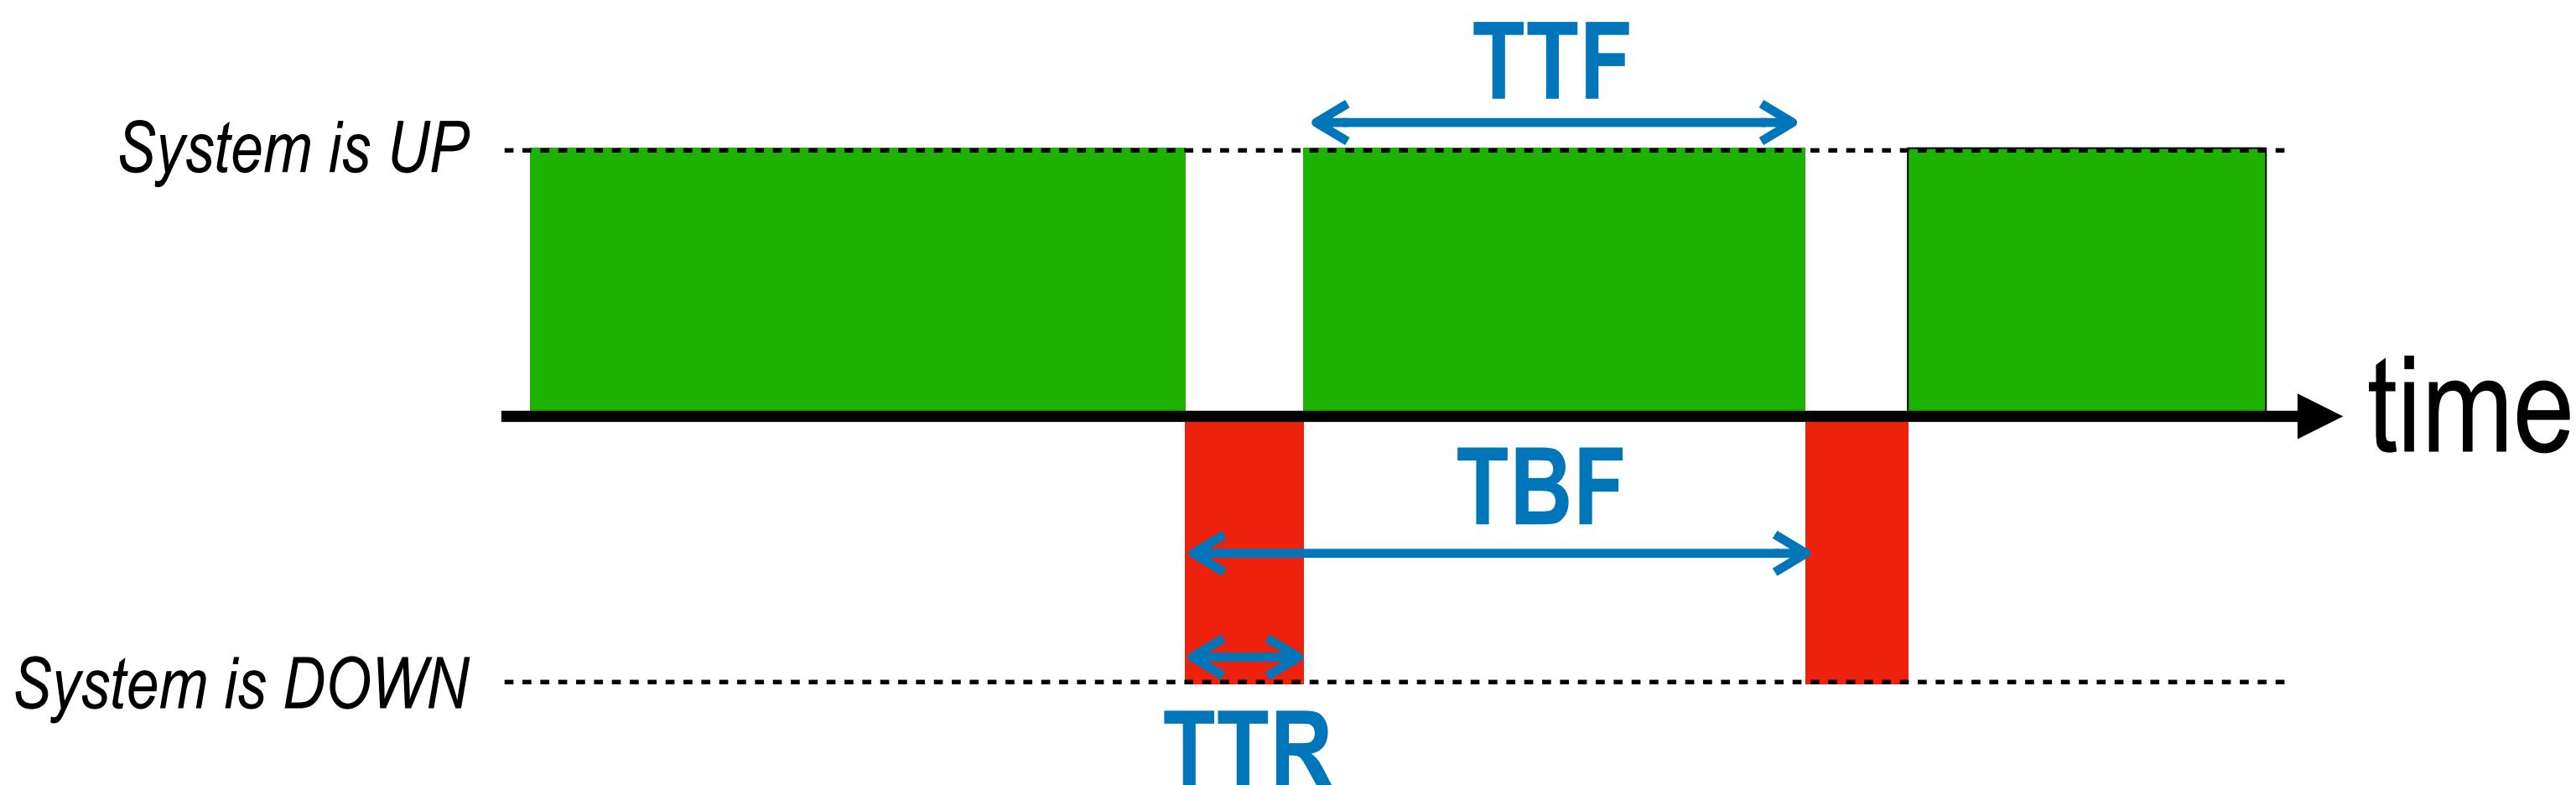
$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$



Reliability

- Reliability = probability of continuous operation
 - *continuous operation* = (correctly) producing outputs in response to inputs

$R(t) = P(\text{module operates correctly at time } t \mid \text{it was operating correctly at } t=0)$



$$\text{MTBF} = \text{MTTF} + \text{MTTR}$$

Measuring reliability

- In general MTBF or MTTF ($MTBF = MTTF + MTTR$)
 - Specifics: Example from SSD spec sheet: *P/E cycles, TBW, GB/day, DWPD, MTBF ...*
- Example: Samsung SSD 850 Pro SATA
 - *Warranty period = 10 years*
 - *MTBF = 2M hours (228 years)*

Measuring reliability

- In general MTBF or MTTF ($MTBF = MTTF + MTTR$)
 - Specifics: Example from SSD spec sheet: P/E cycles, TBW, GB/day, DWPD, MTBF ...
 - Example: Samsung SSD 850 Pro SATA
 - Warranty period = 10 years
 - $MTBF = 2M$ hours (228 years)
- Why different???*

Measuring reliability

- In general MTBF or MTTF ($MTBF = MTTF + MTTR$)
 - Specifics: Example from SSD spec sheet: P/E cycles, TBW, GB/day, DWPD, MTBF ...
- Example: Samsung SSD 850 Pro SATA
 - Warranty period = 10 years
 - $MTBF = 2M$ hours (228 years)
 - assumes operation of 8 hrs/day

Why different???

Measuring reliability

- In general MTBF or MTTF ($MTBF = MTTF + MTTR$)
 - Specifics: Example from SSD spec sheet: P/E cycles, TBW, GB/day, DWPD, MTBF ...
 - Example: Samsung SSD 850 Pro SATA
 - Warranty period = 10 years
 - $MTBF = 2M$ hours (228 years)
 - assumes operation of 8 hrs/day
 - 2.5K SSDs => you'd experience 1 failure every ~100 days ($2M / 8 / 2500$)
- Why different???*

Recap: Reliability

- Dependability = Reliability + Availability + Safety + ...
- Safety-critical vs. reliable
- MTBF = MTTF + MTTR

Availability

- Availability = probability of producing (correct) outputs in response to inputs

Availability

- Availability = probability of producing (correct) outputs in response to inputs

Level of Availability	Percent of Uptime	Downtime per Year	Downtime per Day
1 Nine	90%	36.5 days	2.4 hrs.
2 Nines	99%	3.65 days	14 min.
3 Nines	99.9%	8.76 hrs.	86 sec.
4 Nines	99.99%	52.6 min.	8.6 sec.
5 Nines	99.999%	5.25 min.	.86 sec.
6 Nines	99.9999%	31.5 sec.	8.6 msec

Availability

- Availability = probability of producing (correct) outputs in response to inputs

Level of Availability	Percent of Uptime	Downtime per Year	Downtime per Day
1 Nine	90%	36.5 days	2.4 hrs.
2 Nines	99%	3.65 days	14 min.
3 Nines	99.9%	8.76 hrs.	86 sec.
4 Nines	99.99%	52.6 min.	8.6 sec.
5 Nines	99.999%	5.25 min.	.86 sec.
6 Nines	99.9999%	31.5 sec.	8.6 msec

Availability

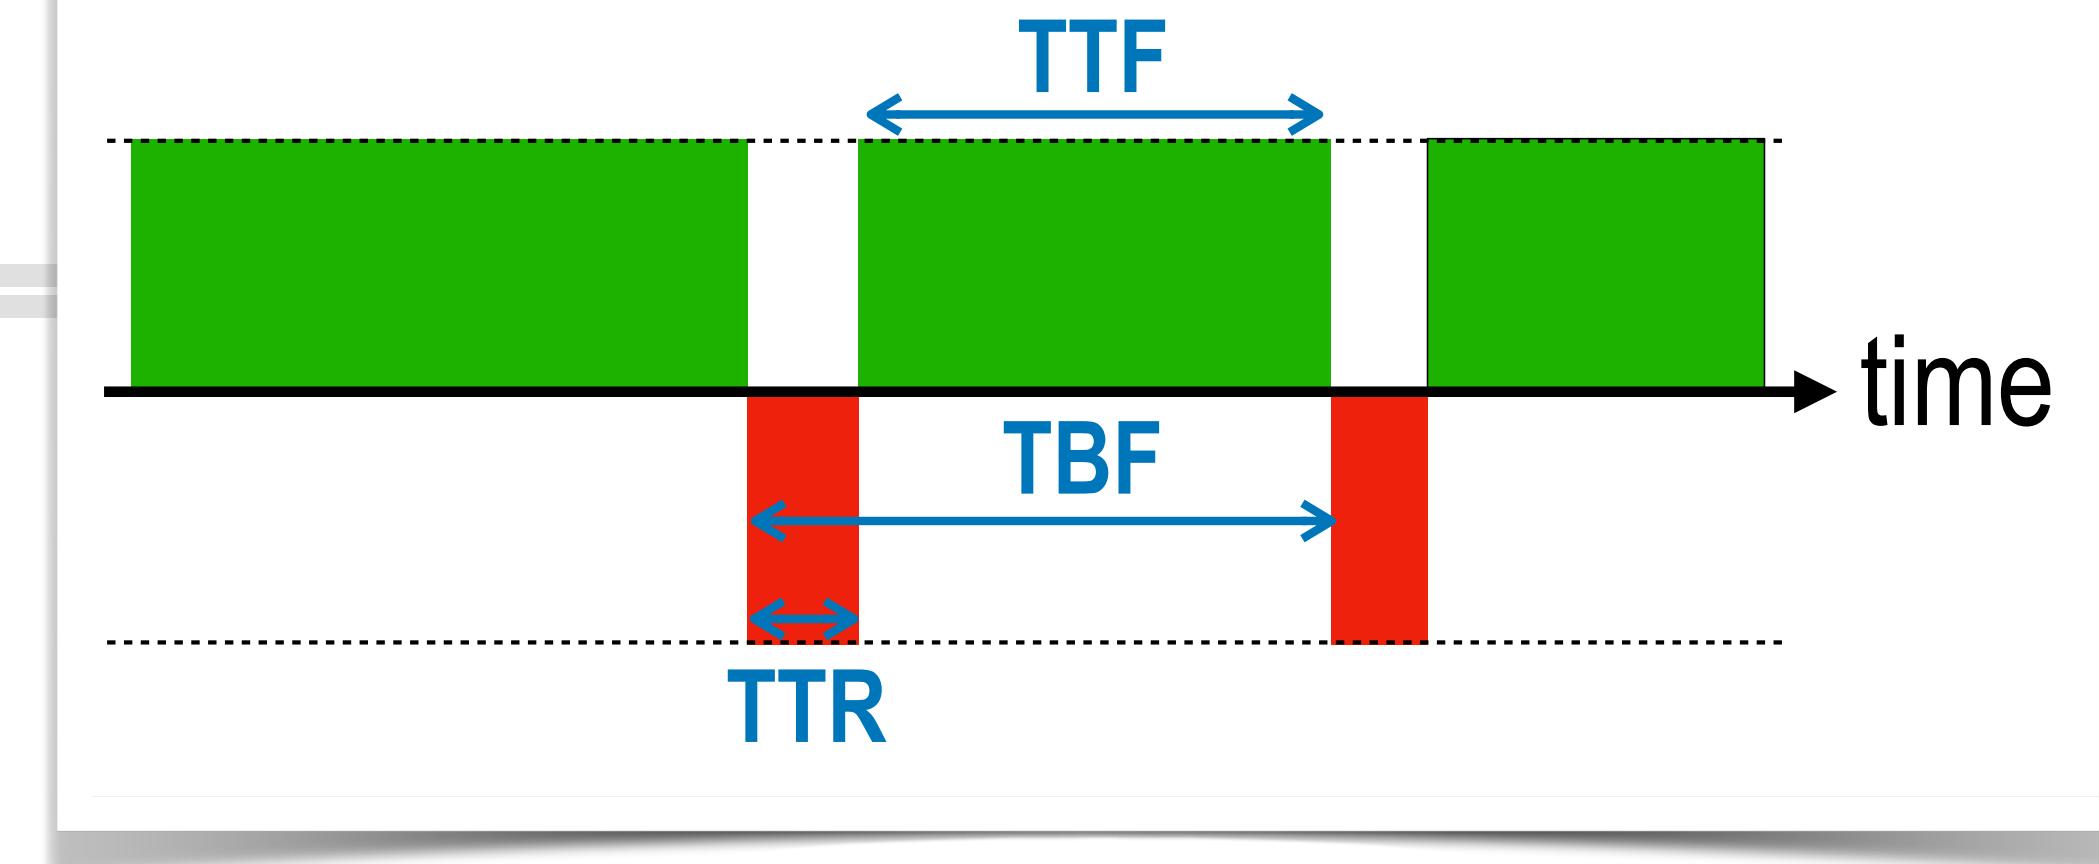
- Availability = probability of producing (correct) outputs in response to inputs

Level of Availability	Percent of Uptime	Downtime per Year	Downtime per Day
1 Nine	90%	36.5 days	2.4 hrs.
2 Nines	99%	3.65 days	14 min.
3 Nines	99.9%	8.76 hrs.	86 sec.
4 Nines	99.99%	52.6 min. ↑×10	.6 sec.
5 Nines	99.999%	5.25 min.	.86 sec.
6 Nines	99.9999%	31.5 sec. ↓÷10	.6 msec

Availability vs. Reliability

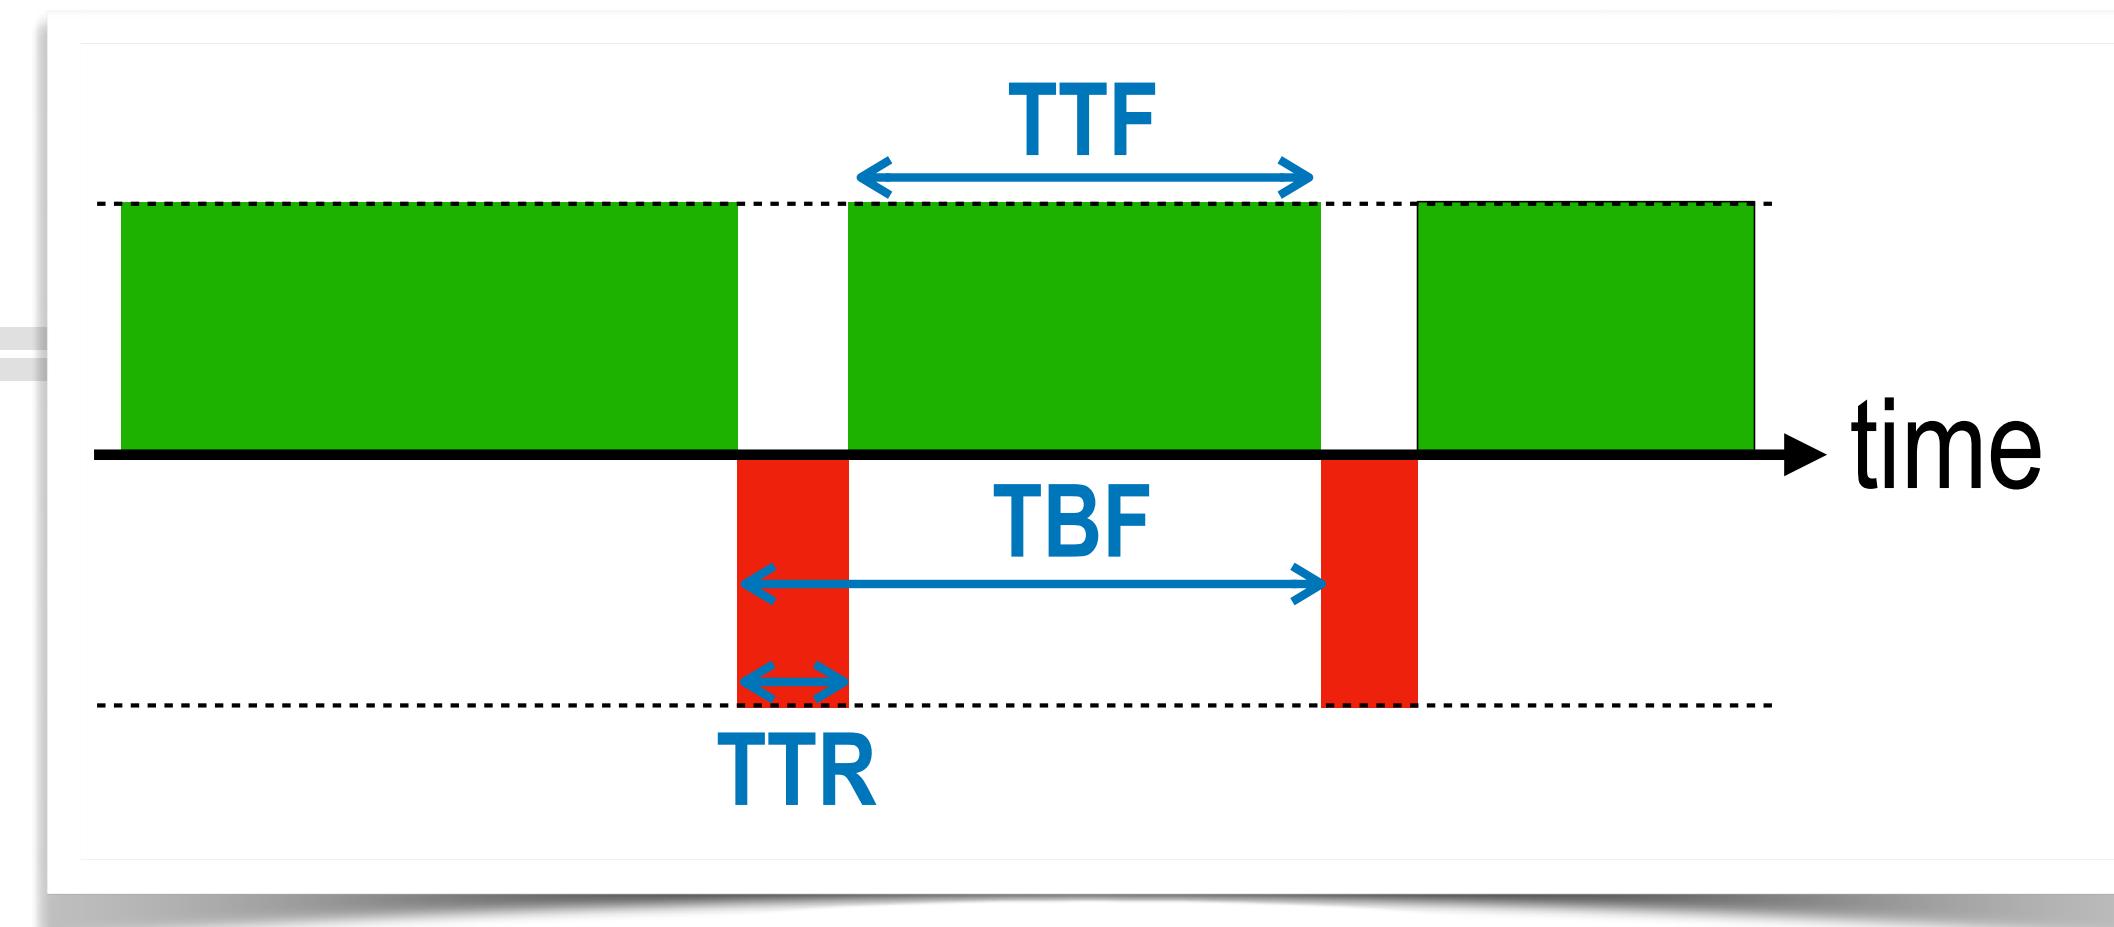
- Continuity of service does not matter (unlike reliability)
 - *In theory: uptime is too strict a measure of availability*
 - *In practice: what's the difference?*
- Uptime => availability but Availability \Rightarrow uptime
- Examples of ...
 - *Highly available systems with poor reliability (and how is redundancy used)*
...
 - *Highly reliable systems with poor availability (and how is redundancy used)*
...

System availability



System availability

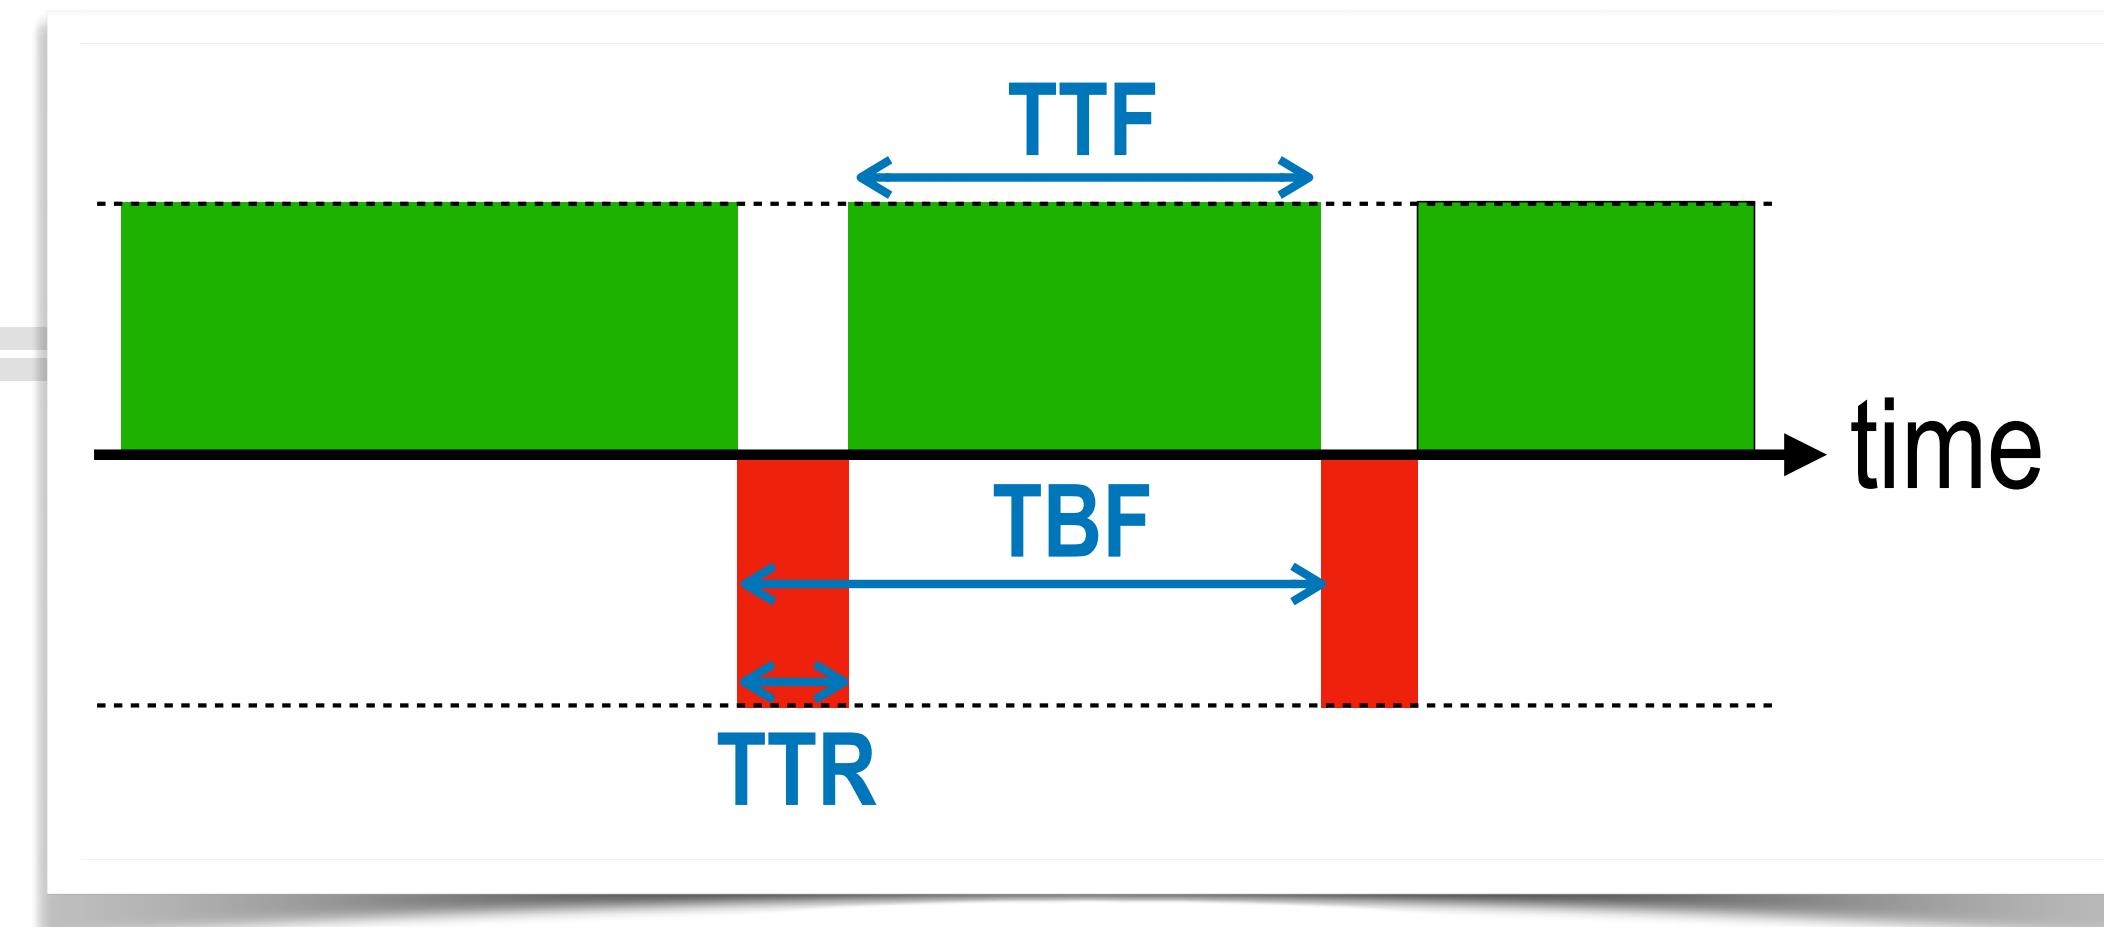
$$\text{Availability} = \frac{\text{MTTF}}{\text{MTBF}}$$



System availability

$$\text{Availability} = \frac{\text{MTTF}}{\text{MTBF}}$$

$$\text{Unavailability} = 1 - \text{Availability} = \frac{\text{MTTR}}{\text{MTBF}}$$

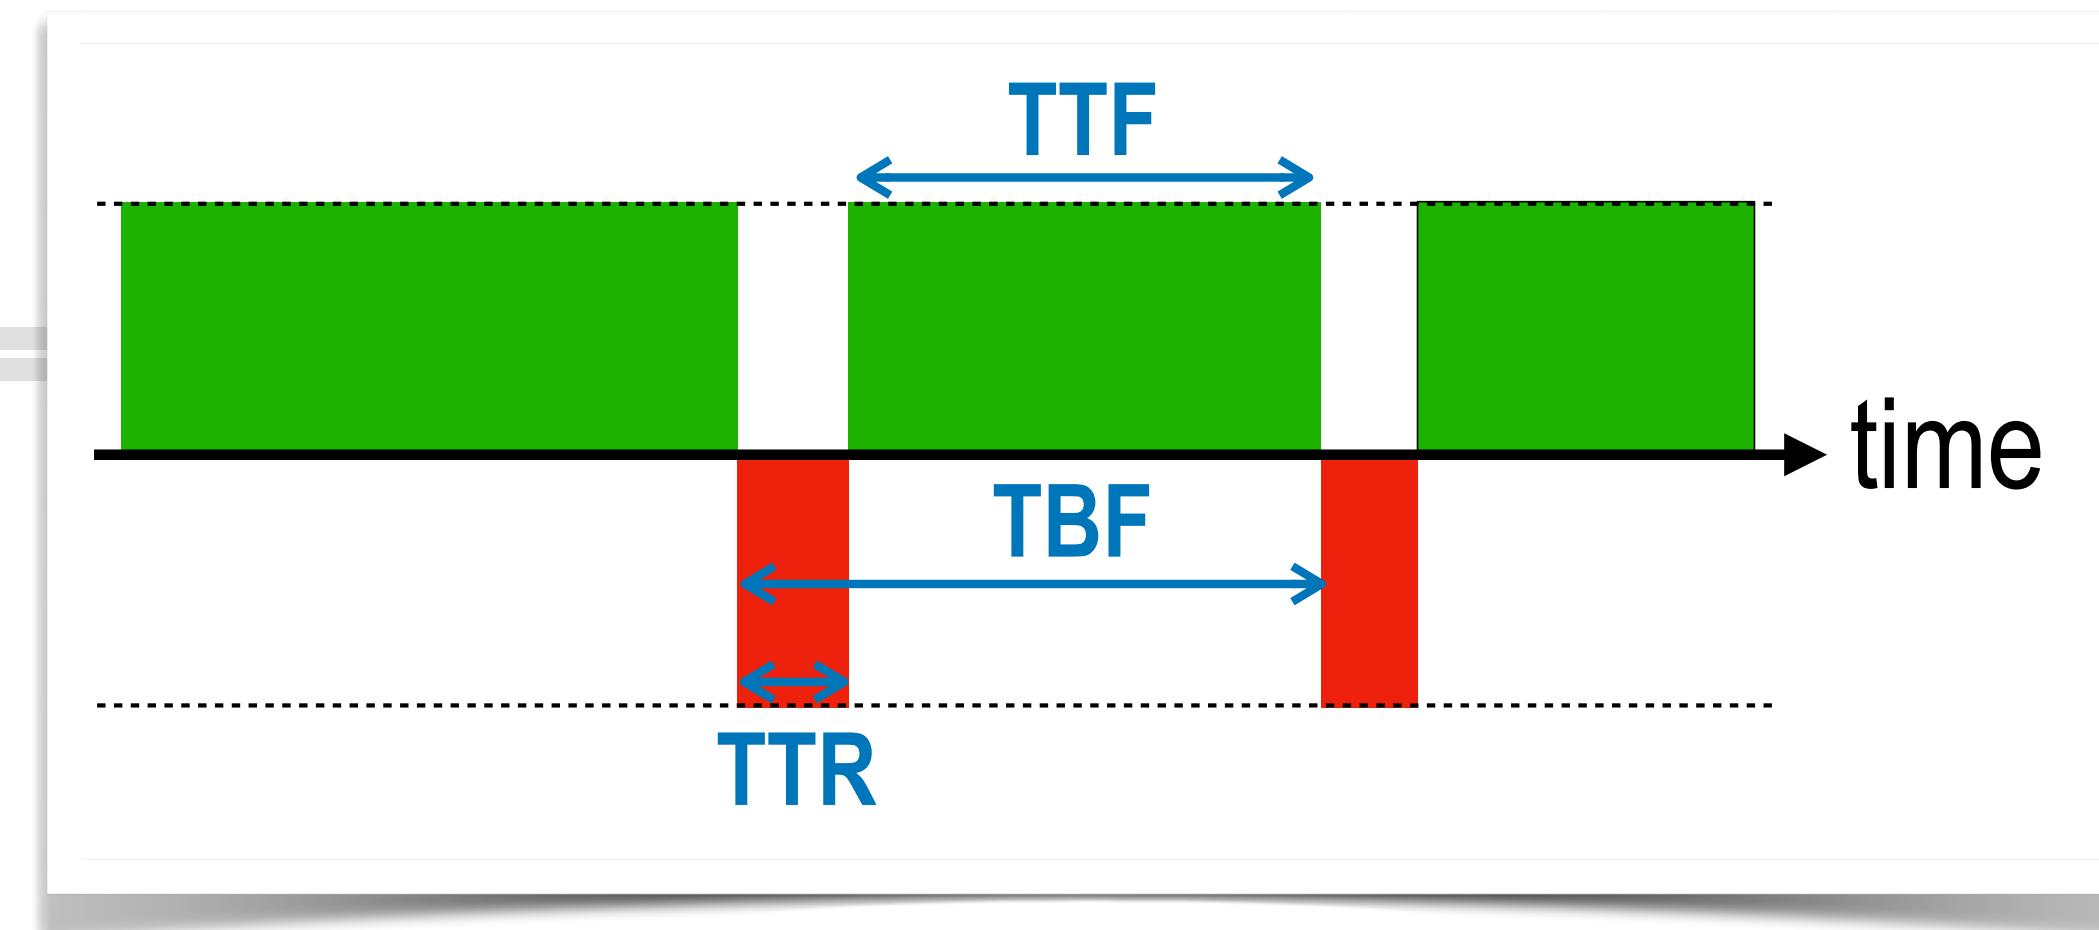


System availability

$$\text{Availability} = \frac{\text{MTTF}}{\text{MTBF}}$$

$$\text{Unavailability} = 1 - \text{Availability} = \frac{\text{MTTR}}{\text{MTBF}}$$

$$\text{MTBF} = \text{MTTF} + \text{MTTR} \approx \text{MTTF} \text{ (if MTTF} \gg \text{MTTR)}$$

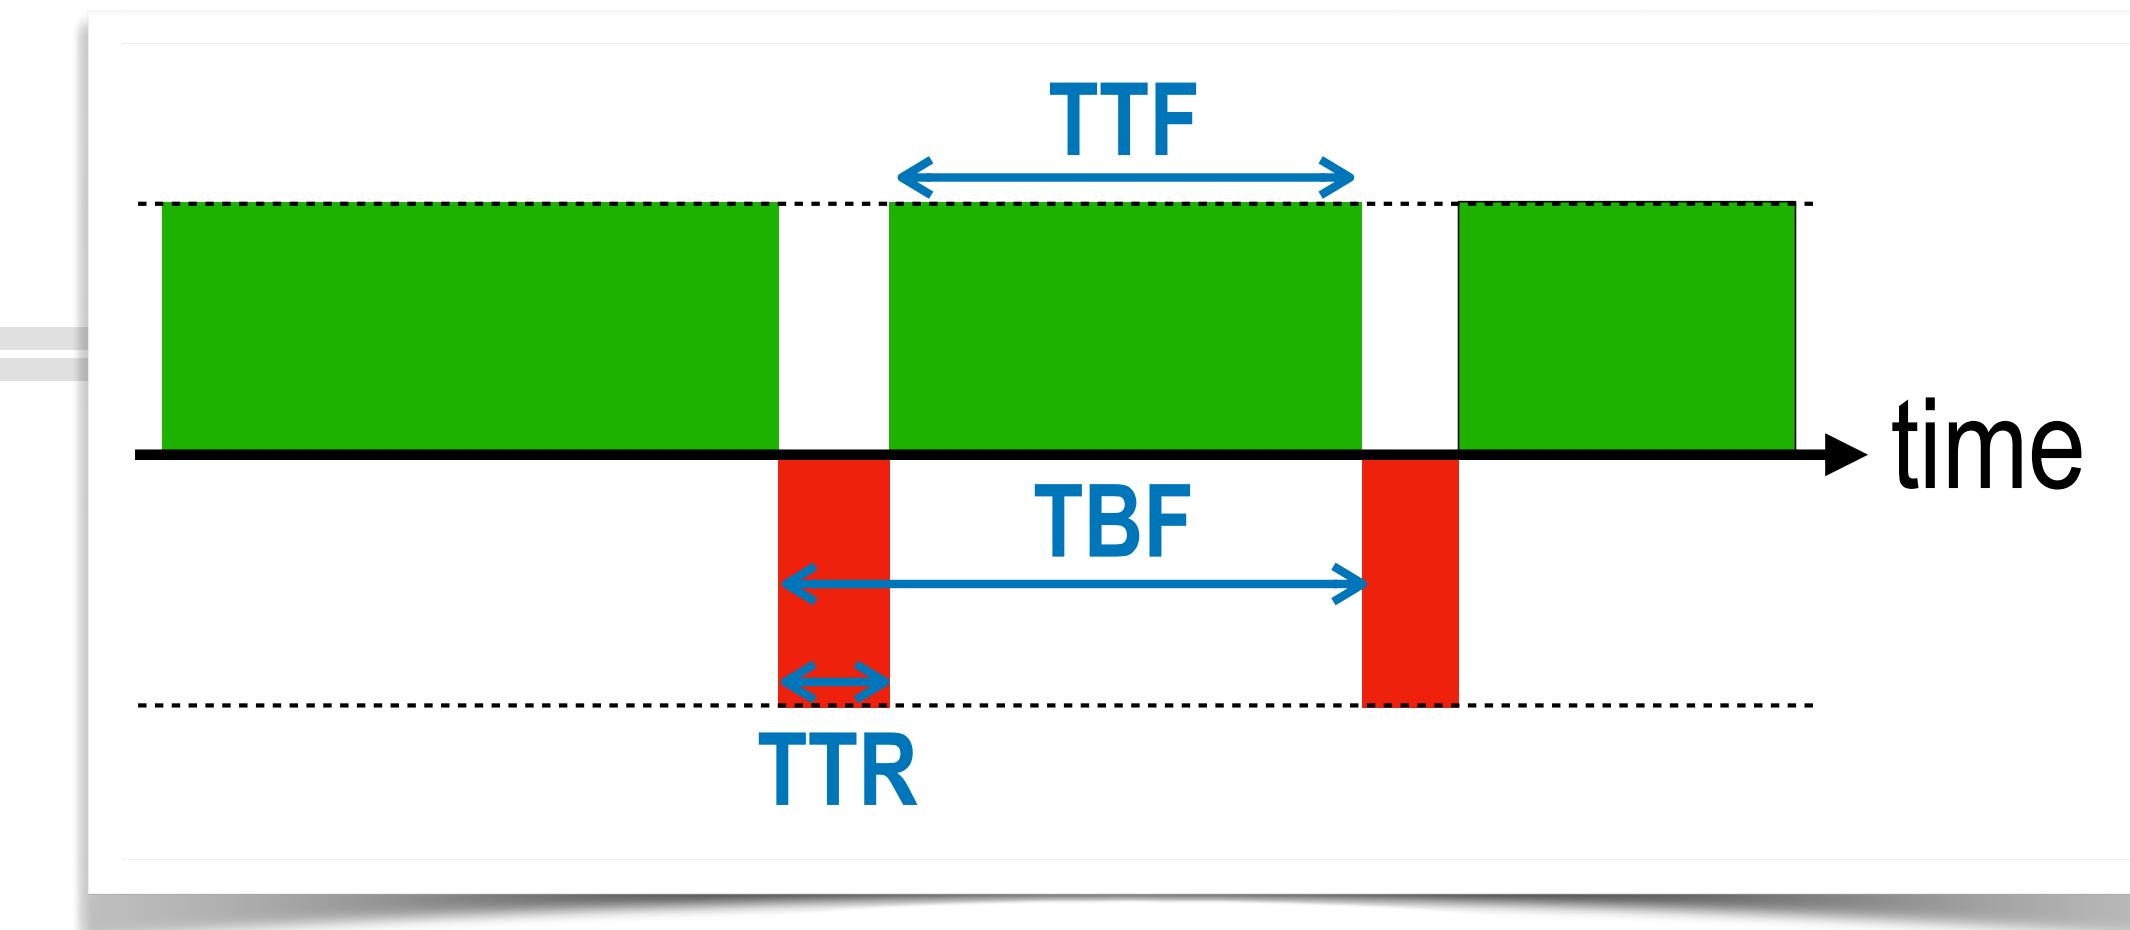


System availability

$$\text{Availability} = \frac{\text{MTTF}}{\text{MTBF}}$$

$$\text{Unavailability} = 1 - \text{Availability} = \frac{\text{MTTR}}{\text{MTBF}}$$

$$\text{MTBF} = \text{MTTF} + \text{MTTR} \approx \text{MTTF} \text{ (if } \text{MTTF} \gg \text{MTTR})$$



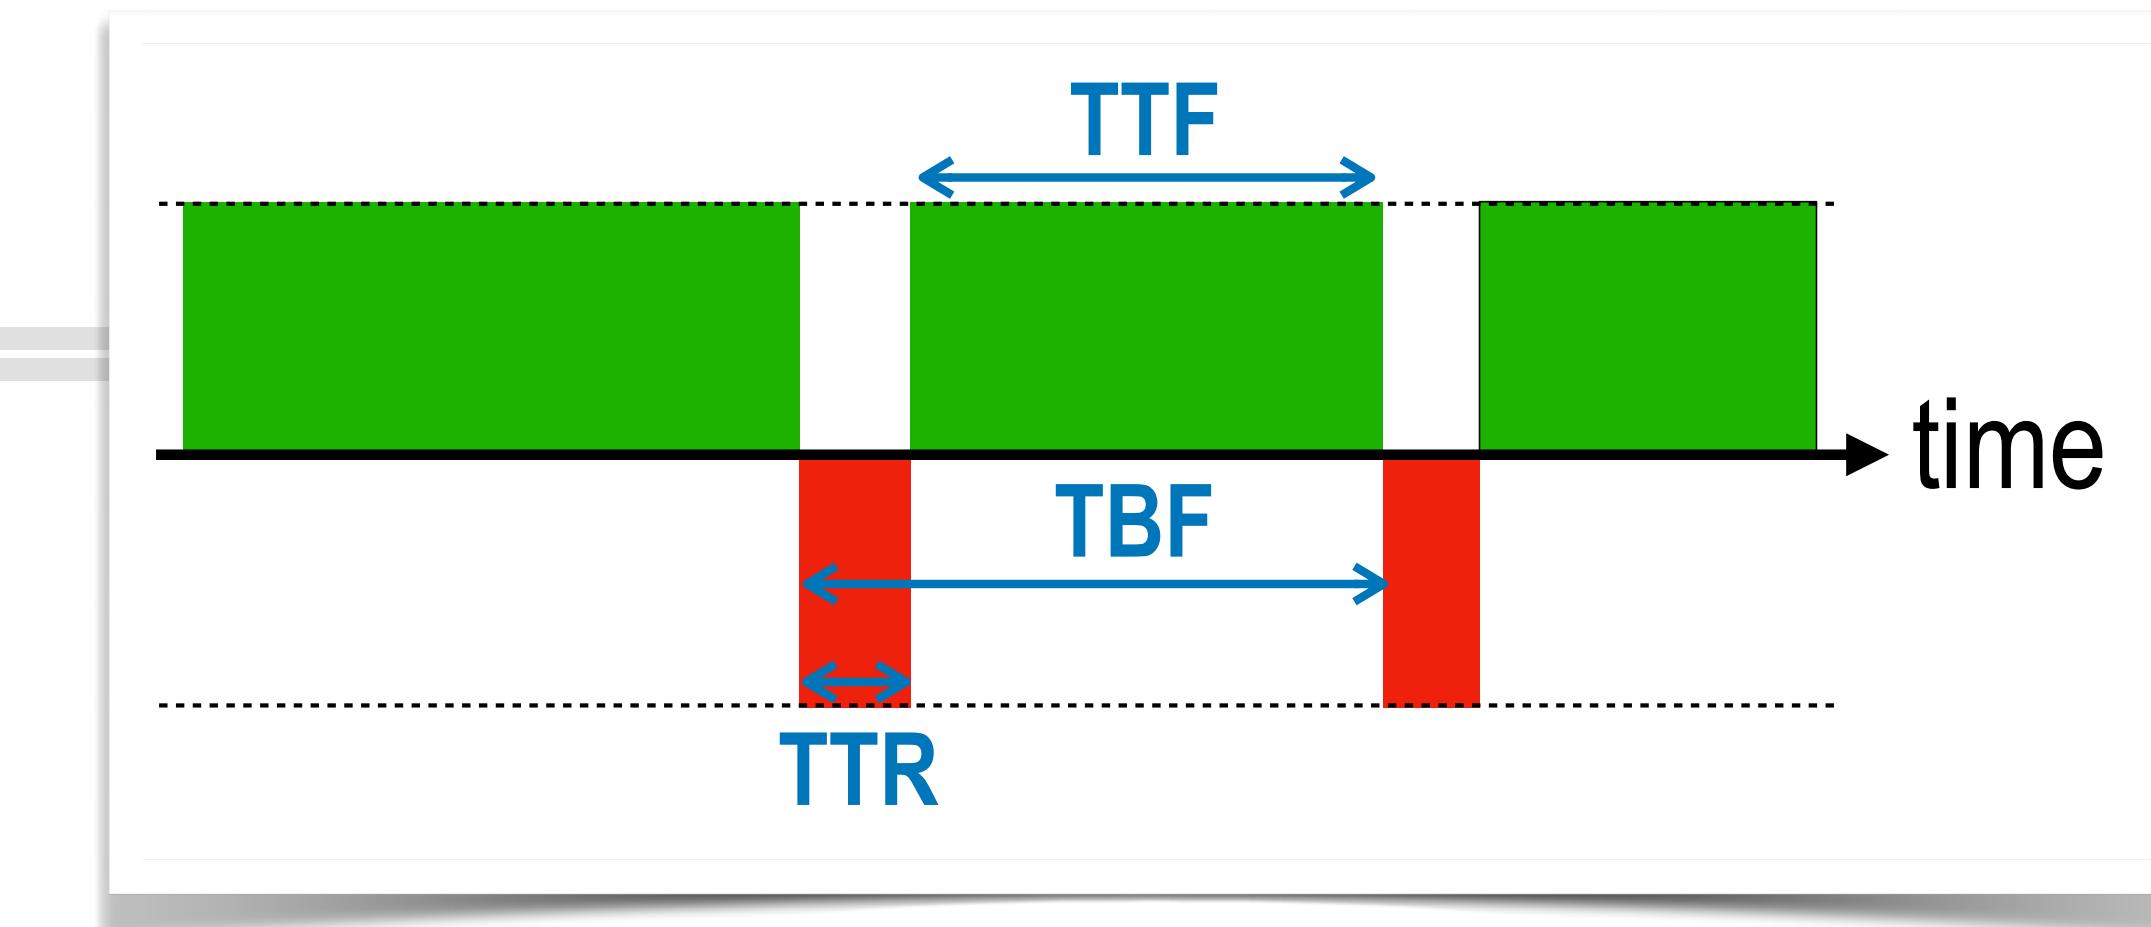
$$\text{Unavailability} \approx \frac{\text{MTTR}}{\text{MTTF}}$$

System availability

$$\text{Availability} = \frac{\text{MTTF}}{\text{MTBF}}$$

$$\text{Unavailability} = 1 - \text{Availability} = \frac{\text{MTTR}}{\text{MTBF}}$$

$$\text{MTBF} = \text{MTTF} + \text{MTTR} \approx \text{MTTF} \text{ (if MTTF} \gg \text{MTTR)}$$



$$\text{Unavailability} \approx \frac{\text{MTTR}}{\text{MTTF}}$$

- Increase availability by
 - *increasing MTTF (higher reliability)*
 - *reducing MTTR (faster recovery)*

Failure modes

Failure modes

- Definition:

When a system fails, how does that failure appear at the interface of a component?
- Four kinds
 - *fail-stop*
 - *fail-fast*
 - *fail-safe*
 - *fail-soft*

Failure mode 1: Fail-stop

- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*

Failure mode 1: Fail-stop

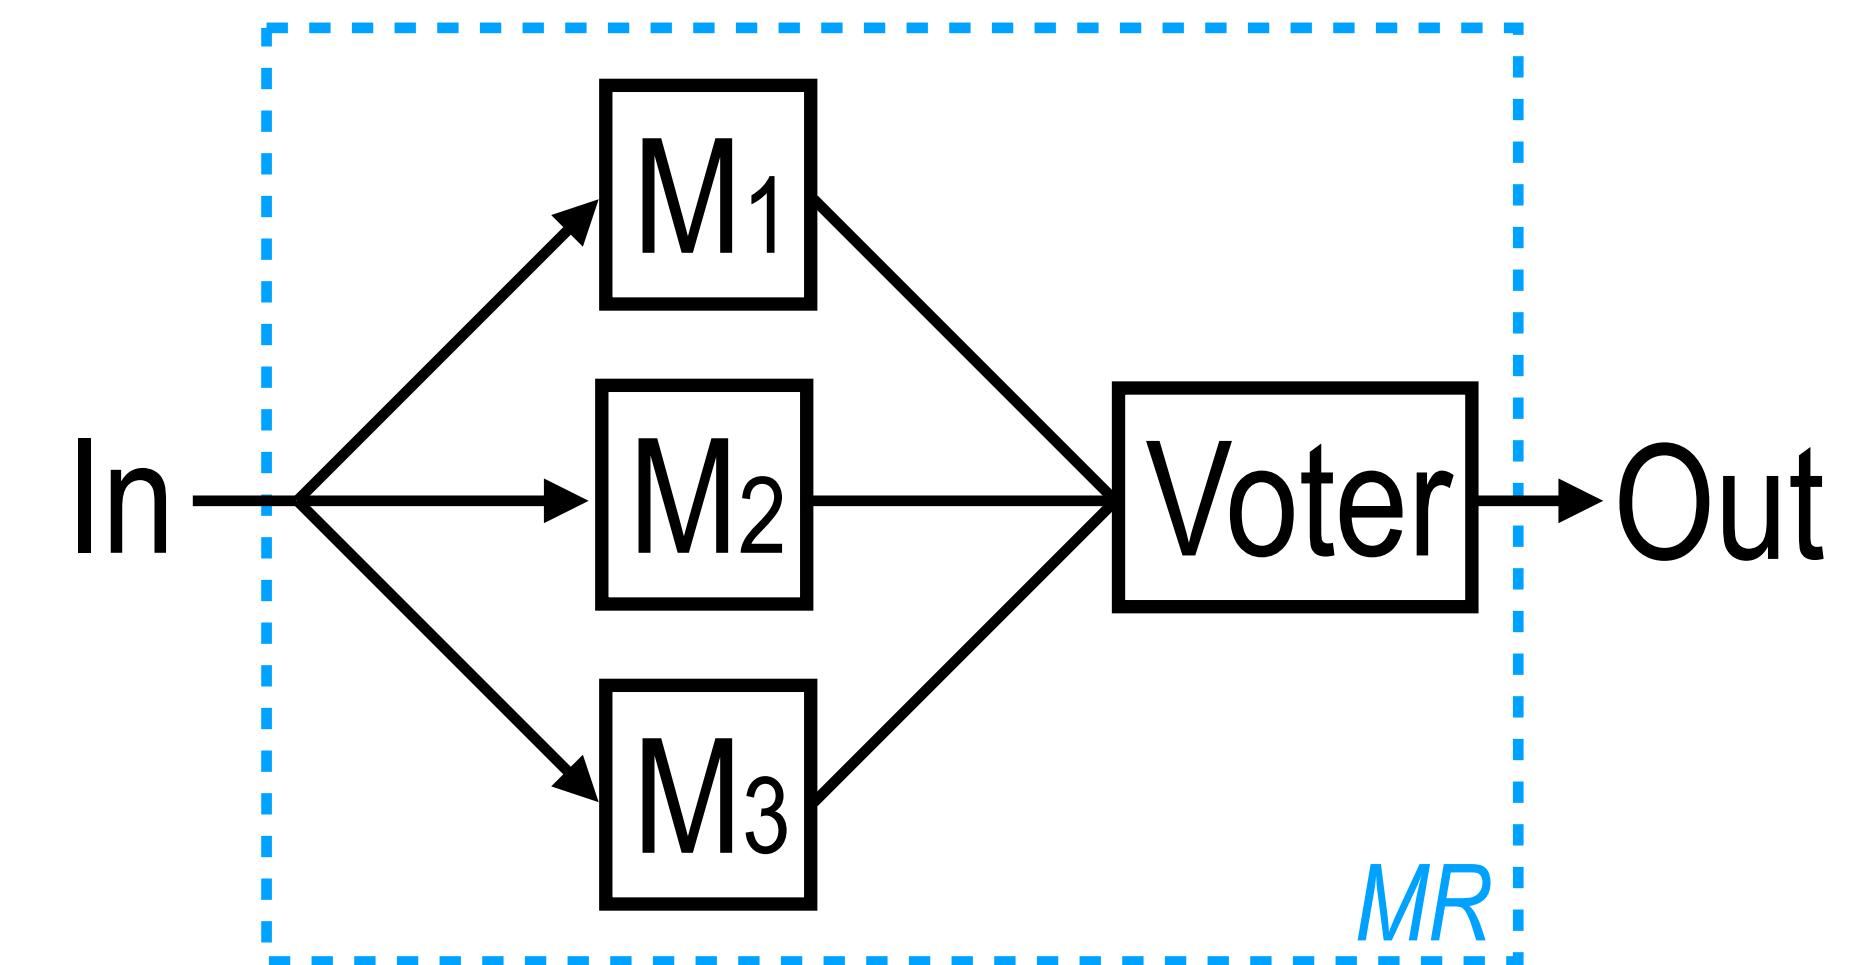
- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*
- Any system can be made fail-stop with simple modular redundancy (MR)

Failure mode 1: Fail-stop

- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*
- Any system can be made fail-stop with simple modular redundancy (MR)
 - *Strict fault model: voter is reliable*

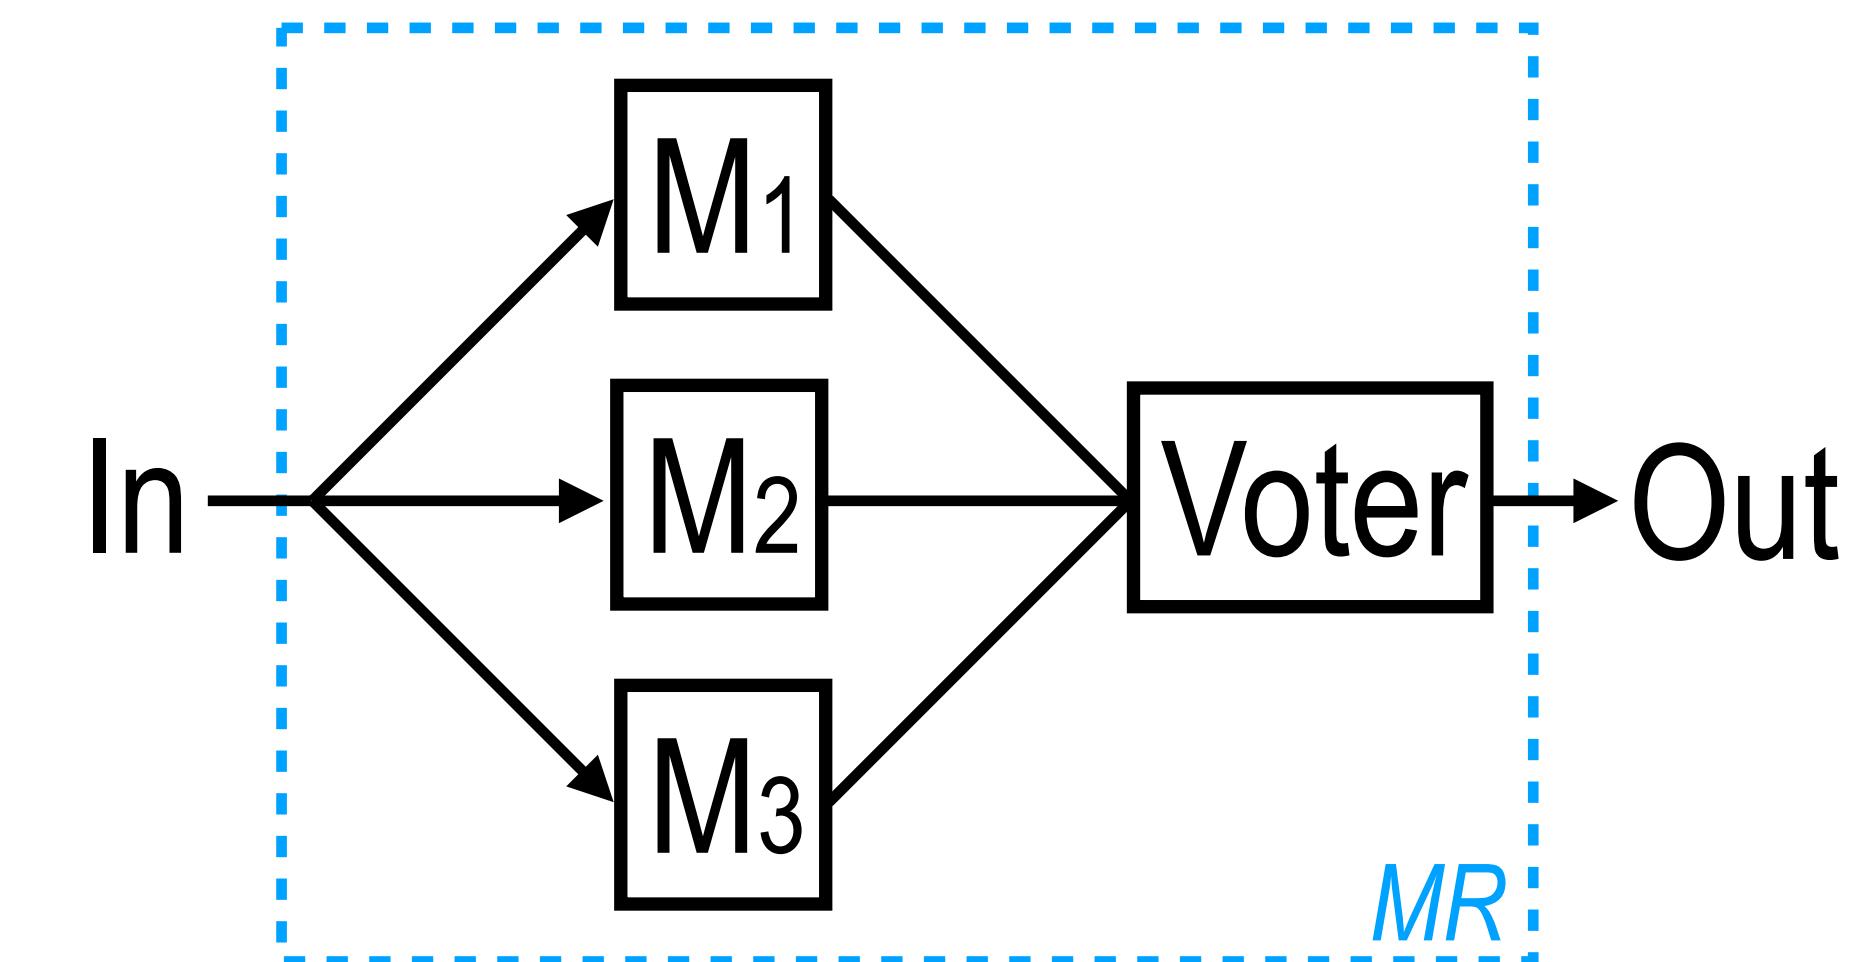
Failure mode 1: Fail-stop

- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*
- Any system can be made fail-stop with simple modular redundancy (MR)
 - *Strict fault model: voter is reliable*
 - *Tolerate failures*
 - level of redundancy: $2f+1$ independent modules



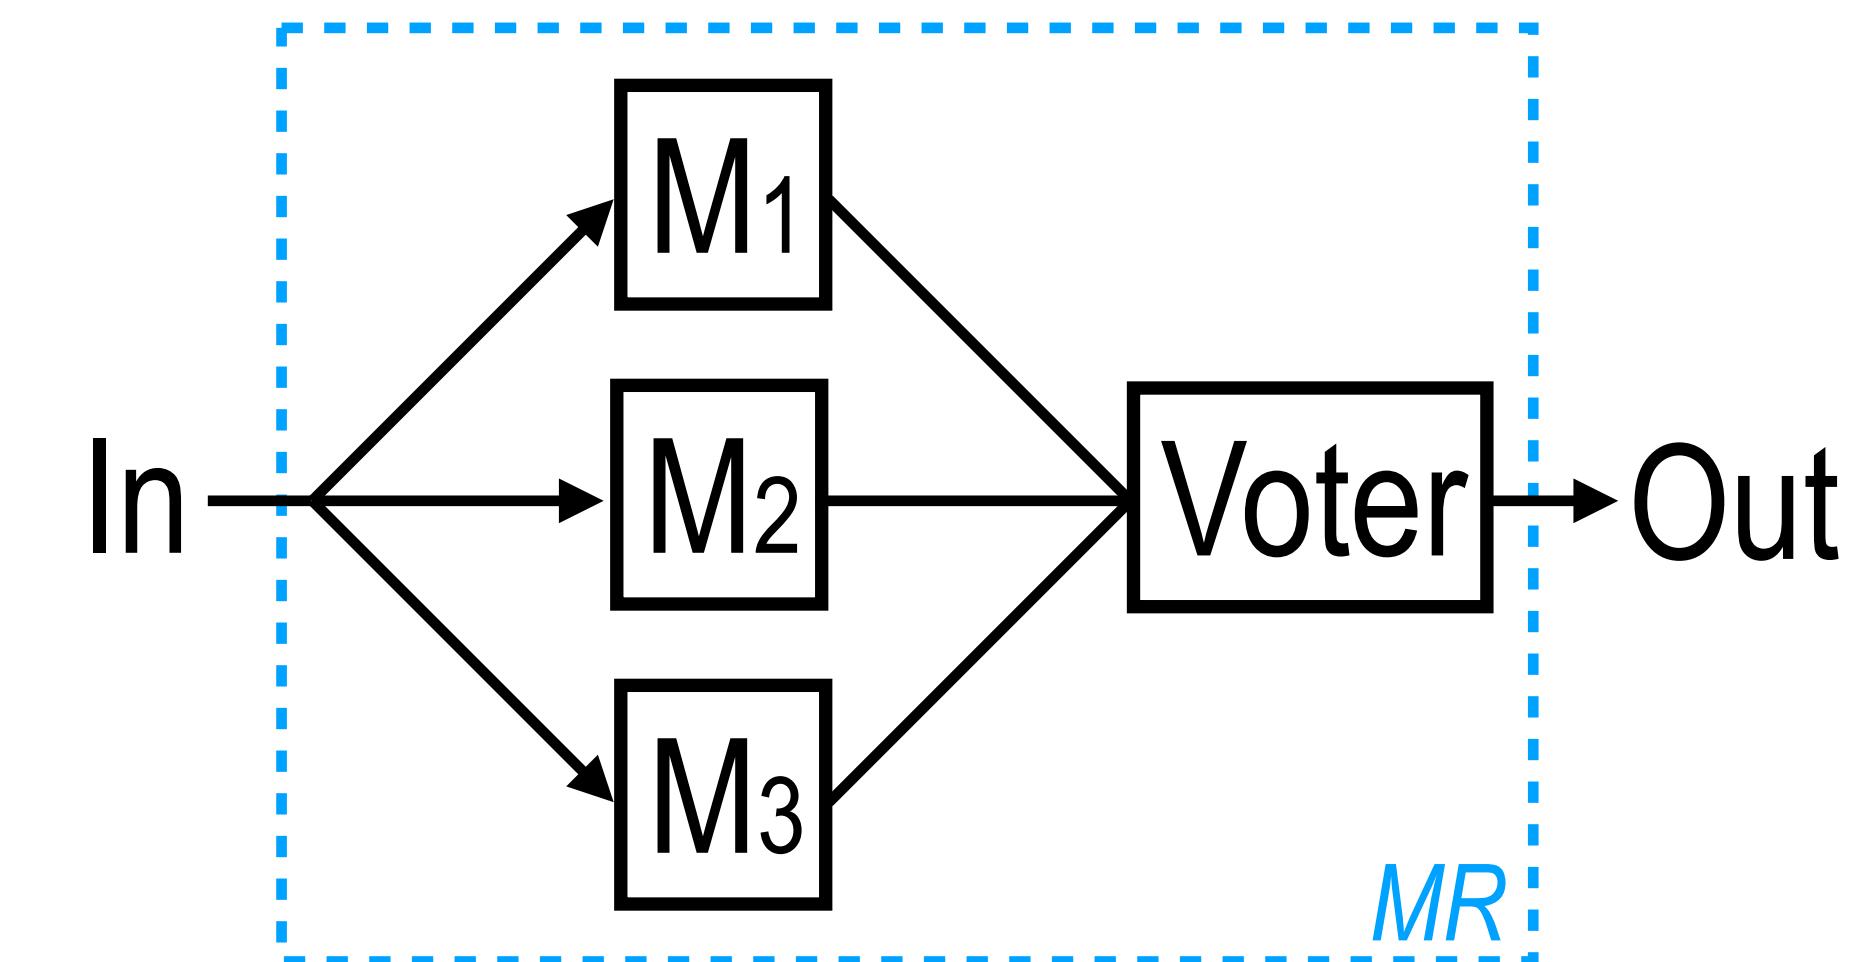
Failure mode 1: Fail-stop

- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*
- Any system can be made fail-stop with simple modular redundancy (MR)
 - *Strict fault model: voter is reliable*
 - *Tolerate failures*
 - level of redundancy: $2f+1$ independent modules
 - can sometimes get away with $f + 1$ (if fail-stop behavior is OK)



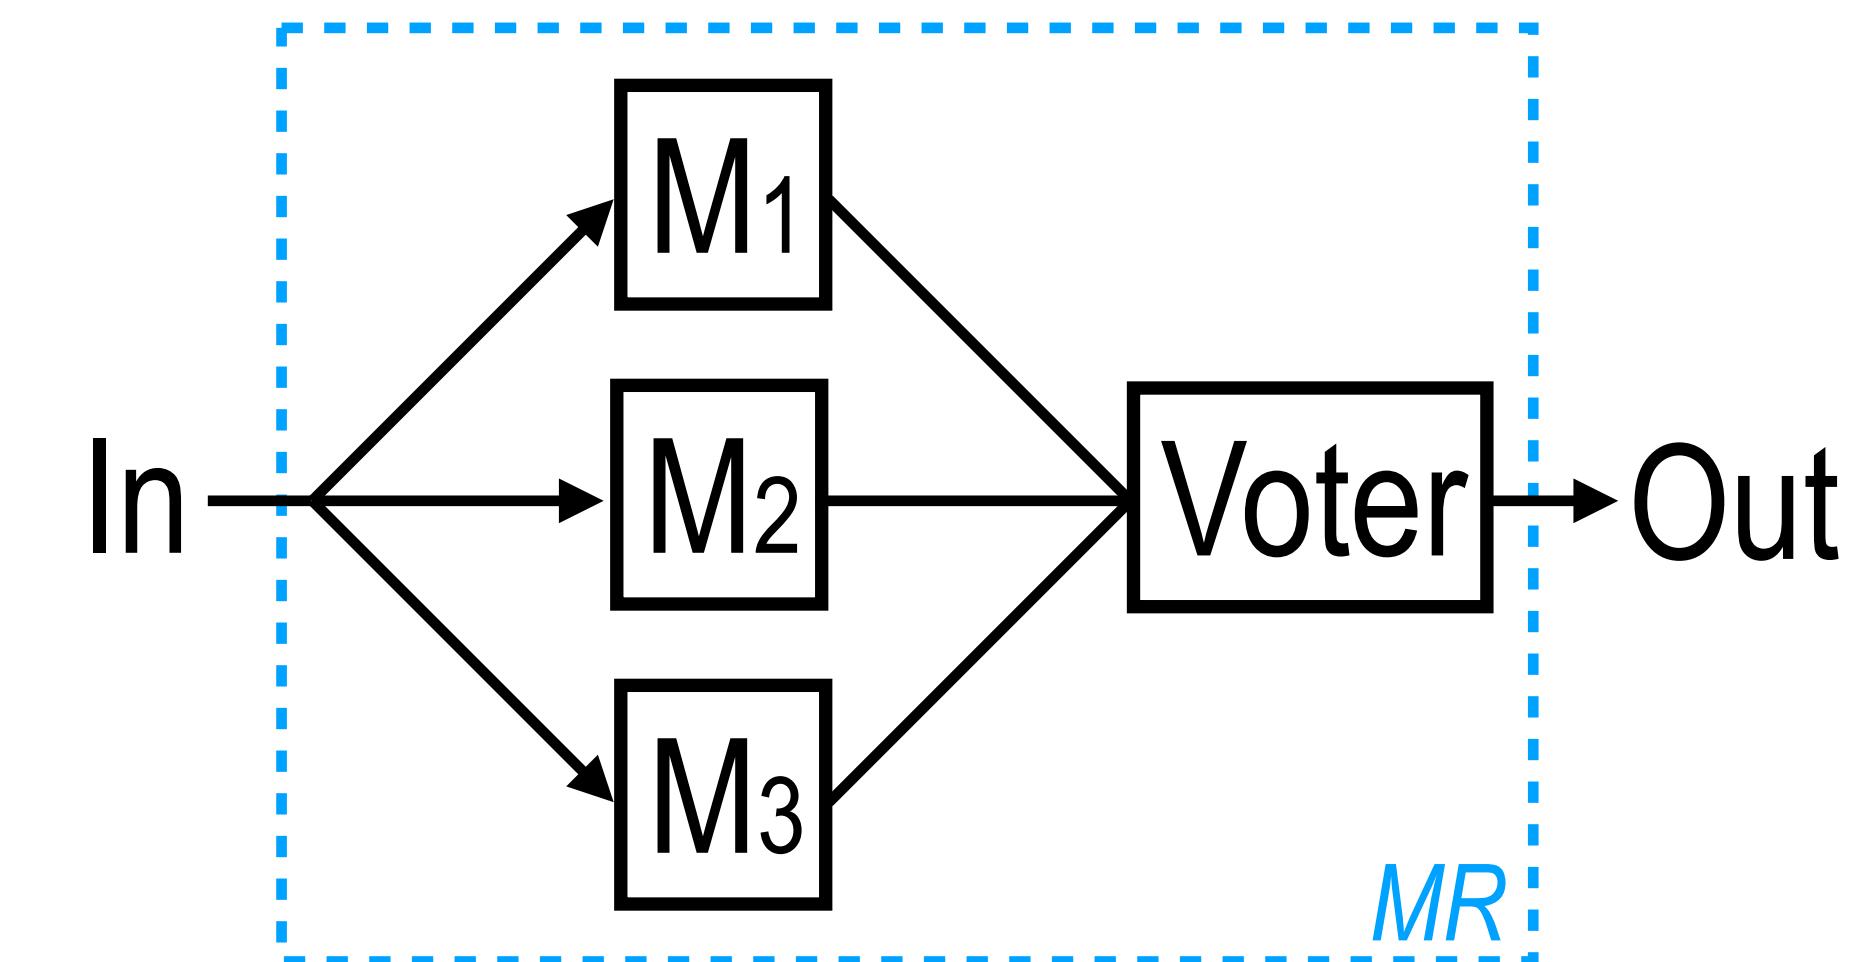
Failure mode 1: Fail-stop

- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*
- Any system can be made fail-stop with simple modular redundancy (MR)
 - *Strict fault model: voter is reliable*
 - *Tolerate failures*
 - level of redundancy: $2f+1$ independent modules
 - can sometimes get away with $f + 1$ (if fail-stop behavior is OK)
 - *Achilles's heel: voter / consensus algorithm*



Failure mode 1: Fail-stop

- a.k.a. "crash failure" mode
- *Definition:* halt in response to any internal error that threatens to turn into a failure, before the failure becomes visible
 - => *never expose arbitrary behavior*
- Any system can be made fail-stop with simple modular redundancy (MR)
 - *Strict fault model: voter is reliable*
 - *Tolerate failures*
 - level of redundancy: $2f+1$ independent modules
 - can sometimes get away with $f + 1$ (if fail-stop behavior is OK)
 - *Achilles's heel: voter / consensus algorithm*



Failure mode 2: Fail-fast

Failure mode 2: Fail-fast

- *Definition:* immediately report at interface any situation that could lead to failure
 - *Can stop immediately after detection or delay (if expect recovery)*
 - *Must stop before failure manifests externally*

Failure mode 2: Fail-fast

- *Definition:* immediately report at interface any situation that could lead to failure
 - *Can stop immediately after detection or delay (if expect recovery)*
 - *Must stop before failure manifests externally*
- Requires frequent checks of state invariants

Failure mode 2: Fail-fast

- *Definition:* immediately report at interface any situation that could lead to failure
 - *Can stop immediately after detection or delay (if expect recovery)*
 - *Must stop before failure manifests externally*
- Requires frequent checks of state invariants
- Get auditability of error propagation

Failure mode 3: Fail-safe

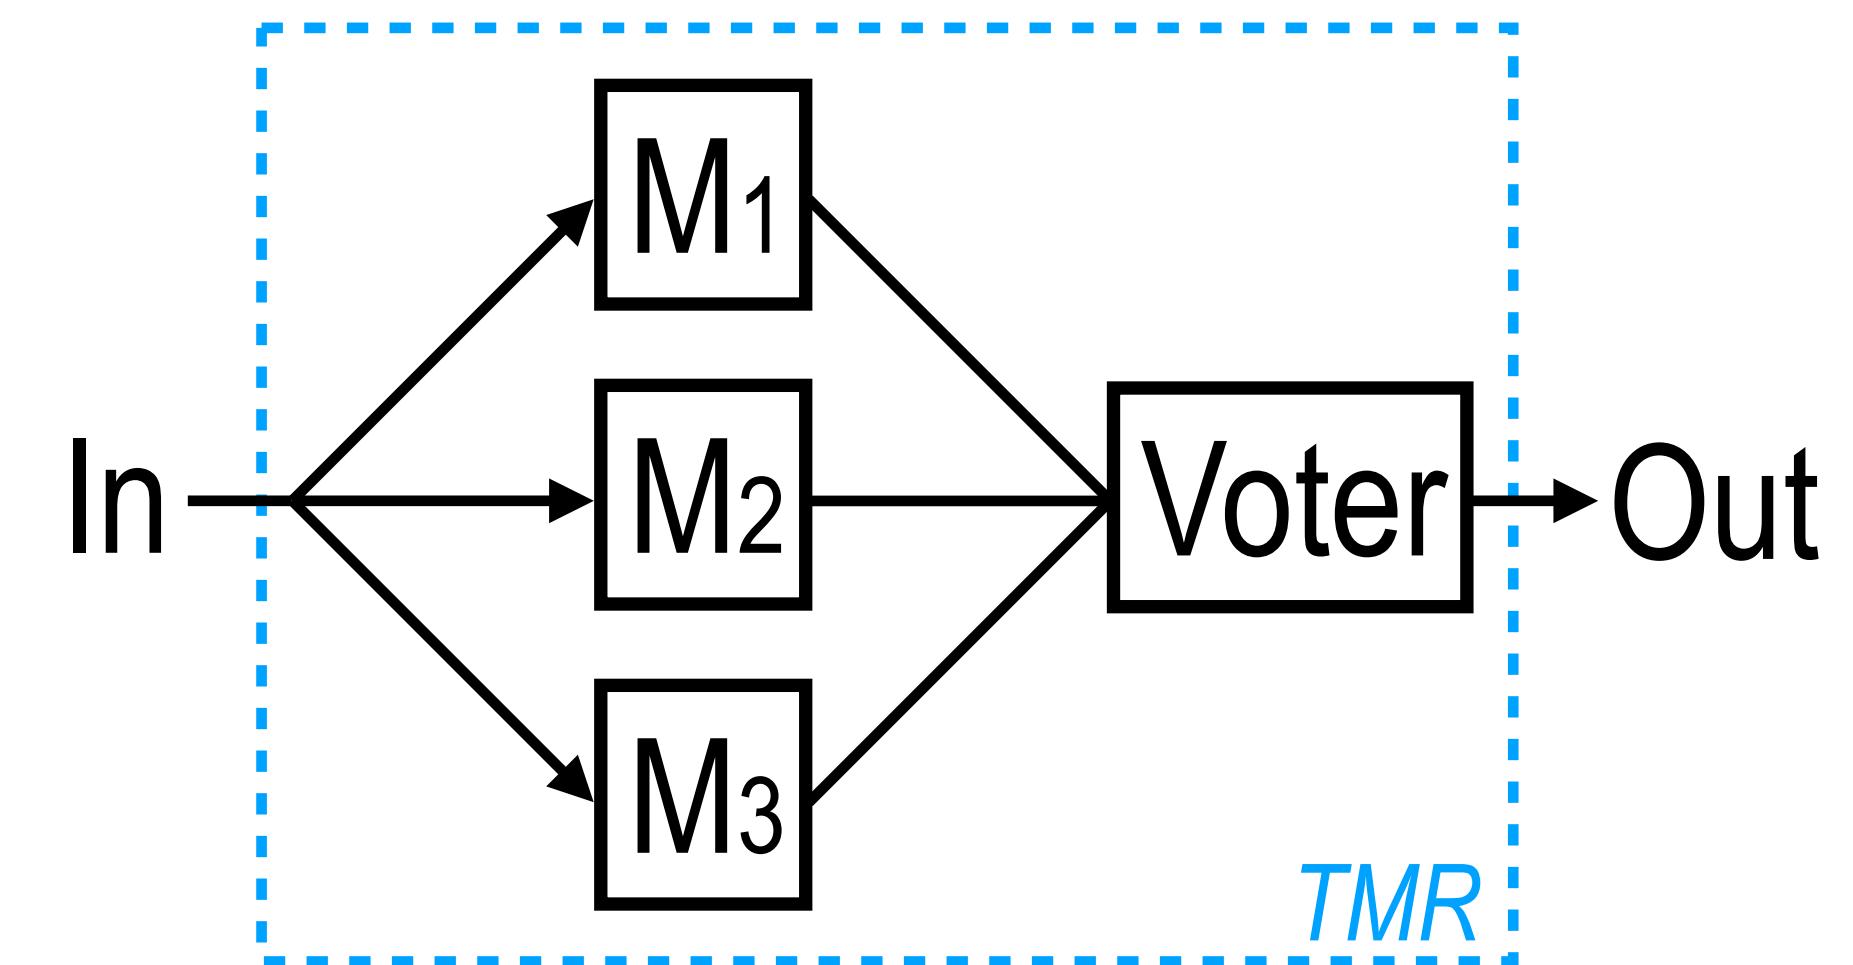
- *Definition:* the component remains safe in the face of failure
 - *but possibly degraded functionality or performance*

Failure mode 3: Fail-safe

- *Definition:* the component remains safe in the face of failure
 - *but possibly degraded functionality or performance*
- "Safety" is context-dependent

Failure mode 3: Fail-safe

- *Definition:* the component remains safe in the face of failure
 - *but possibly degraded functionality or performance*
- "Safety" is context-dependent
- "Controlled" failure

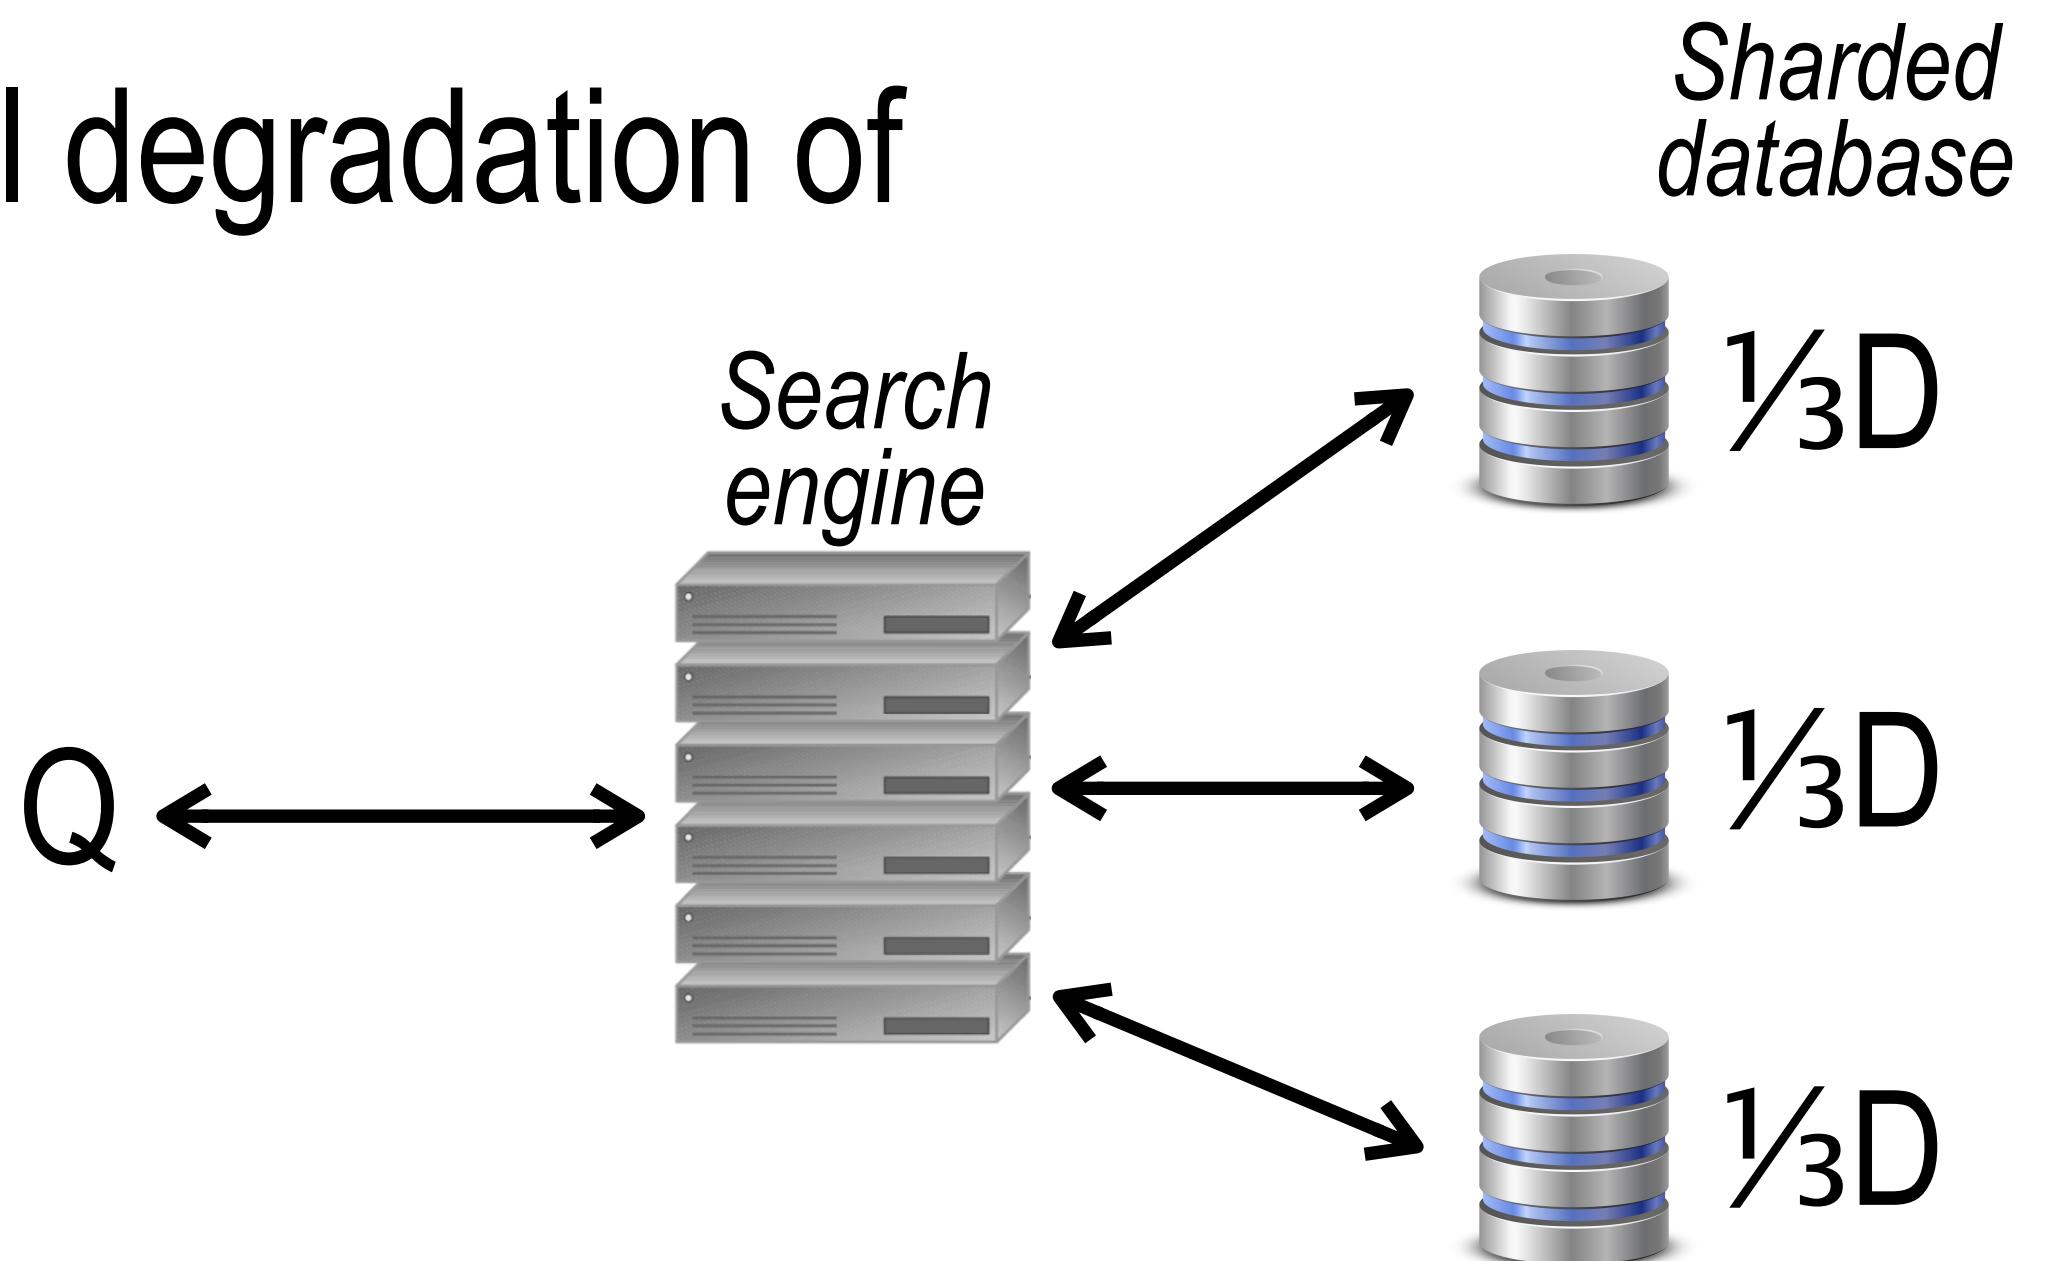


Failure mode 4: Fail-soft

- Definition: internal failures lead to graceful degradation of functionality instead of outright failure

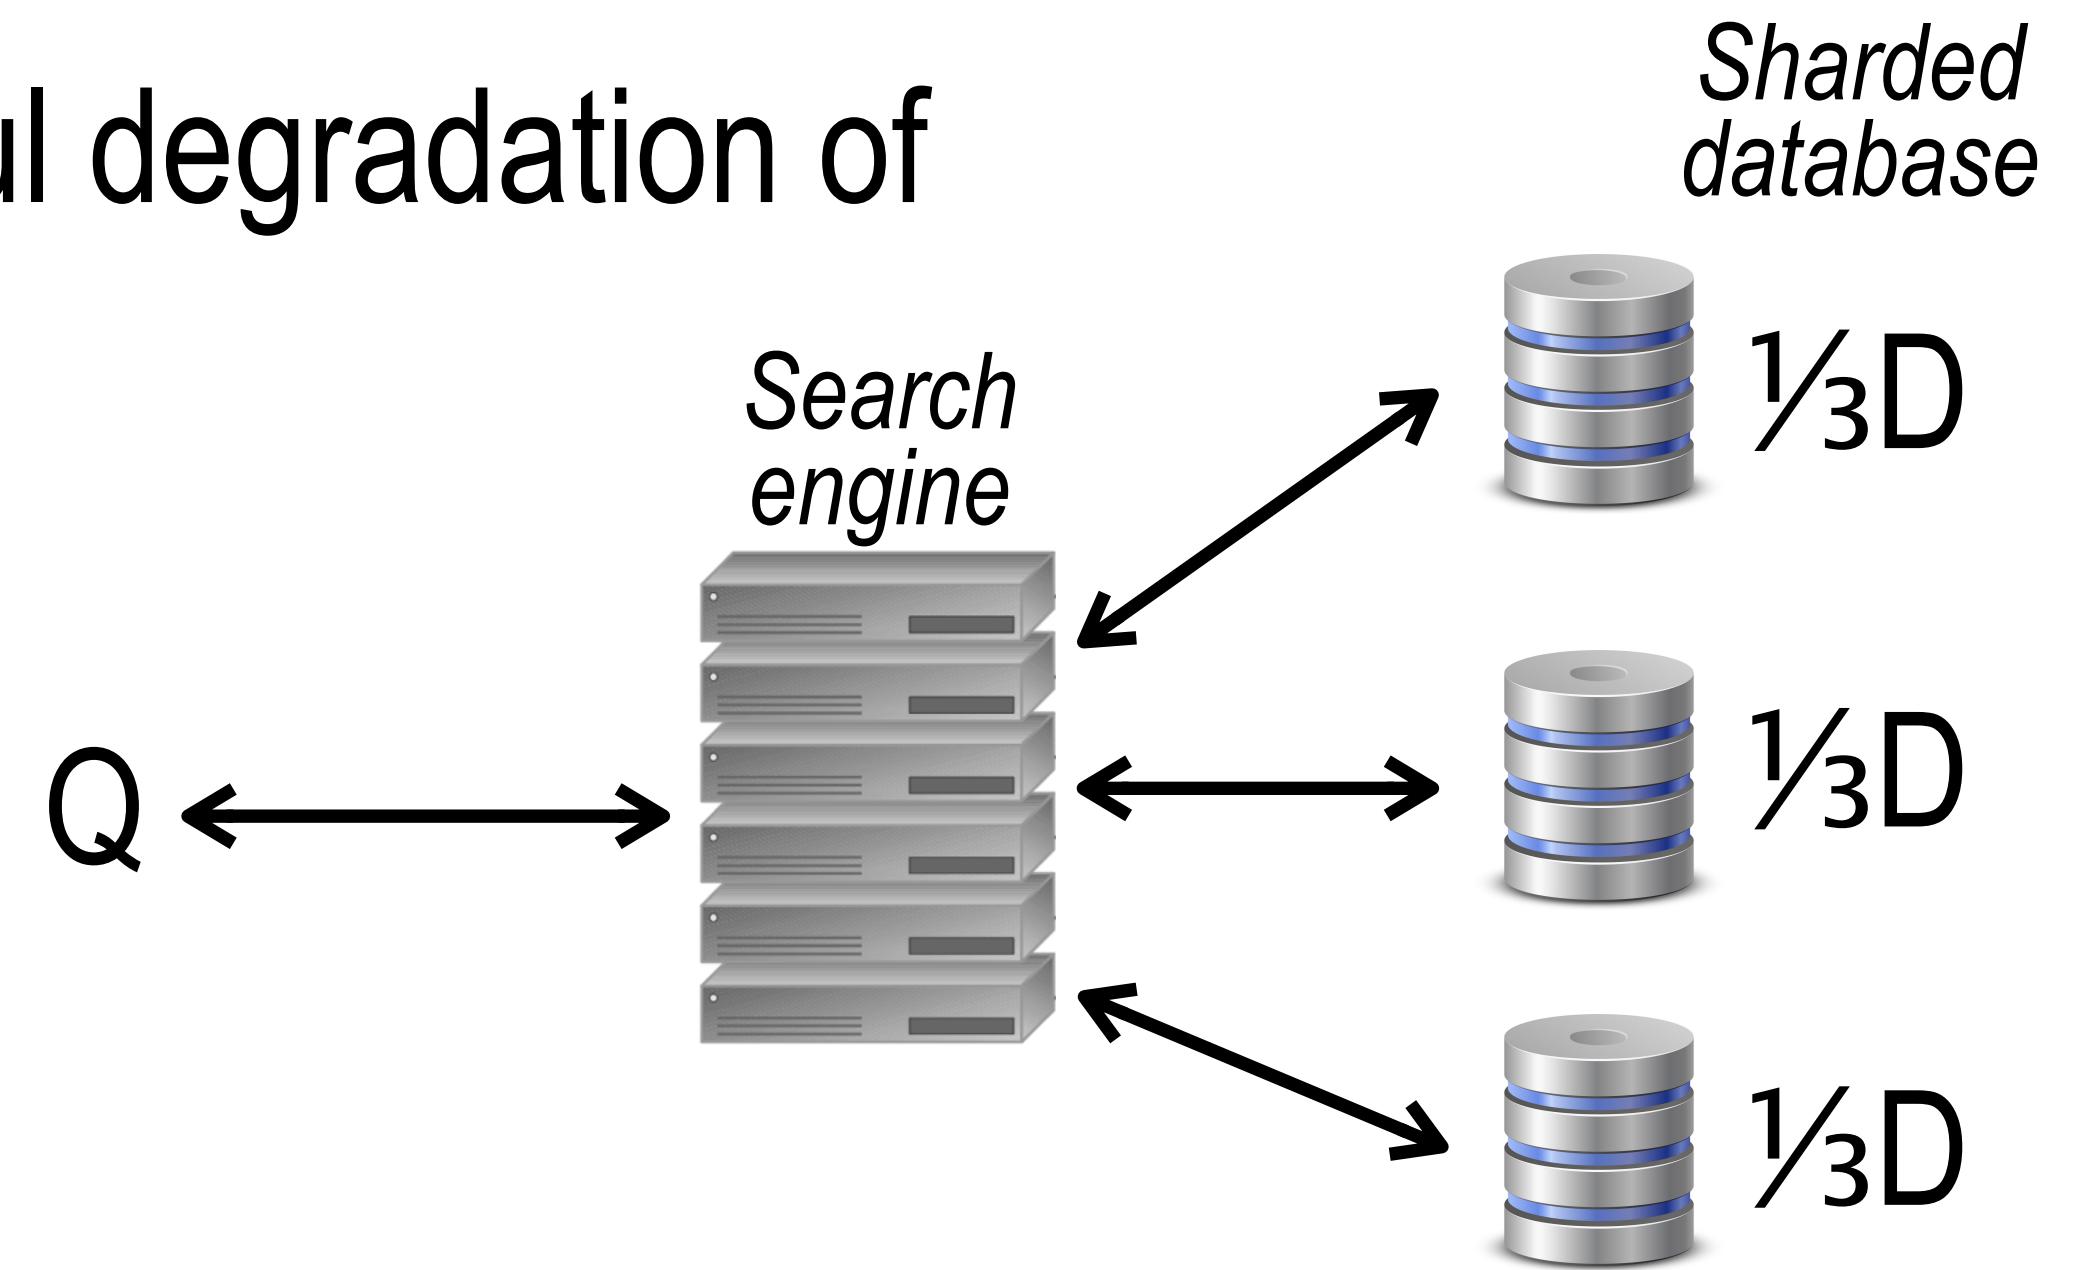
Failure mode 4: Fail-soft

- Definition: internal failures lead to graceful degradation of functionality instead of outright failure
- Example: simple search engine
 - *system has redundancy at every level*



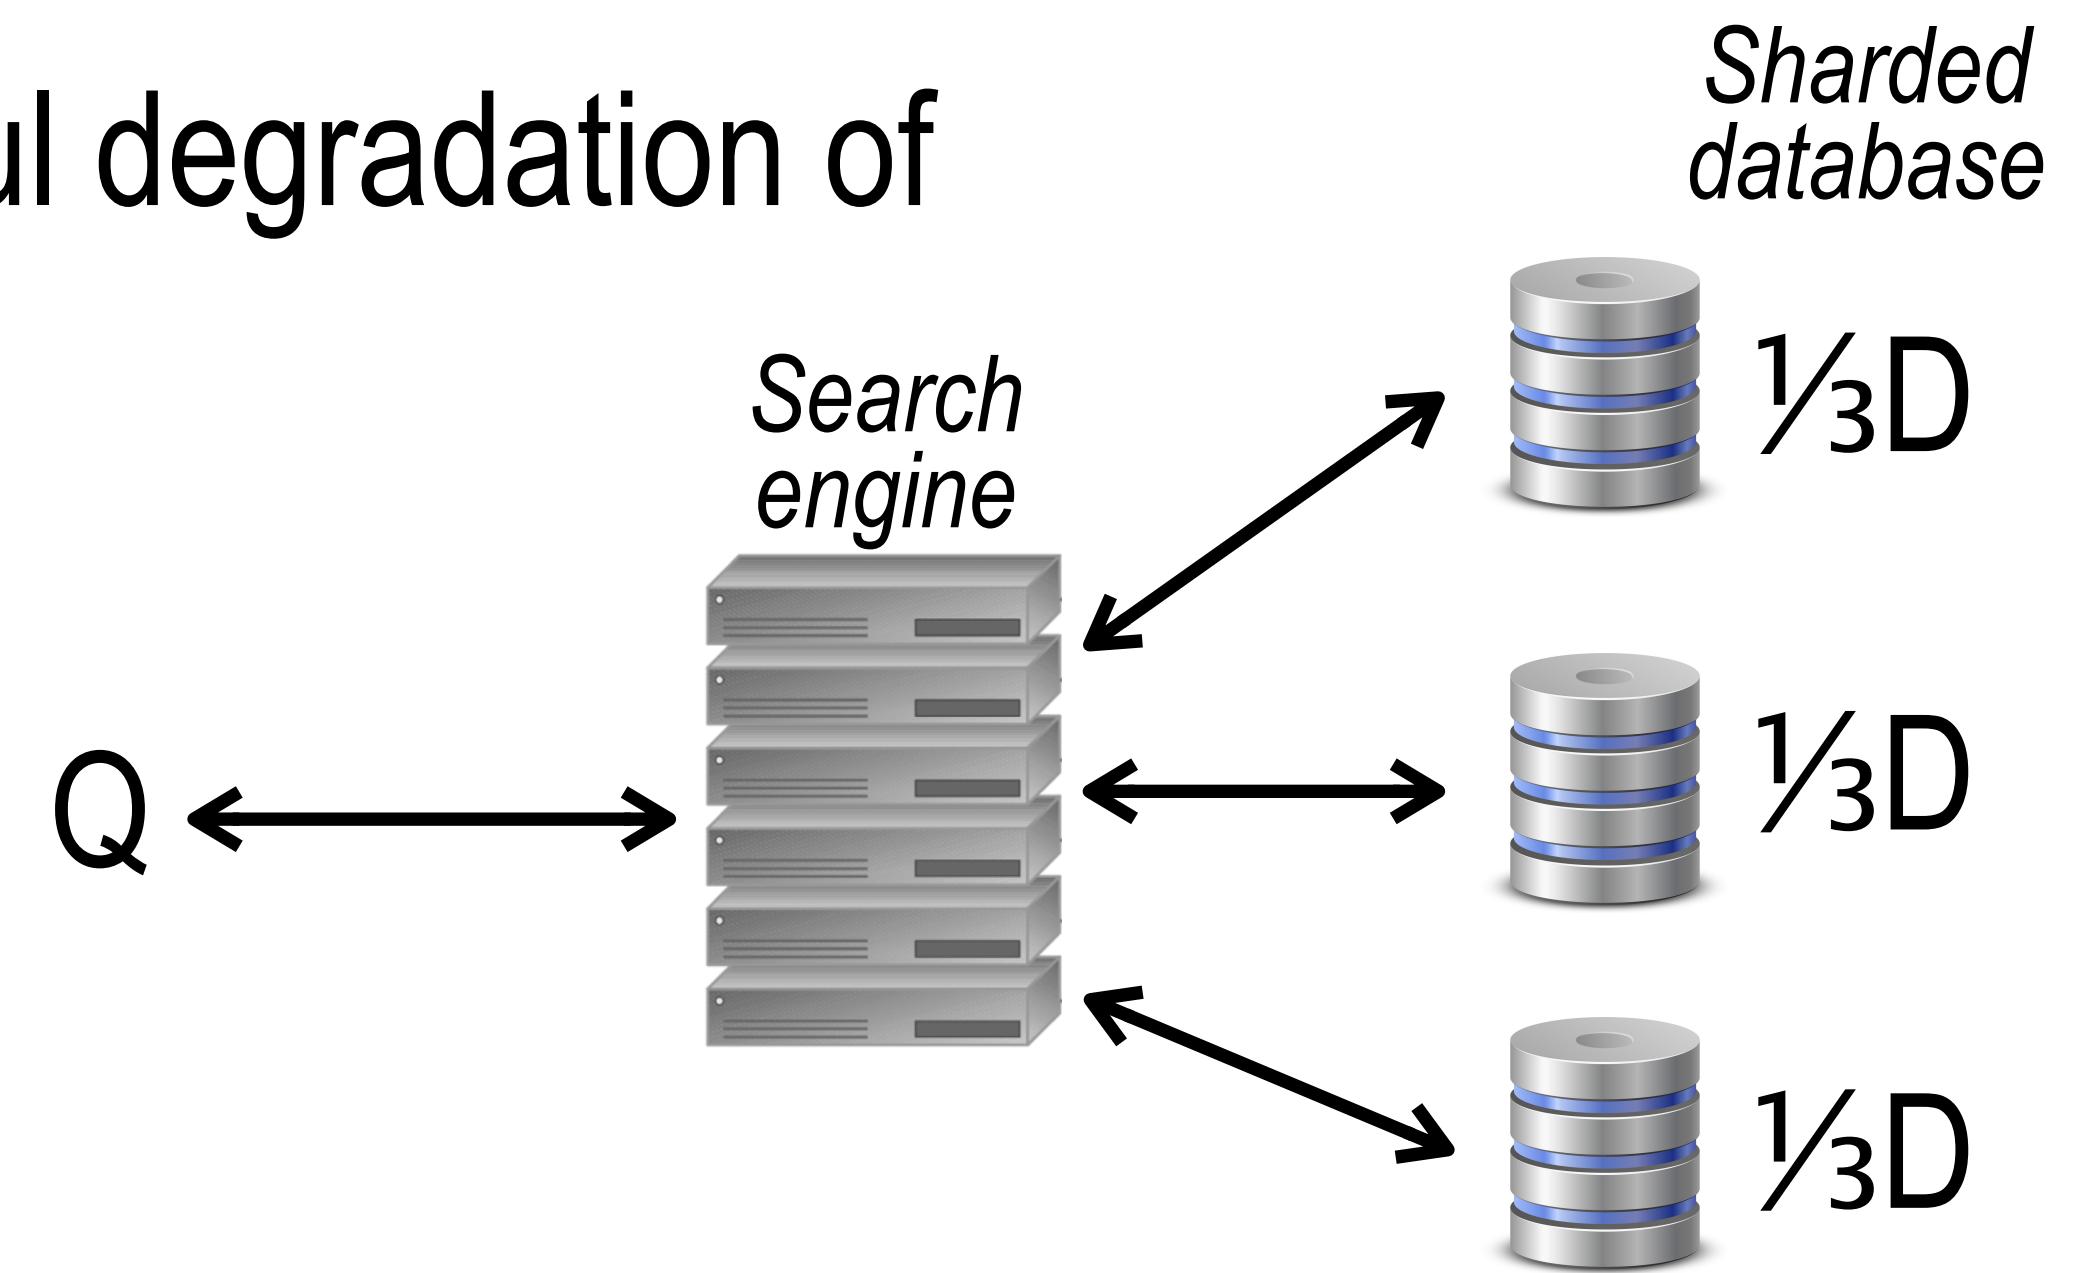
Failure mode 4: Fail-soft

- Definition: internal failures lead to graceful degradation of functionality instead of outright failure
- Example: simple search engine
 - *system has redundancy at every level*
- Intuition
 - *Functionality is typically bottlenecked on data movement (disks, network switches)*
 - => Functionality tied to how much data can be moved per unit of time



Failure mode 4: Fail-soft

- Definition: internal failures lead to graceful degradation of functionality instead of outright failure
- Example: simple search engine
 - *system has redundancy at every level*
- Intuition
 - *Functionality is typically bottlenecked on data movement (disks, network switches)*
 - => Functionality tied to how much data can be moved per unit of time
 - ***Harvest (completeness of responses) vs. yield (fraction of requests served)***



Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)

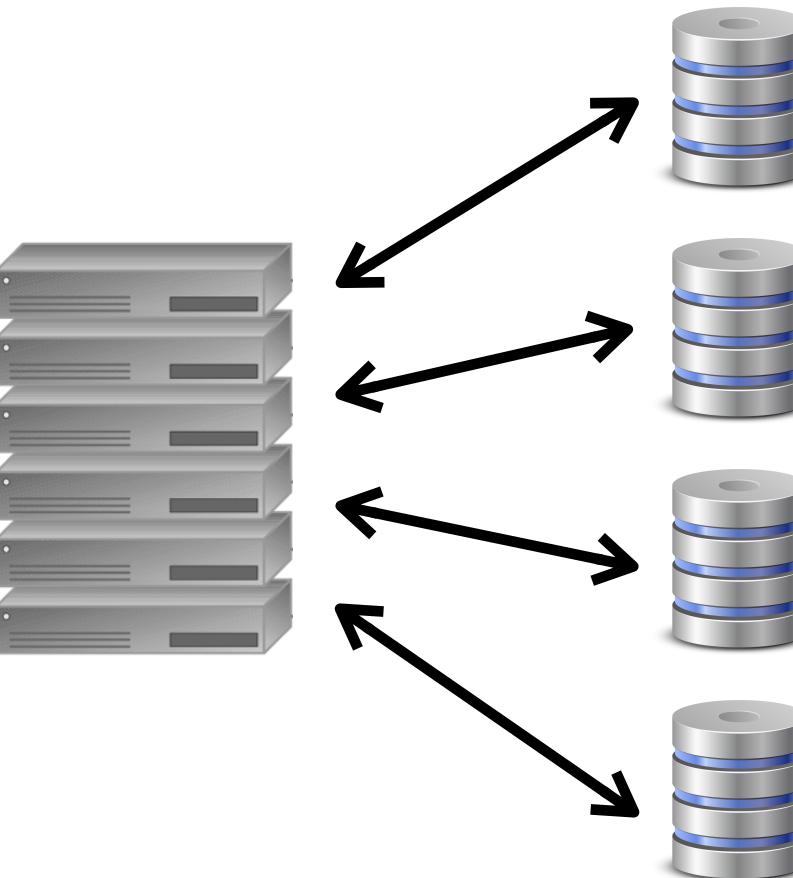
Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)



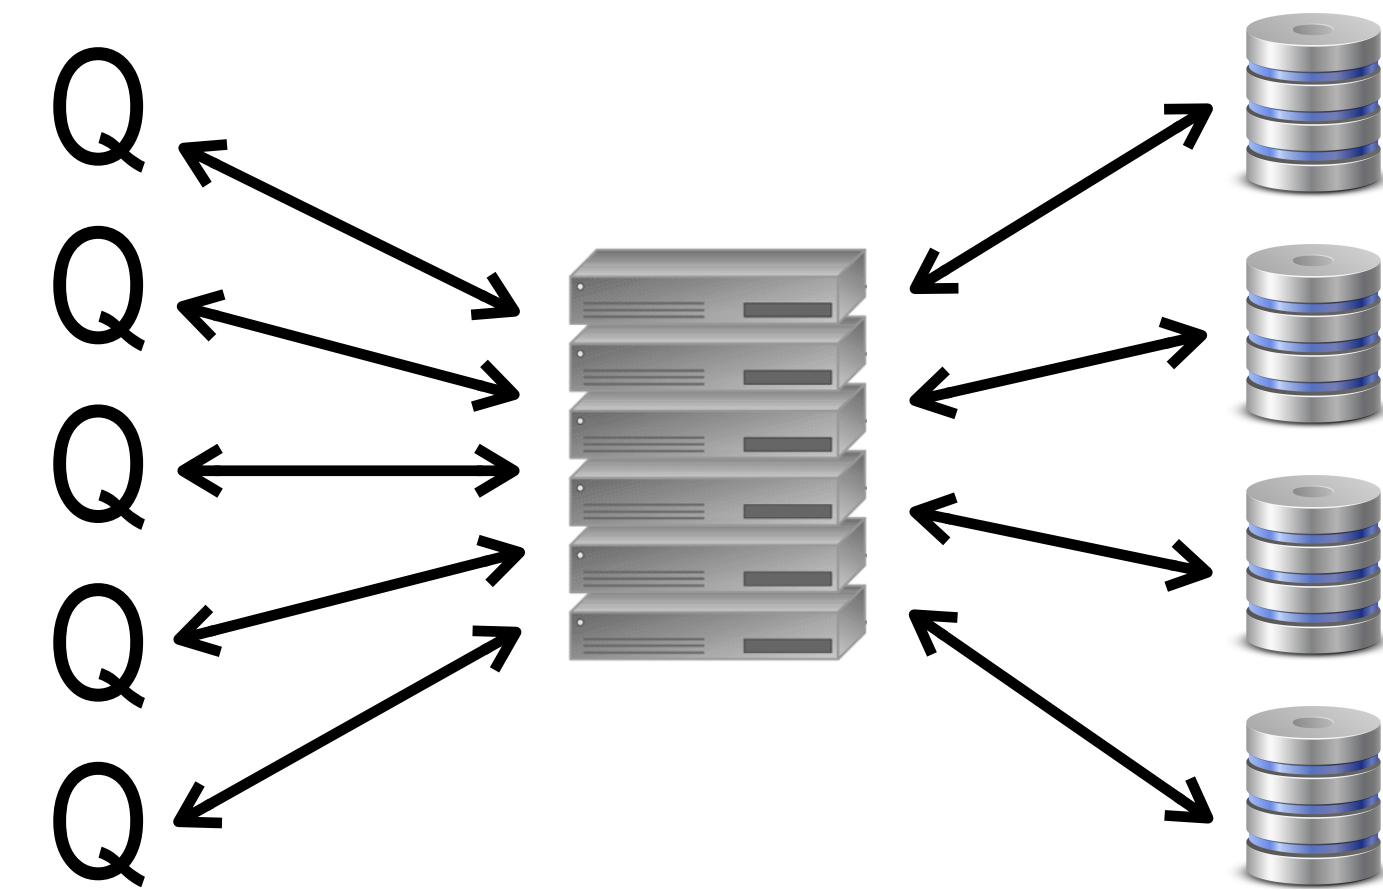
Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)



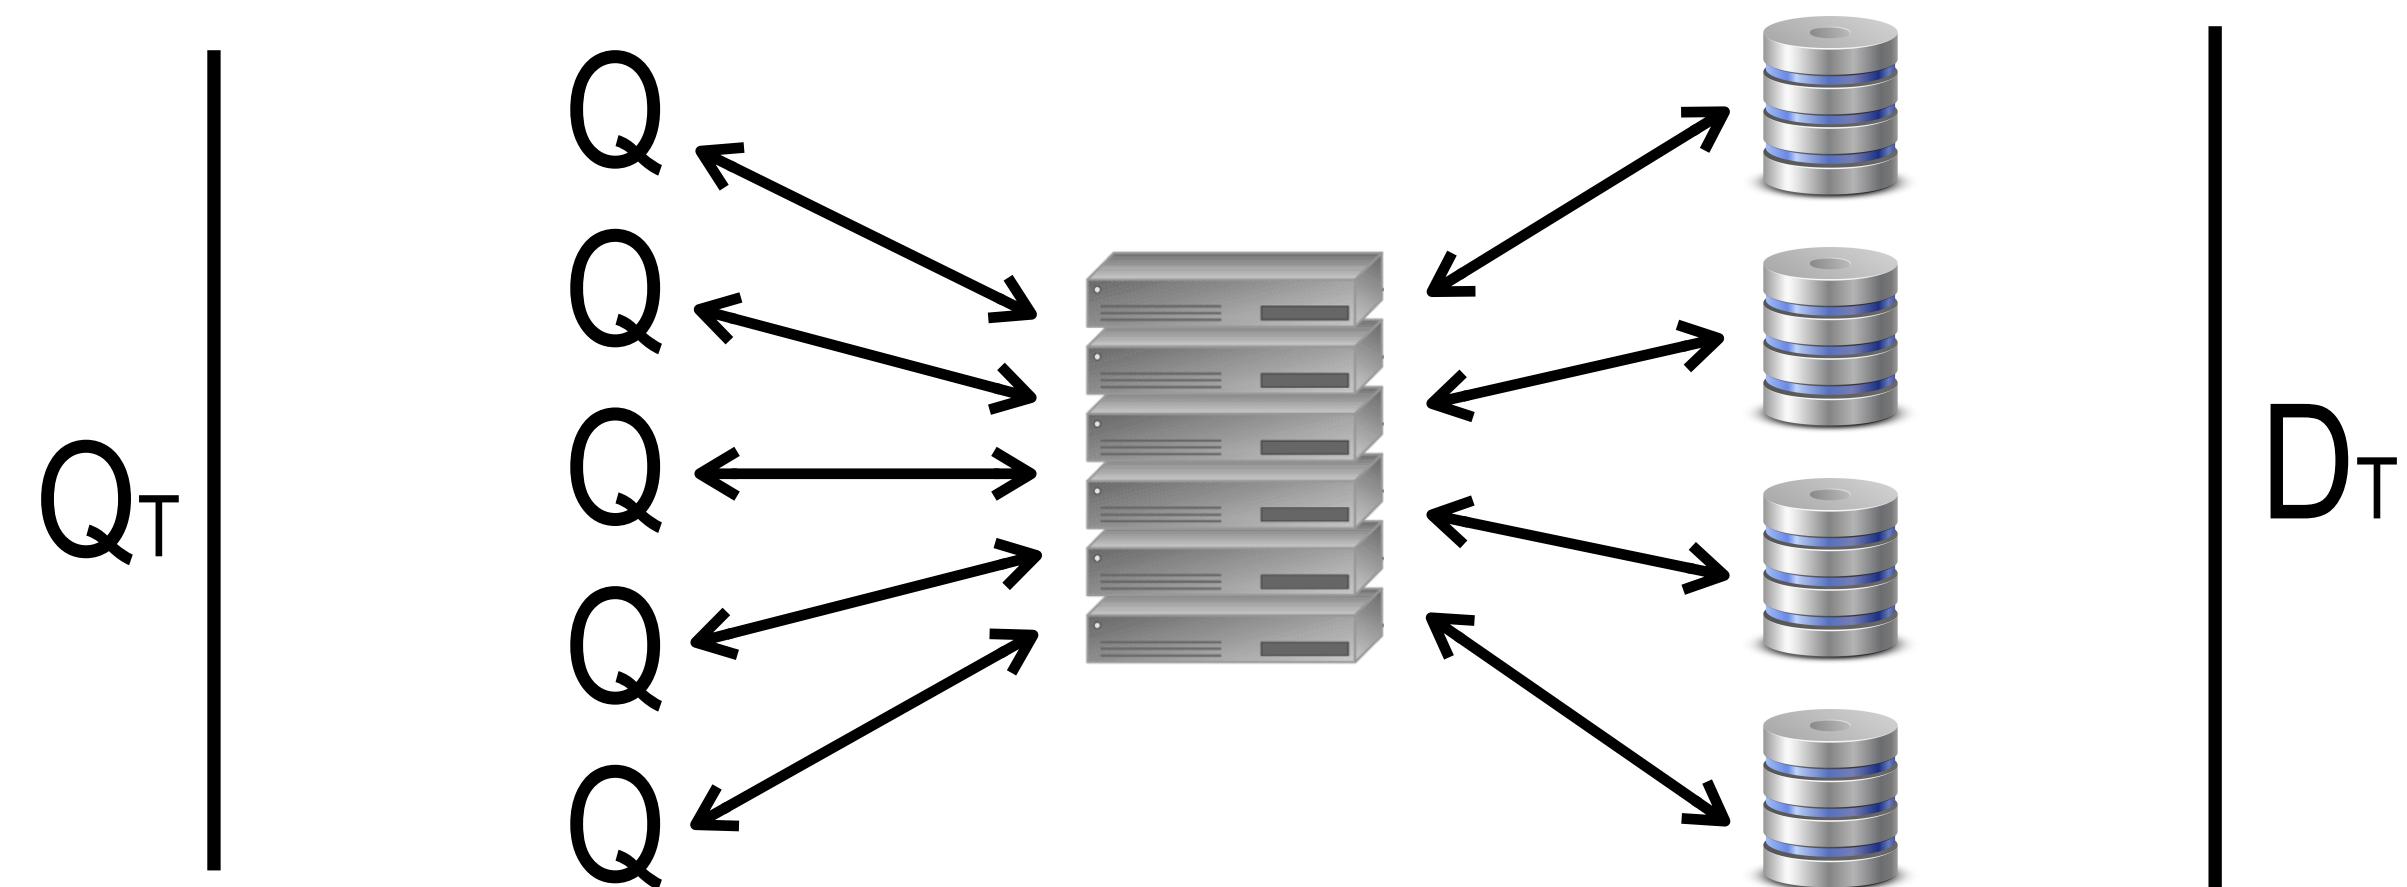
Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

DQ Principle: " $D \times Q$ is constant"

(DQ value ρ determined by system configuration)



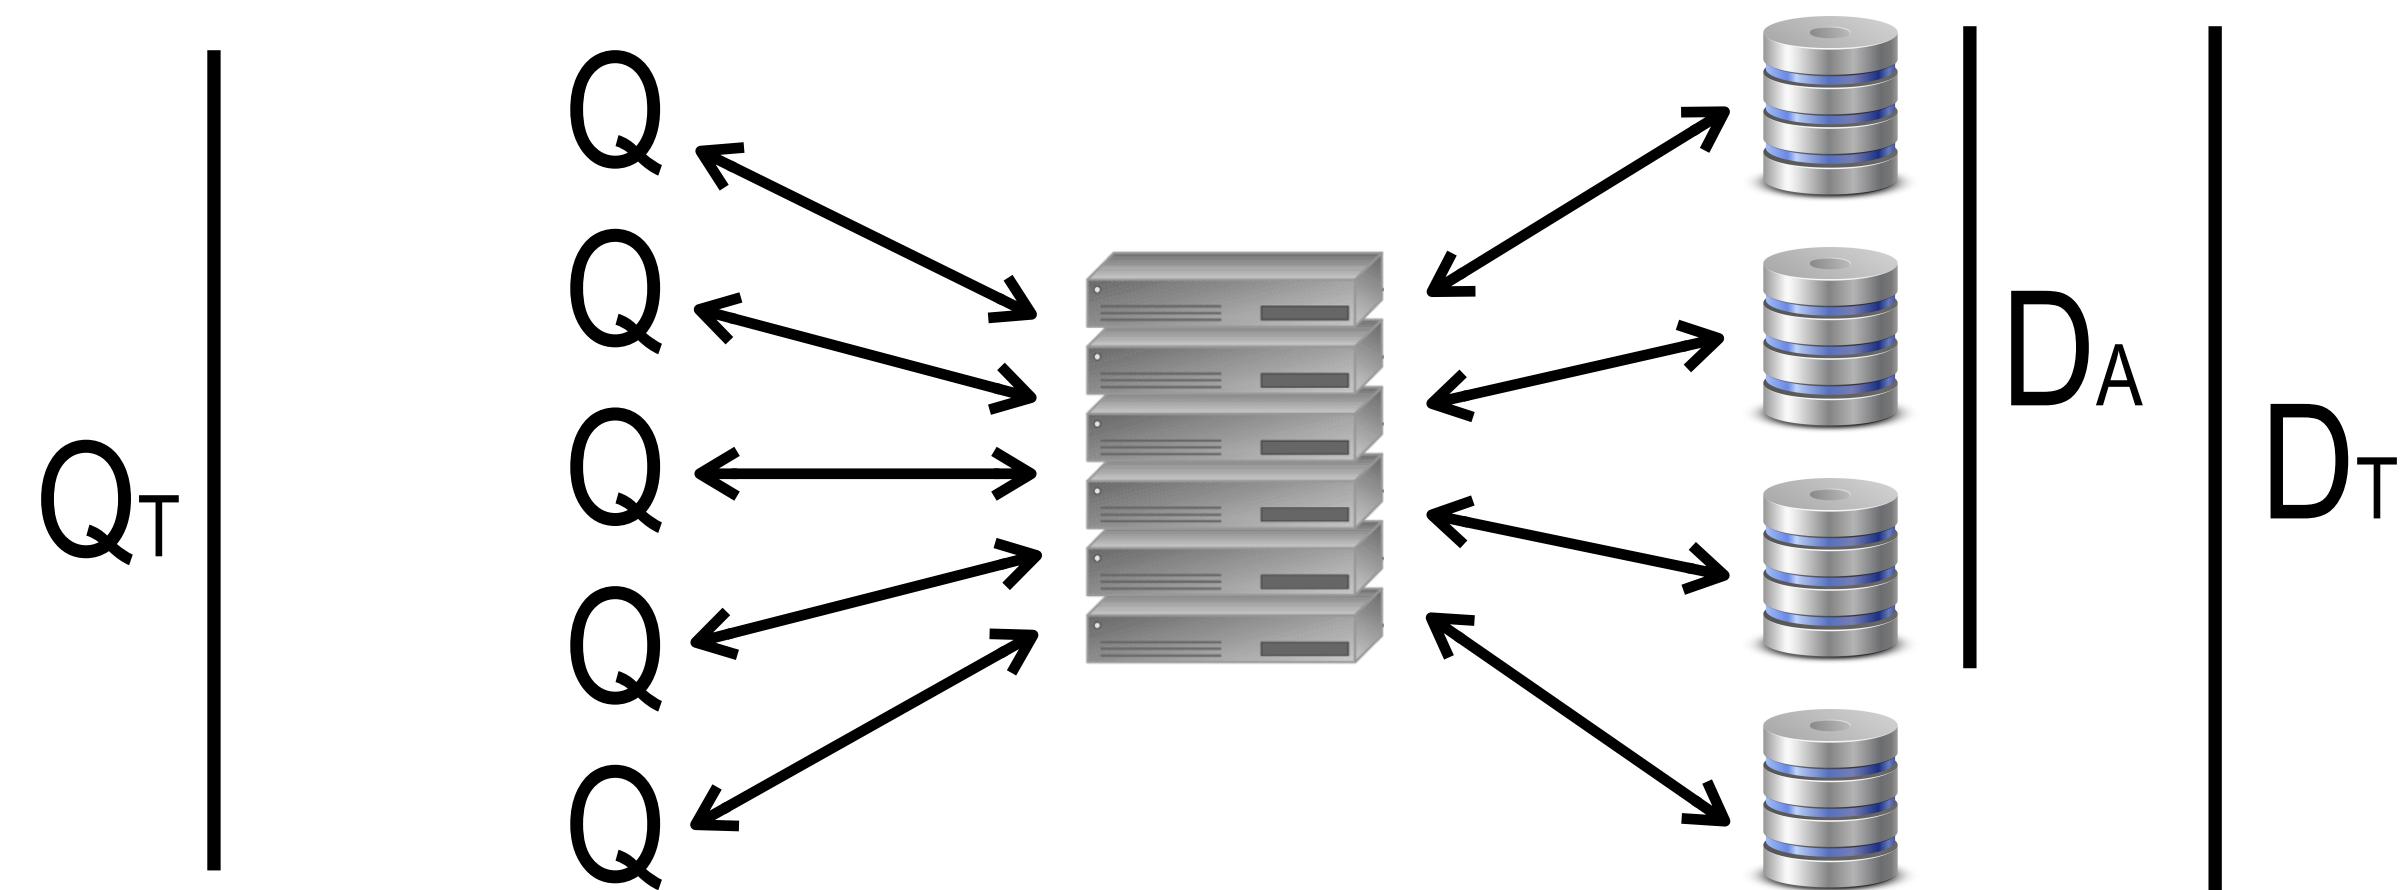
Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

DQ Principle: " $D \times Q$ is constant"

(DQ value ρ determined by system configuration)



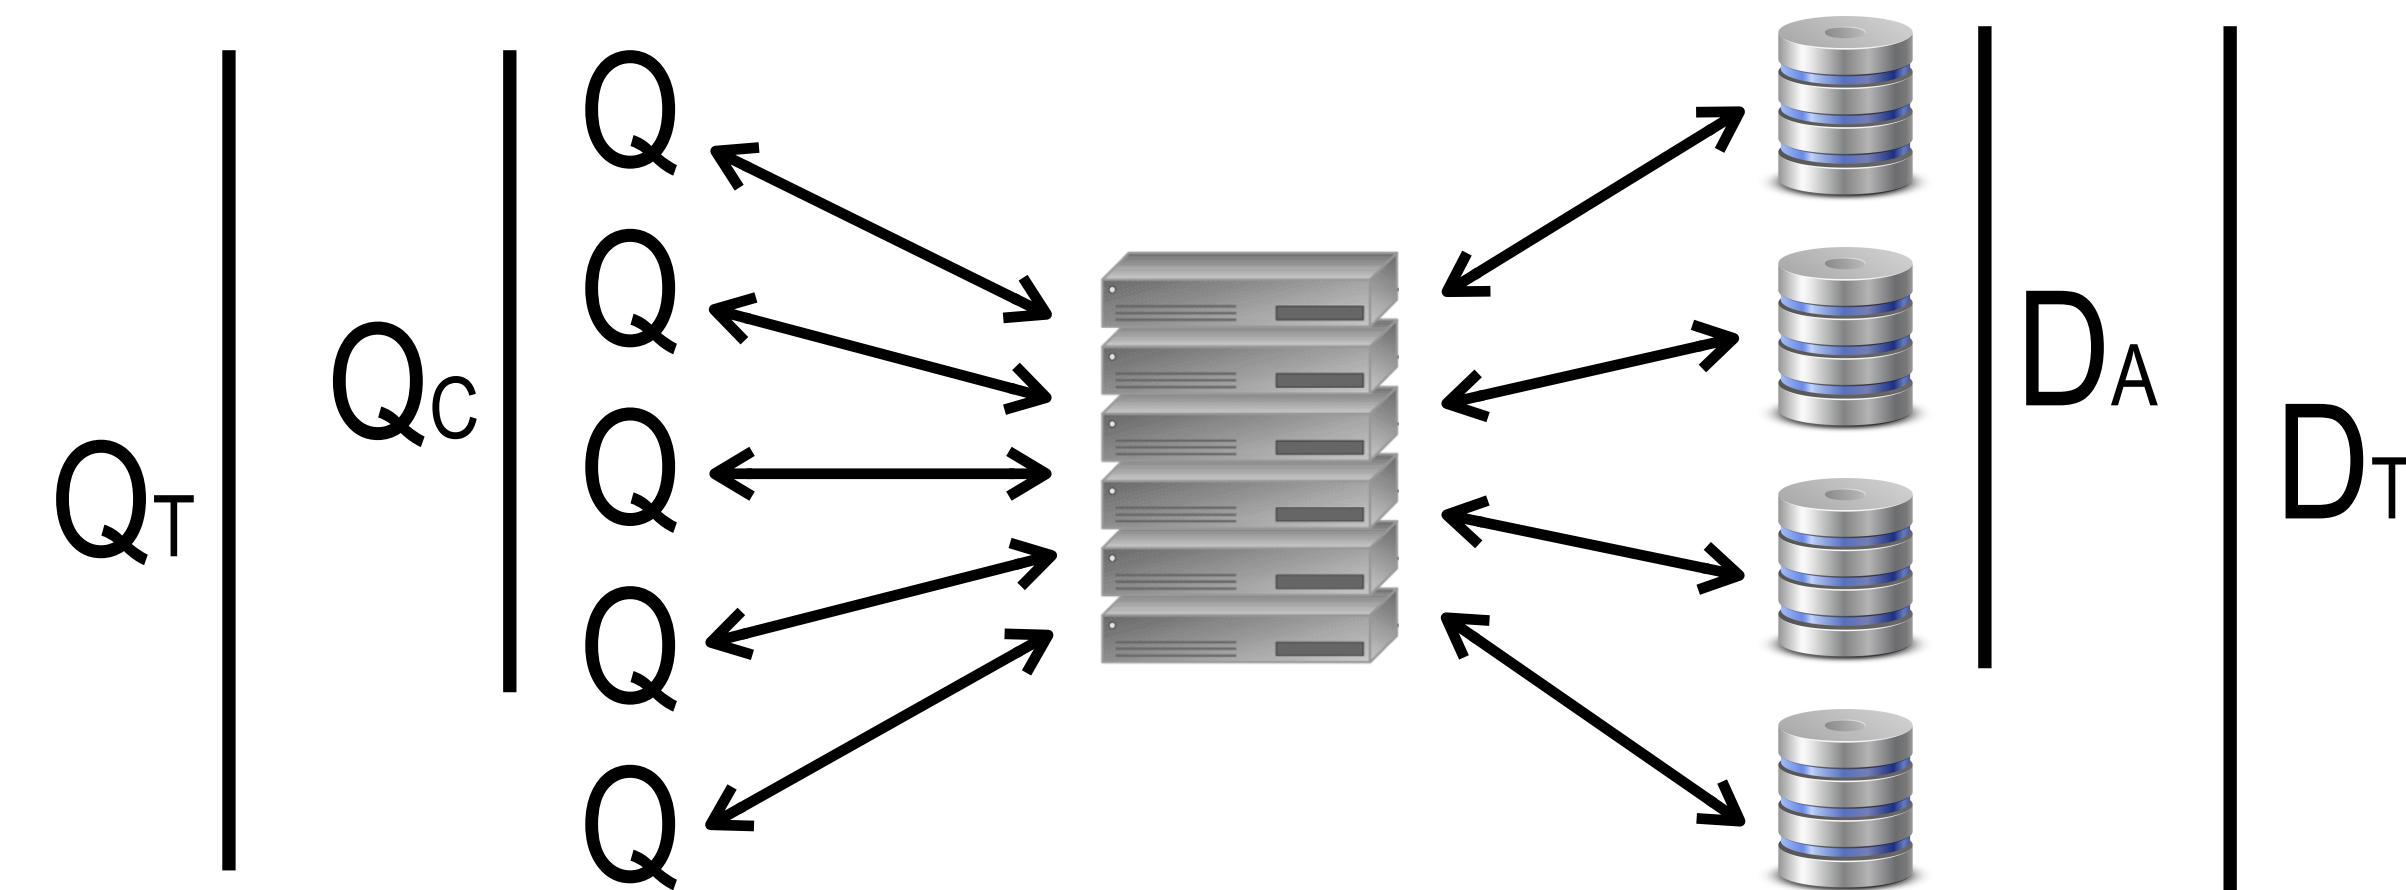
Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

DQ Principle: " $D \times Q$ is constant"

(DQ value ρ determined by system configuration)



Failure mode 4: Fail-soft: DQ Principle

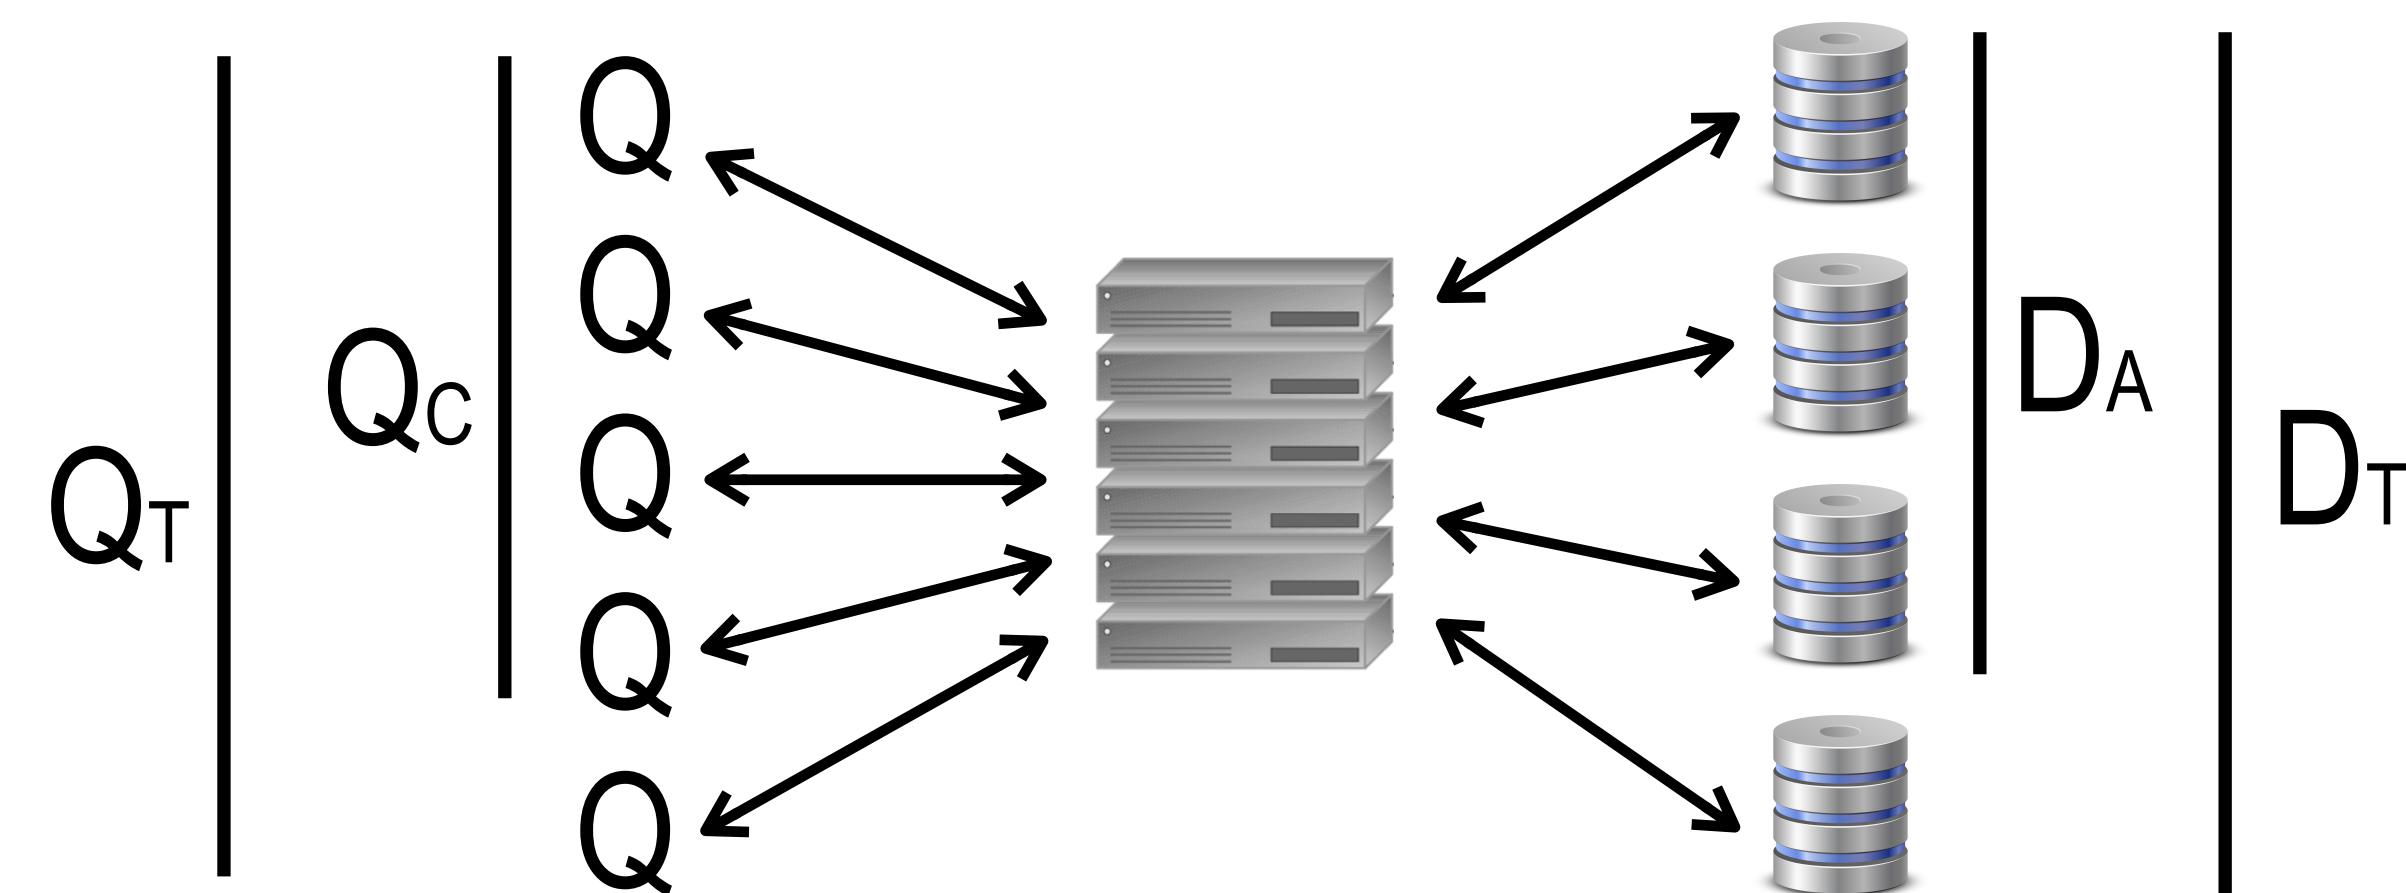
D = data/query

Q = queries/sec

DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)

$$\text{Harvest } H = \frac{D_A}{D_T}$$



Failure mode 4: Fail-soft: DQ Principle

D = data/query

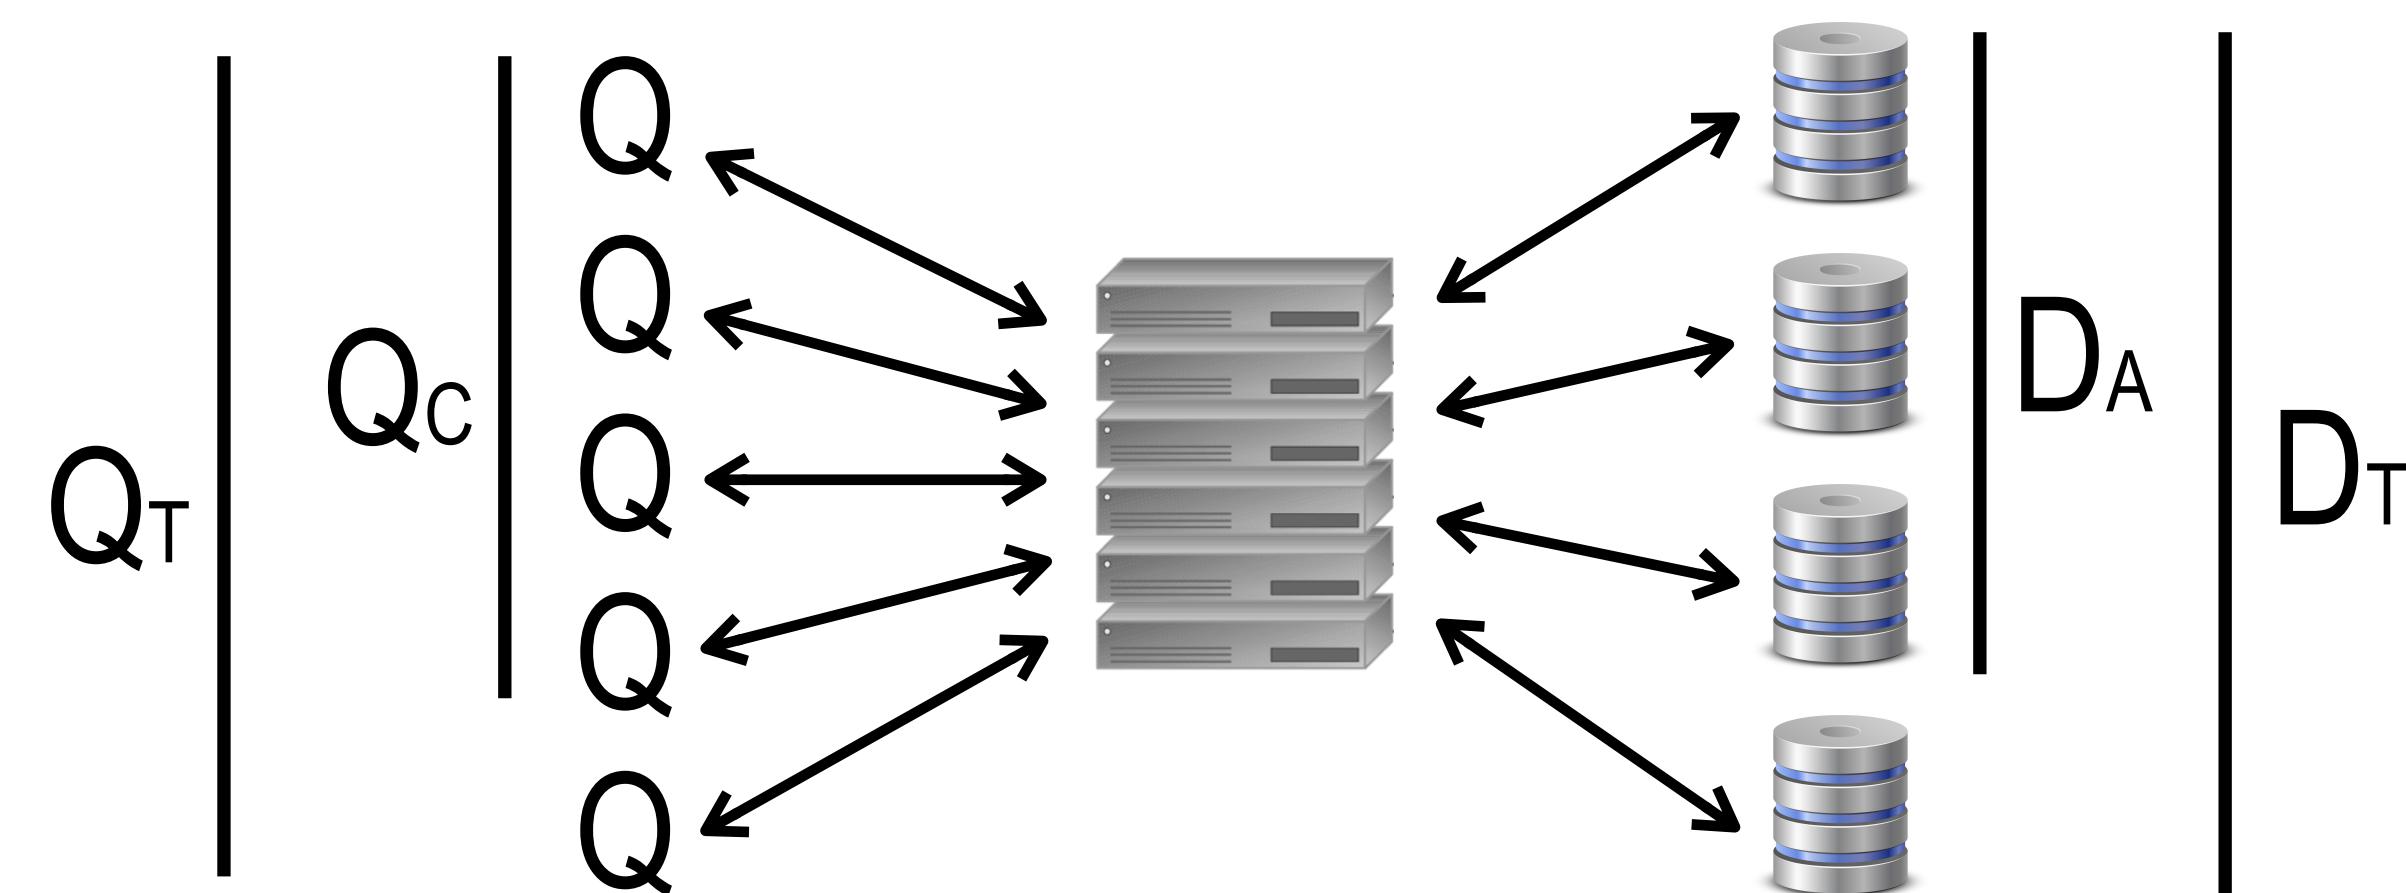
Q = queries/sec

DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)

$$\text{Harvest } H = \frac{D_A}{D_T}$$

$$\text{Yield } Y = \frac{Q_C}{Q_T}$$



Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

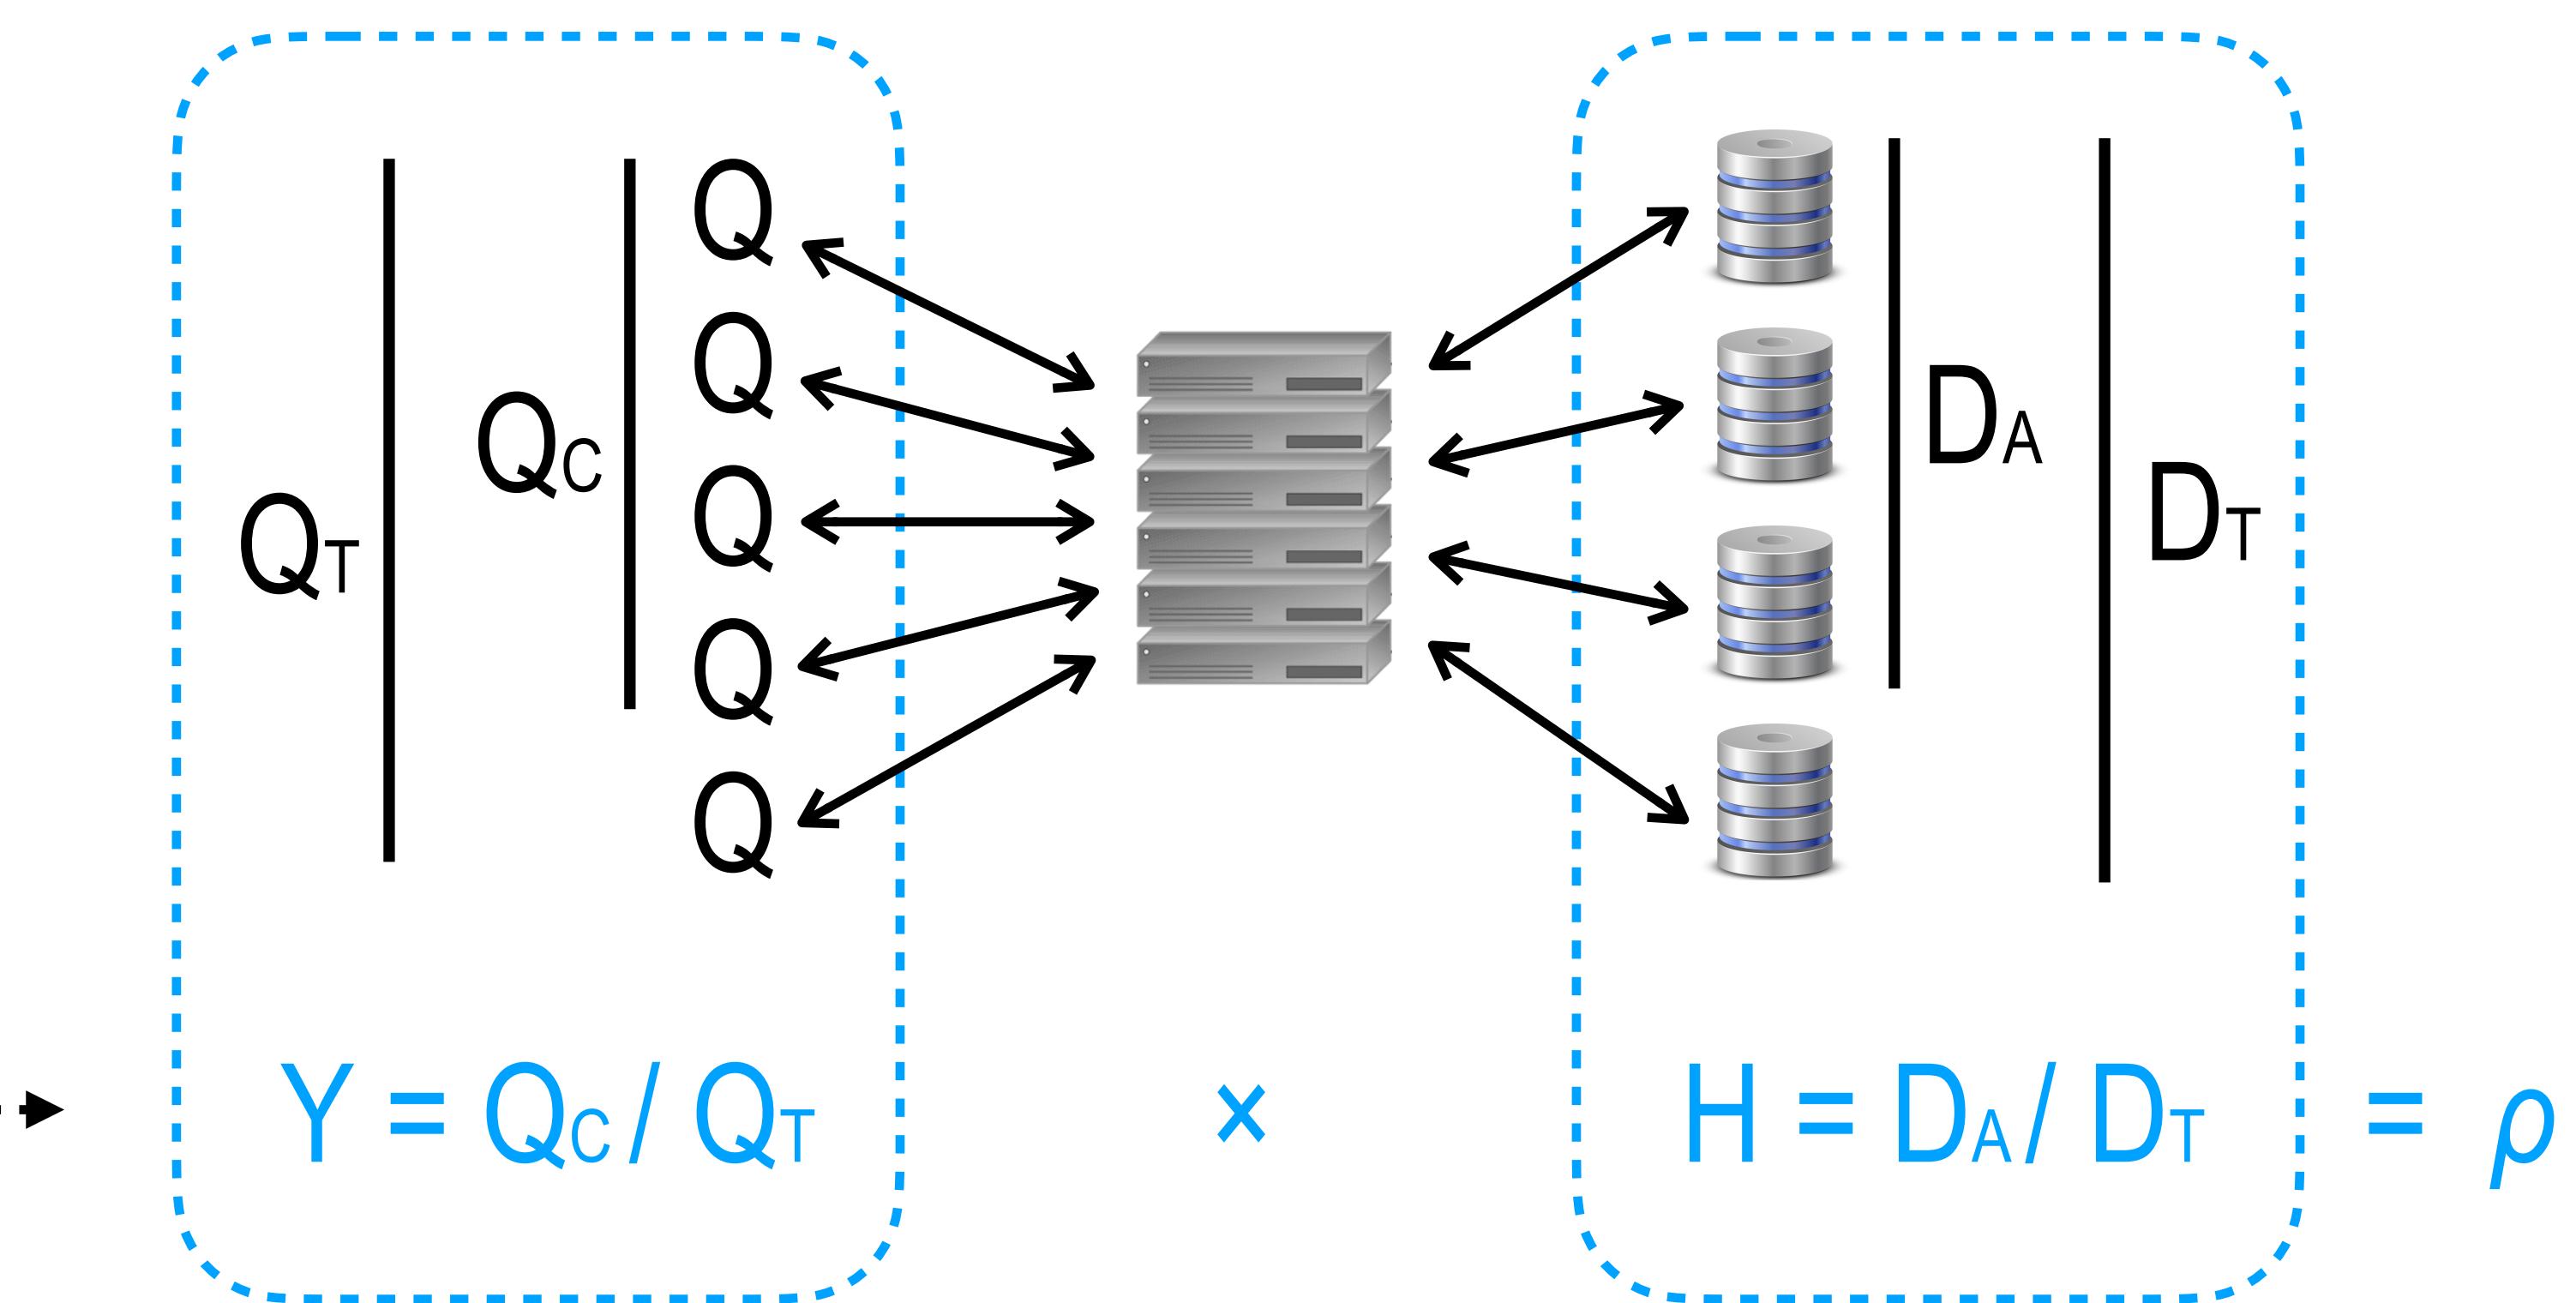
DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)

$$\text{Harvest } H = \frac{D_A}{D_T}$$

$$\text{Yield } Y = \frac{Q_c}{Q_T}$$

$$\text{DQ Principle: } H \times Y = \rho \longrightarrow$$



¹ Assumption: Disk is the only bottleneck. All the other resources (e.g., network bandwidth) are plentiful.

Failure mode 4: Fail-soft: DQ Principle

D = data/query

Q = queries/sec

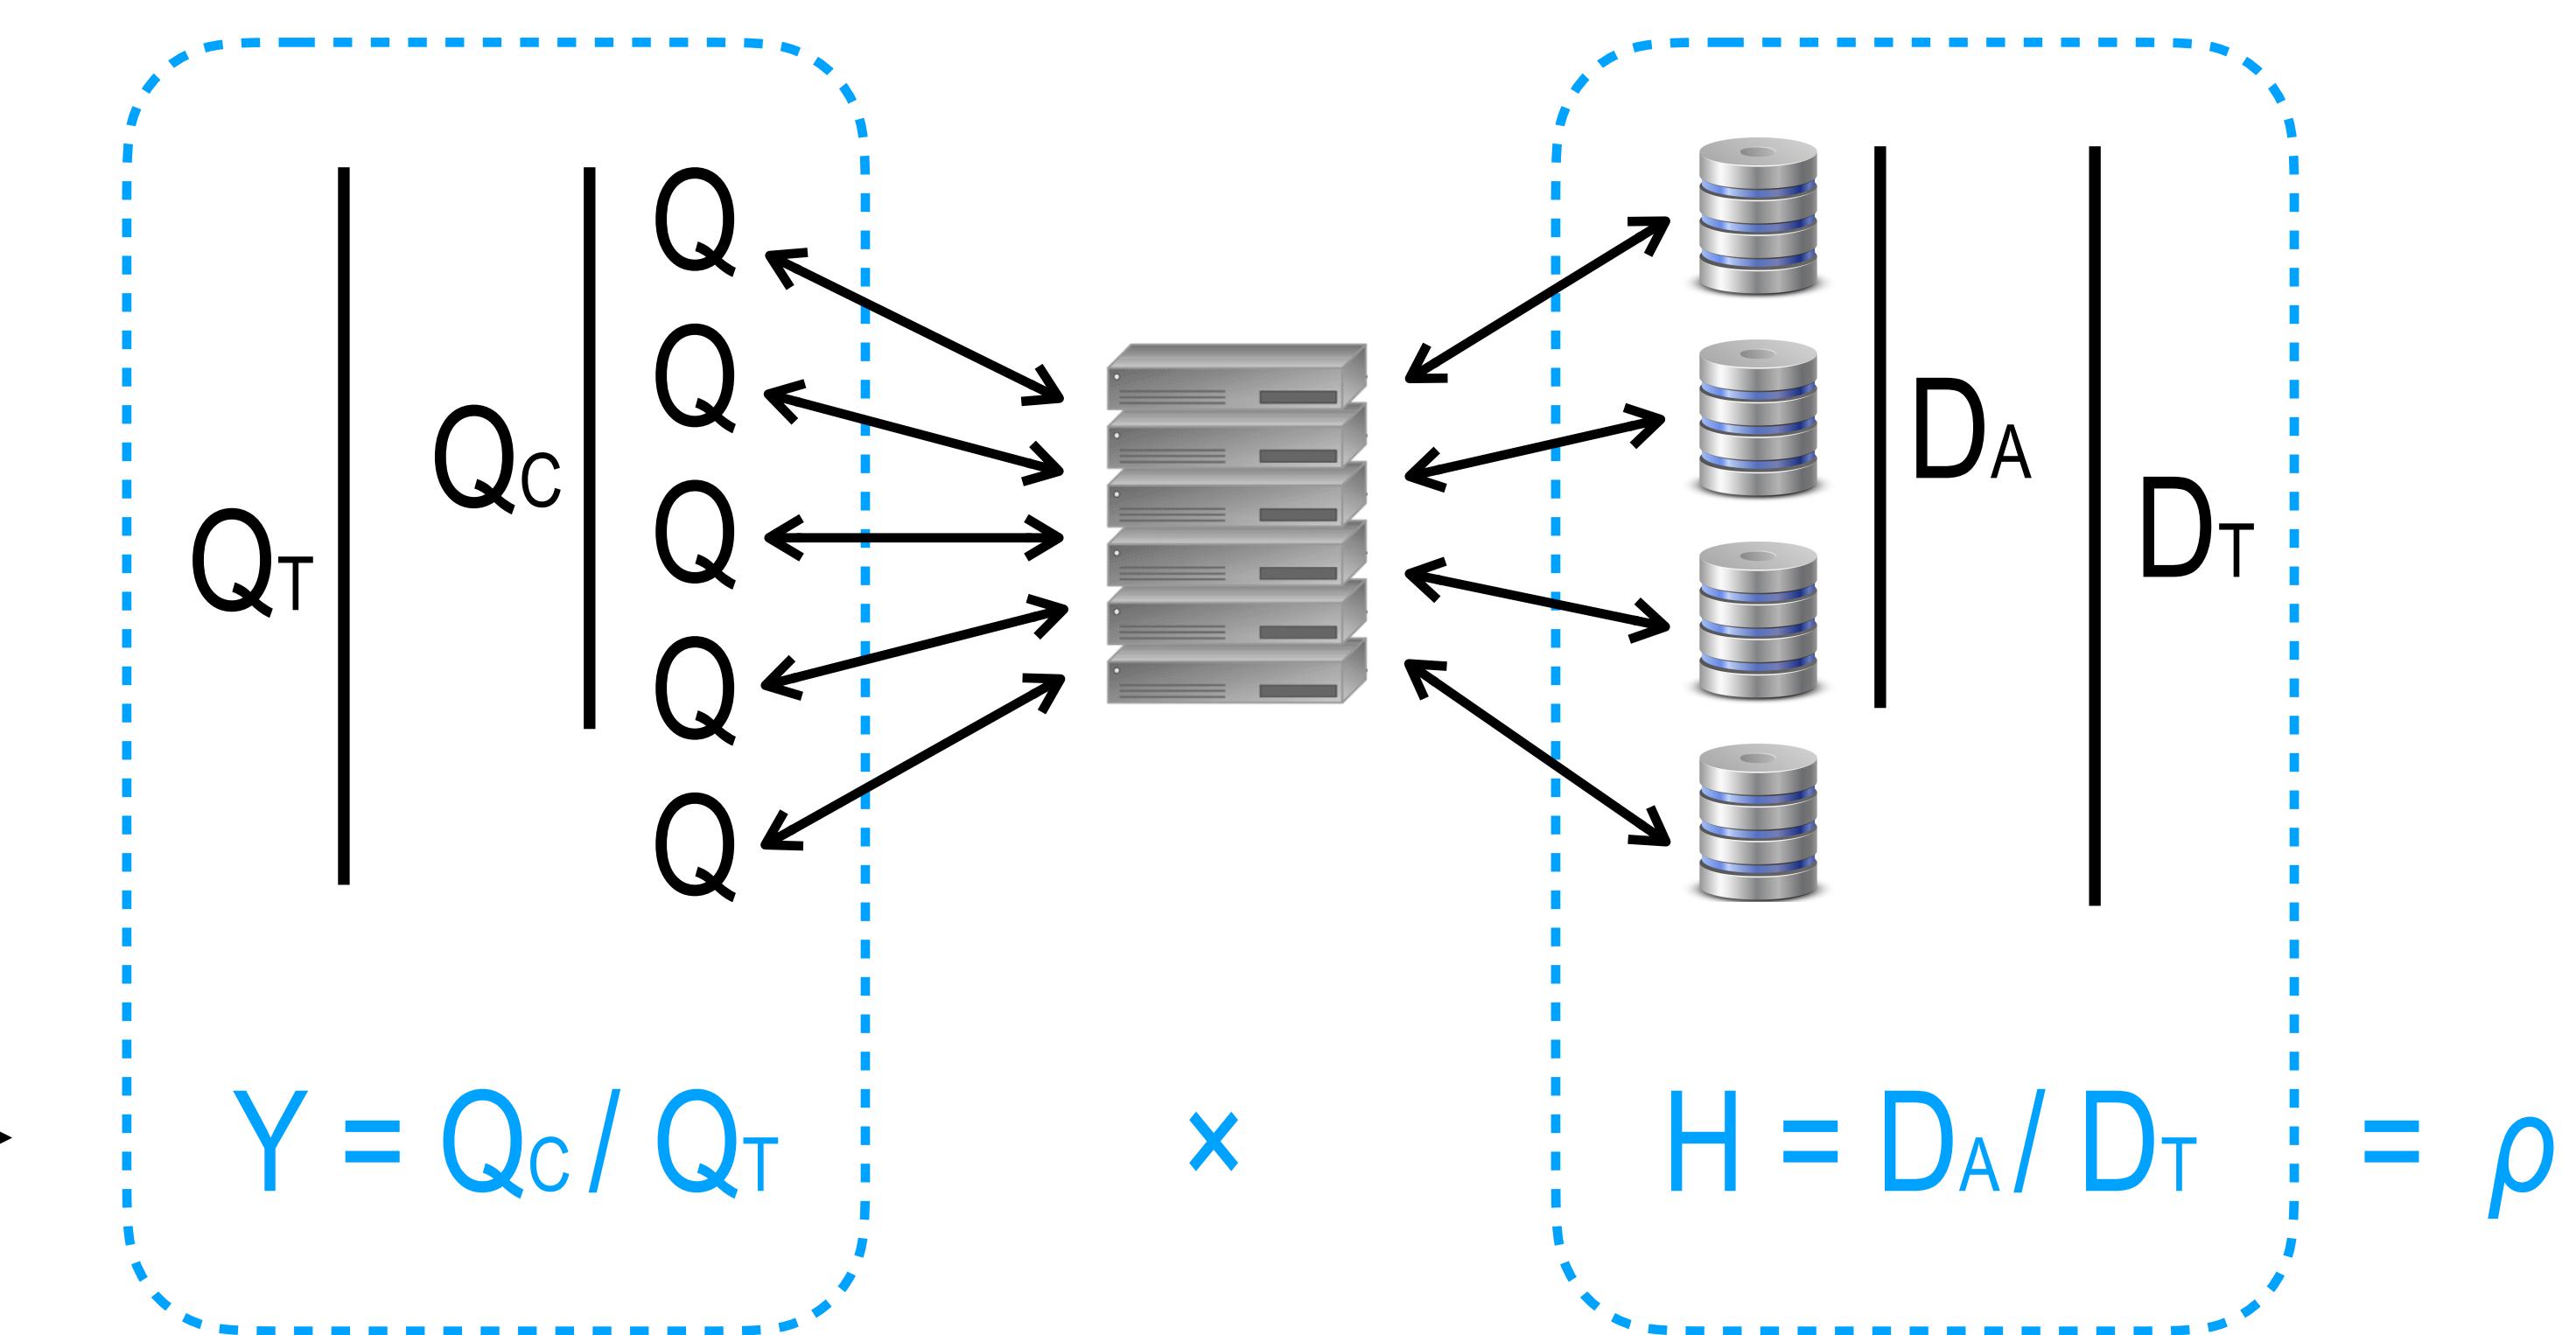
DQ Principle: "D×Q is constant"

(DQ value ρ determined by system configuration)

$$\text{Harvest } H = \frac{D_A}{D_T}$$

$$\text{Yield } Y = \frac{Q_c}{Q_T}$$

$$\text{DQ Principle: } H \times Y = \rho \longrightarrow$$



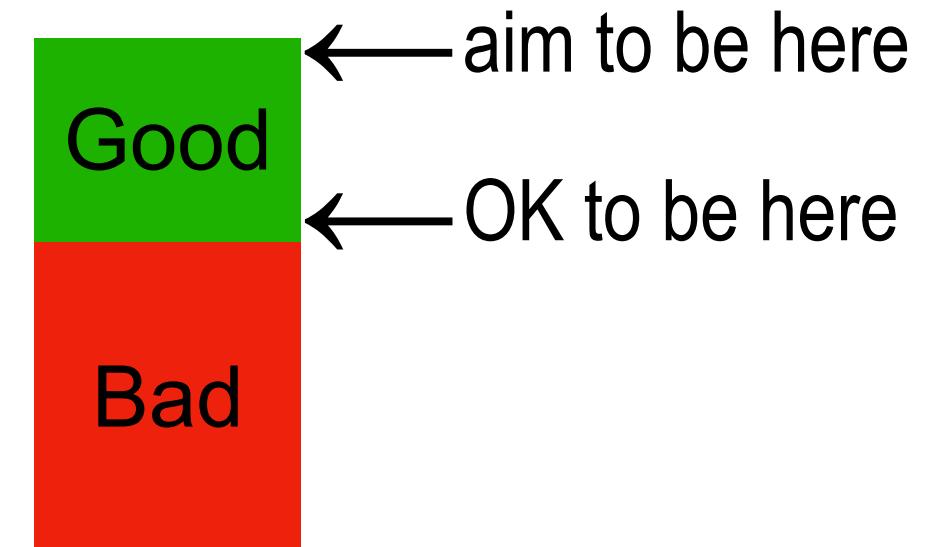
¹ Assumption: Disk is the only bottleneck. All the other resources (e.g., network bandwidth) are plentiful.

Recap: Failure modes

- Fail-stop (TMR)
- Fail-fast (Redundant invariant checks)
- Fail-safe
 - *OK to fail, as long as safety is not compromised*
- Fail-soft (Weaker spec)
 - *Redundant resources for top band of acceptable system behavior*
 - *Harvest/yield and the DQ principle in data-intensive parallel systems*

Recap: Failure modes

- Fail-stop (TMR)
- Fail-fast (Redundant invariant checks)
- Fail-safe
 - *OK to fail, as long as safety is not compromised*
- Fail-soft (Weaker spec)
 - *Redundant resources for top band of acceptable system behavior*
 - *Harvest/yield and the DQ principle in data-intensive parallel systems*





How to reduce unavailability by $10 \times$?

How to reduce unavailability by 10× ?

$$\text{Availability} = \frac{\text{MTTF}}{\text{MTBF}}$$

How to reduce unavailability by 10 \times ?

$$\text{Unavailability} \cong \frac{\text{MTTR}}{\text{MTTF}}$$

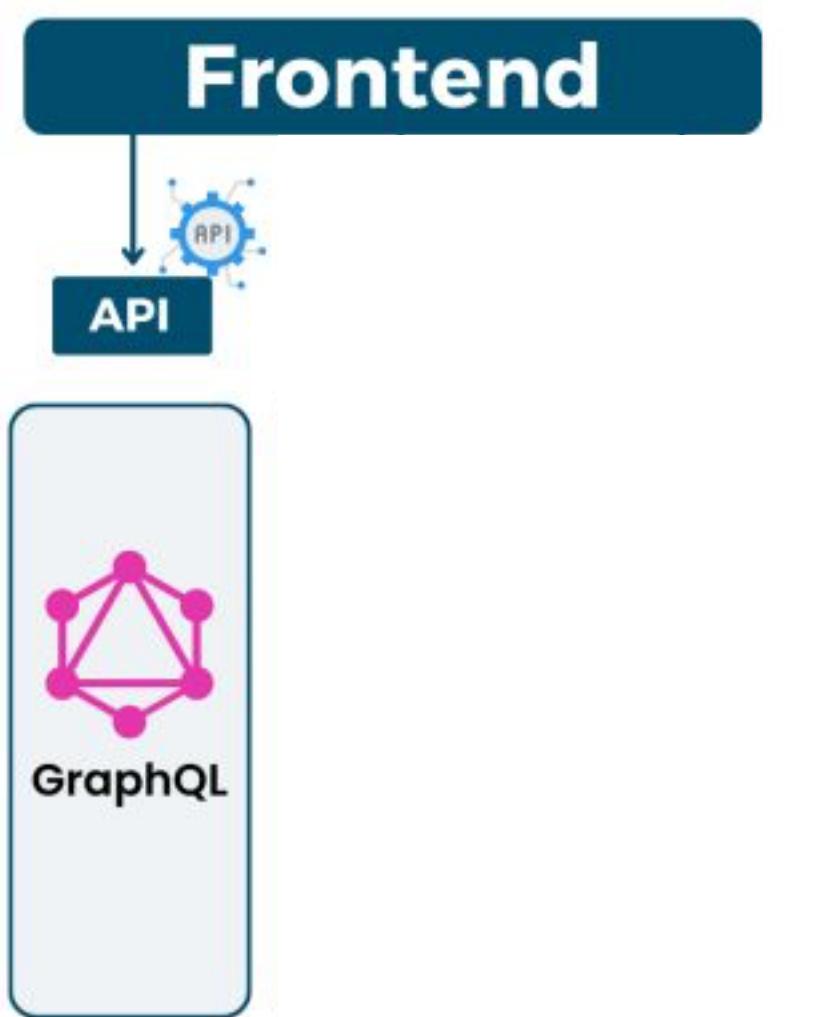
How to reduce unavailability by 10× ?

$$\text{Unavailability} \cong \frac{\text{MTTR}}{\text{MTTF}} \uparrow \times 10$$

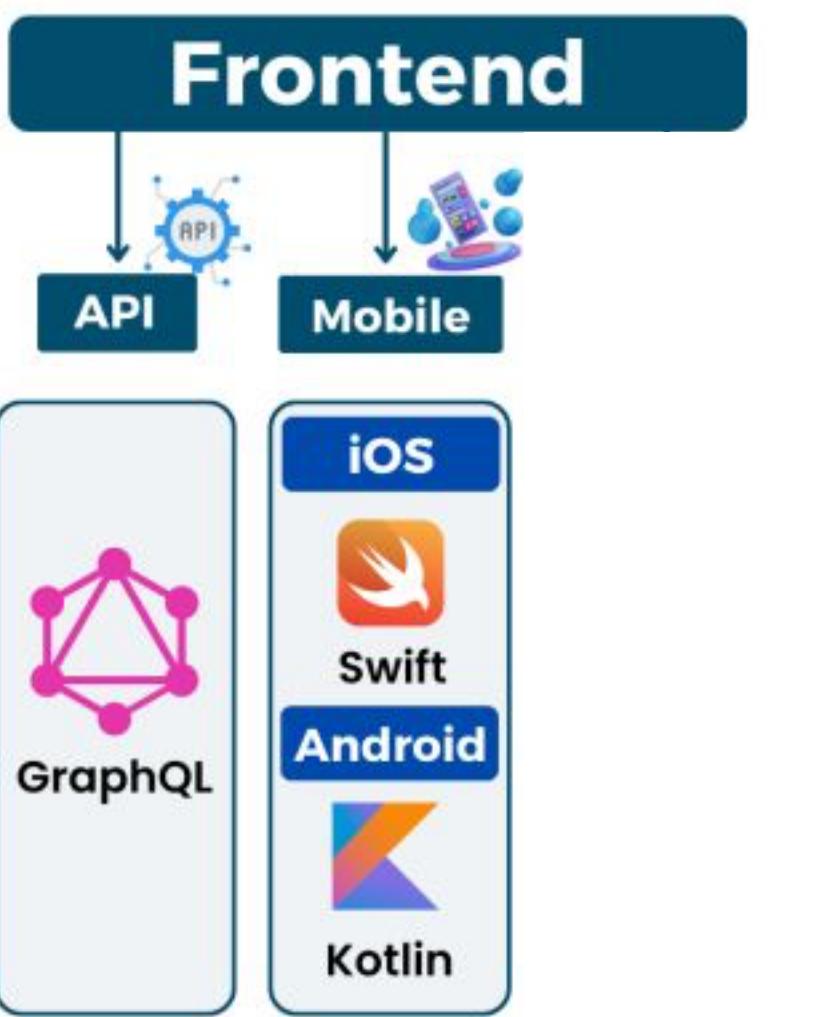
NETFLIX



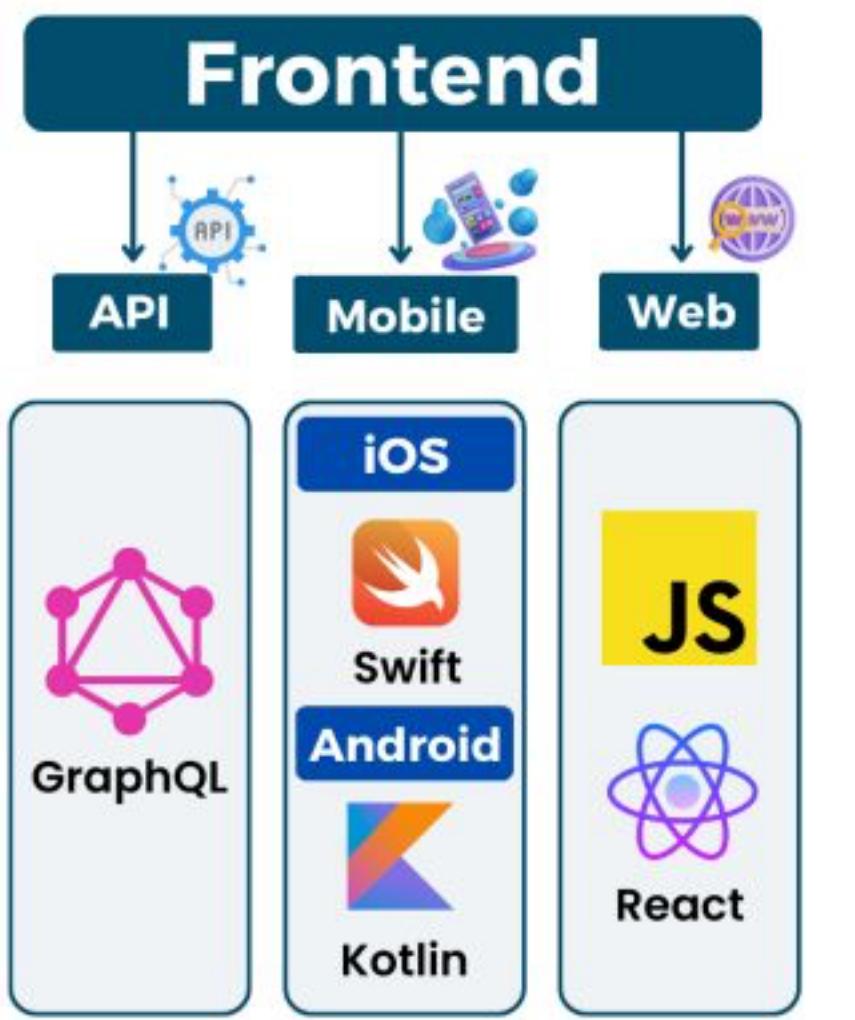
NETFLIX



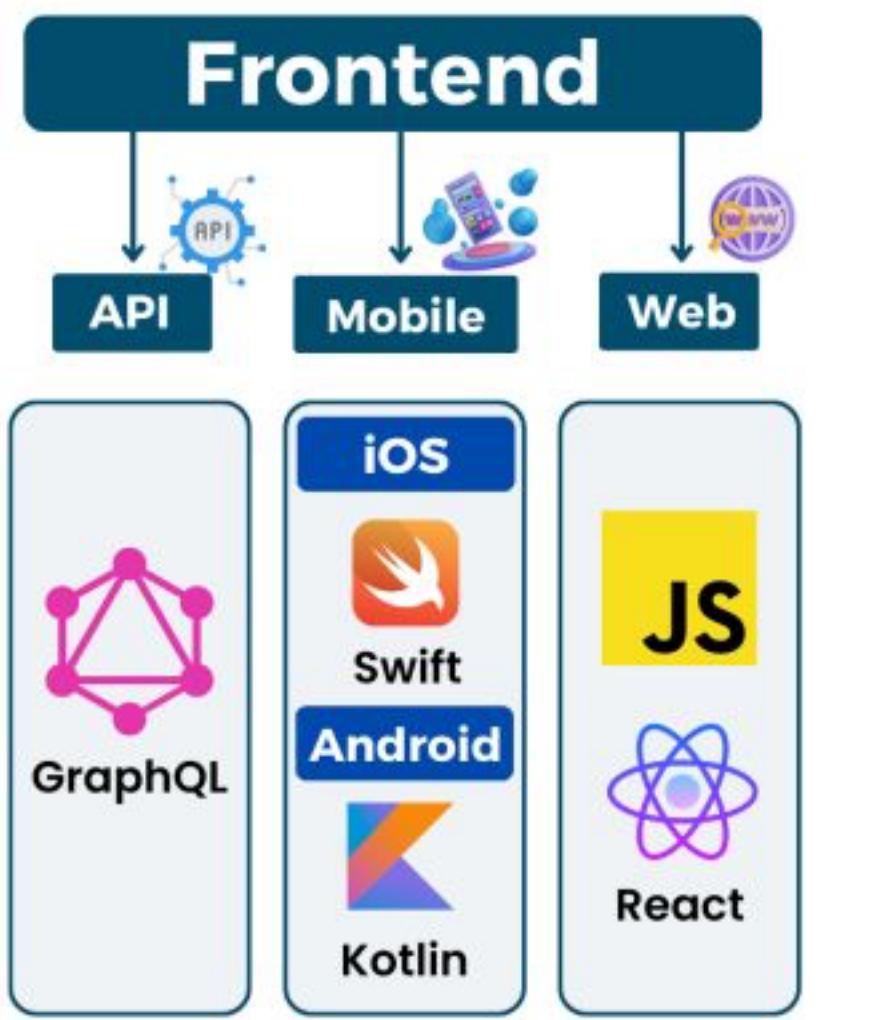
NETFLIX



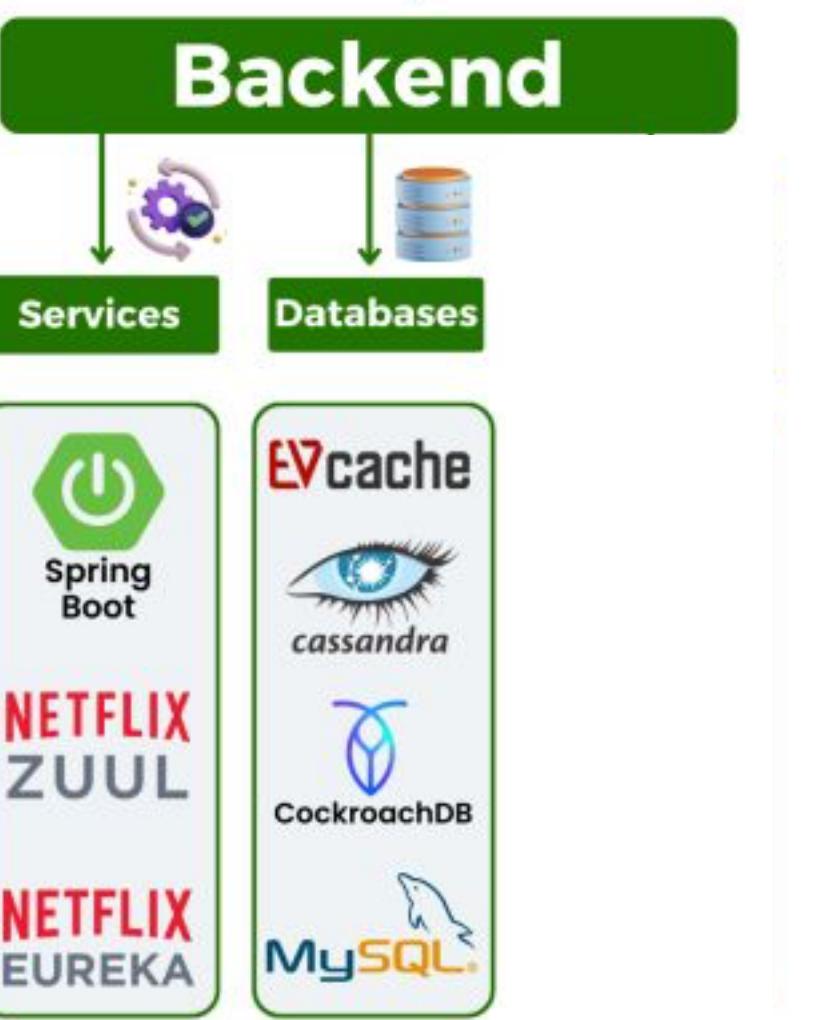
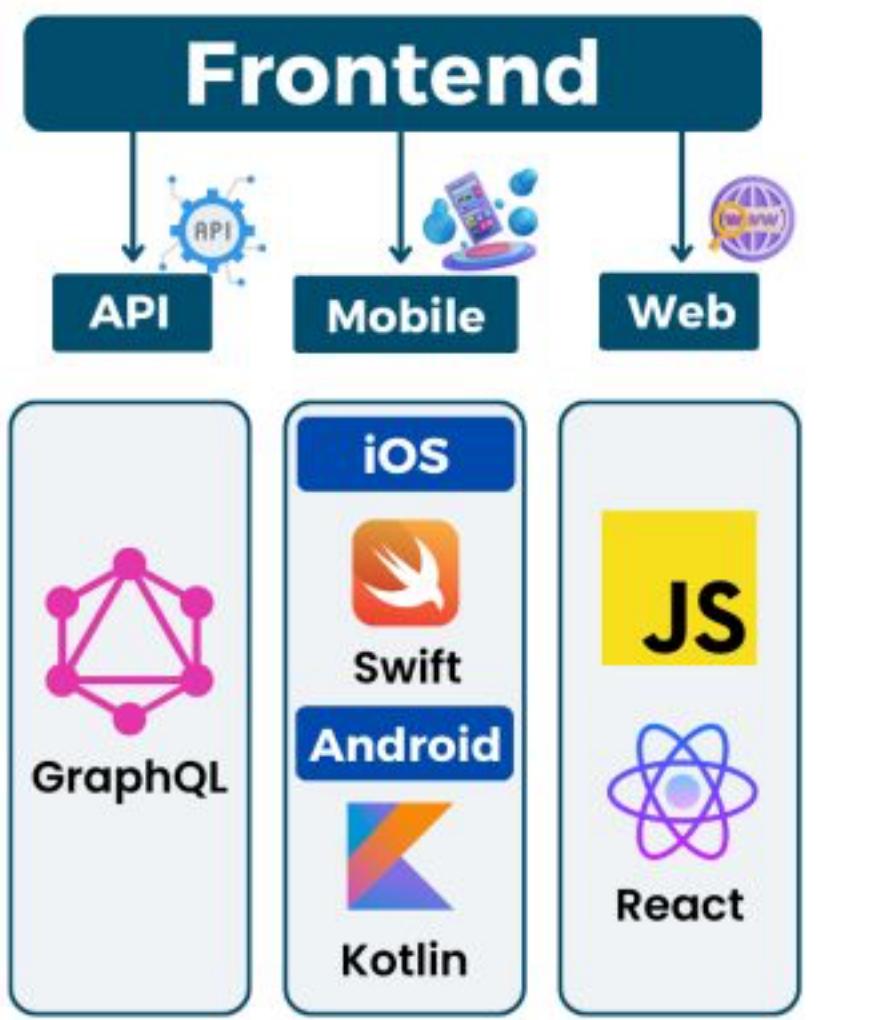
NETFLIX



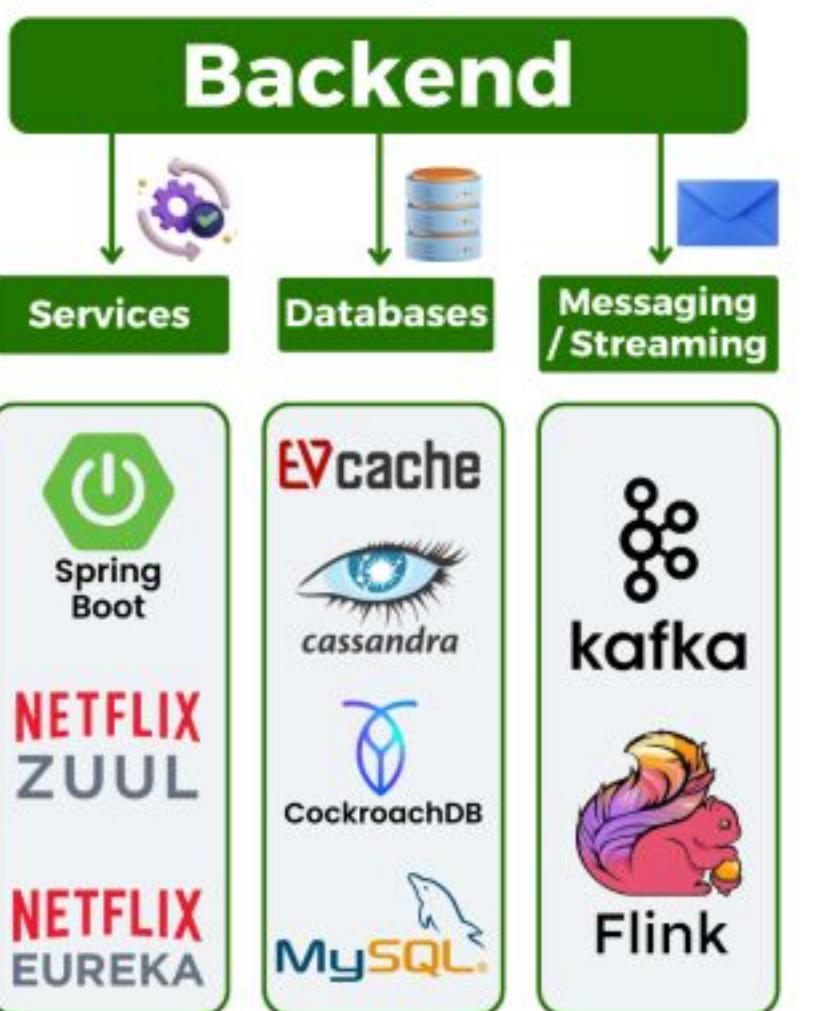
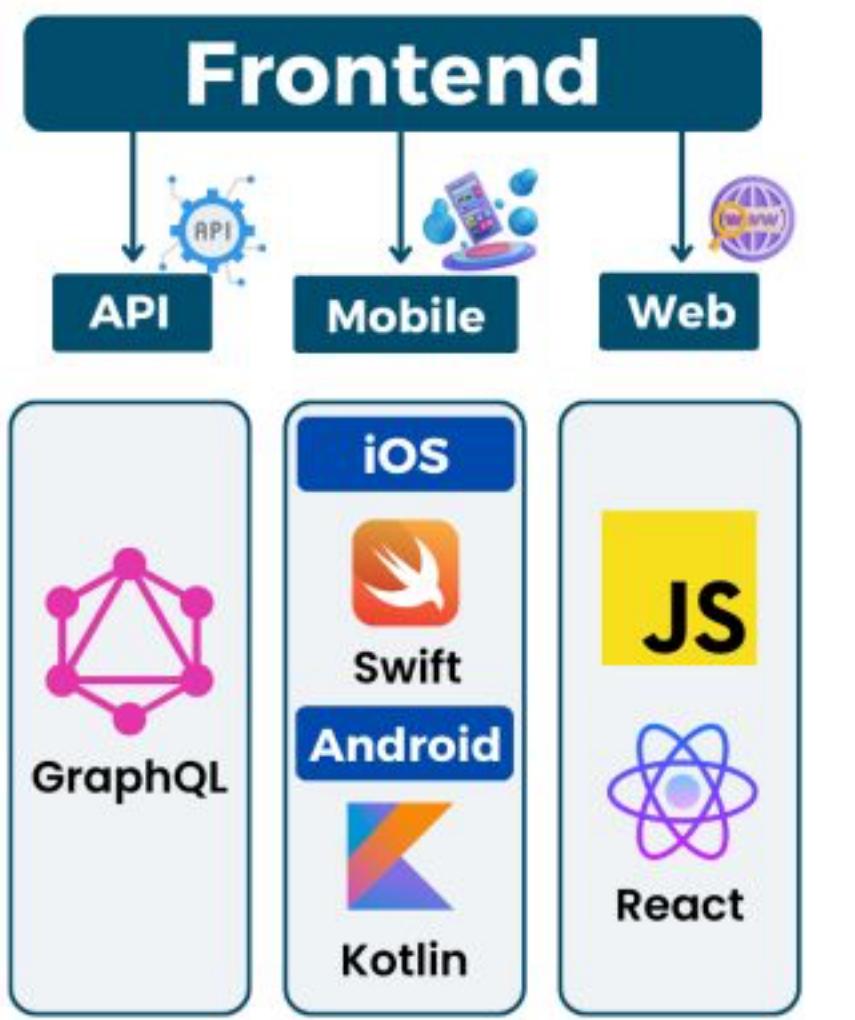
NETFLIX



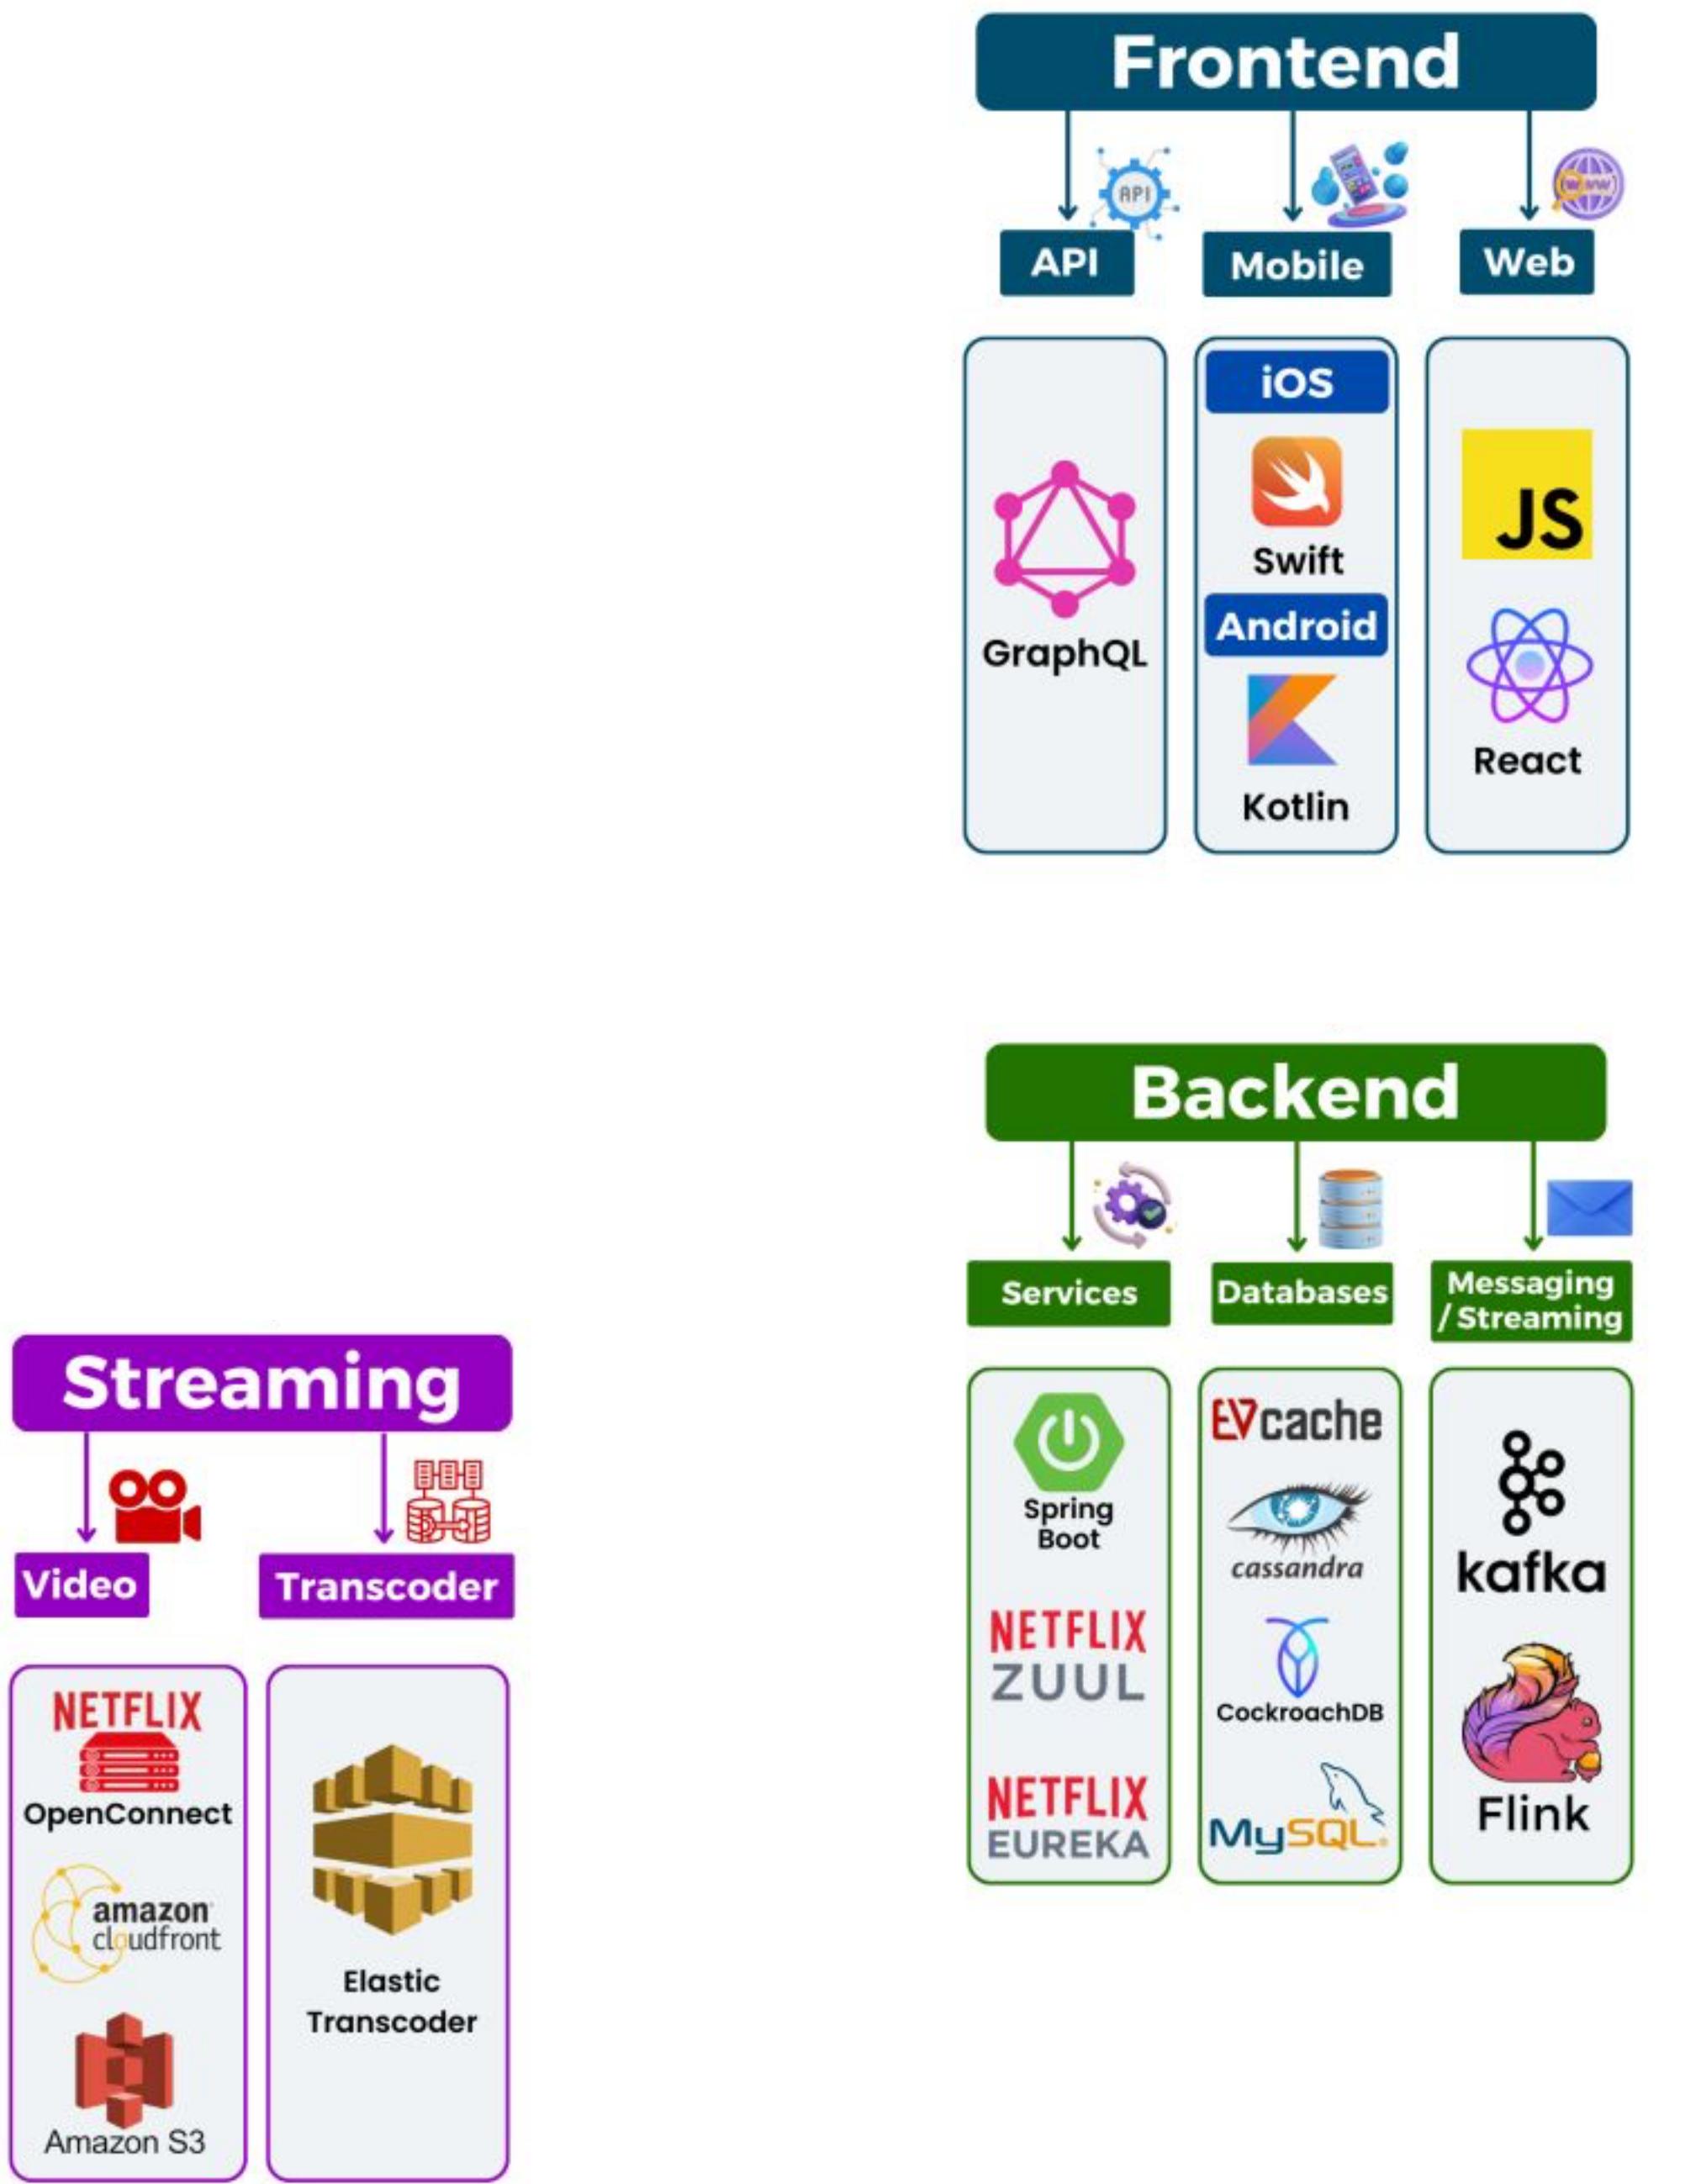
NETFLIX



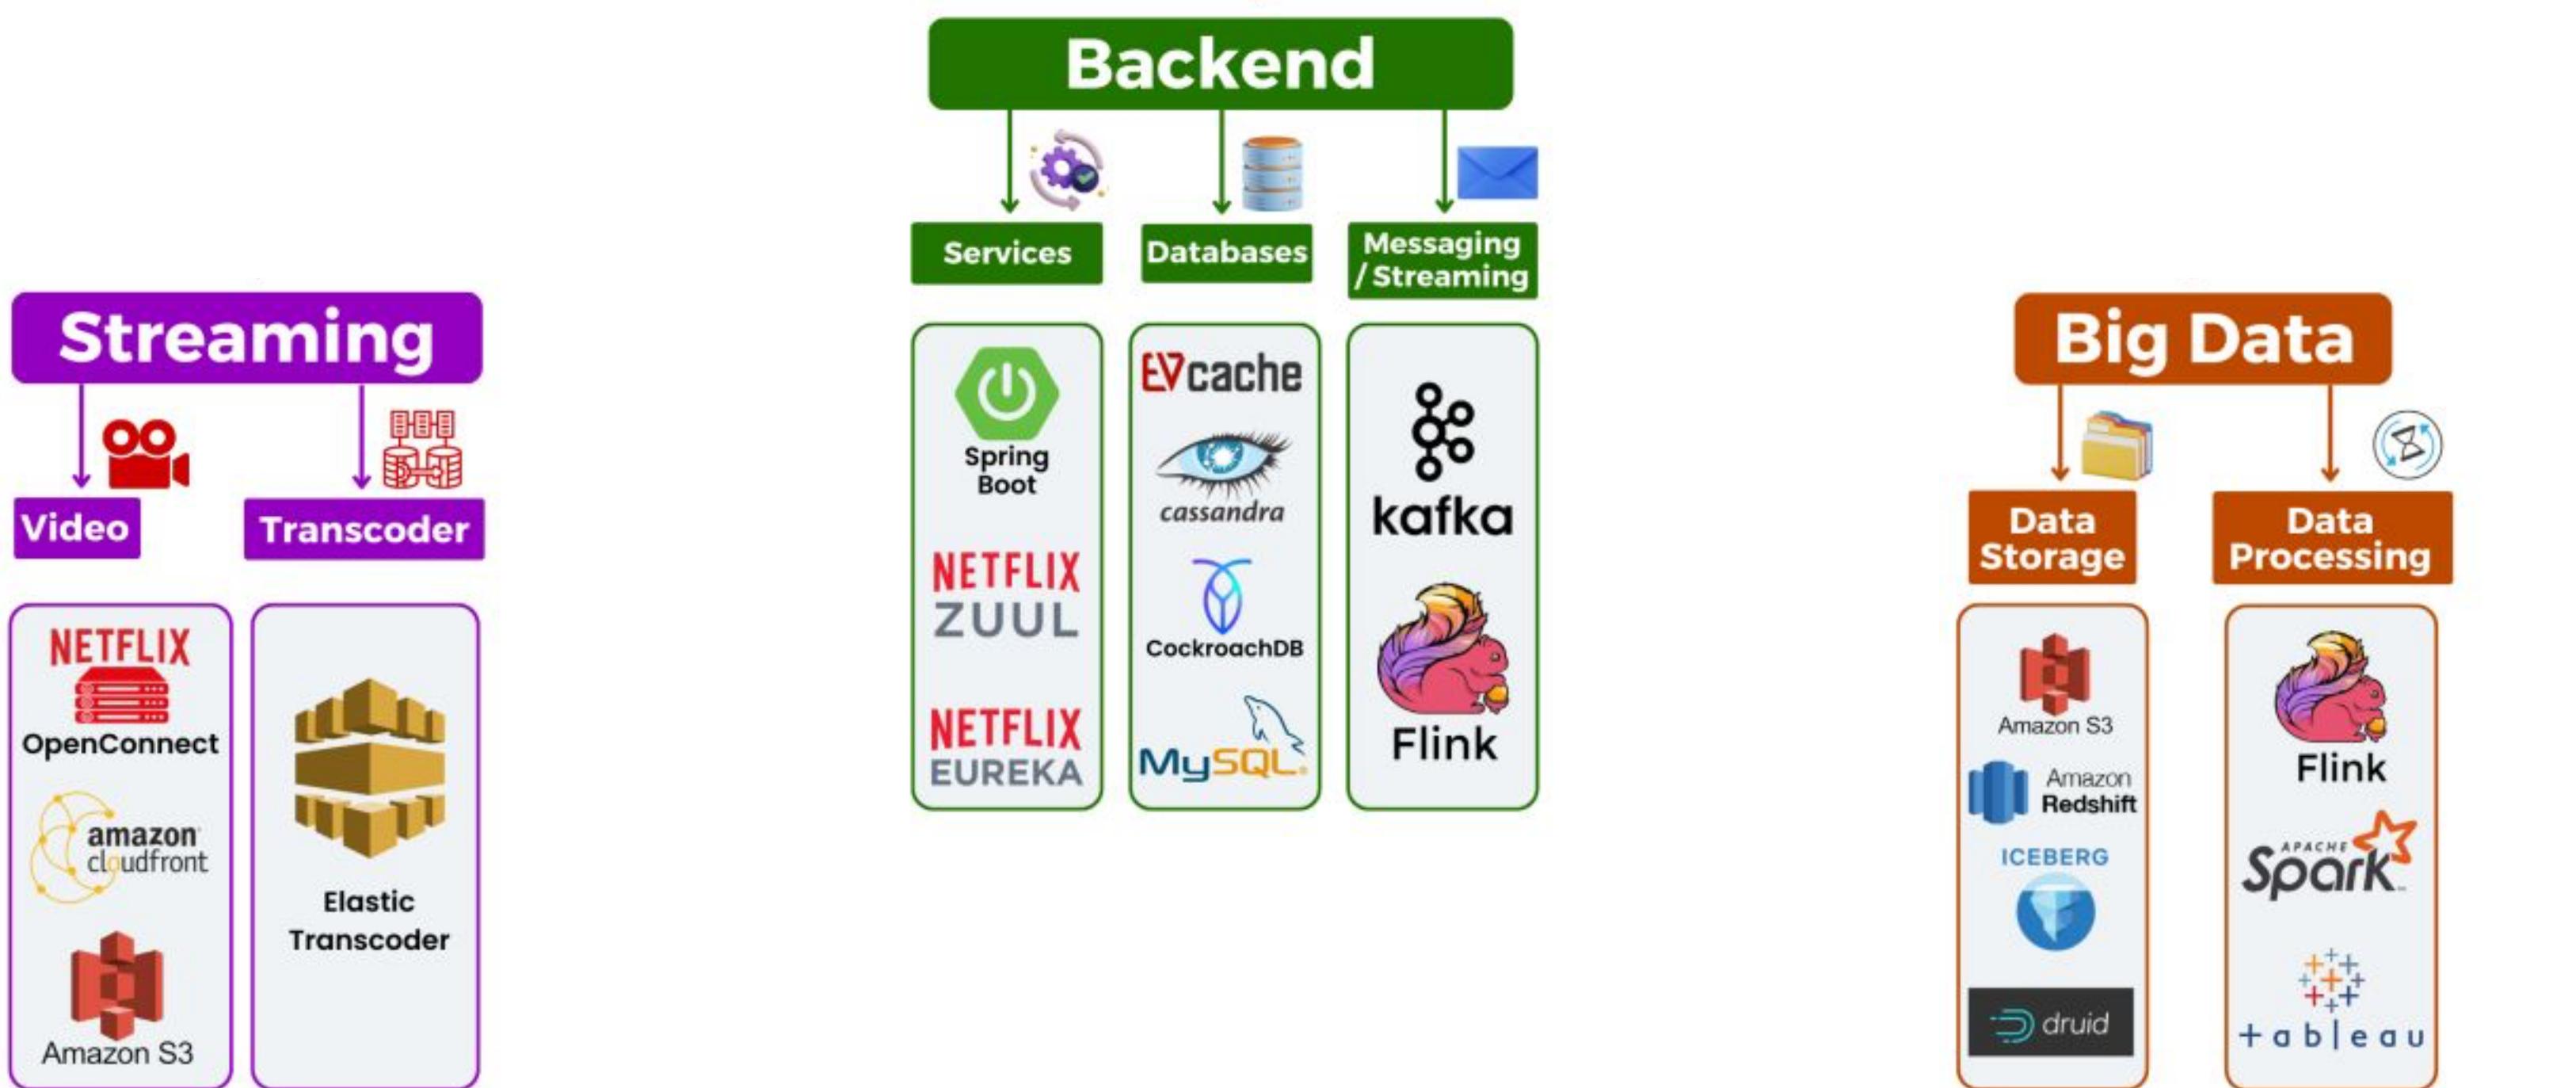
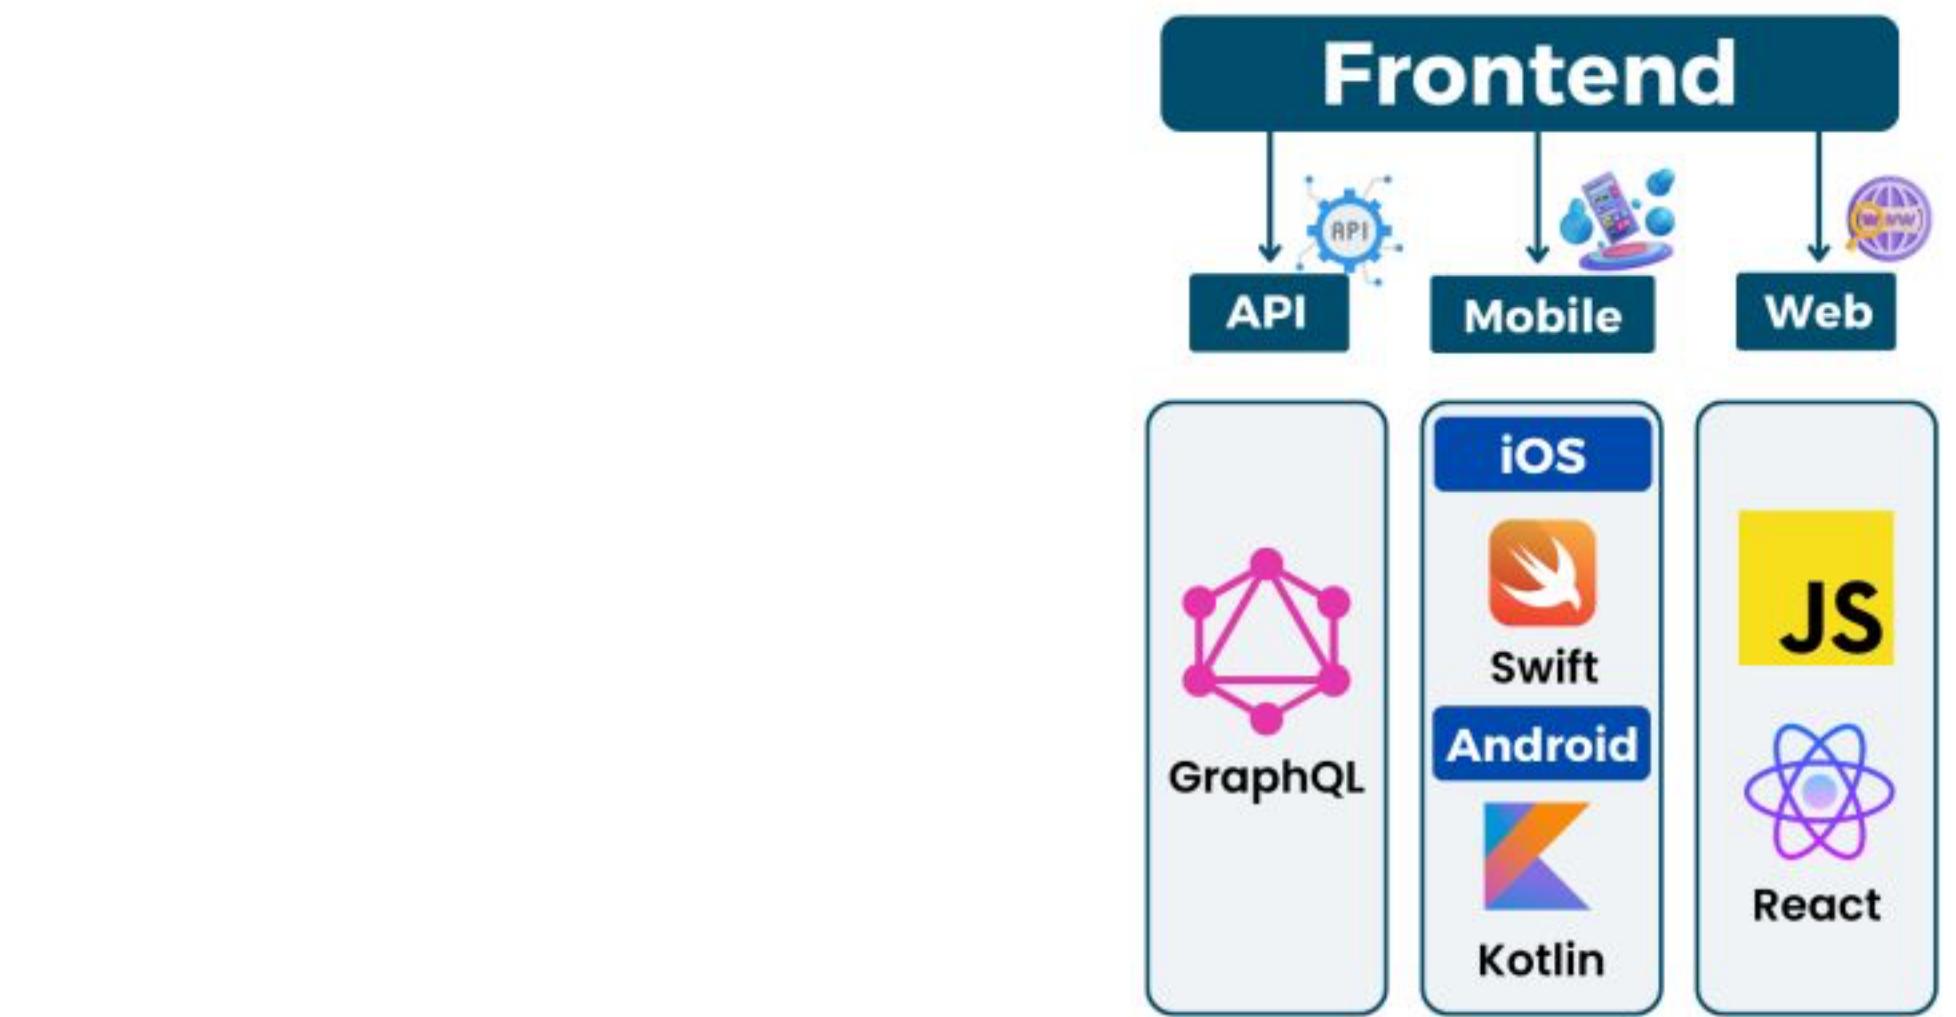
NETFLIX



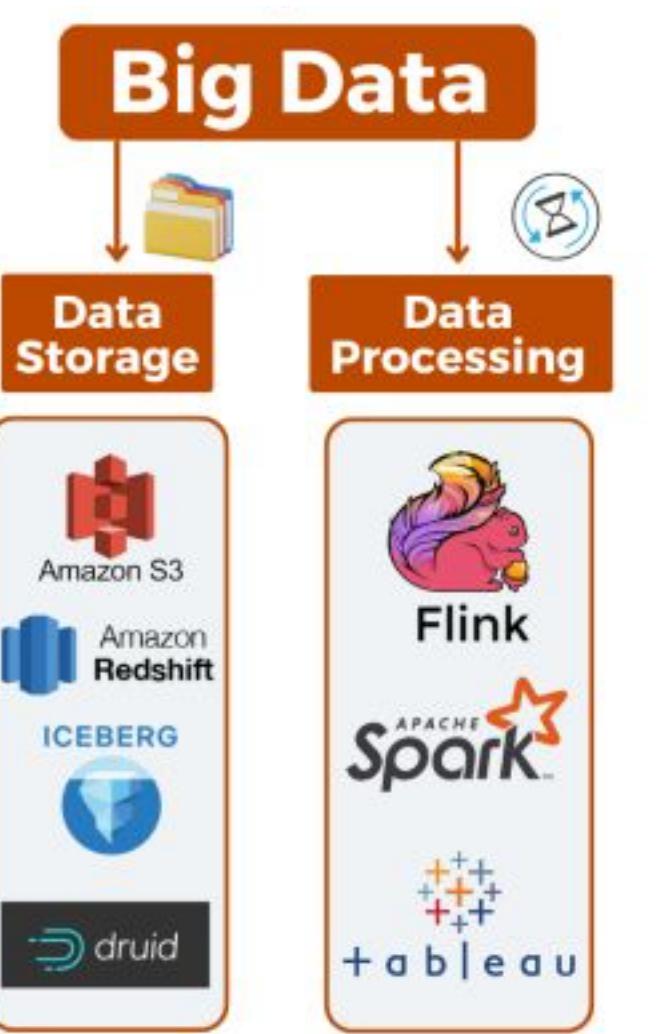
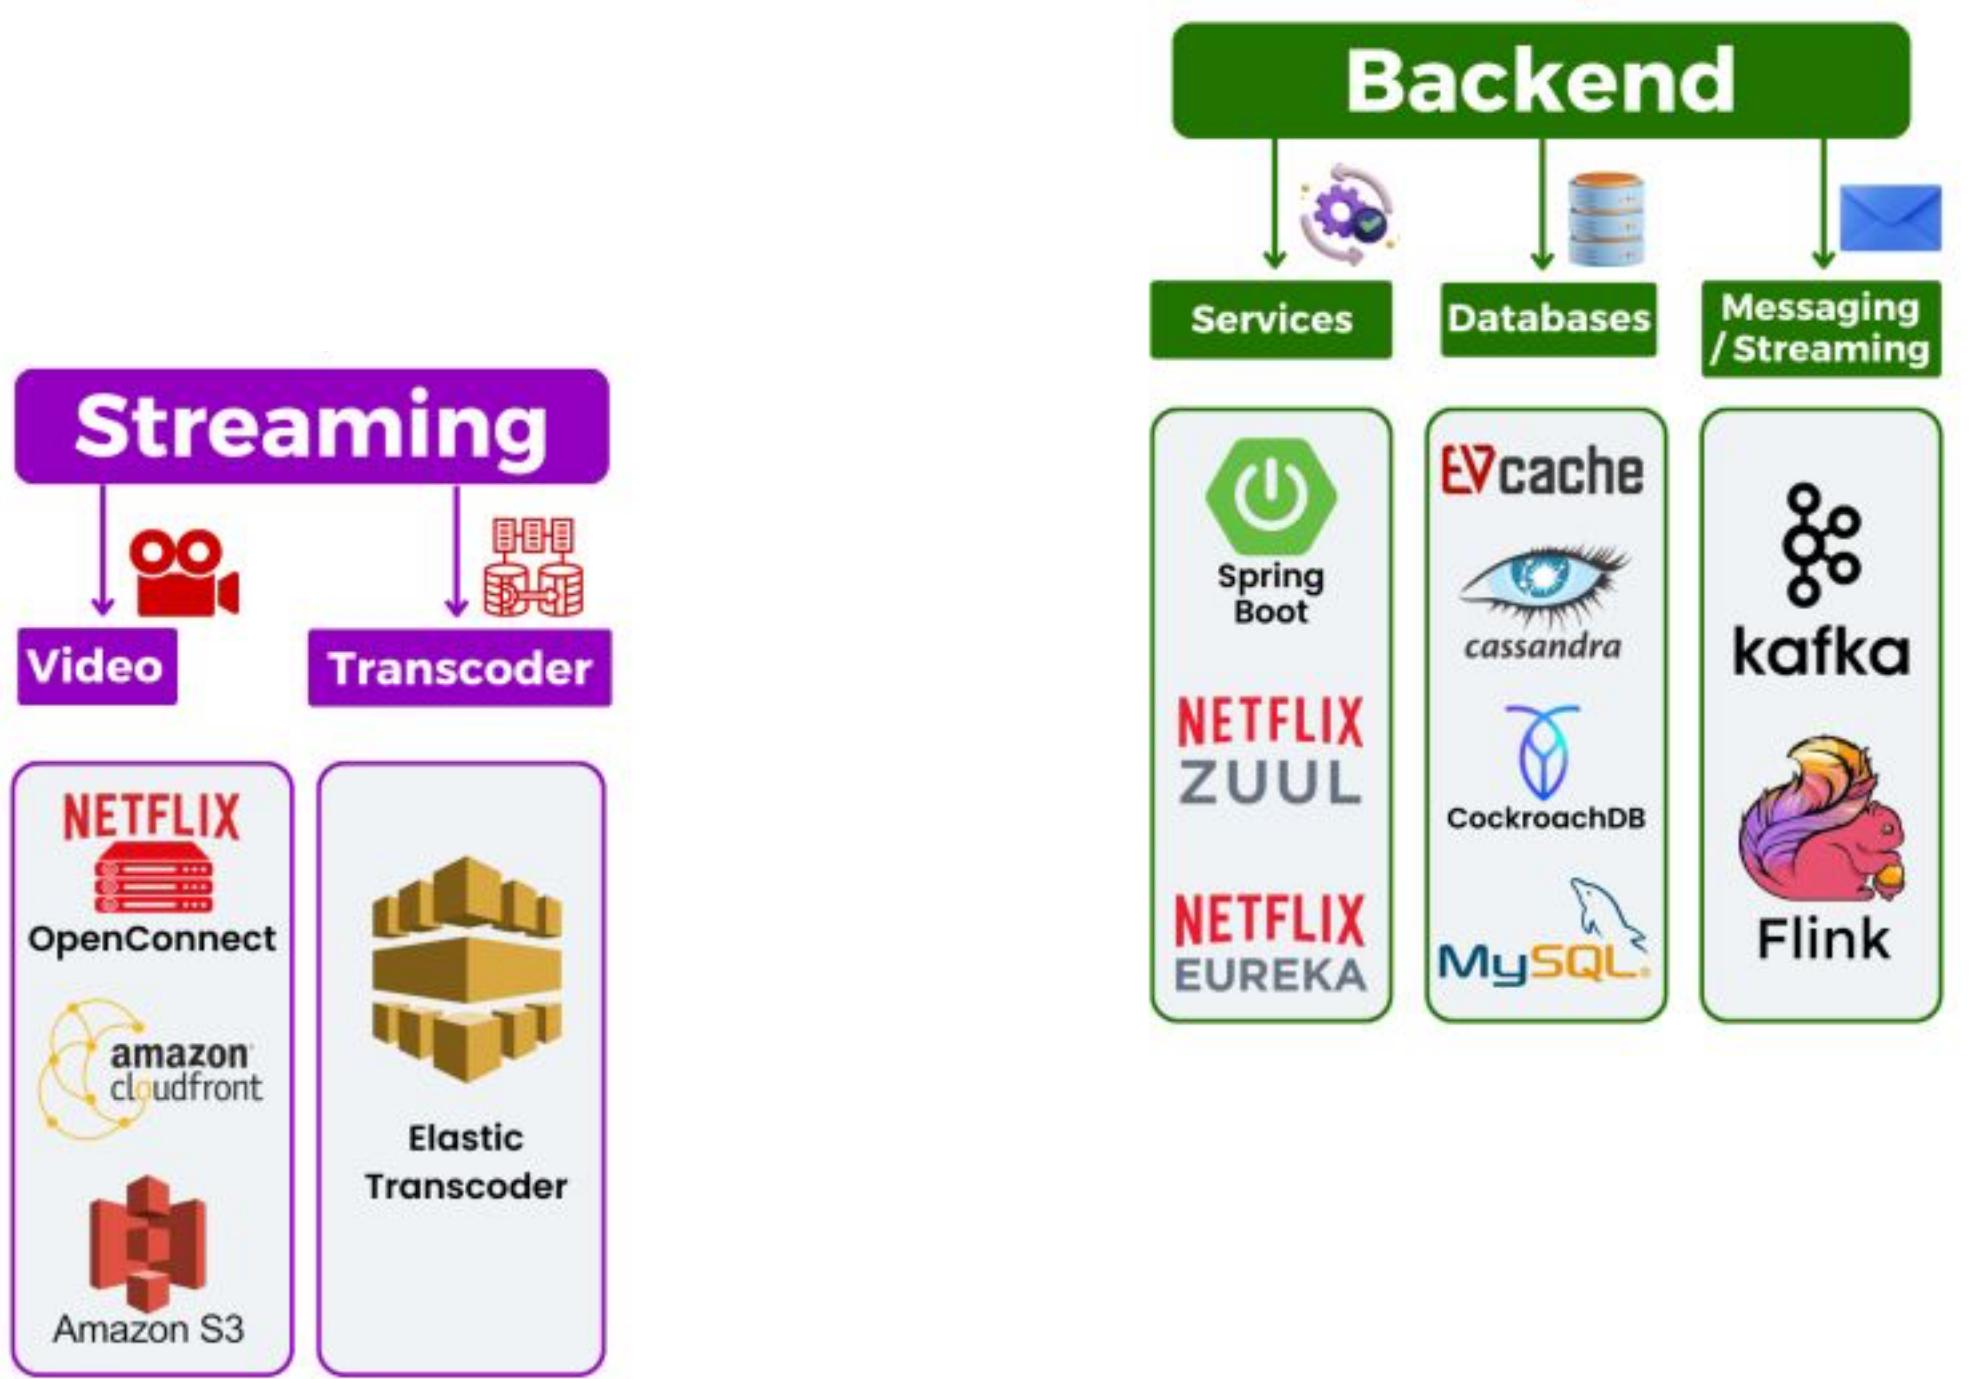
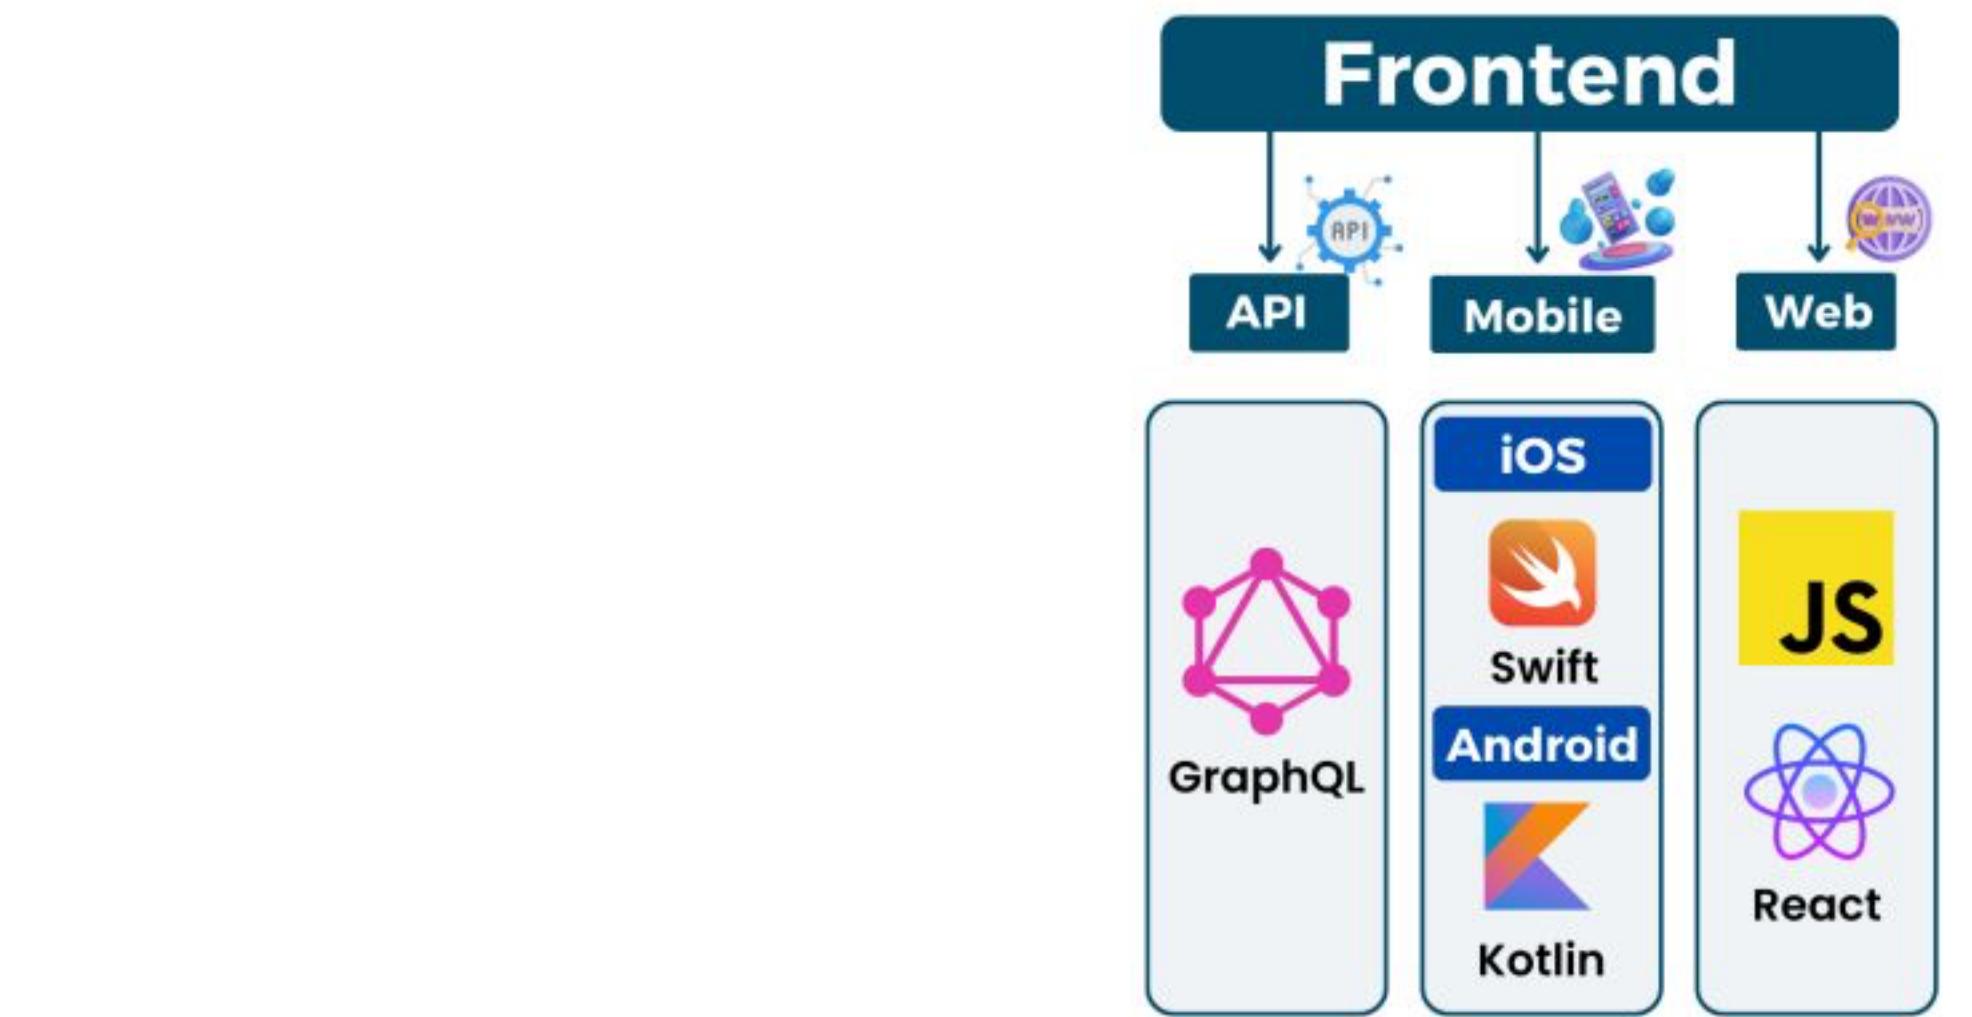
NETFLIX

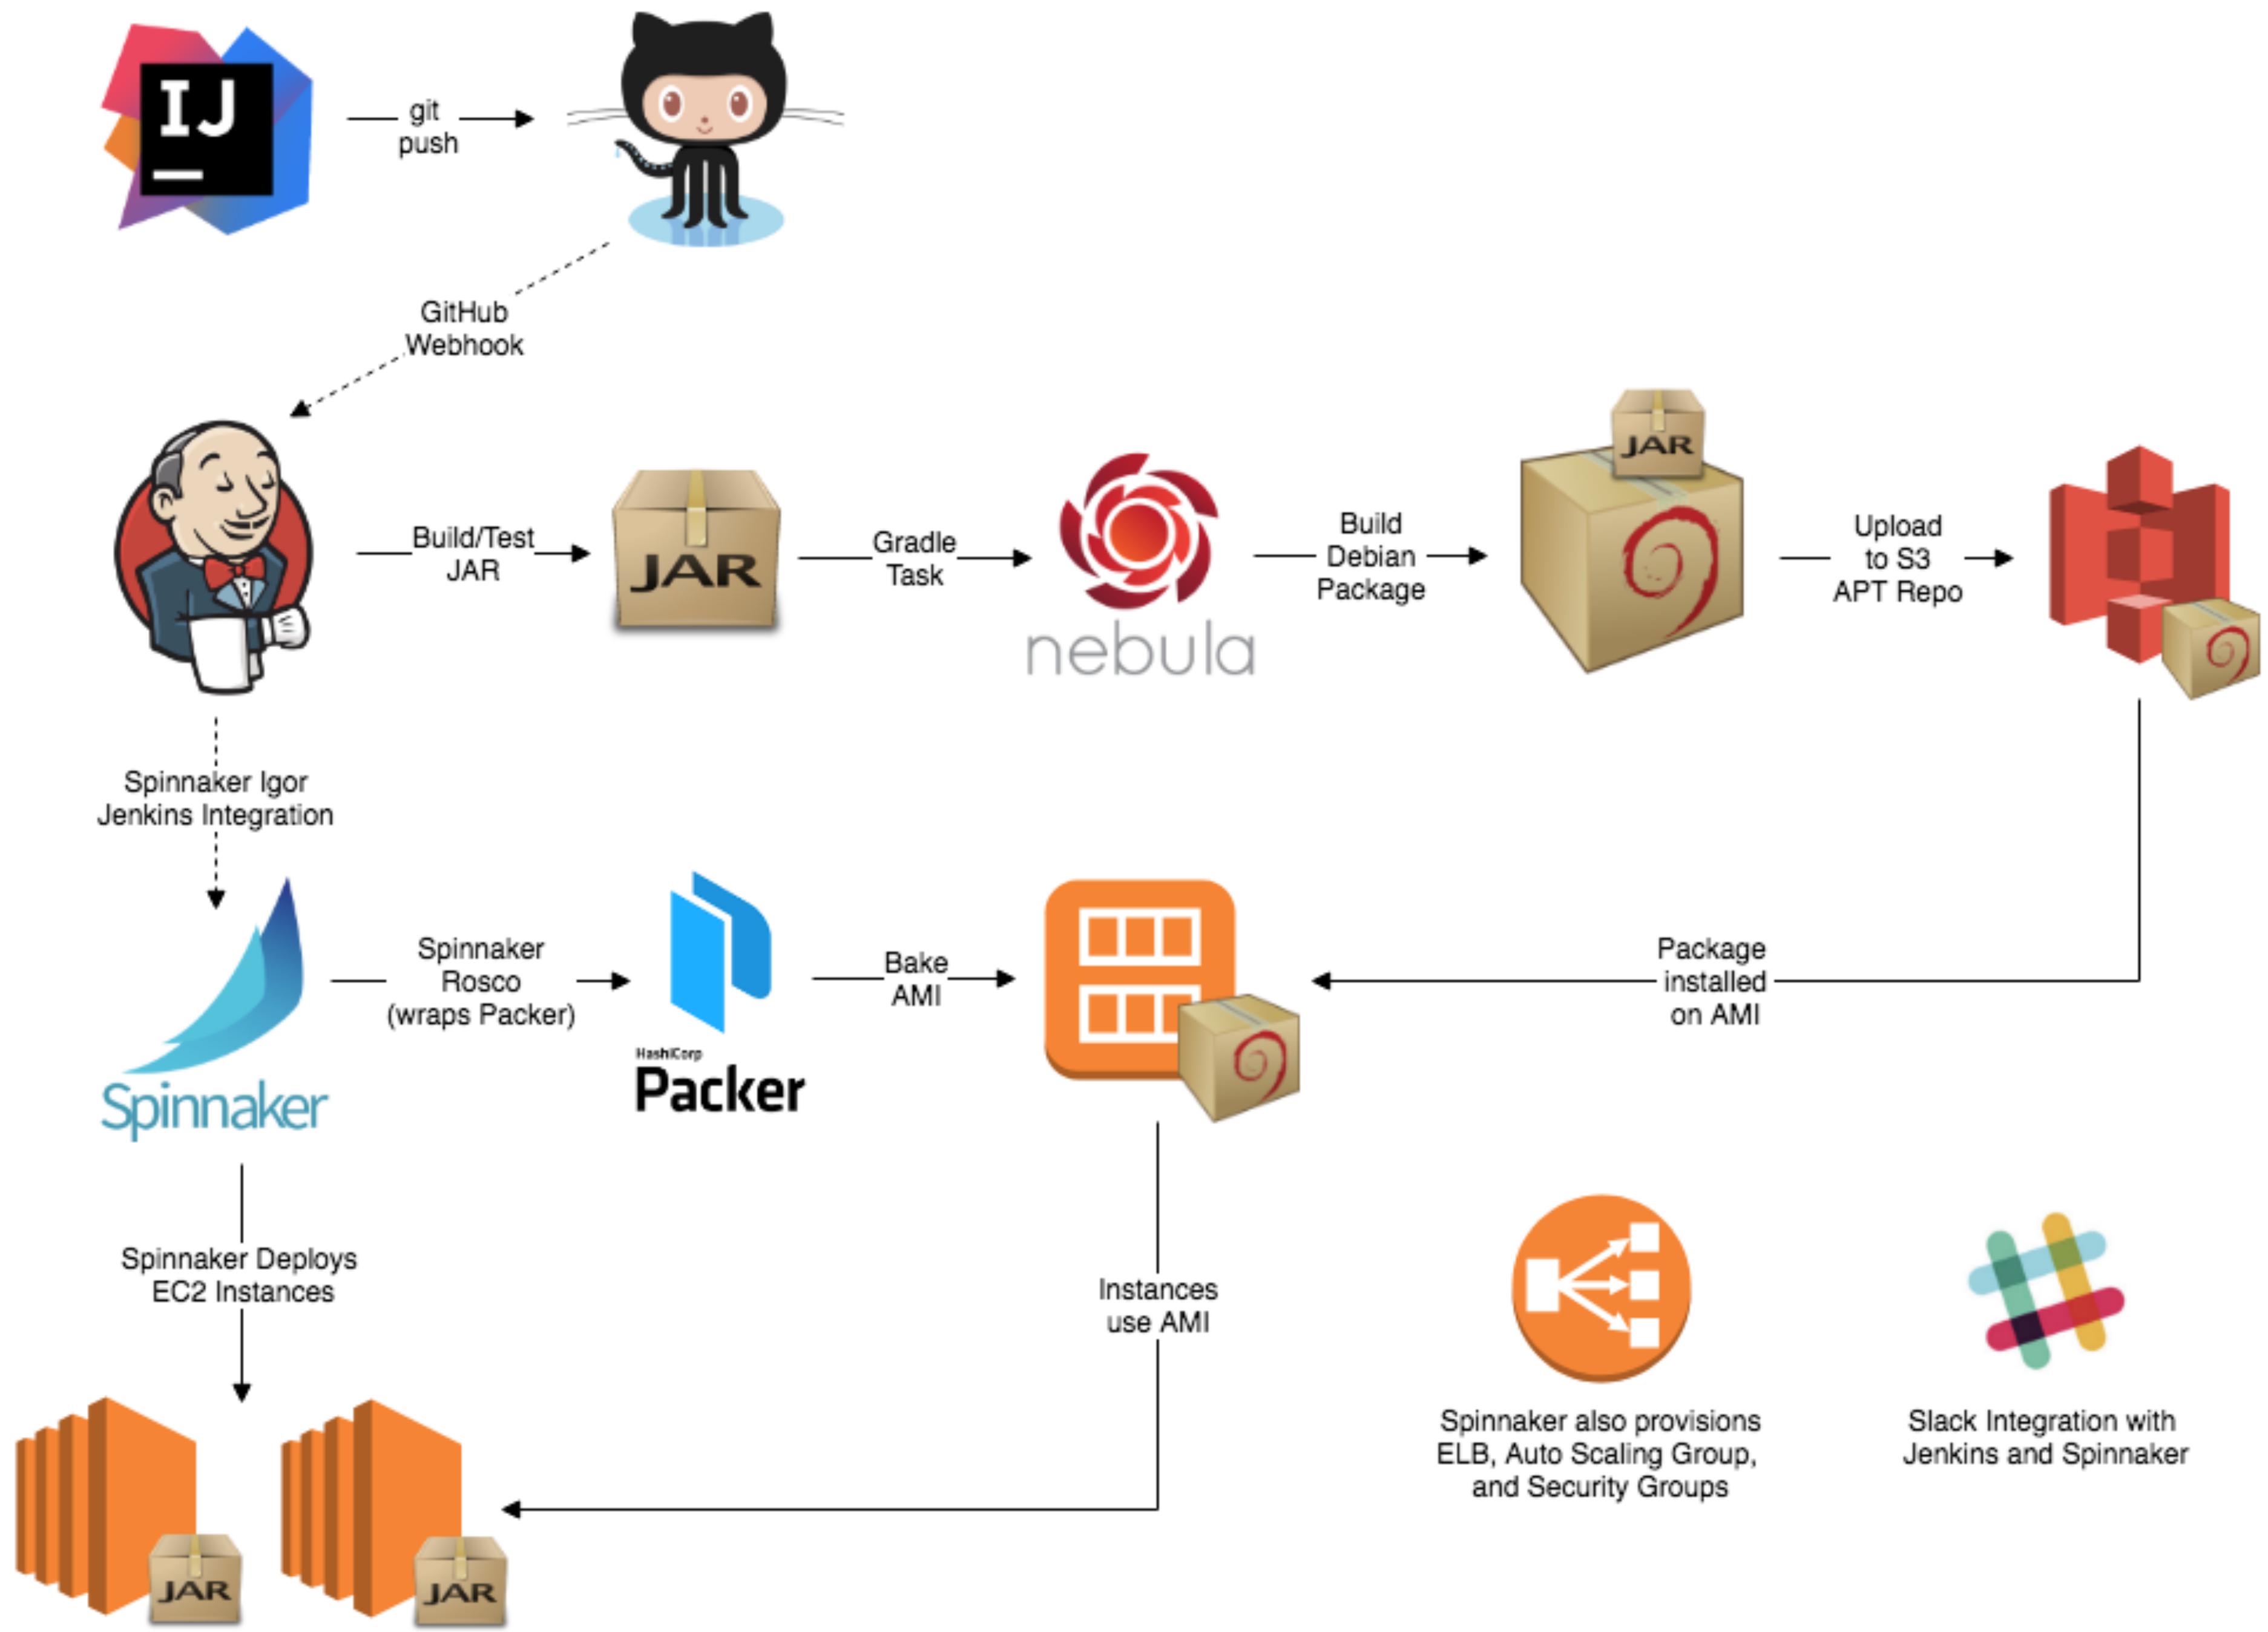


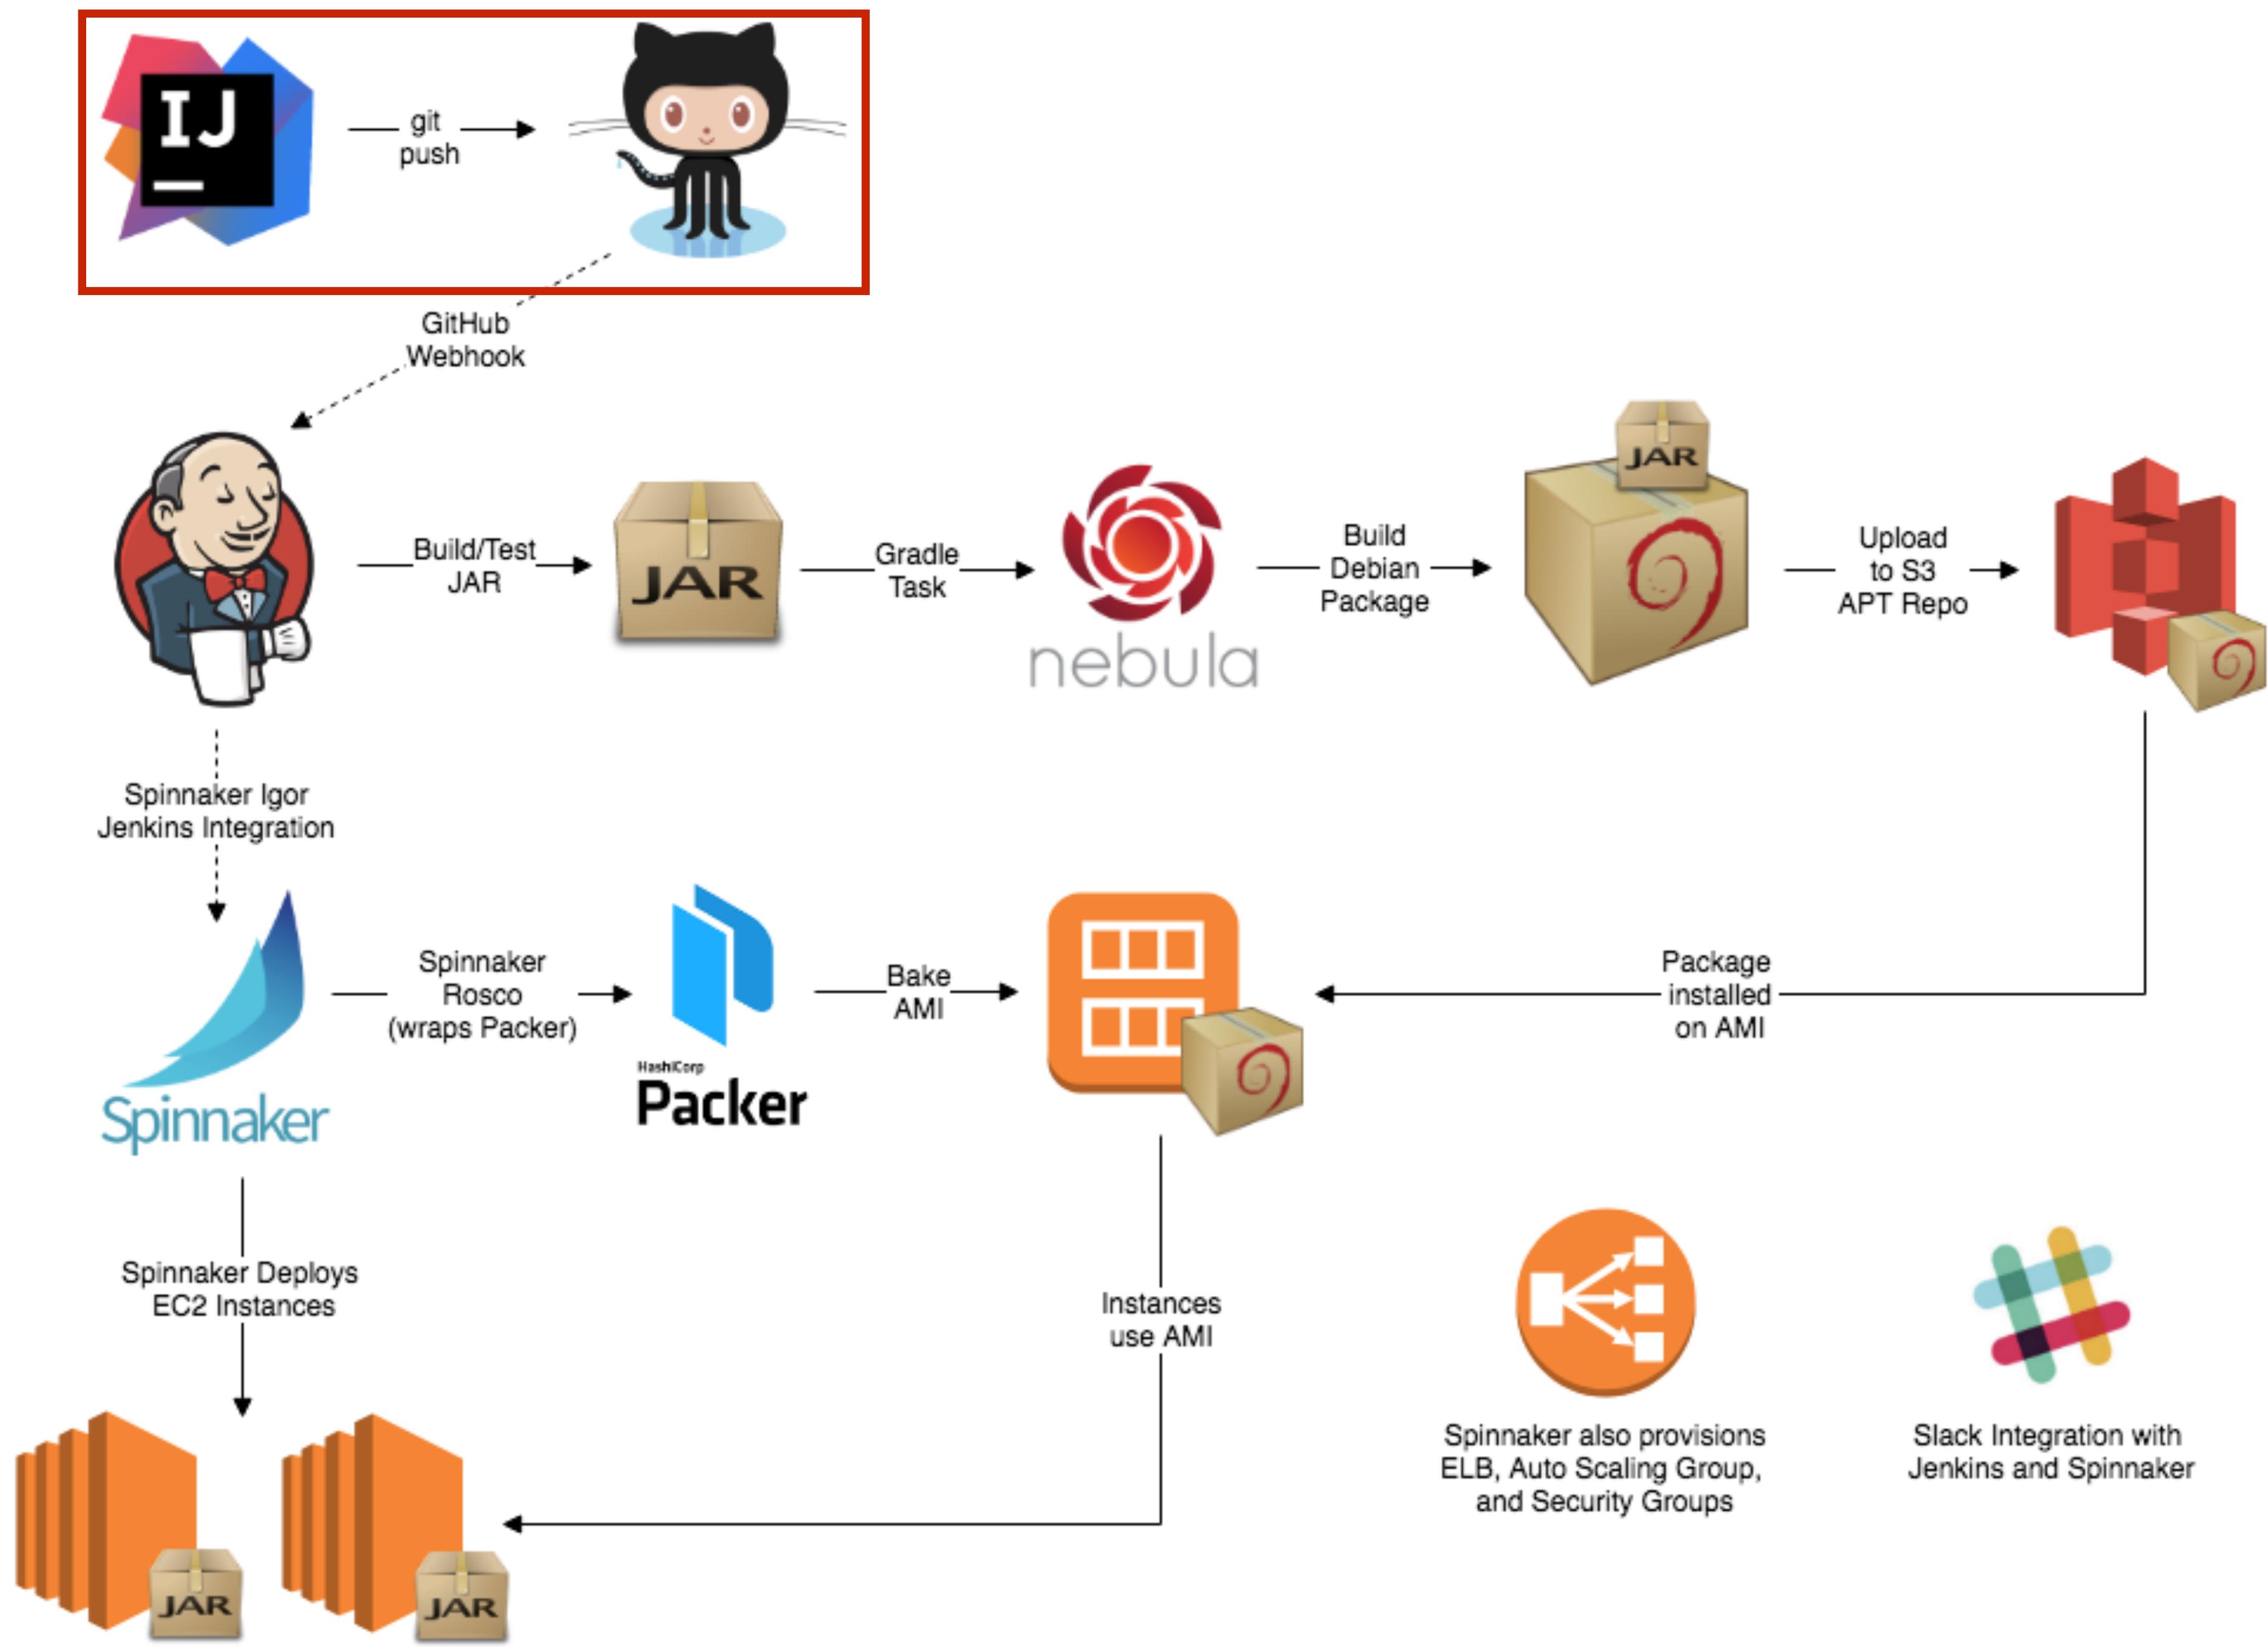
NETFLIX

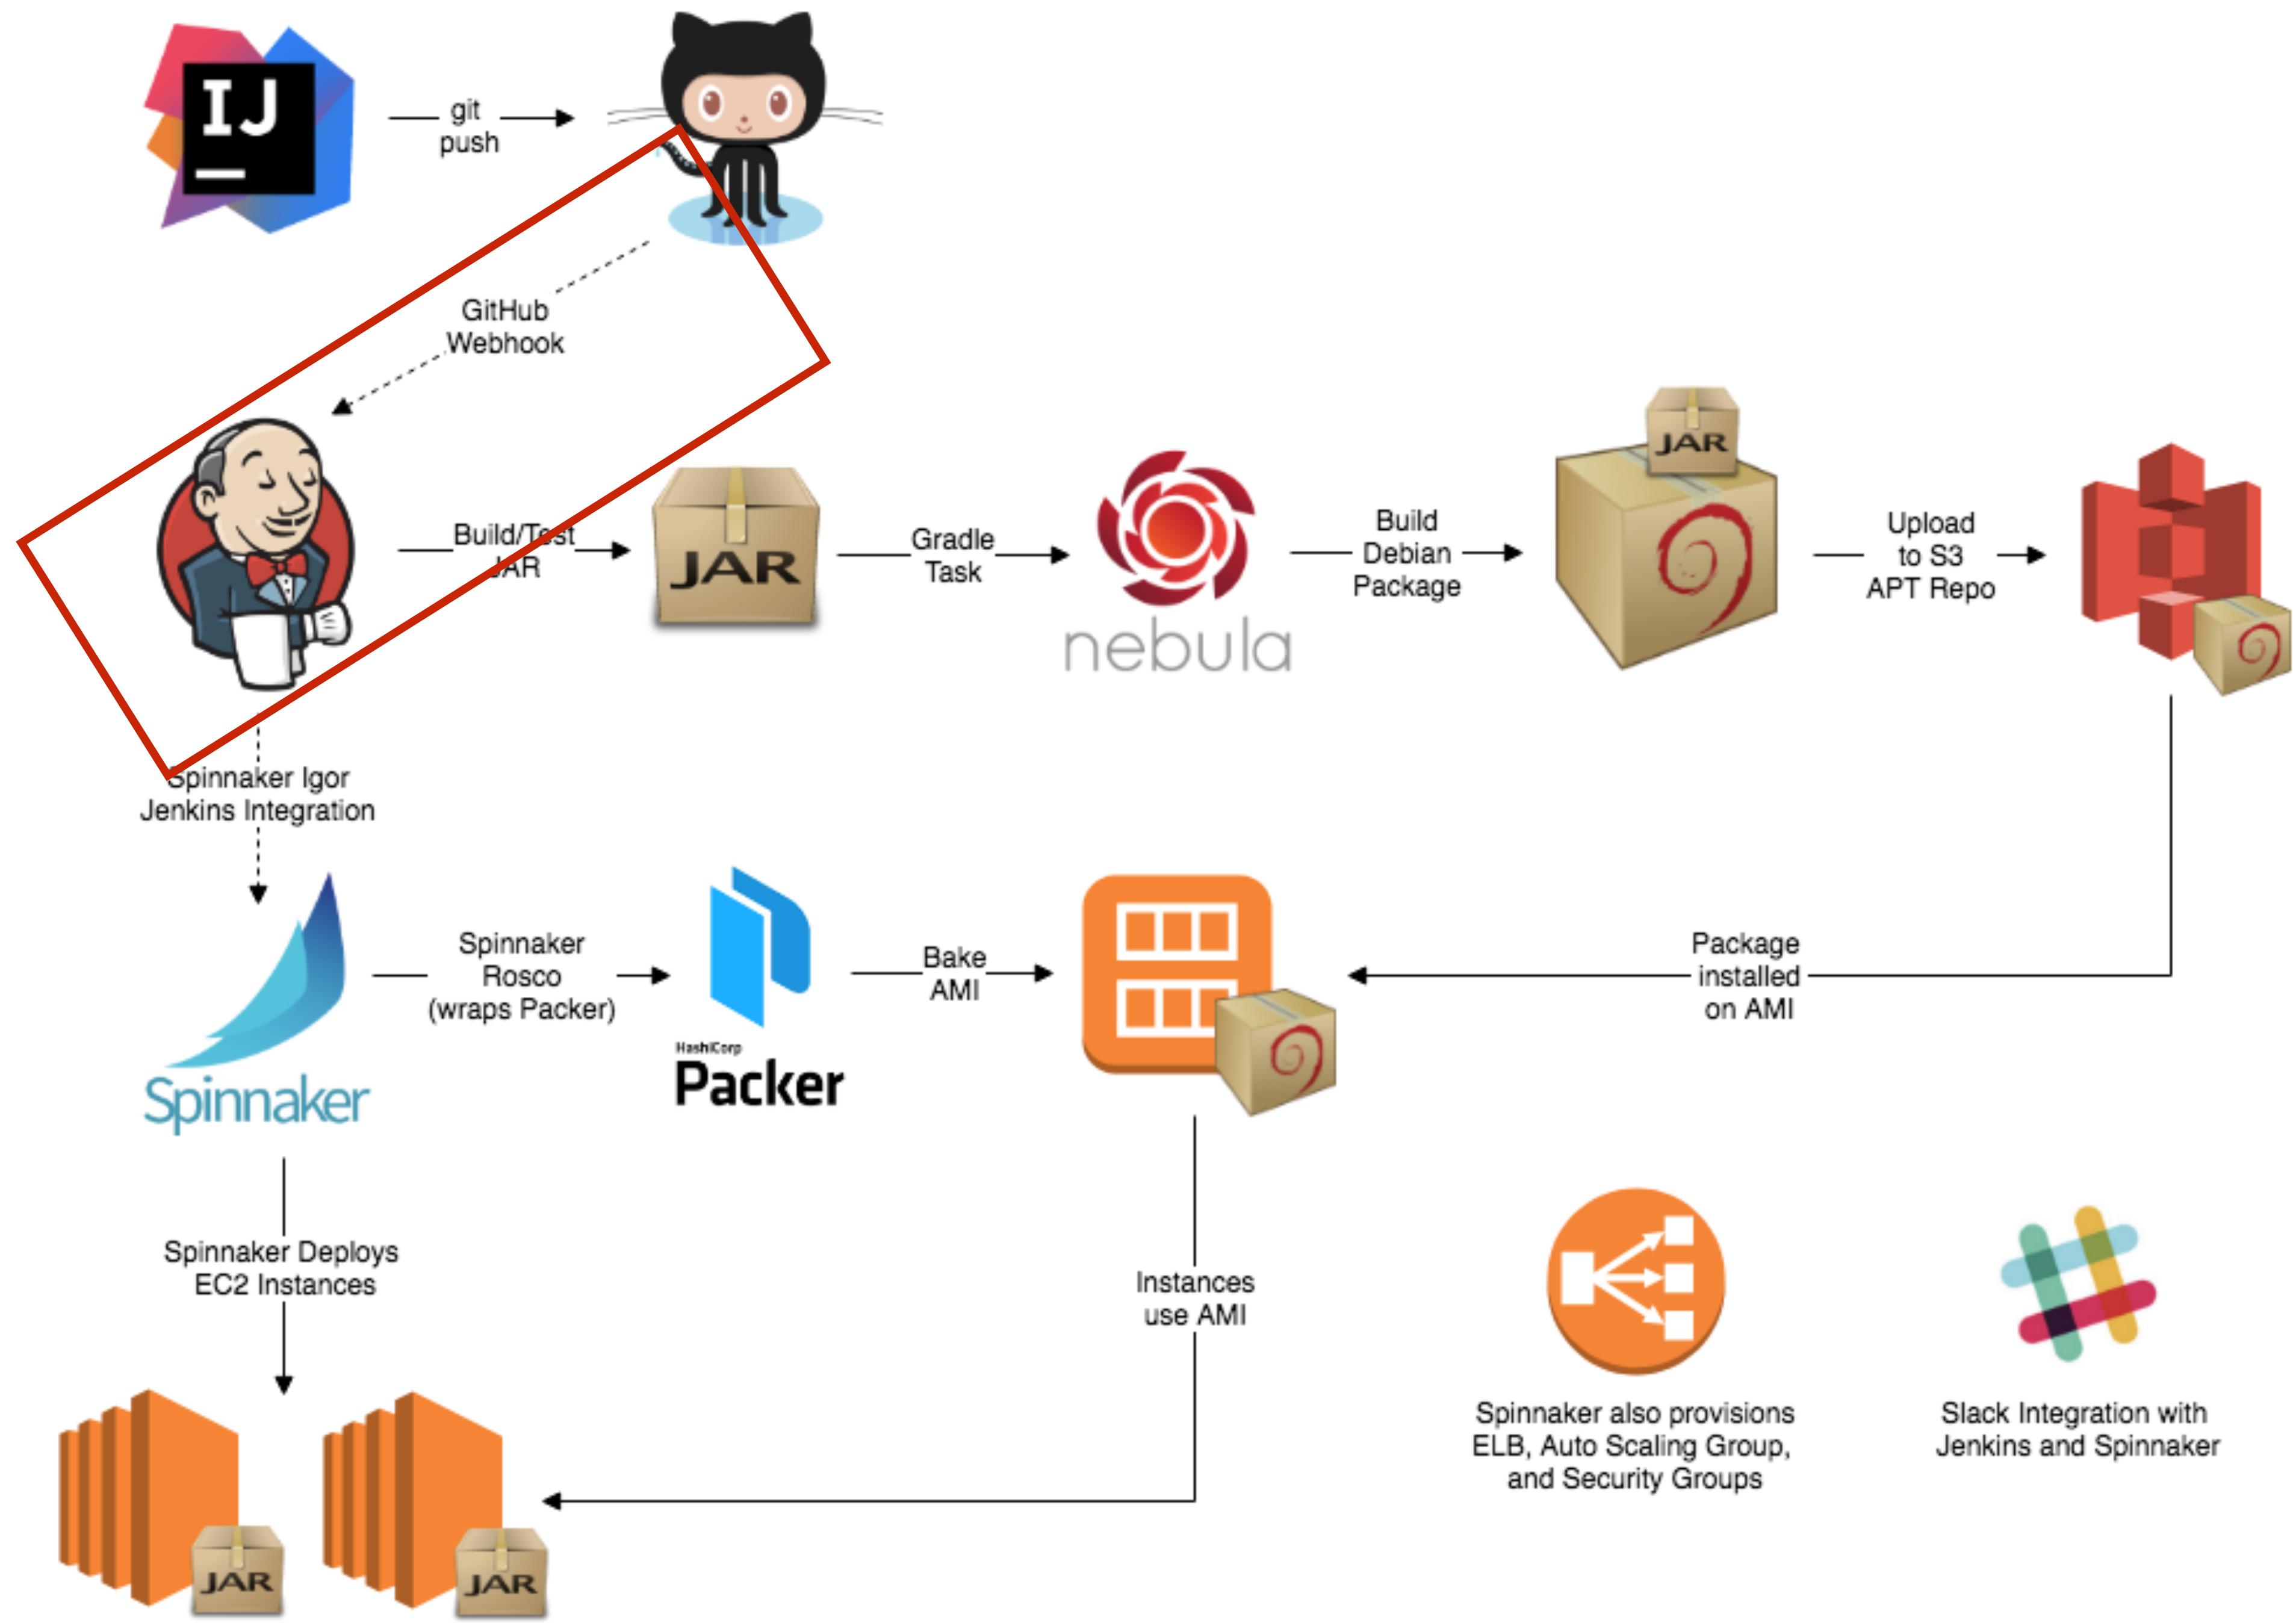


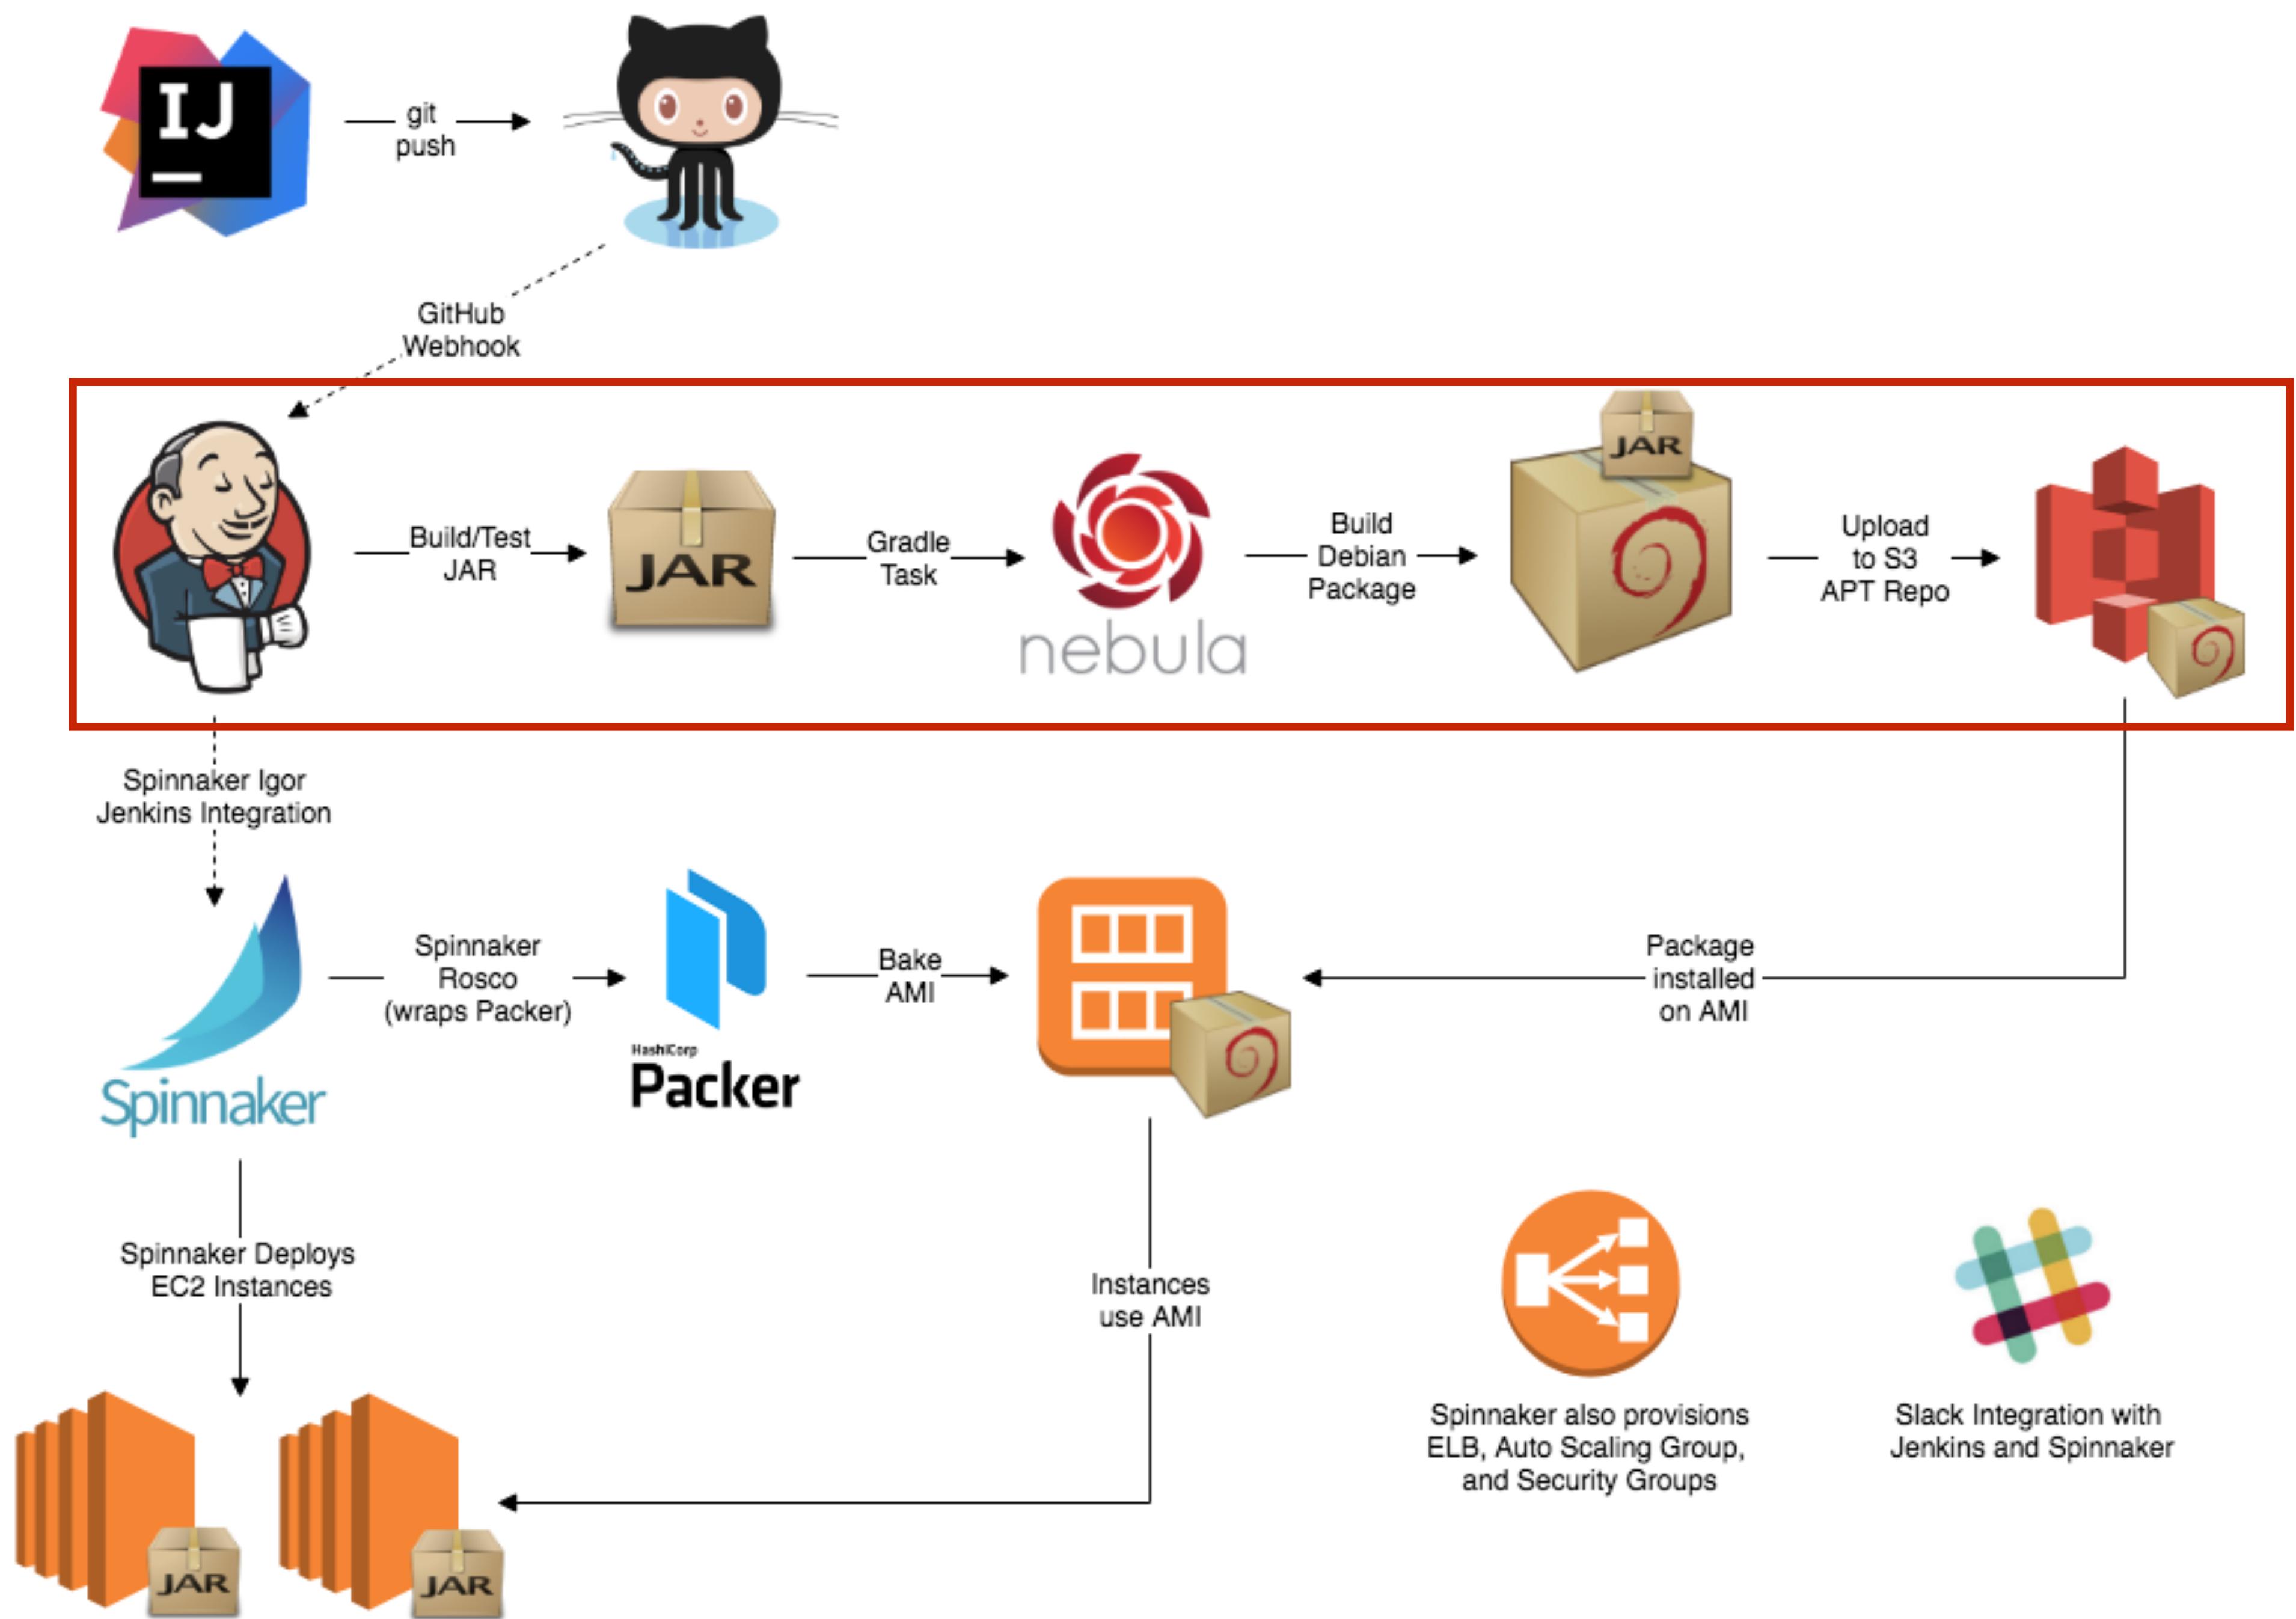
NETFLIX

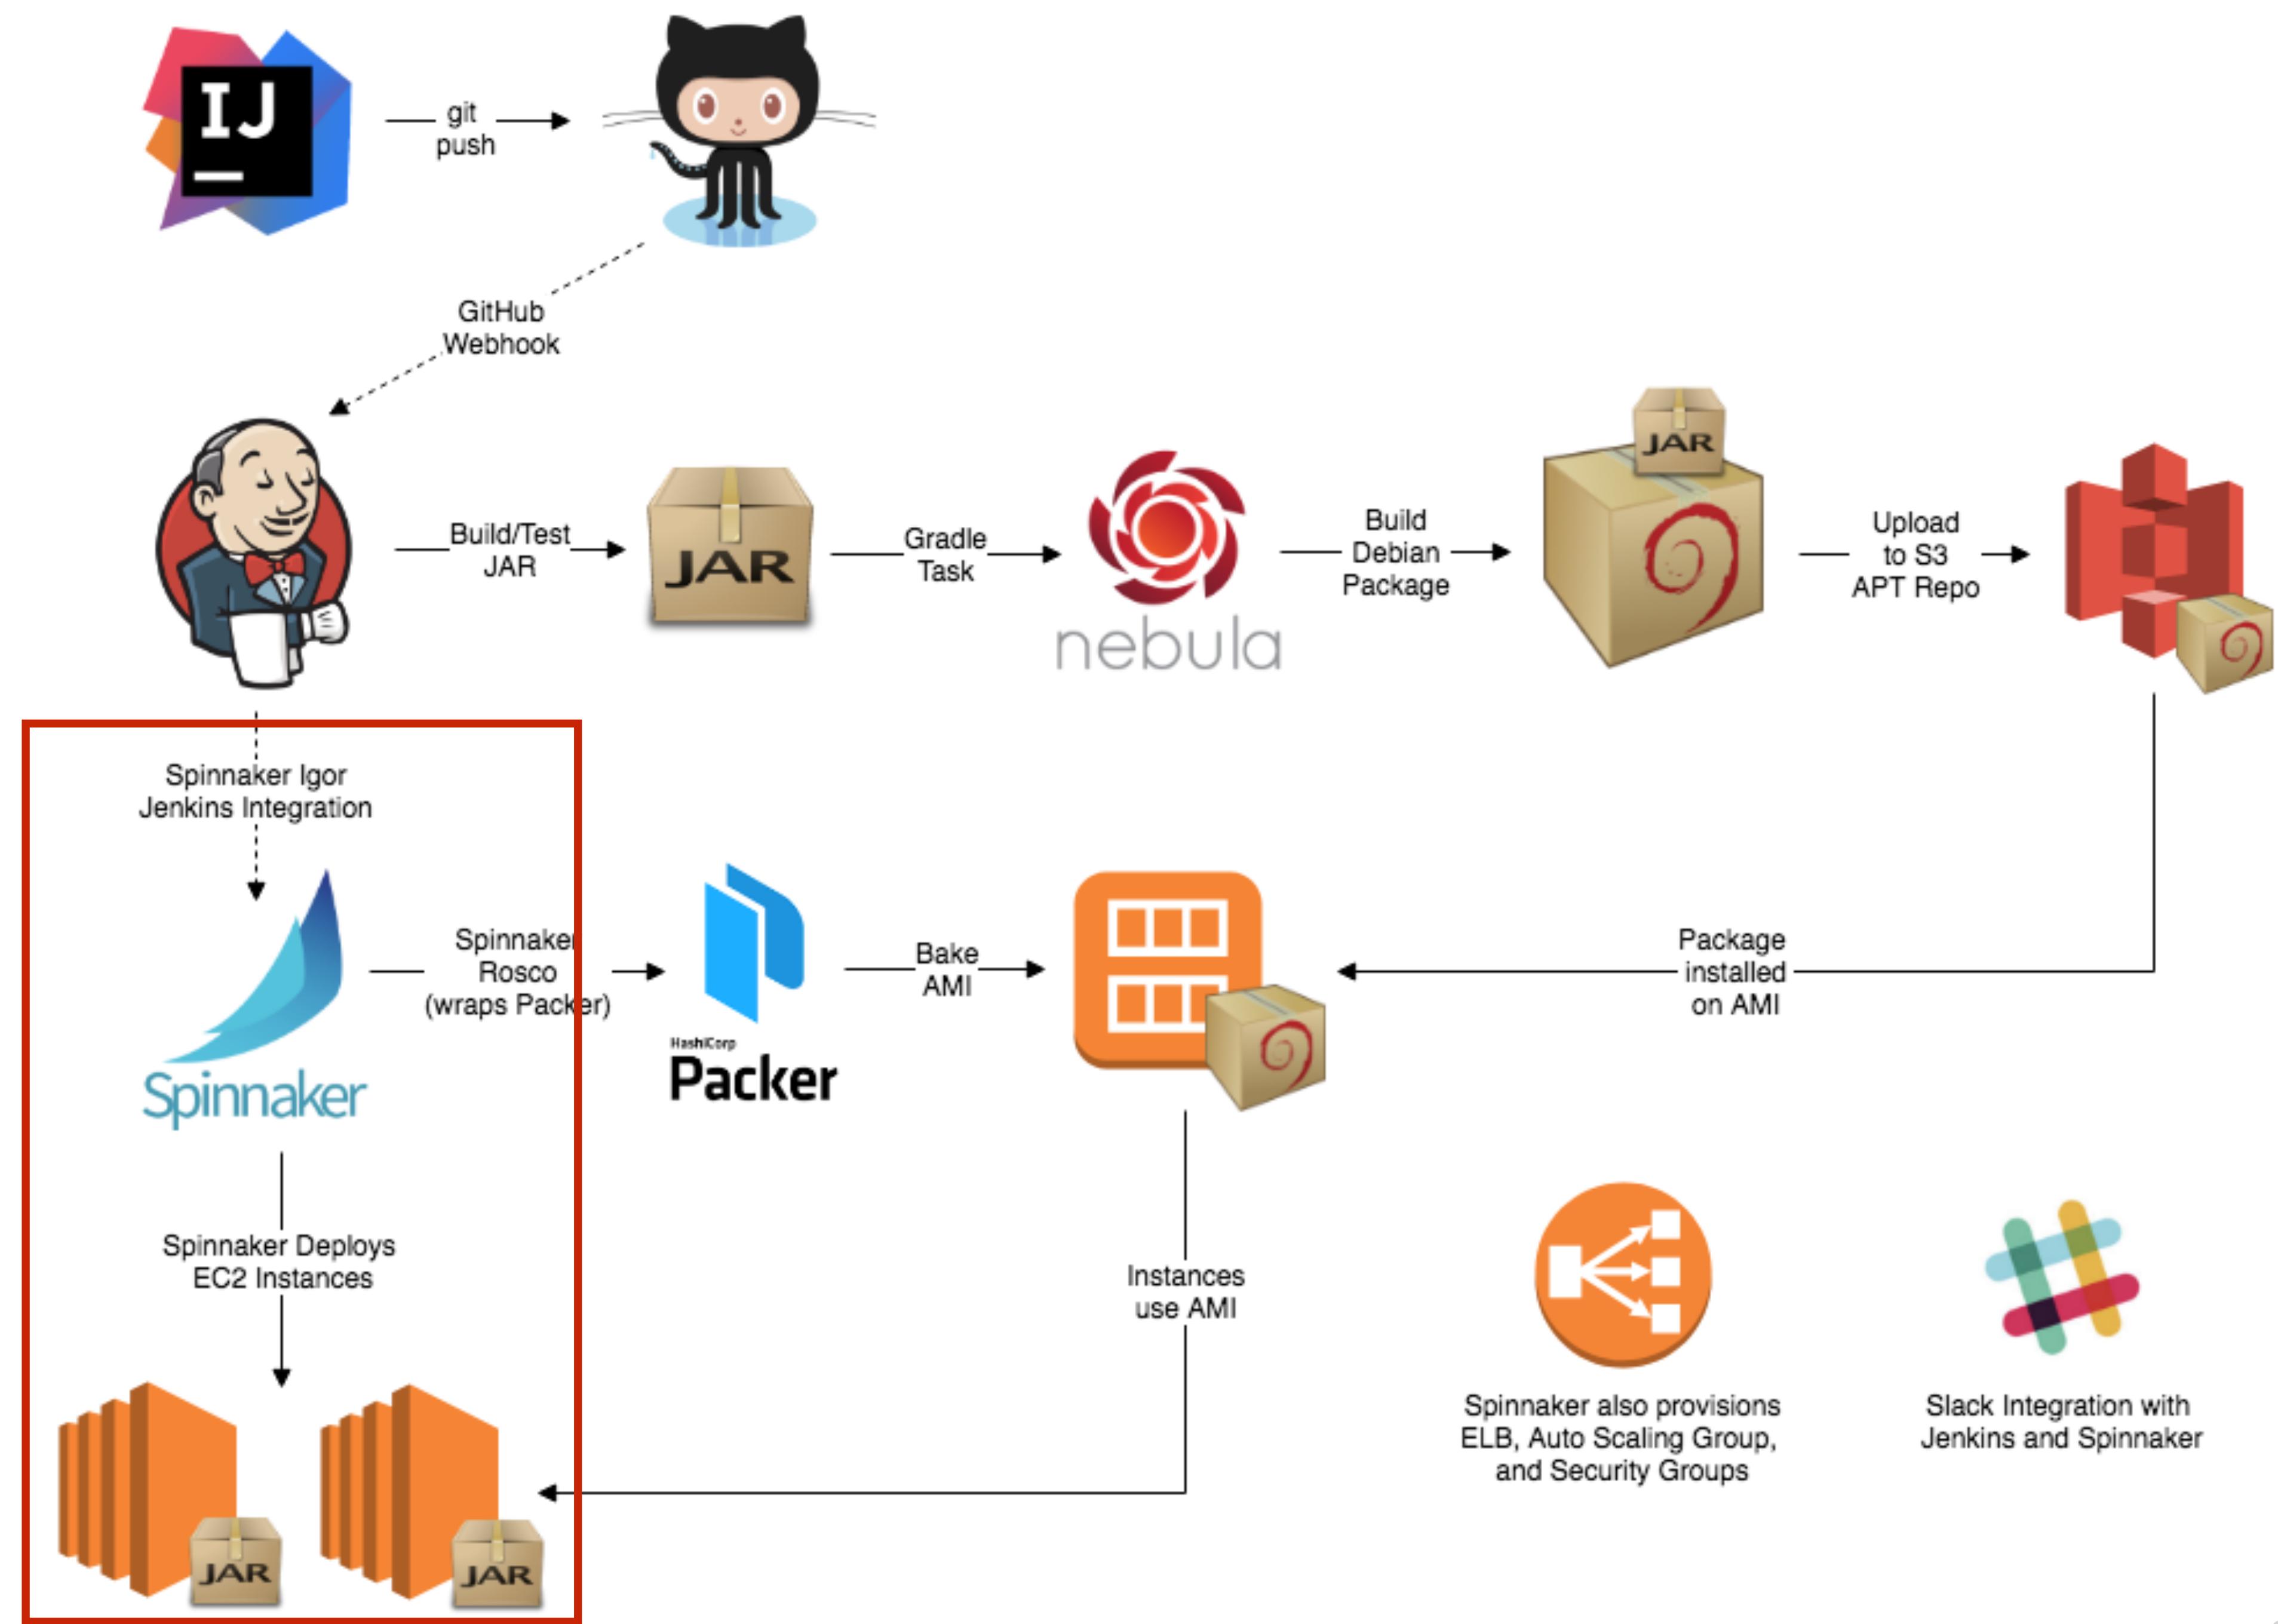


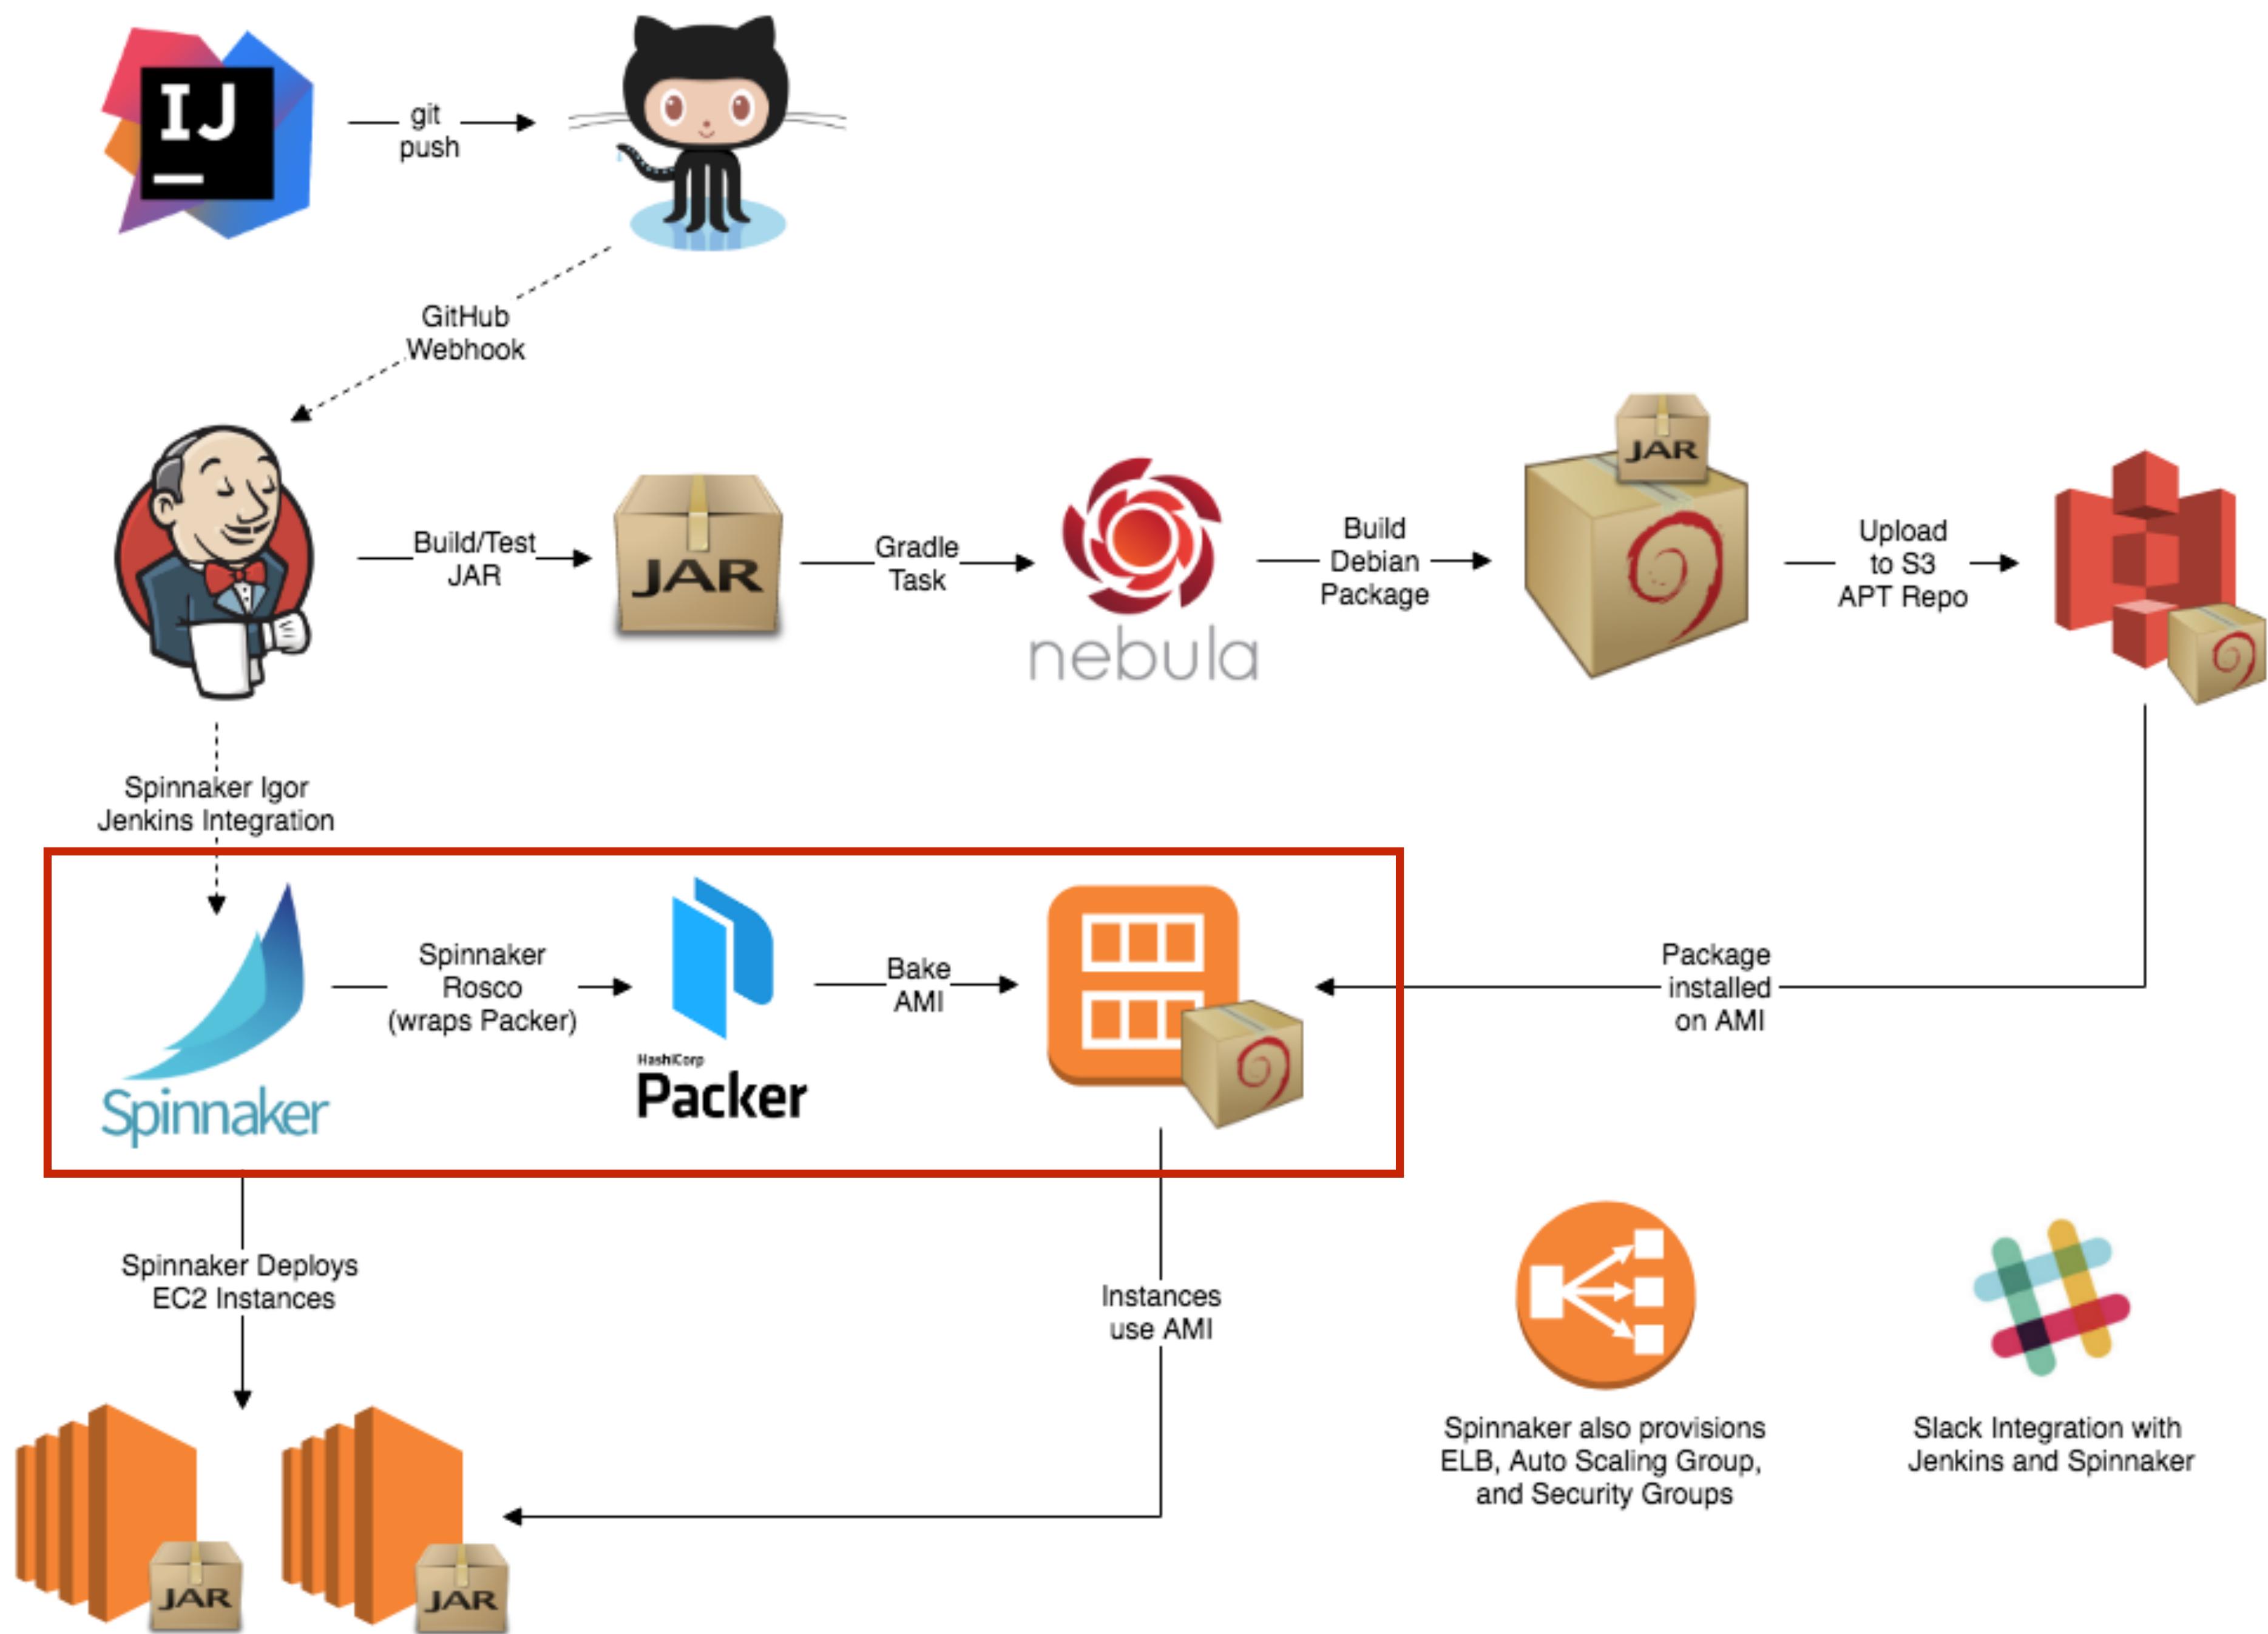


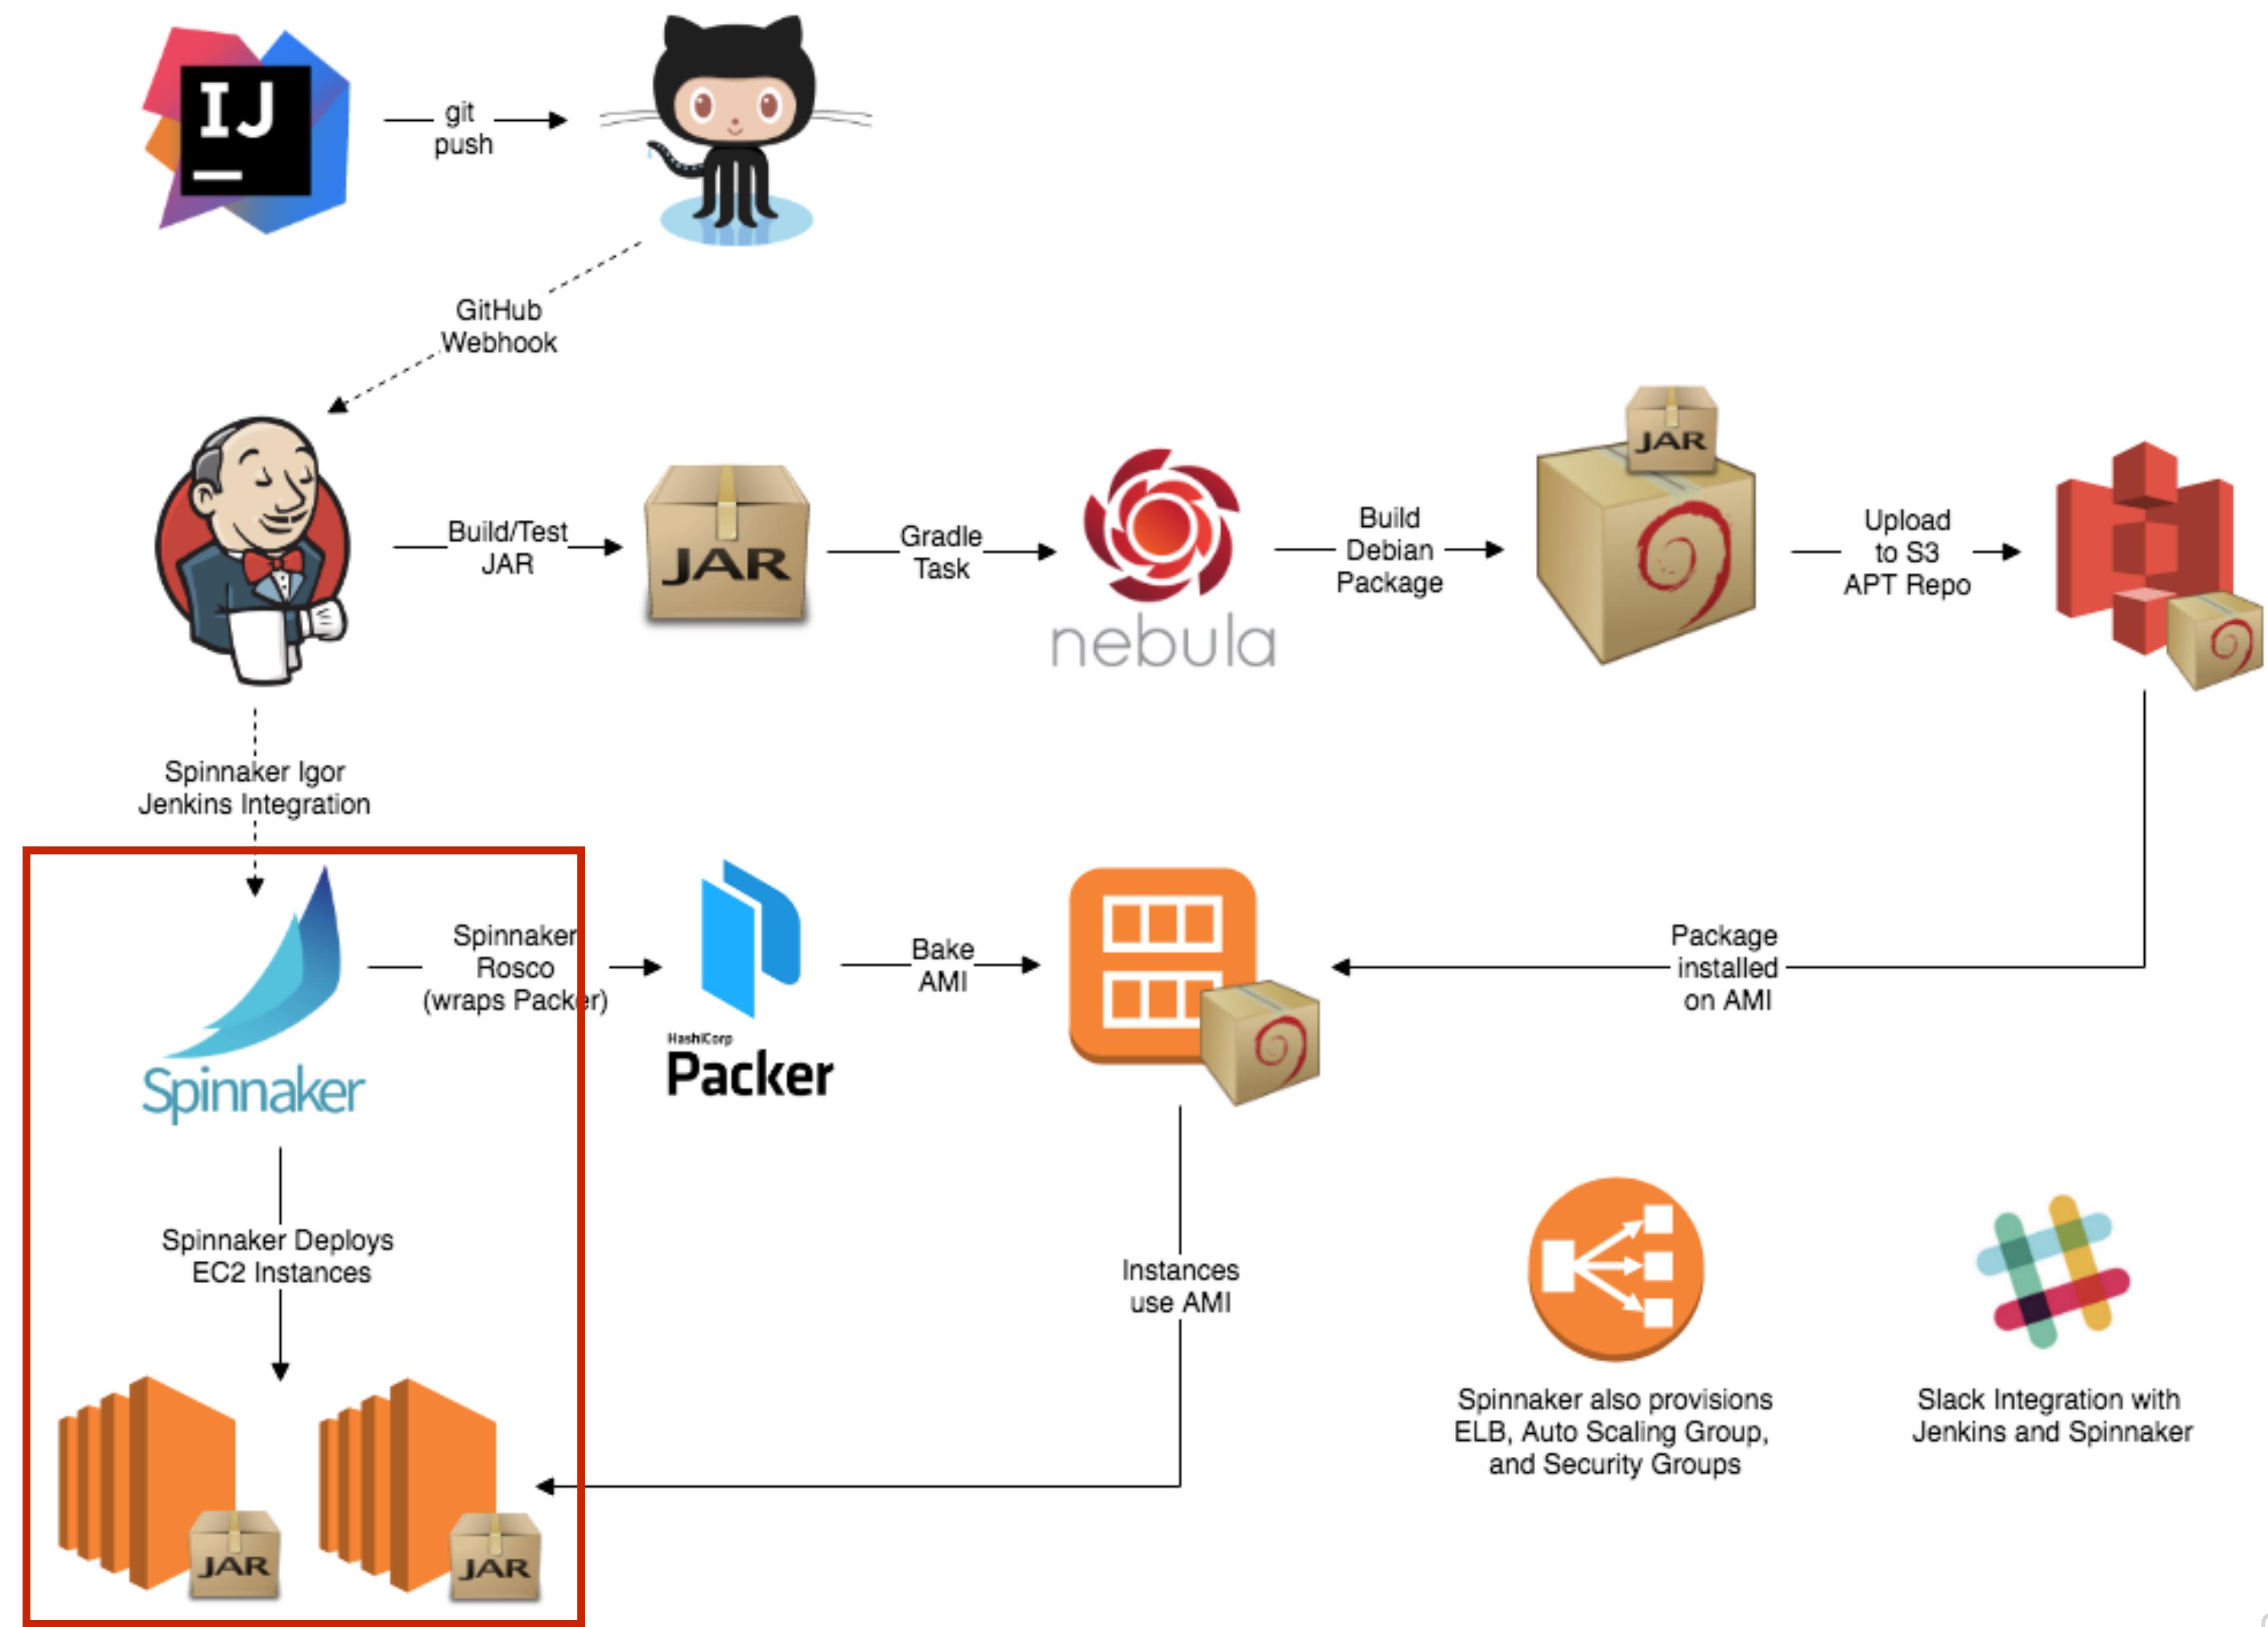


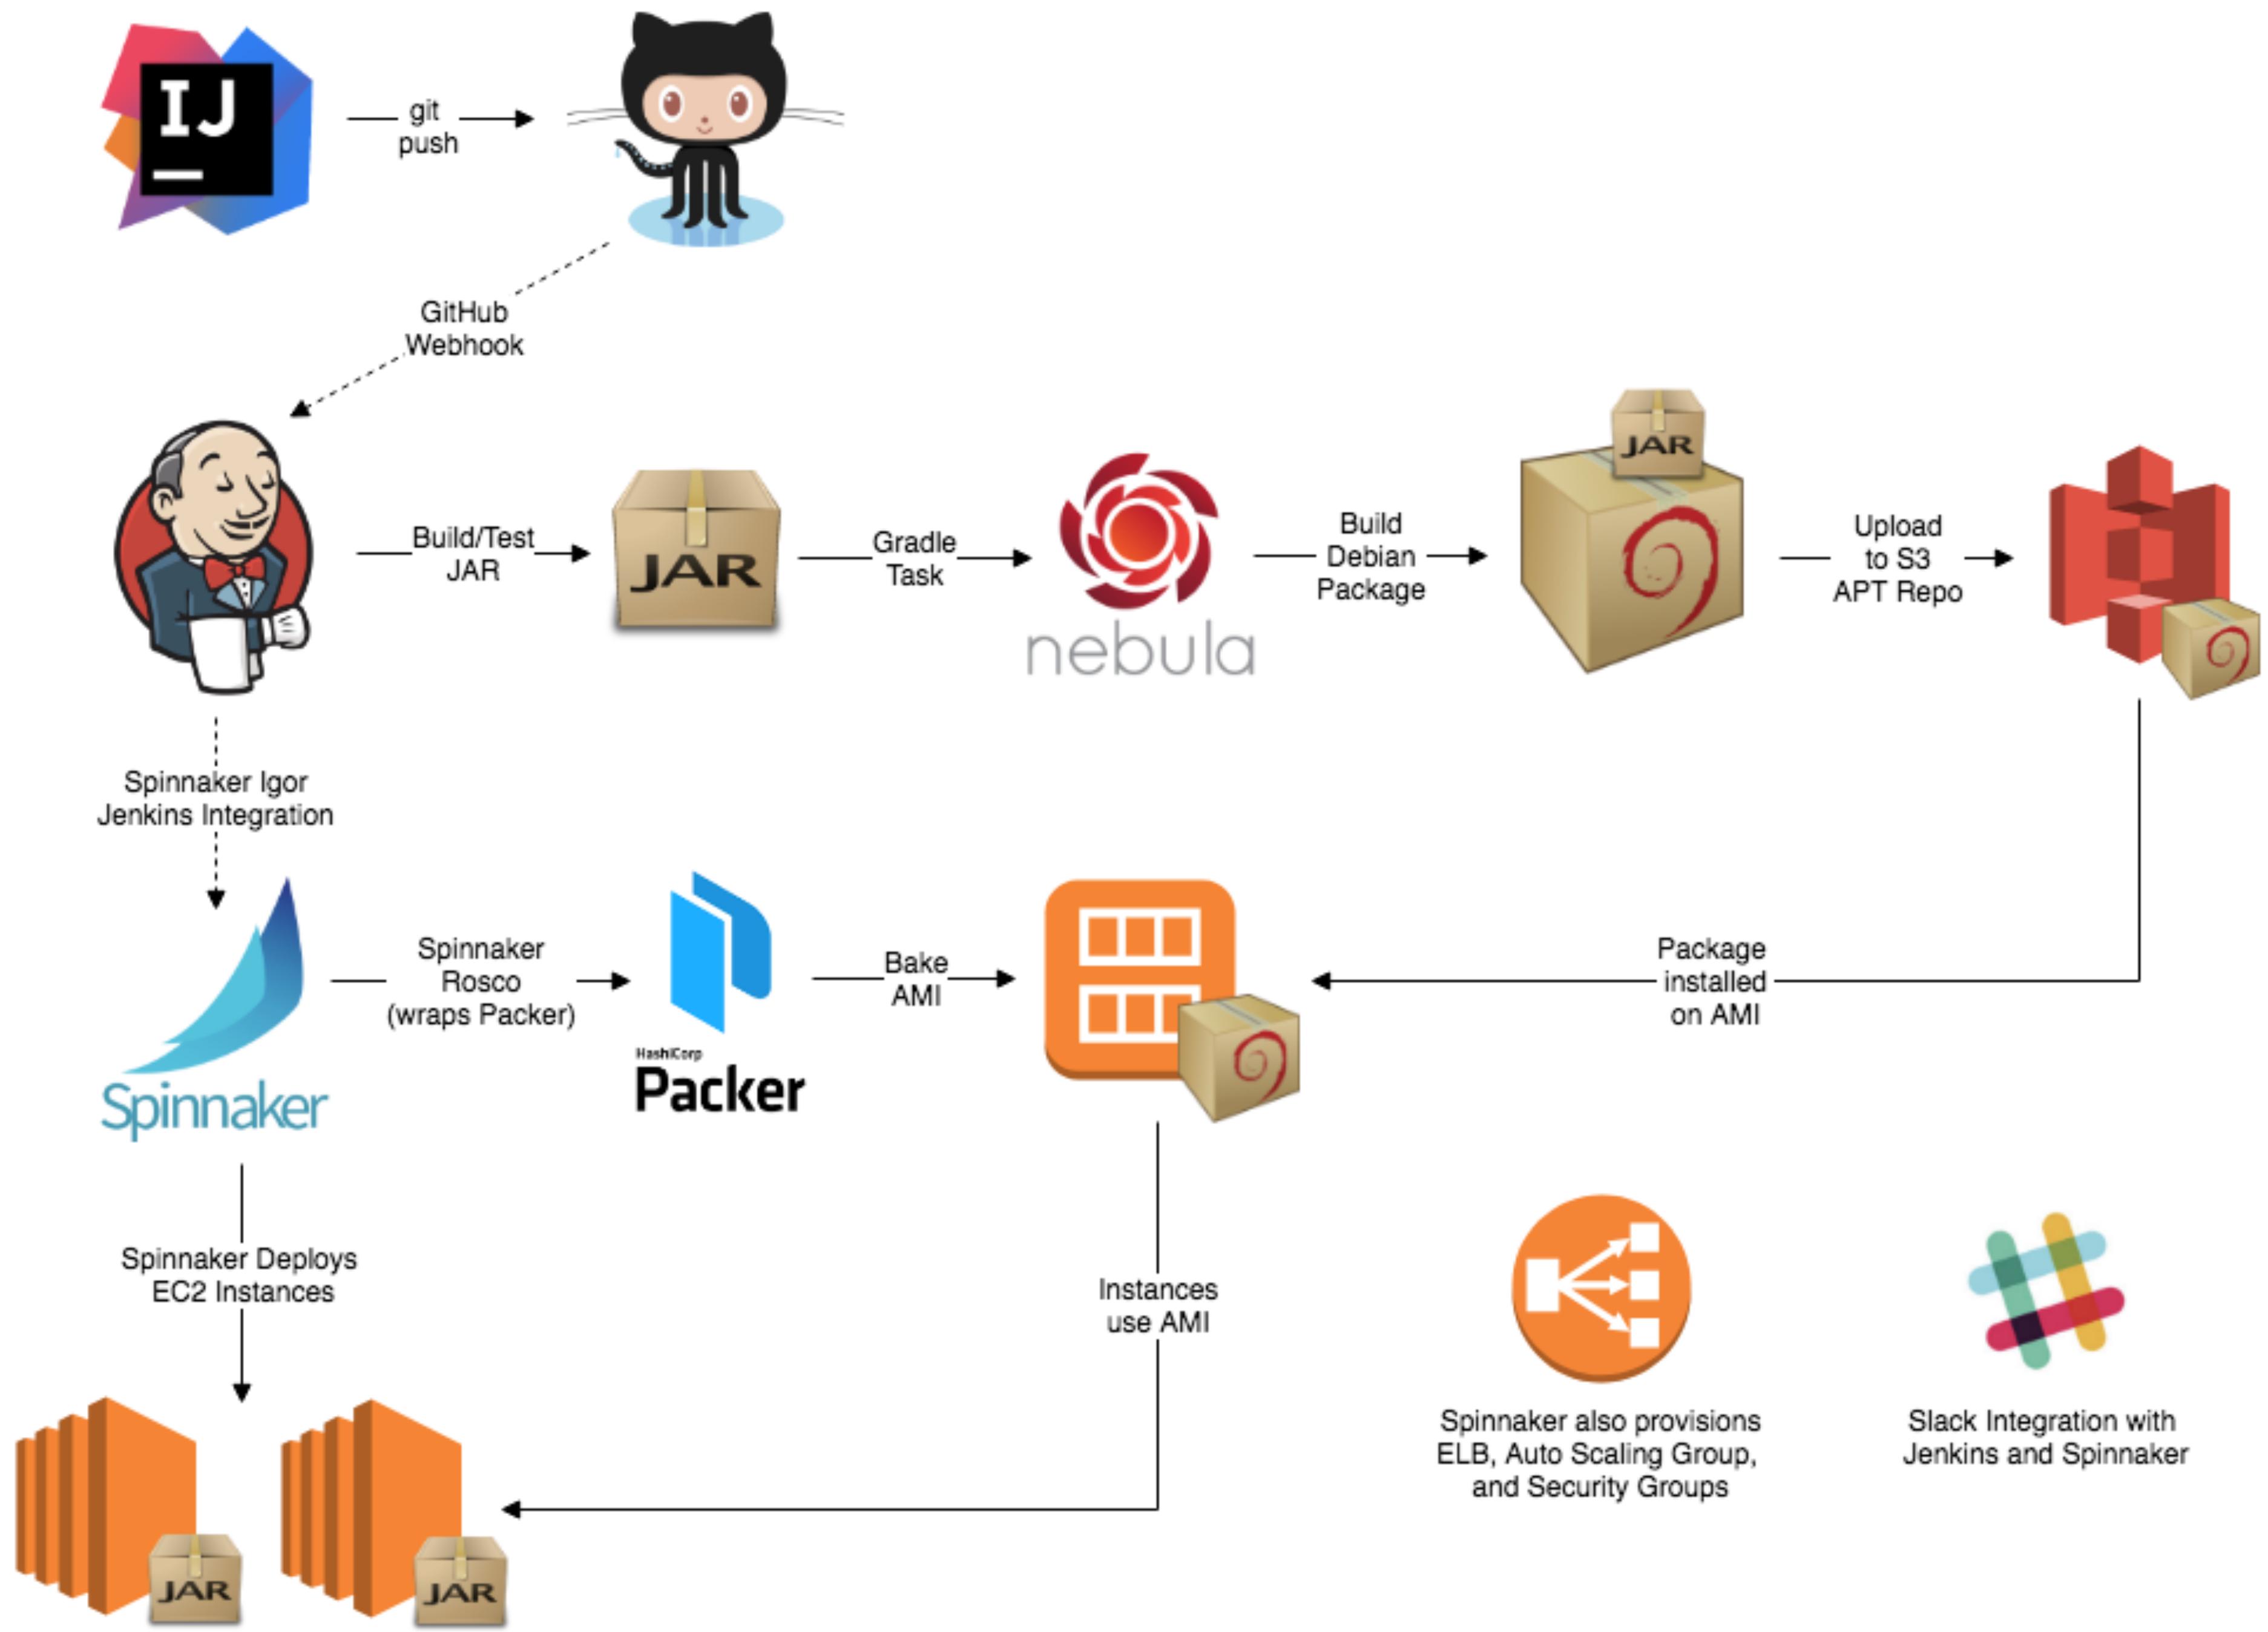














Chaos Monkey

Randomly disables production instances



Chaos Gorilla

Outage of entire Amazon Availability Zone



Janitor Monkey

Identifies and disposes unused resources



Chaos Kong

Drops a full AWS Region

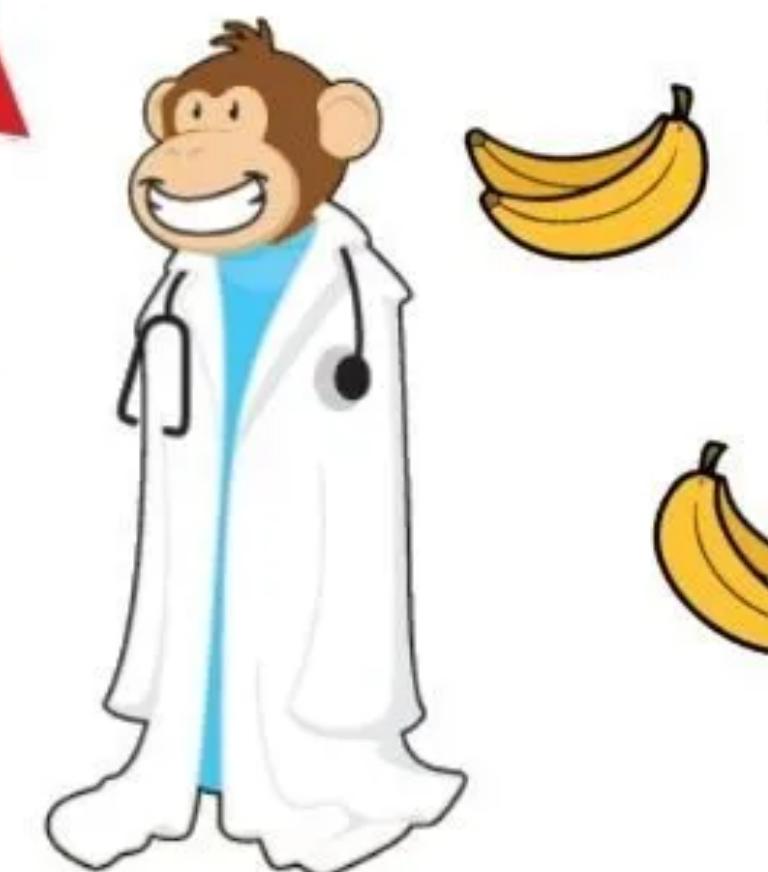


Conformity Monkey

Shuts down instances not adhering to best-practices

NETFLIX

SIMIAN ARMY



Doctor Monkey

Taps into health checks and fixes unhealthy resources



Latency Monkey

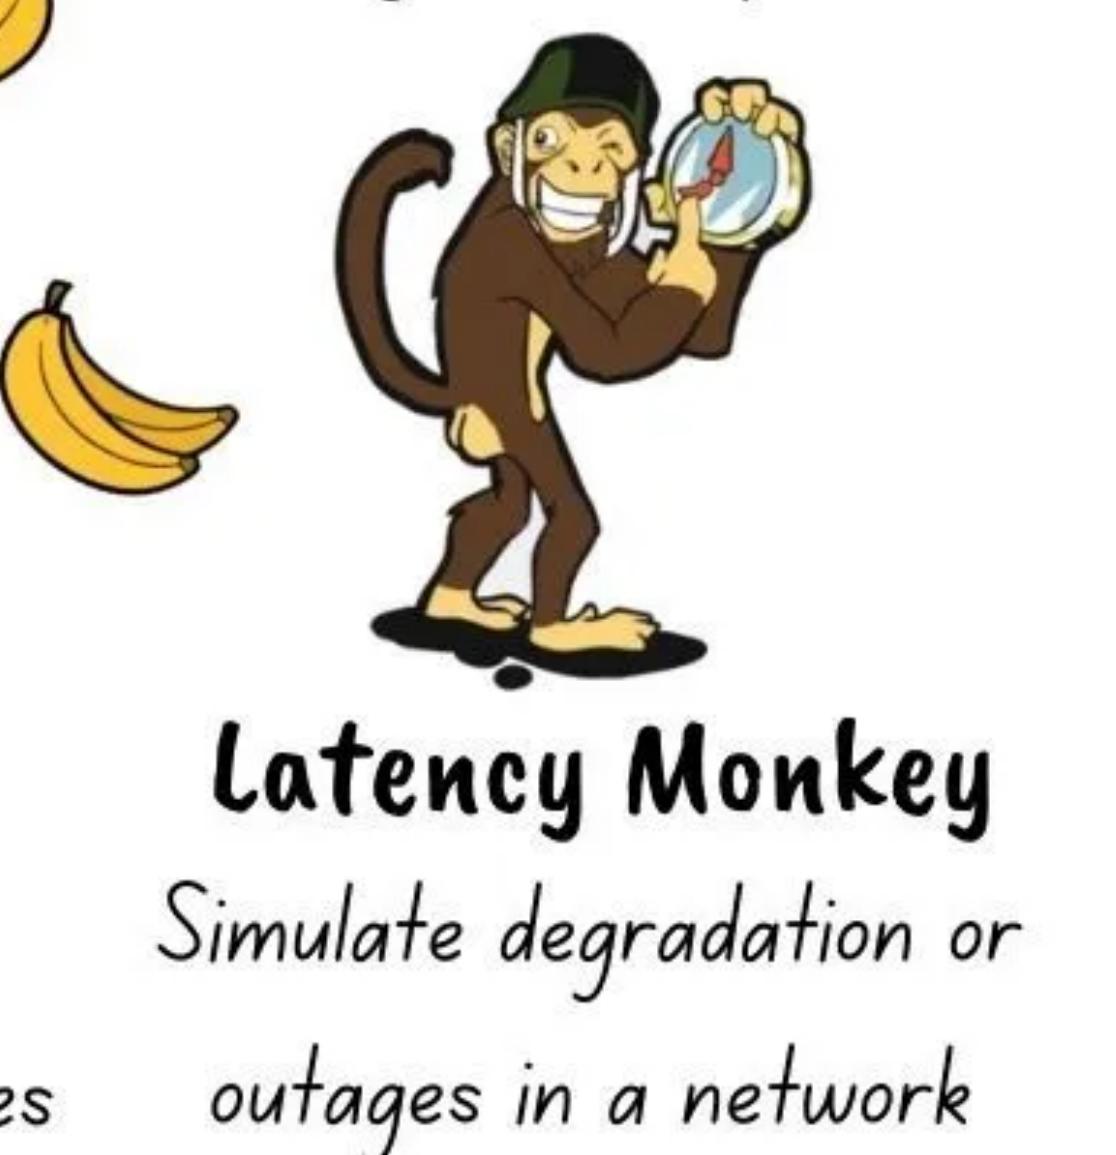
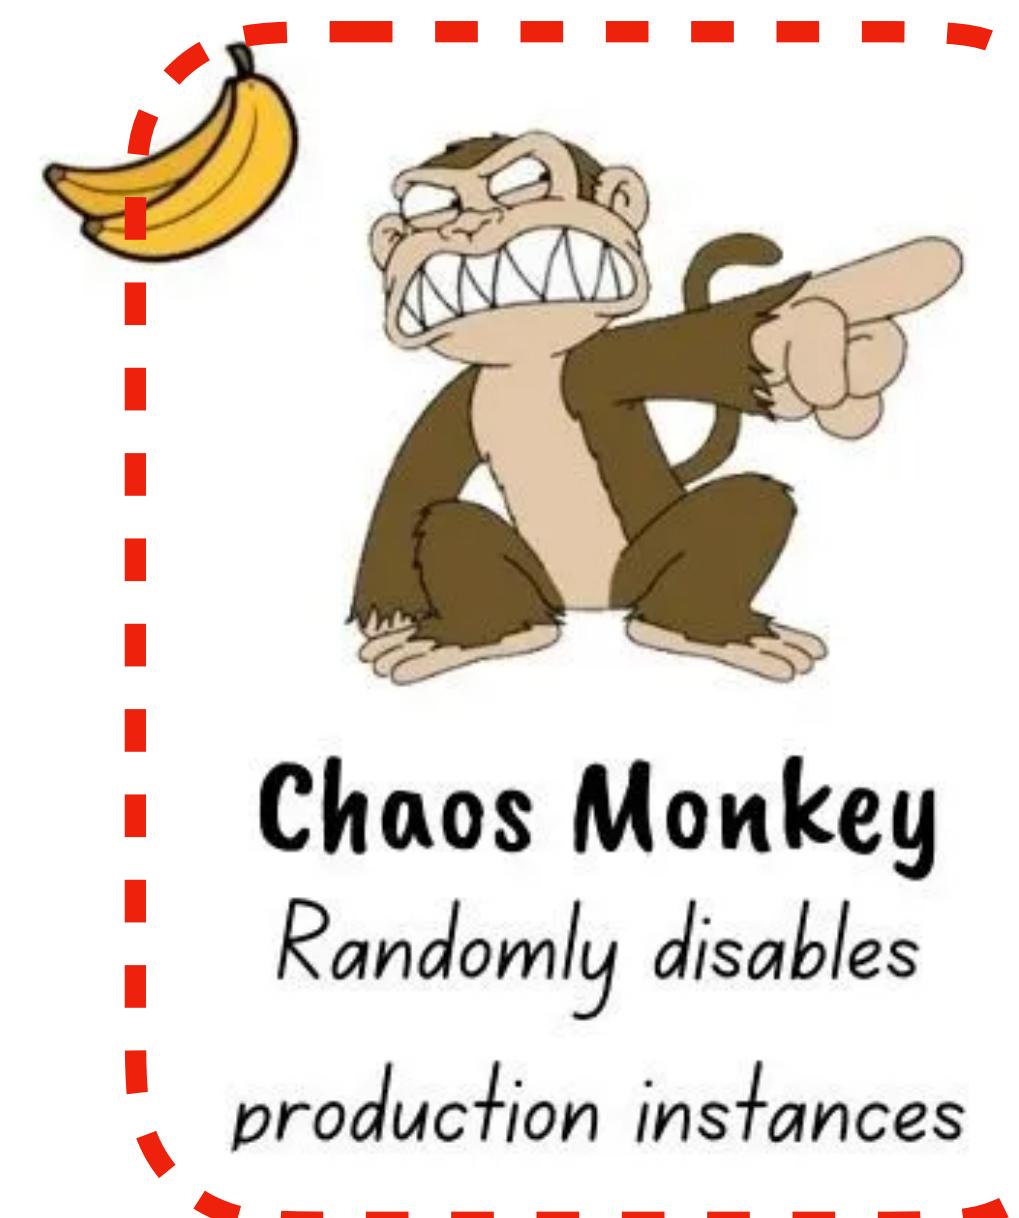
Simulate degradation or outages in a network



Security Monkey

Finds security violations and vulnerability







Chaos Monkey

Randomly disables production instances



Chaos Gorilla

Outage of entire Amazon Availability Zone



Janitor Monkey

Identifies and disposes unused resources



Chaos Kong

Drops a full AWS Region



Conformity Monkey

Shuts down instances not adhering to best-practices



Security Monkey

Finds security violations and vulnerability



Doctor Monkey

Taps into health checks and fixes unhealthy resources



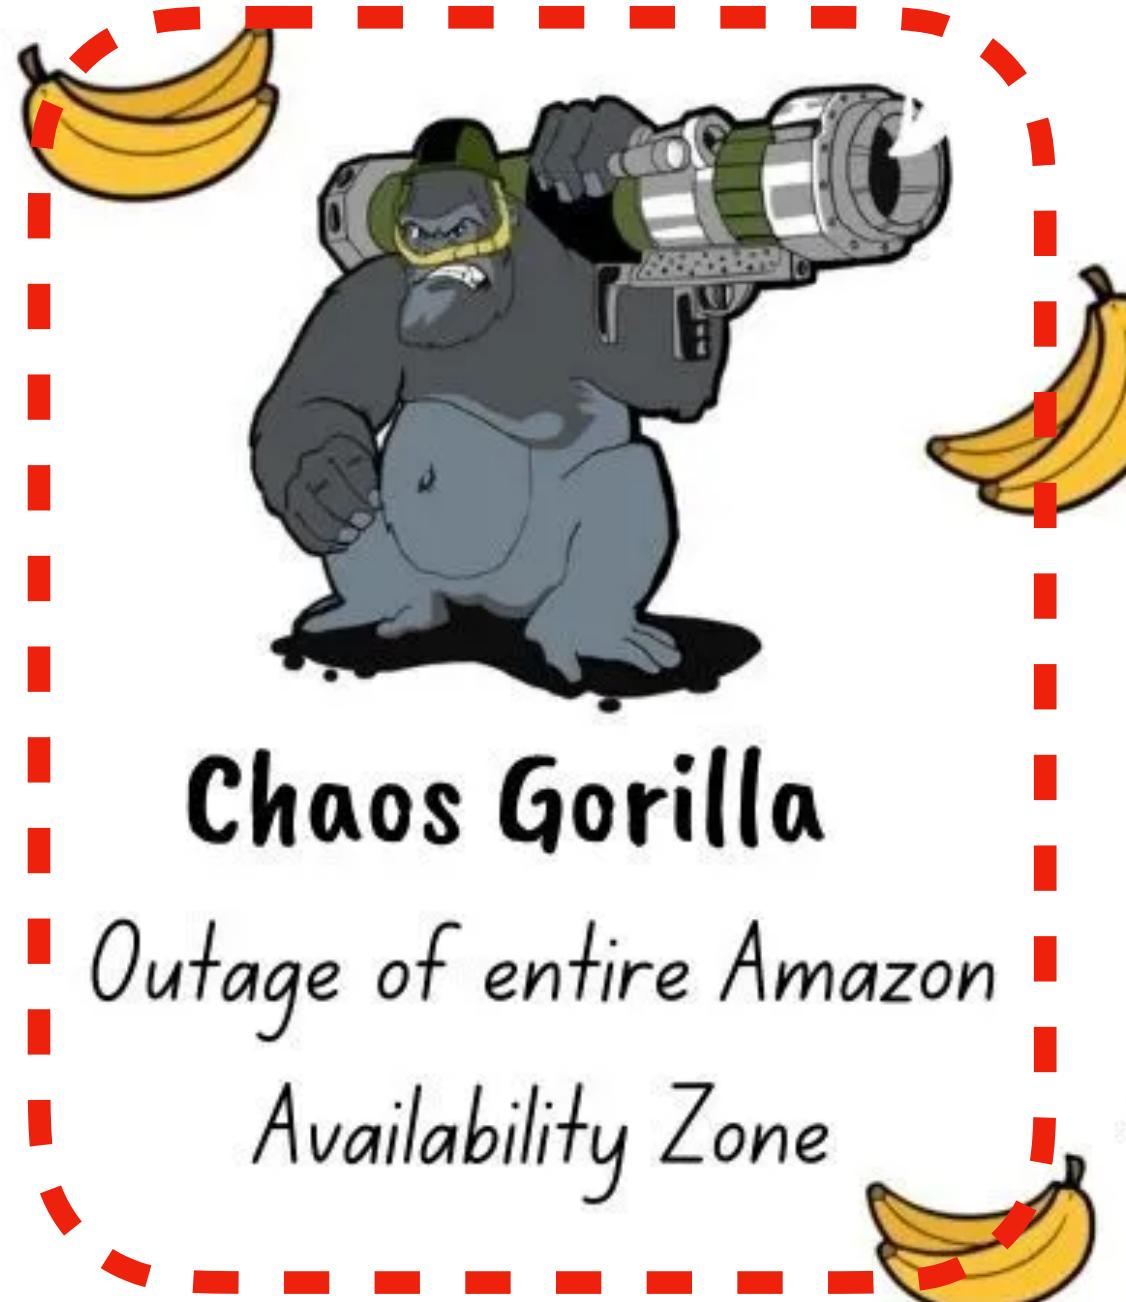
Latency Monkey

Simulate degradation or outages in a network



Chaos Monkey

Randomly disables production instances



Chaos Gorilla

Outage of entire Amazon Availability Zone



Janitor Monkey

Identifies and disposes unused resources



Chaos Kong

Drops a full AWS Region



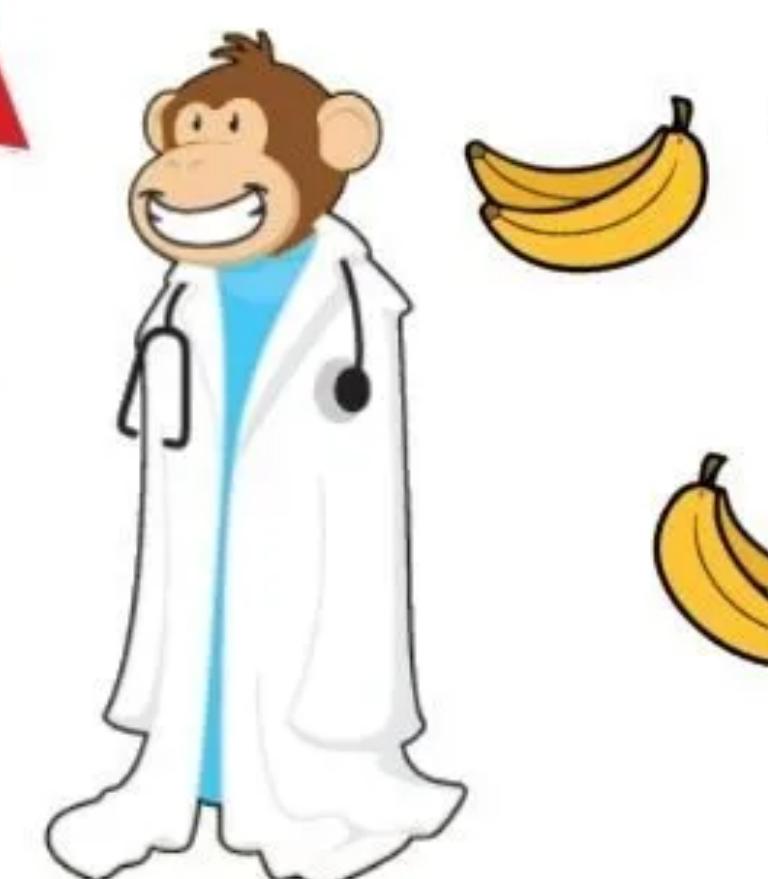
Conformity Monkey

Shuts down instances not adhering to best-practices



Security Monkey

Finds security violations and vulnerability



Doctor Monkey

Taps into health checks and fixes unhealthy resources



Latency Monkey

Simulate degradation or outages in a network



@geosley



Chaos Monkey
Randomly disables production instances



Chaos Gorilla
Outage of entire Amazon Availability Zone



Janitor Monkey
Identifies and disposes unused resources



Chaos Kong
Drops a full AWS Region



Conformity Monkey
Shuts down instances not adhering to best-practices



Security Monkey
Finds security violations and vulnerability

Doctor Monkey
Taps into health checks and fixes unhealthy resources



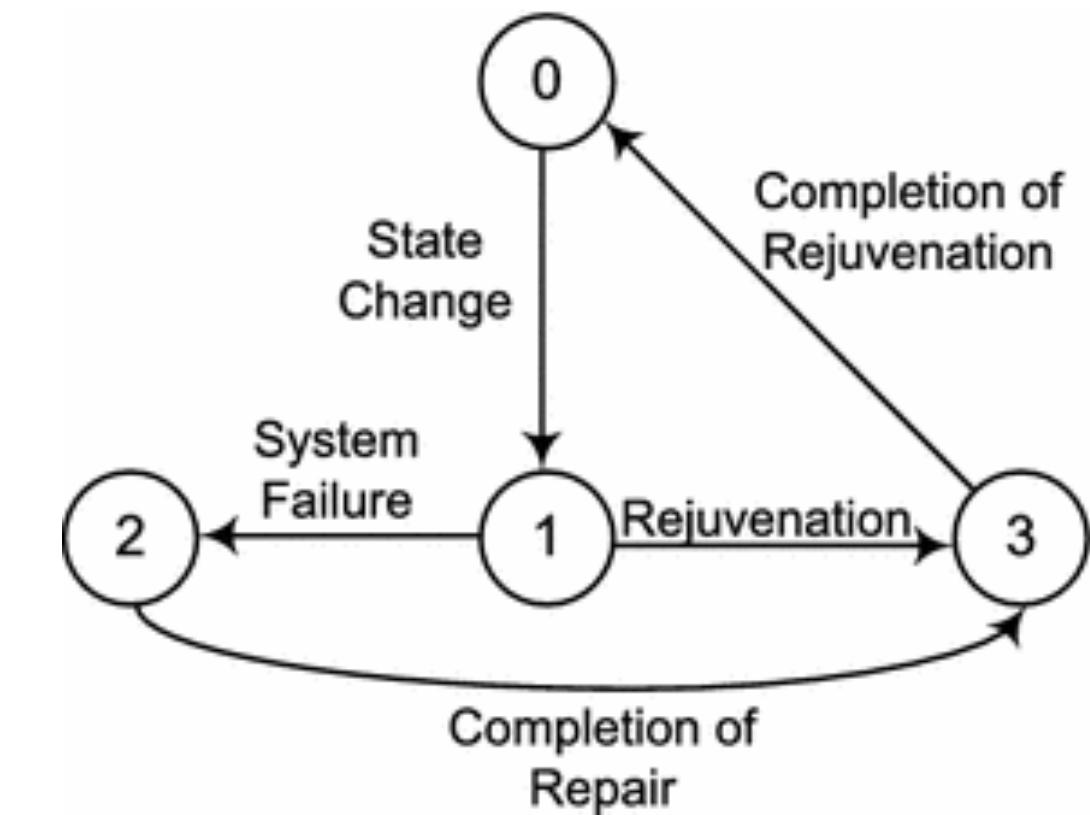
Latency Monkey
Simulate degradation or outages in a network

Software rejuvenation

- Goal: clean up state to prevent accumulation of errors
 - *Insight: Reboot as a prophylactic*
 - *Does nothing about defects, but reduces probability of turning errors into failures*

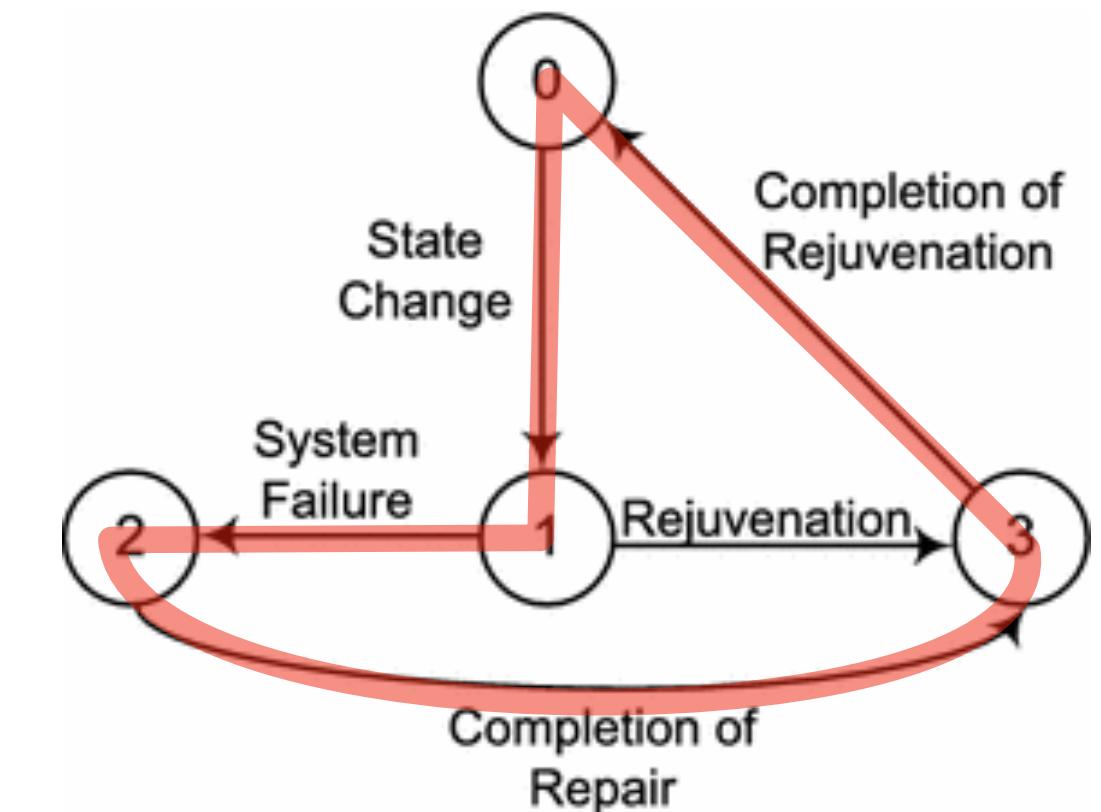
Software rejuvenation

- Goal: clean up state to prevent accumulation of errors
 - *Insight: Reboot as a prophylactic*
 - *Does nothing about defects, but reduces probability of turning errors into failures*
- Turns unplanned downtime into planned downtime
 - *Dynamic version of "preventive maintenance"*
 - *Release leaked resources, wipe out data corruption, ...*



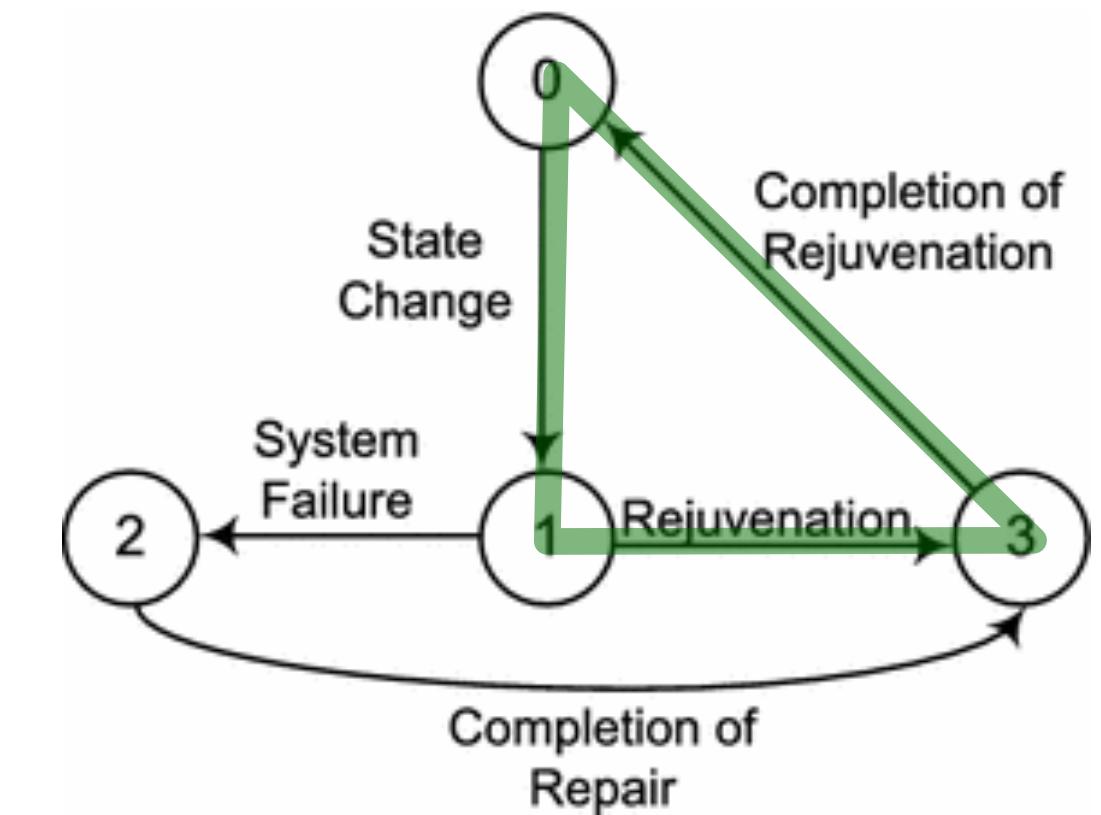
Software rejuvenation

- Goal: clean up state to prevent accumulation of errors
 - *Insight: Reboot as a prophylactic*
 - *Does nothing about defects, but reduces probability of turning errors into failures*
- Turns unplanned downtime into planned downtime
 - *Dynamic version of "preventive maintenance"*
 - *Release leaked resources, wipe out data corruption, ...*



Software rejuvenation

- Goal: clean up state to prevent accumulation of errors
 - *Insight: Reboot as a prophylactic*
 - *Does nothing about defects, but reduces probability of turning errors into failures*
- Turns unplanned downtime into planned downtime
 - *Dynamic version of "preventive maintenance"*
 - *Release leaked resources, wipe out data corruption, ...*



How to reduce unavailability by 10× ?

$$\text{Unavailability} \cong \frac{\text{MTTR}}{\text{MTTF}} \uparrow \times 10$$

How to reduce unavailability by 10× ?

$$\text{Unavailability} \approx \frac{\text{MTTR}}{\text{MTTF}} \downarrow \div 10$$

Components of recovery time

- $T_{recover} = T_{detect} + T_{diagnose} + T_{repair}$

Components of recovery time

- $T_{recover} = T_{detect} + T_{diagnose} + T_{repair}$
- How to reduce T_{detect} ?
 - Automation
 - Prediction/anticipation
 - Trade-offs between FPs and FNs

Components of recovery time

- $T_{recover} = T_{detect} + T_{diagnose} + T_{repair}$
- How to reduce T_{detect} ?
 - Automation
 - Prediction/anticipation
 - Trade-offs between FPs and FNs

Detection/Prediction says...

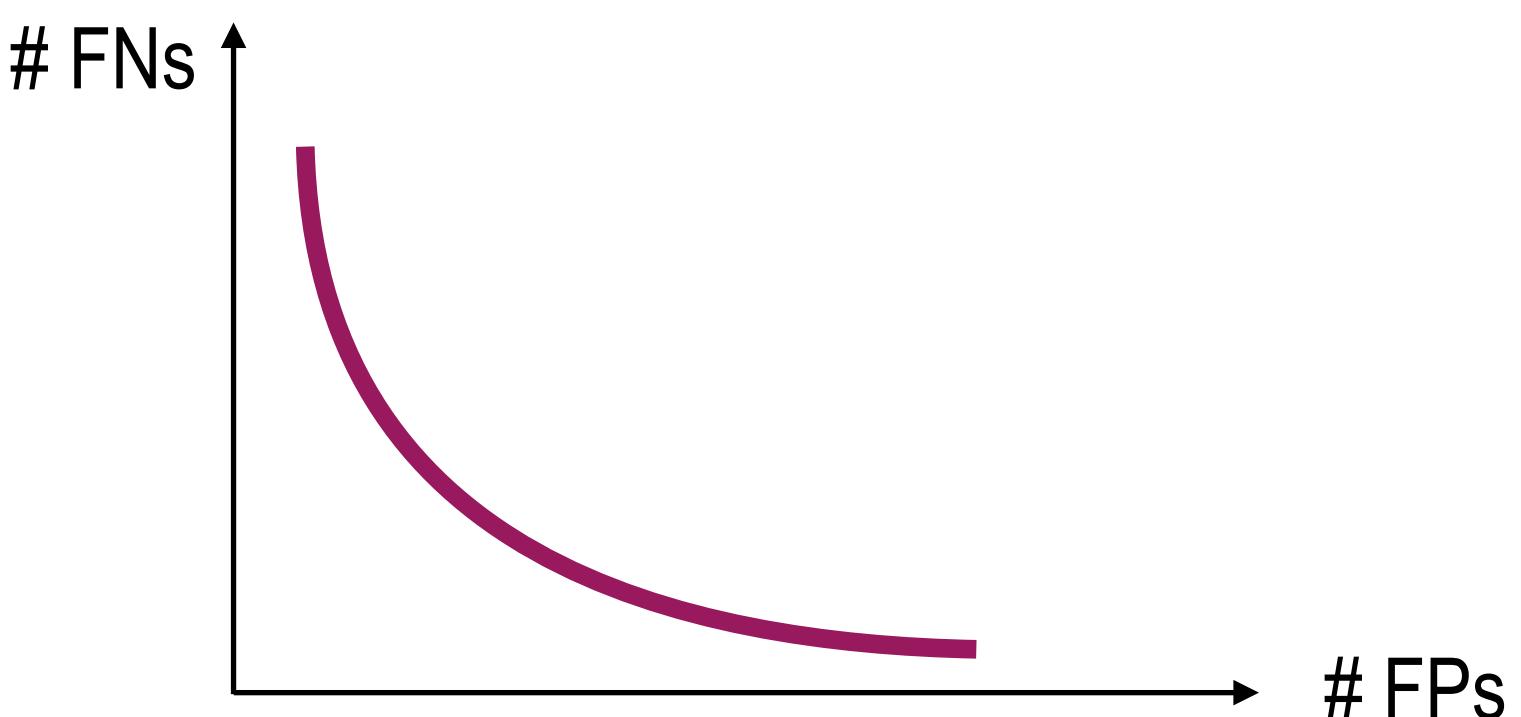
		Failure	No Failure
Truth is...	Failure	FP	TN
	No Failure	TP	FN

Components of recovery time

- $T_{recover} = T_{detect} + T_{diagnose} + T_{repair}$
- How to reduce T_{detect} ?
 - Automation
 - Prediction/anticipation
 - Trade-offs between FPs and FNs

Detection/Prediction says...

		Failure	No Failure
Truth is...	Failure	FP	TN
	No Failure	TP	FN

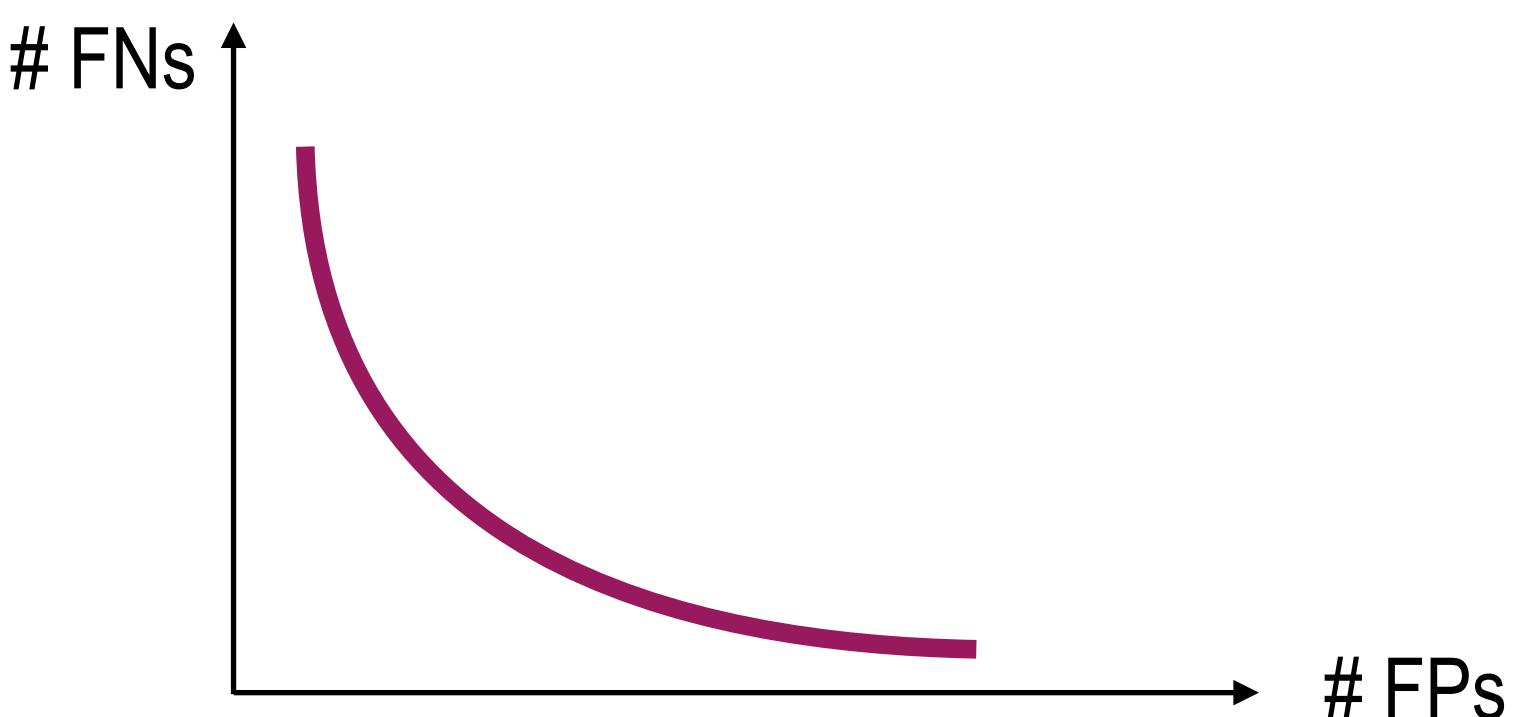


Components of recovery time

- $T_{recover} = T_{detect} + T_{diagnose} + T_{repair}$
- How to reduce T_{detect} ?
 - Automation
 - Prediction/anticipation
 - Trade-offs between FPs and FNs
- How to reduce $T_{diagnose}$?
 - Lots of instrumentation, ML, ...
 - Also a function of what recovery mechanism have available

Detection/Prediction says...

		Failure	No Failure
Truth is...	Failure	FP	TN
	No Failure	TP	FN

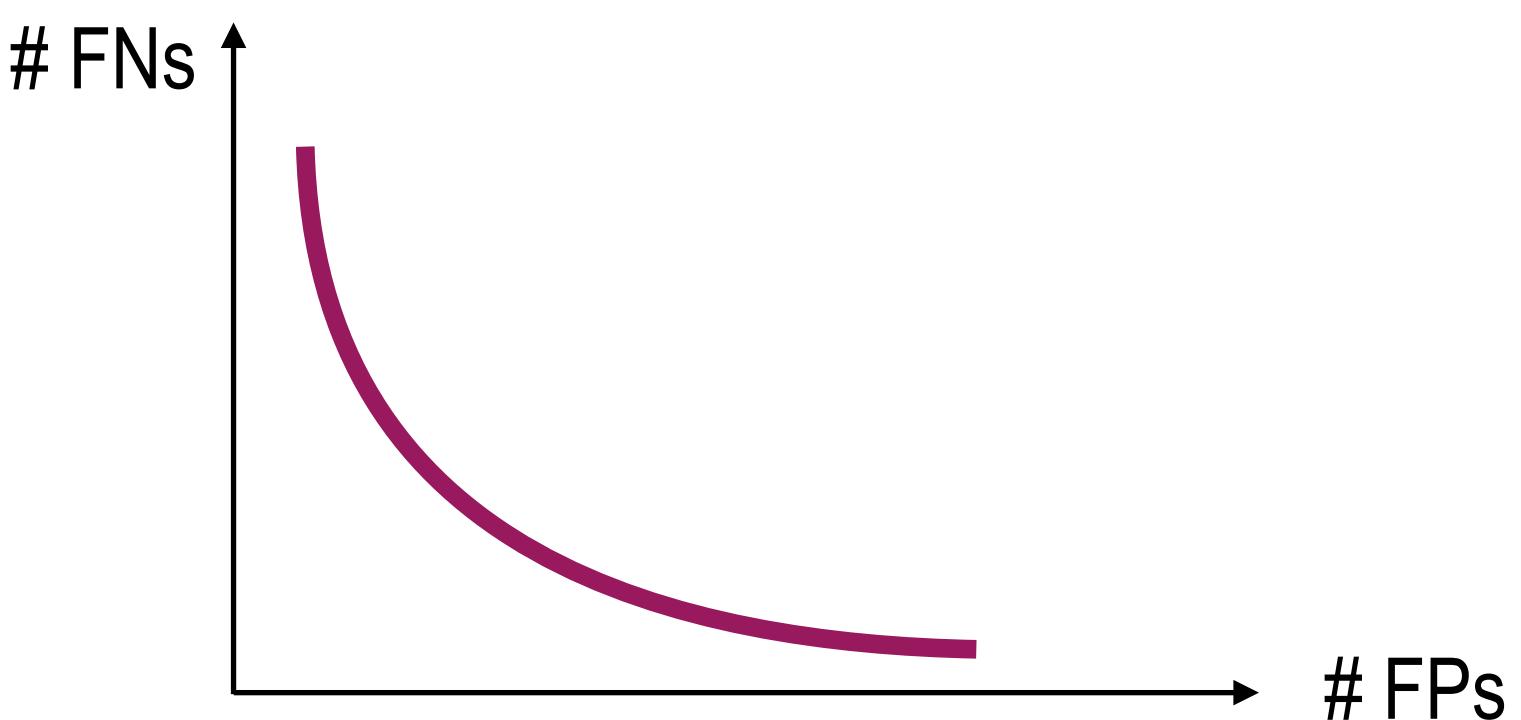


Components of recovery time

- $T_{recover} = T_{detect} + T_{diagnose} + T_{repair}$
- How to reduce T_{detect} ?
 - Automation
 - Prediction/anticipation
 - Trade-offs between FPs and FNs
- How to reduce $T_{diagnose}$?
 - Lots of instrumentation, ML, ...
 - Also a function of what recovery mechanism have available
- How to reduce T_{repair} ?
 - Mostly app-specific
 - Reboot is universal

Detection/Prediction says...

		Failure	No Failure
Truth is...	Failure	FP	TN
	No Failure	TP	FN



Reboot-based Recovery



Reboot-based Recovery





Reboot-based Recovery

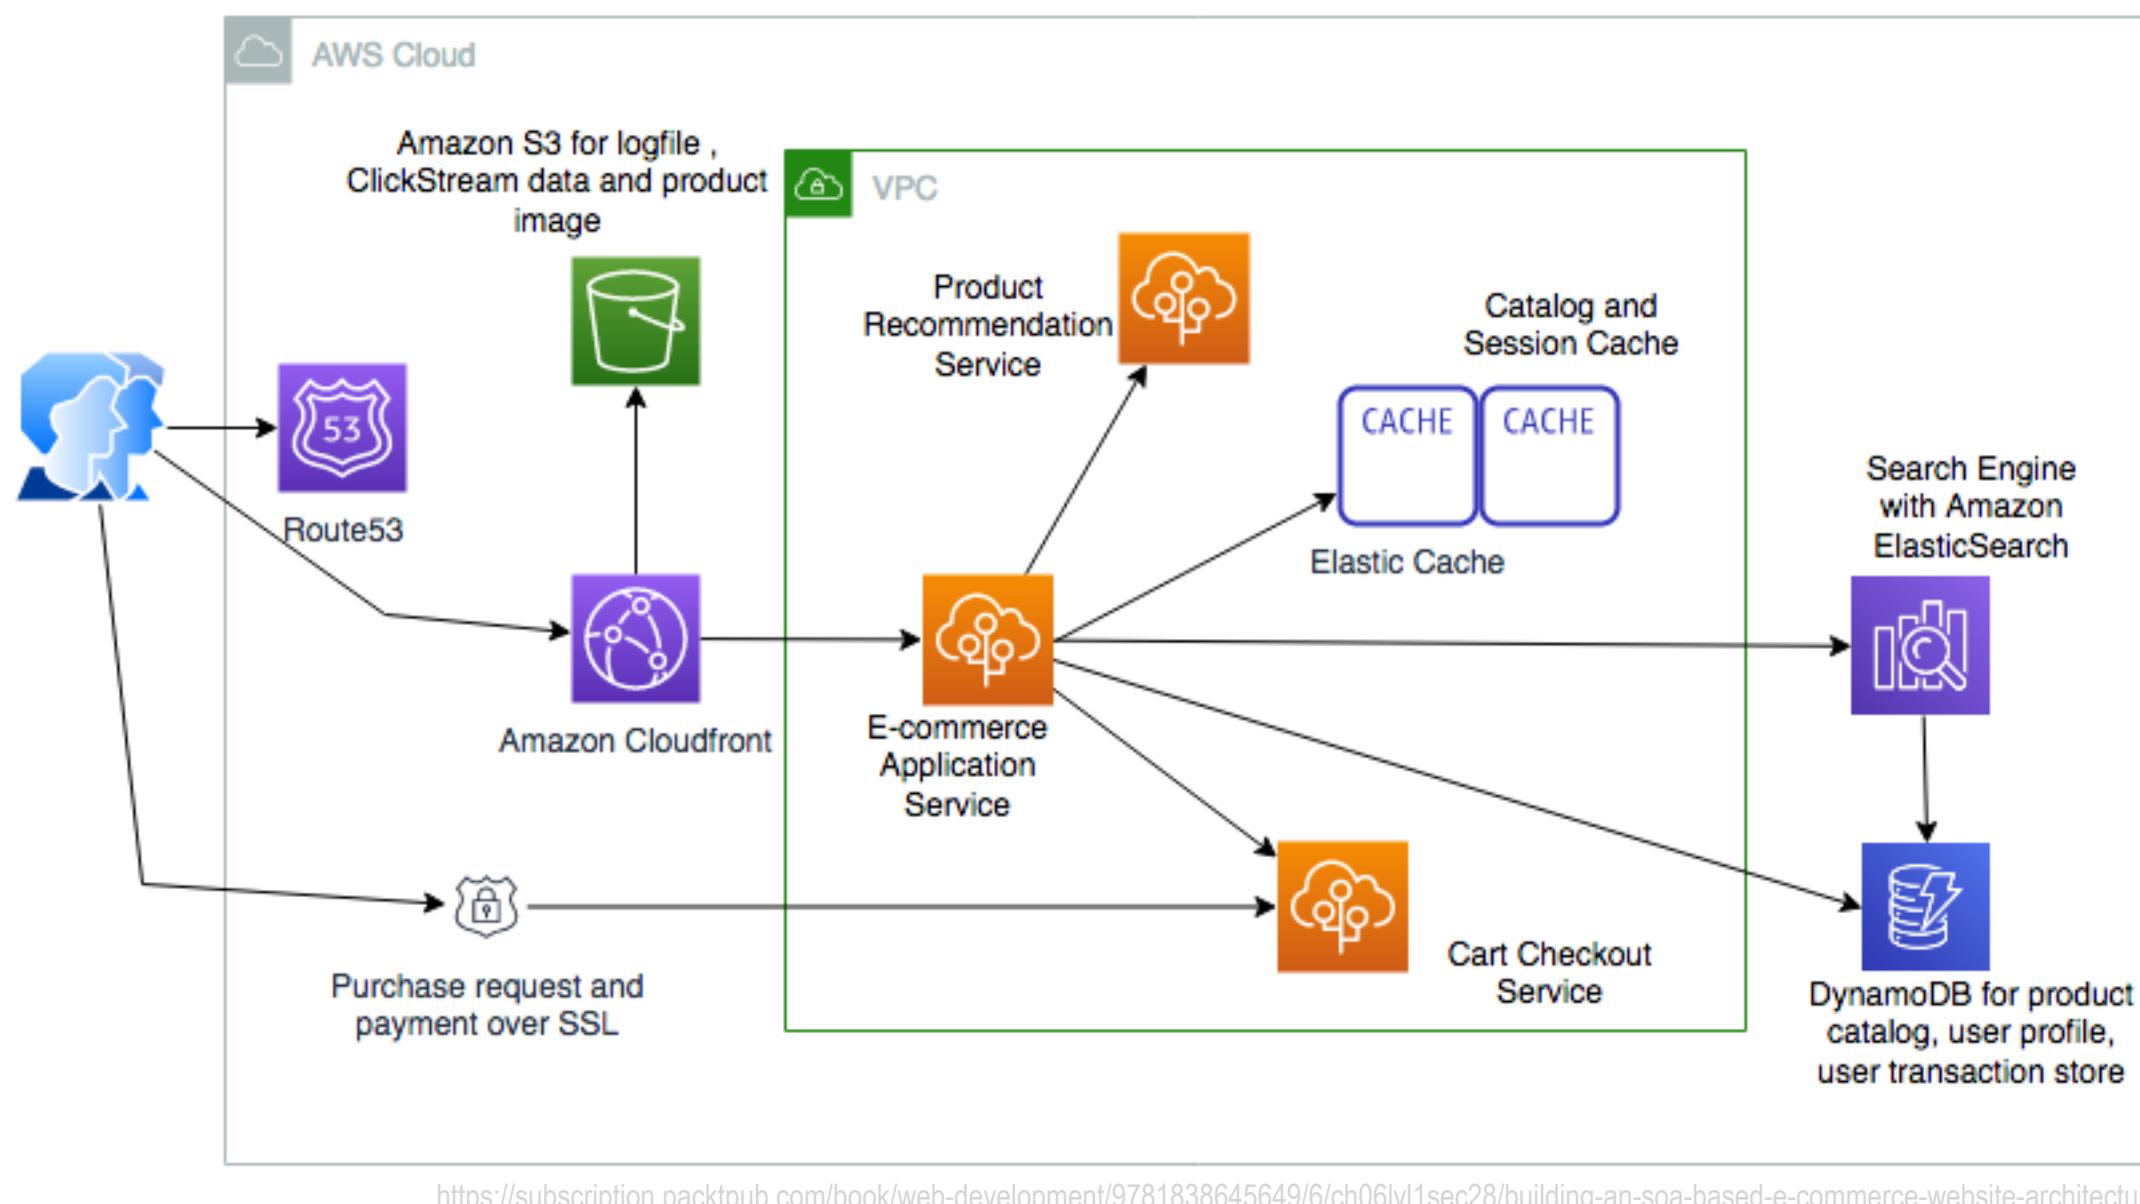


Step 1: Modularize system into fine-grained components

- Components with individual loci of control
 - *Well defined interfaces*
 - *Small in terms of program logic and startup time*

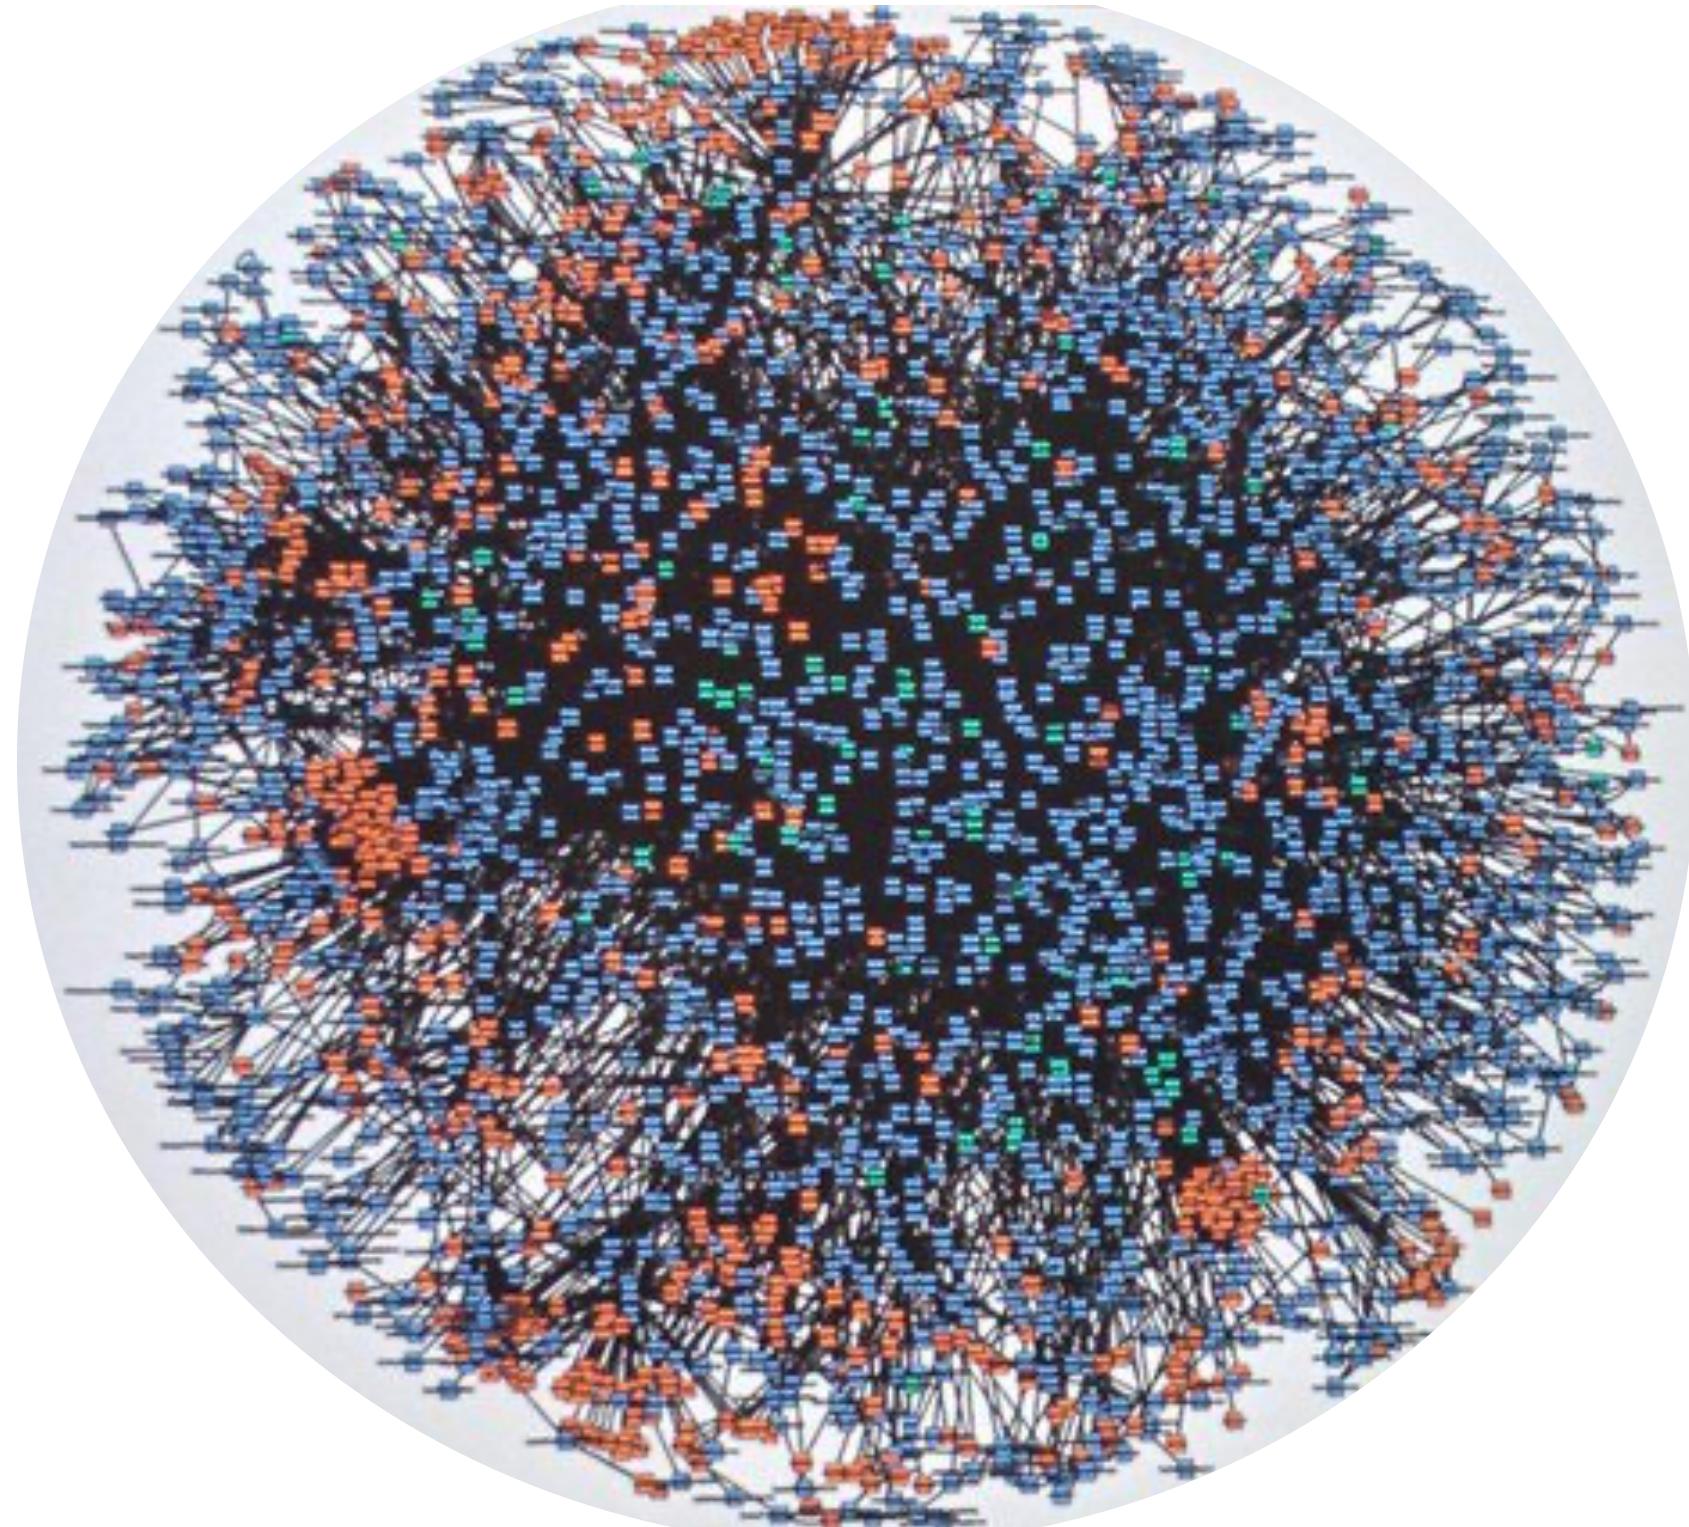
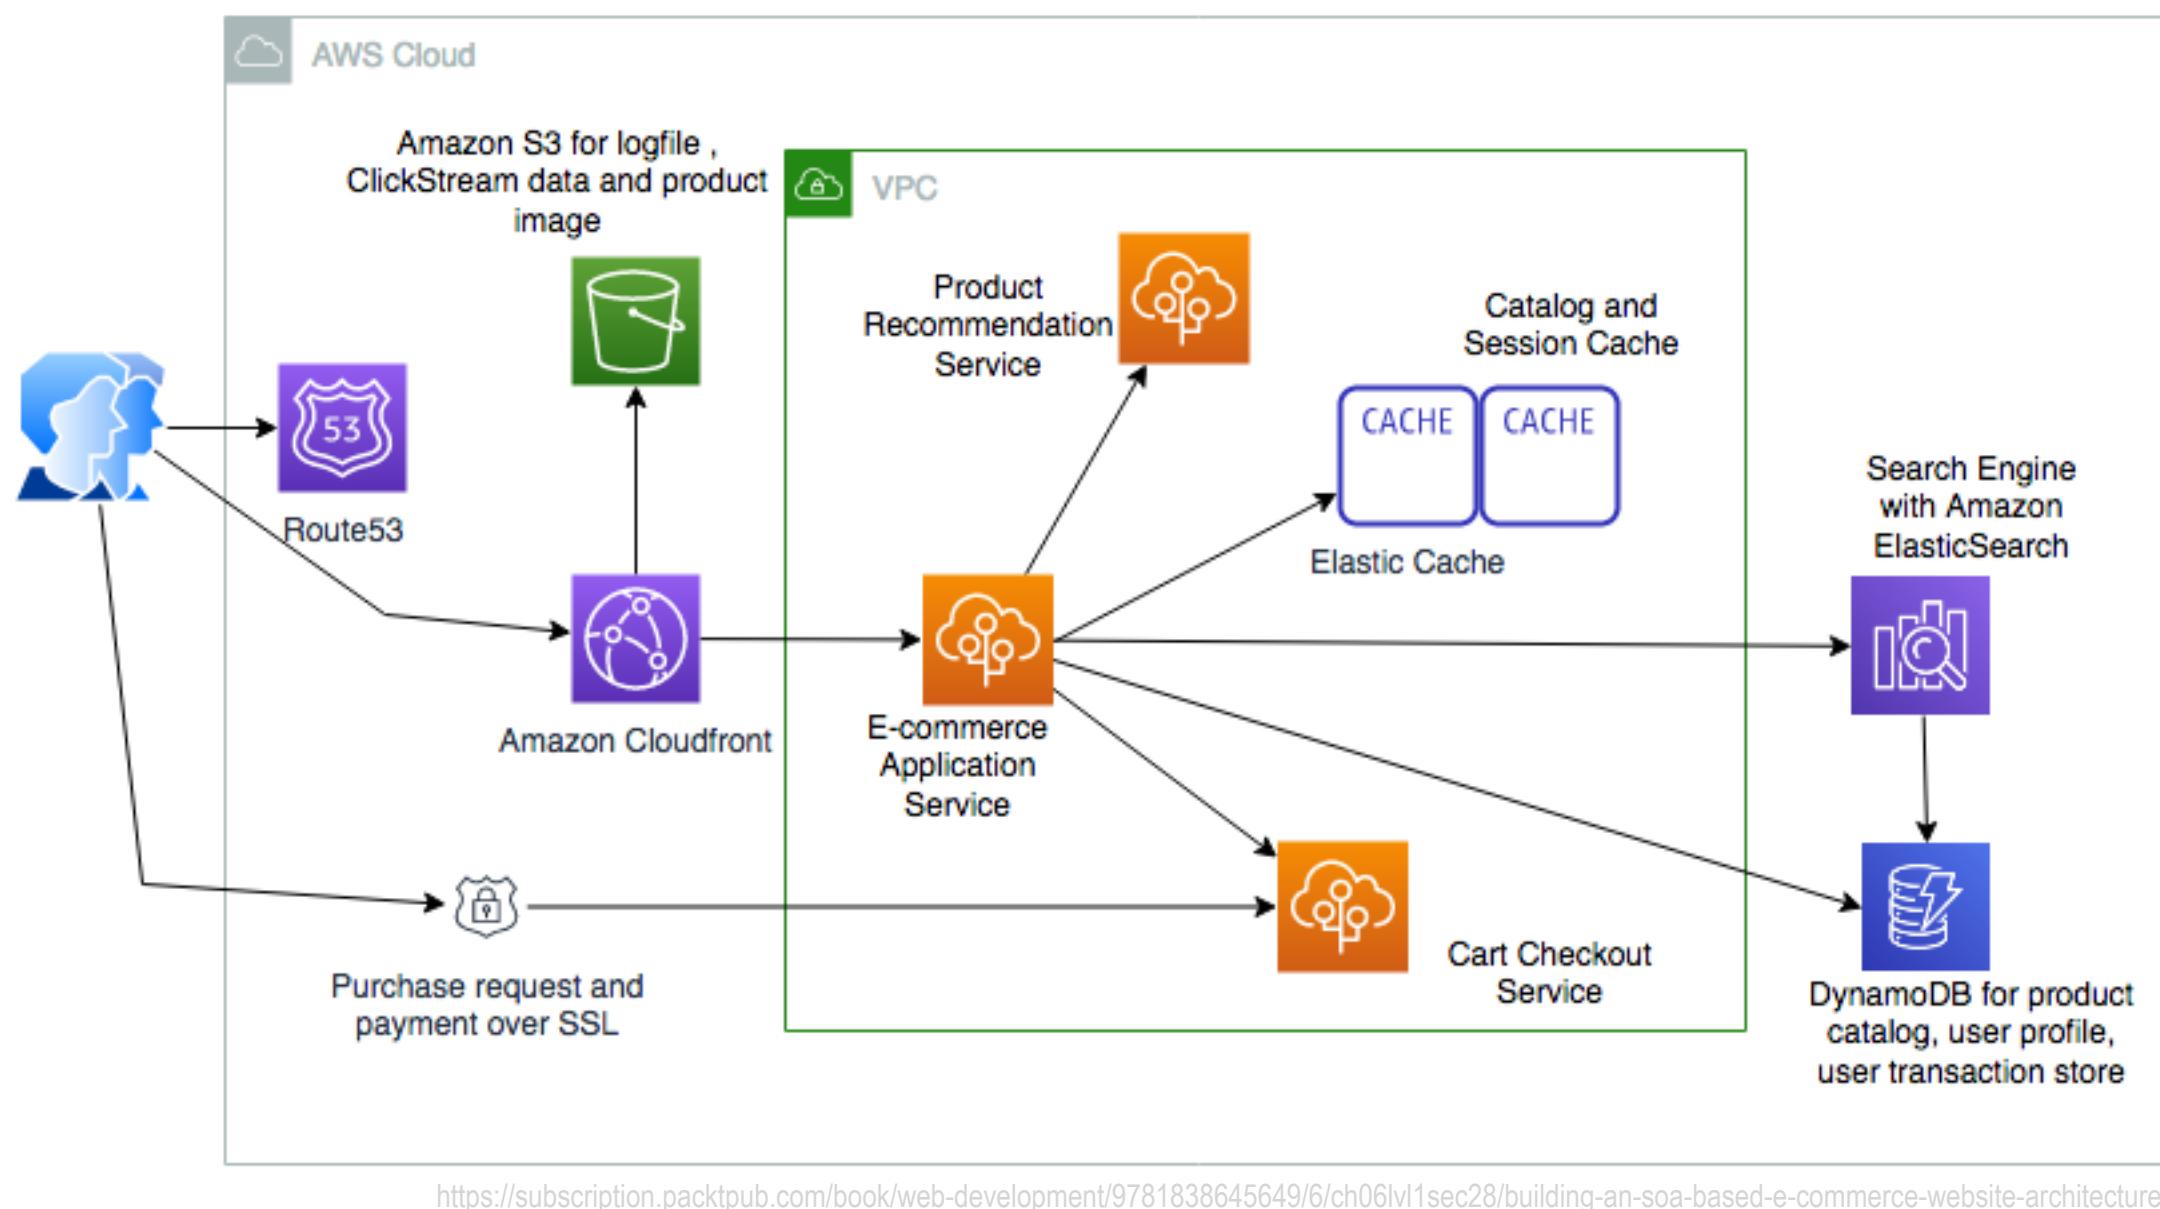
Step 1: Modularize system into fine-grained components

- Components with individual loci of control
 - *Well defined interfaces*
 - *Small in terms of program logic and startup time*



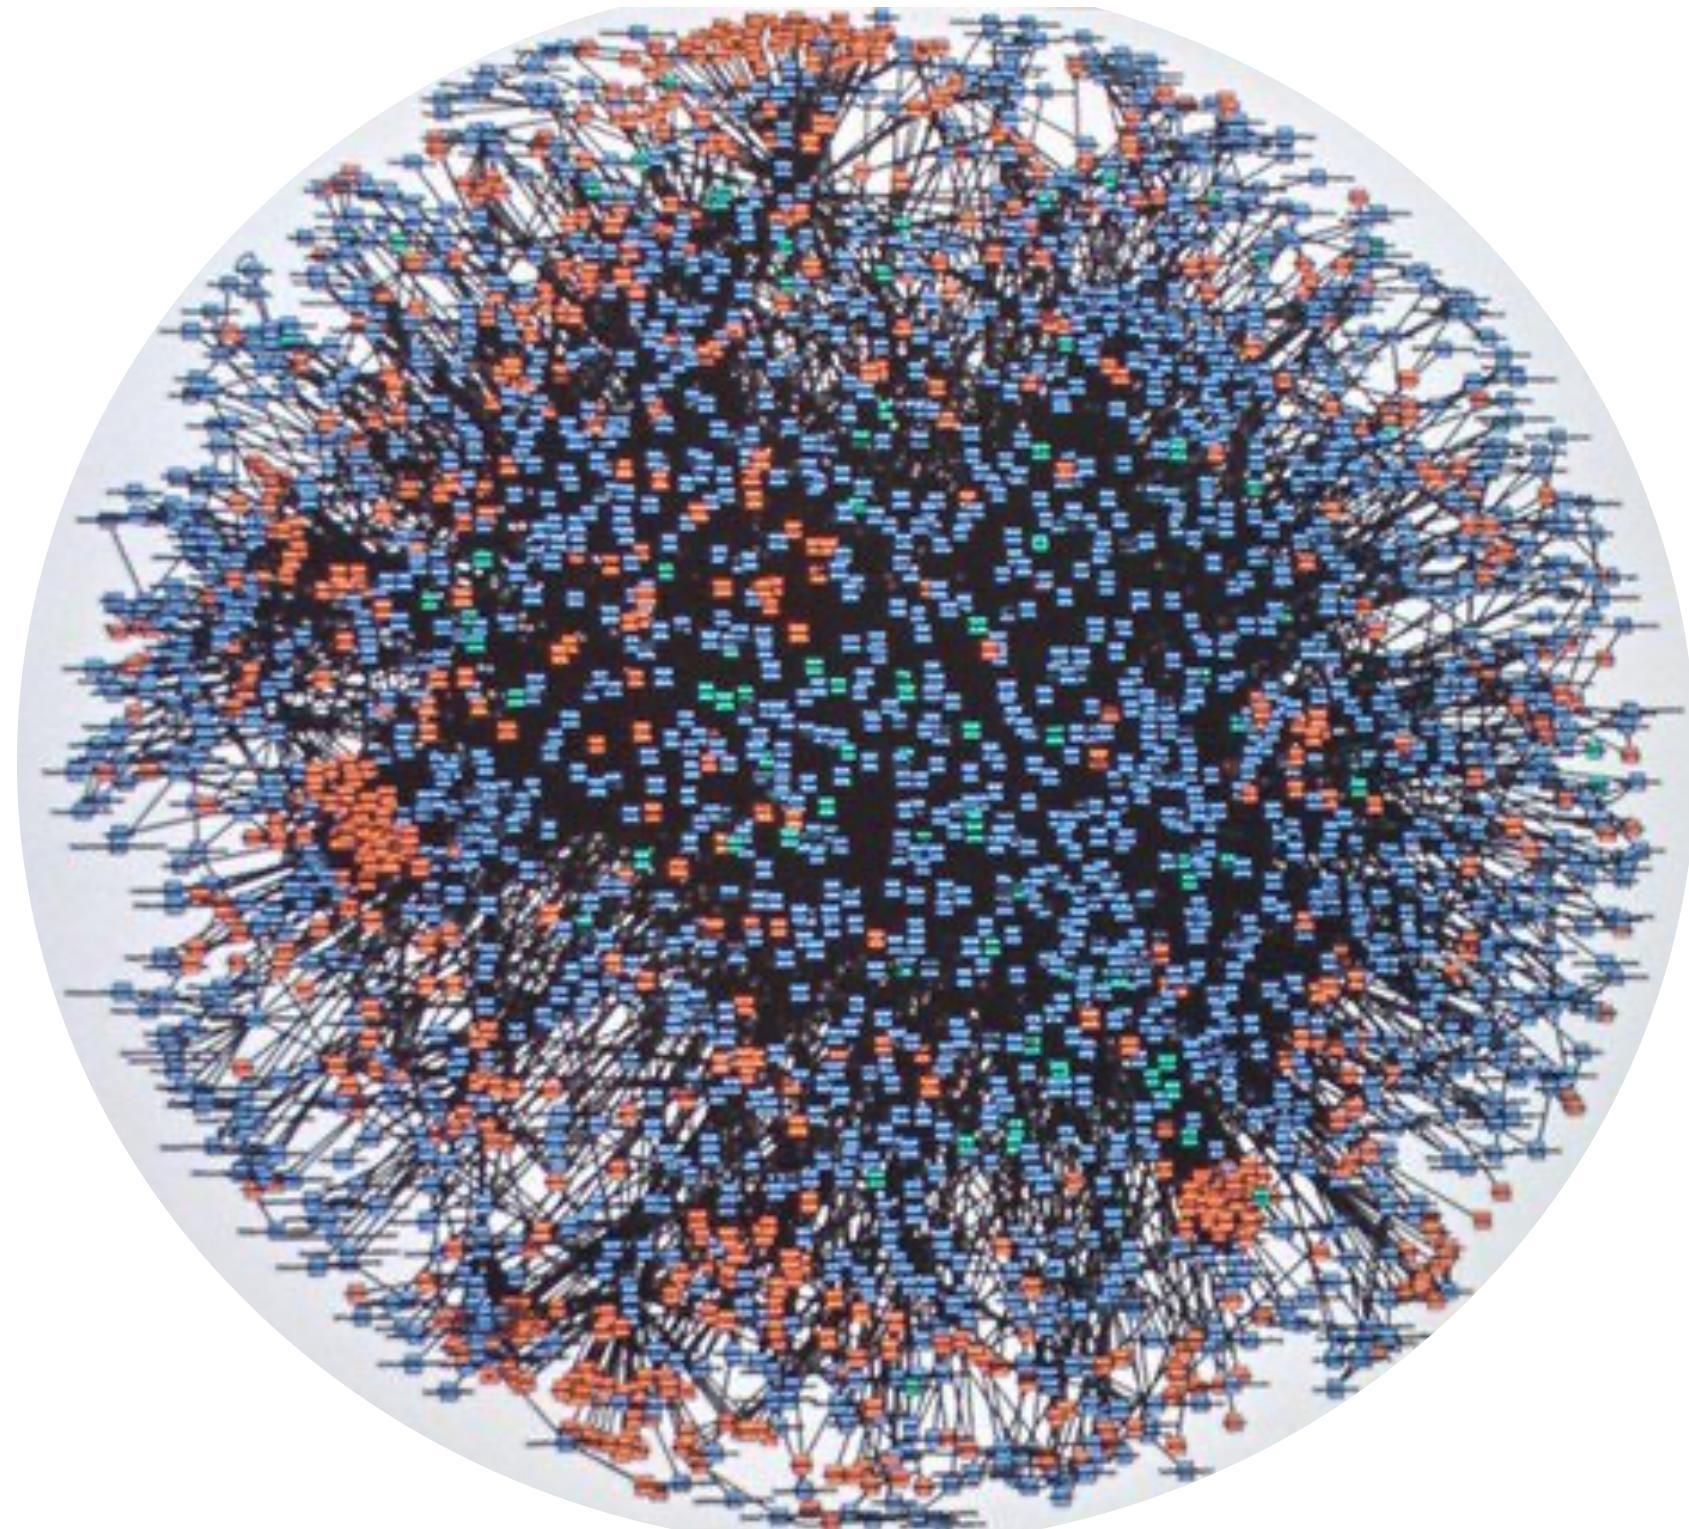
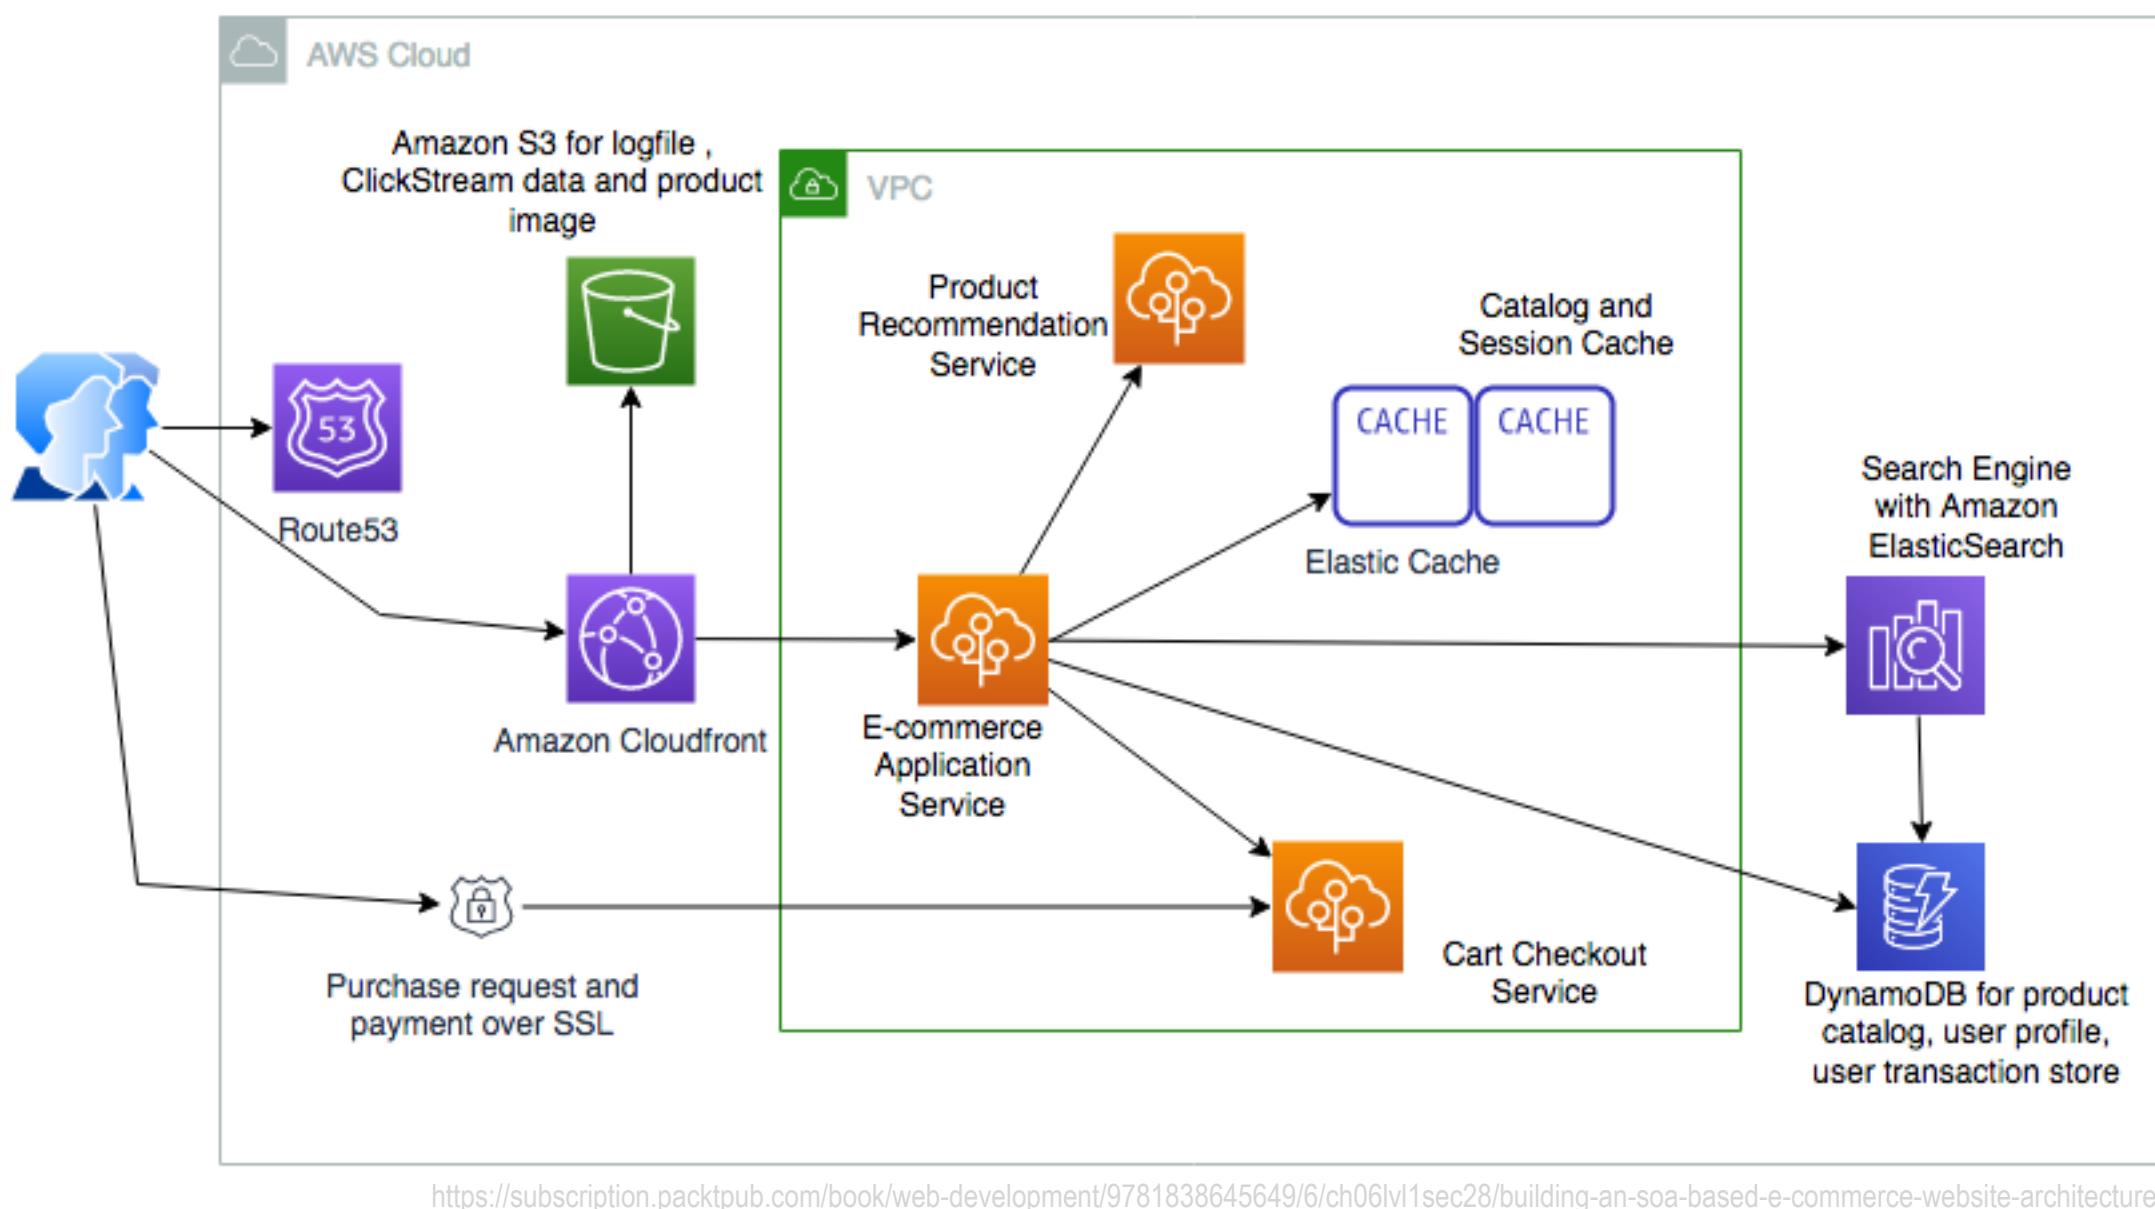
Step 1: Modularize system into fine-grained components

- Components with individual loci of control
 - *Well defined interfaces*
 - *Small in terms of program logic and startup time*



Step 1: Modularize system into fine-grained components

- Components with individual loci of control
 - *Well defined interfaces*
 - *Small in terms of program logic and startup time*
- $T_{reboot} = T_{restart} + T_{initialization}$



Problems with microrebooting

1. Component reboot can induce state corruption/inconsistency that persists across microrebooting
2. A component I depend on (i.e., need to call) is microrebooting when I need it
3. How to avoid resource leakage after arbitrary microrebooting?
4. How does a component re-integrate after microrebooting?

...

State segregation

- Goal: prevent microreboot from inducing corruption or state inconsistency

State segregation

- Goal: prevent microreboot from inducing corruption or state inconsistency
- Keep all state that should survive a reboot in dedicated state stores
 - *stores located outside the application ...*
 - *... behind strongly-enforced high-level APIs (e.g., DBs, KV stores)*

State segregation

- Goal: prevent microreboot from inducing corruption or state inconsistency
- Keep all state that should survive a reboot in dedicated state stores
 - *stores located outside the application ...*
 - *... behind strongly-enforced high-level APIs (e.g., DBs, KV stores)*
- Segment the state by lifetime
 - *apply modularization idea to all state: session state vs. persistent state*

State segregation

- Goal: prevent microreboot from inducing corruption or state inconsistency
- Keep all state that should survive a reboot in dedicated state stores
 - *stores located outside the application ...*
 - *... behind strongly-enforced high-level APIs (e.g., DBs, KV stores)*
- Segment the state by lifetime
 - *apply modularization idea to all state: session state vs. persistent state*
- Separate data recovery from app recovery => do each one better

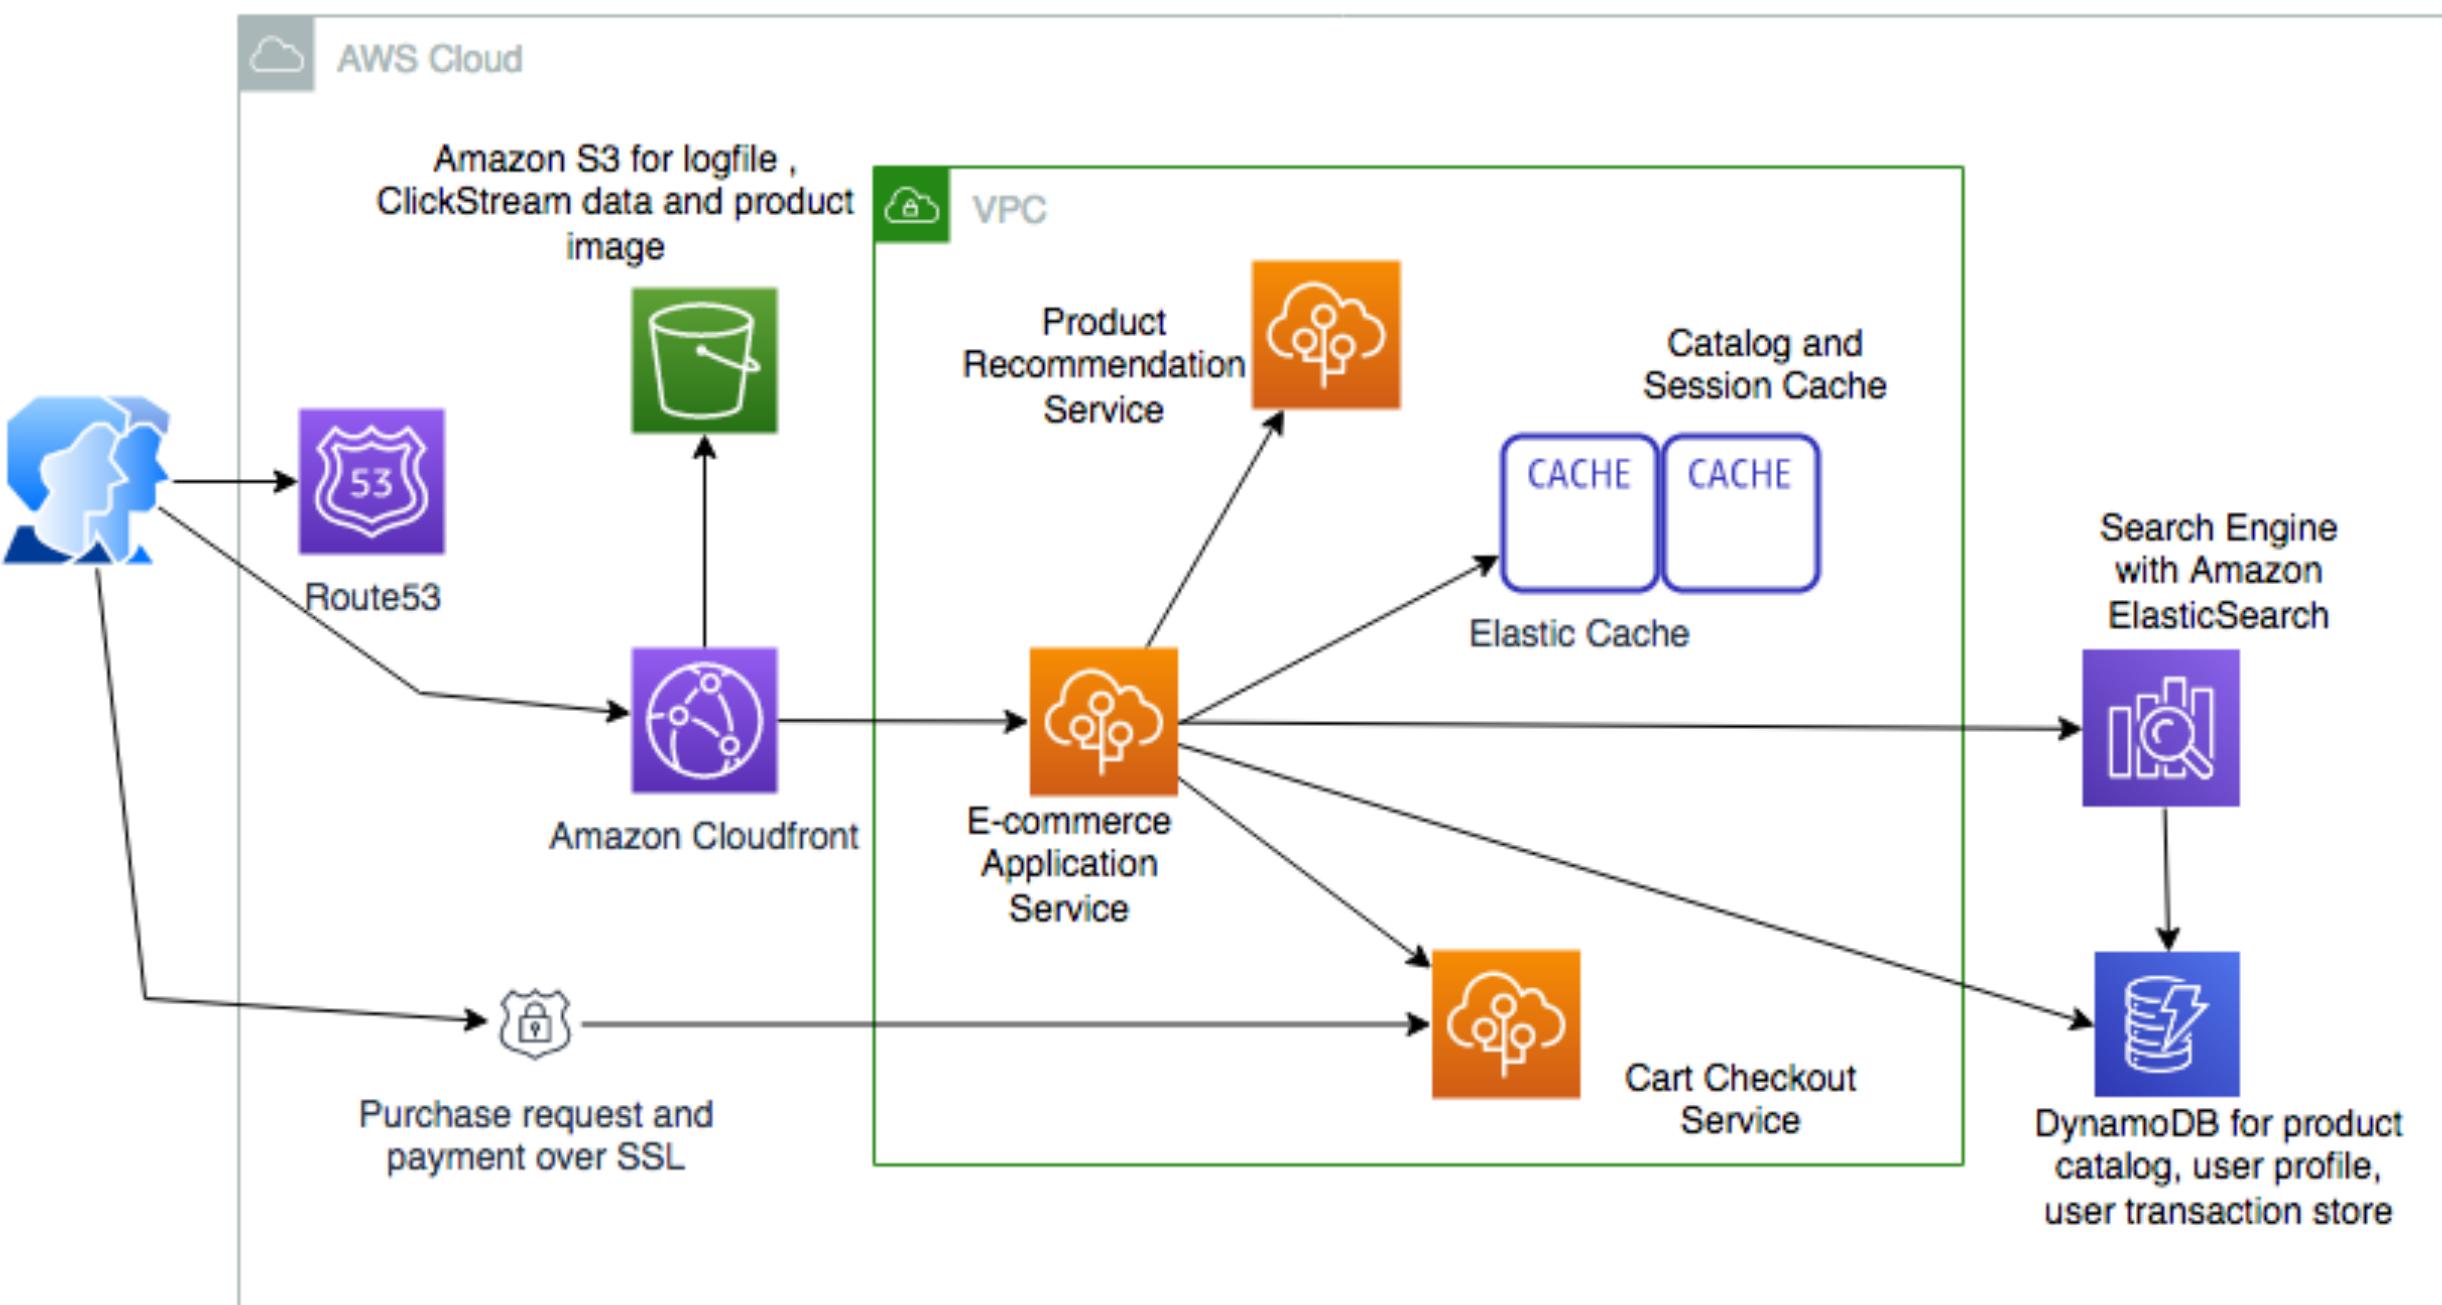
Problems with microrebooting

1. Component reboot can induce state corruption/inconsistency that persists across microrebooting
2. A component I depend on (i.e., need to call) is microrebooting when I need it
3. How to avoid resource leakage after arbitrary microrebooting?
4. How does a component re-integrate after microrebooting?

...

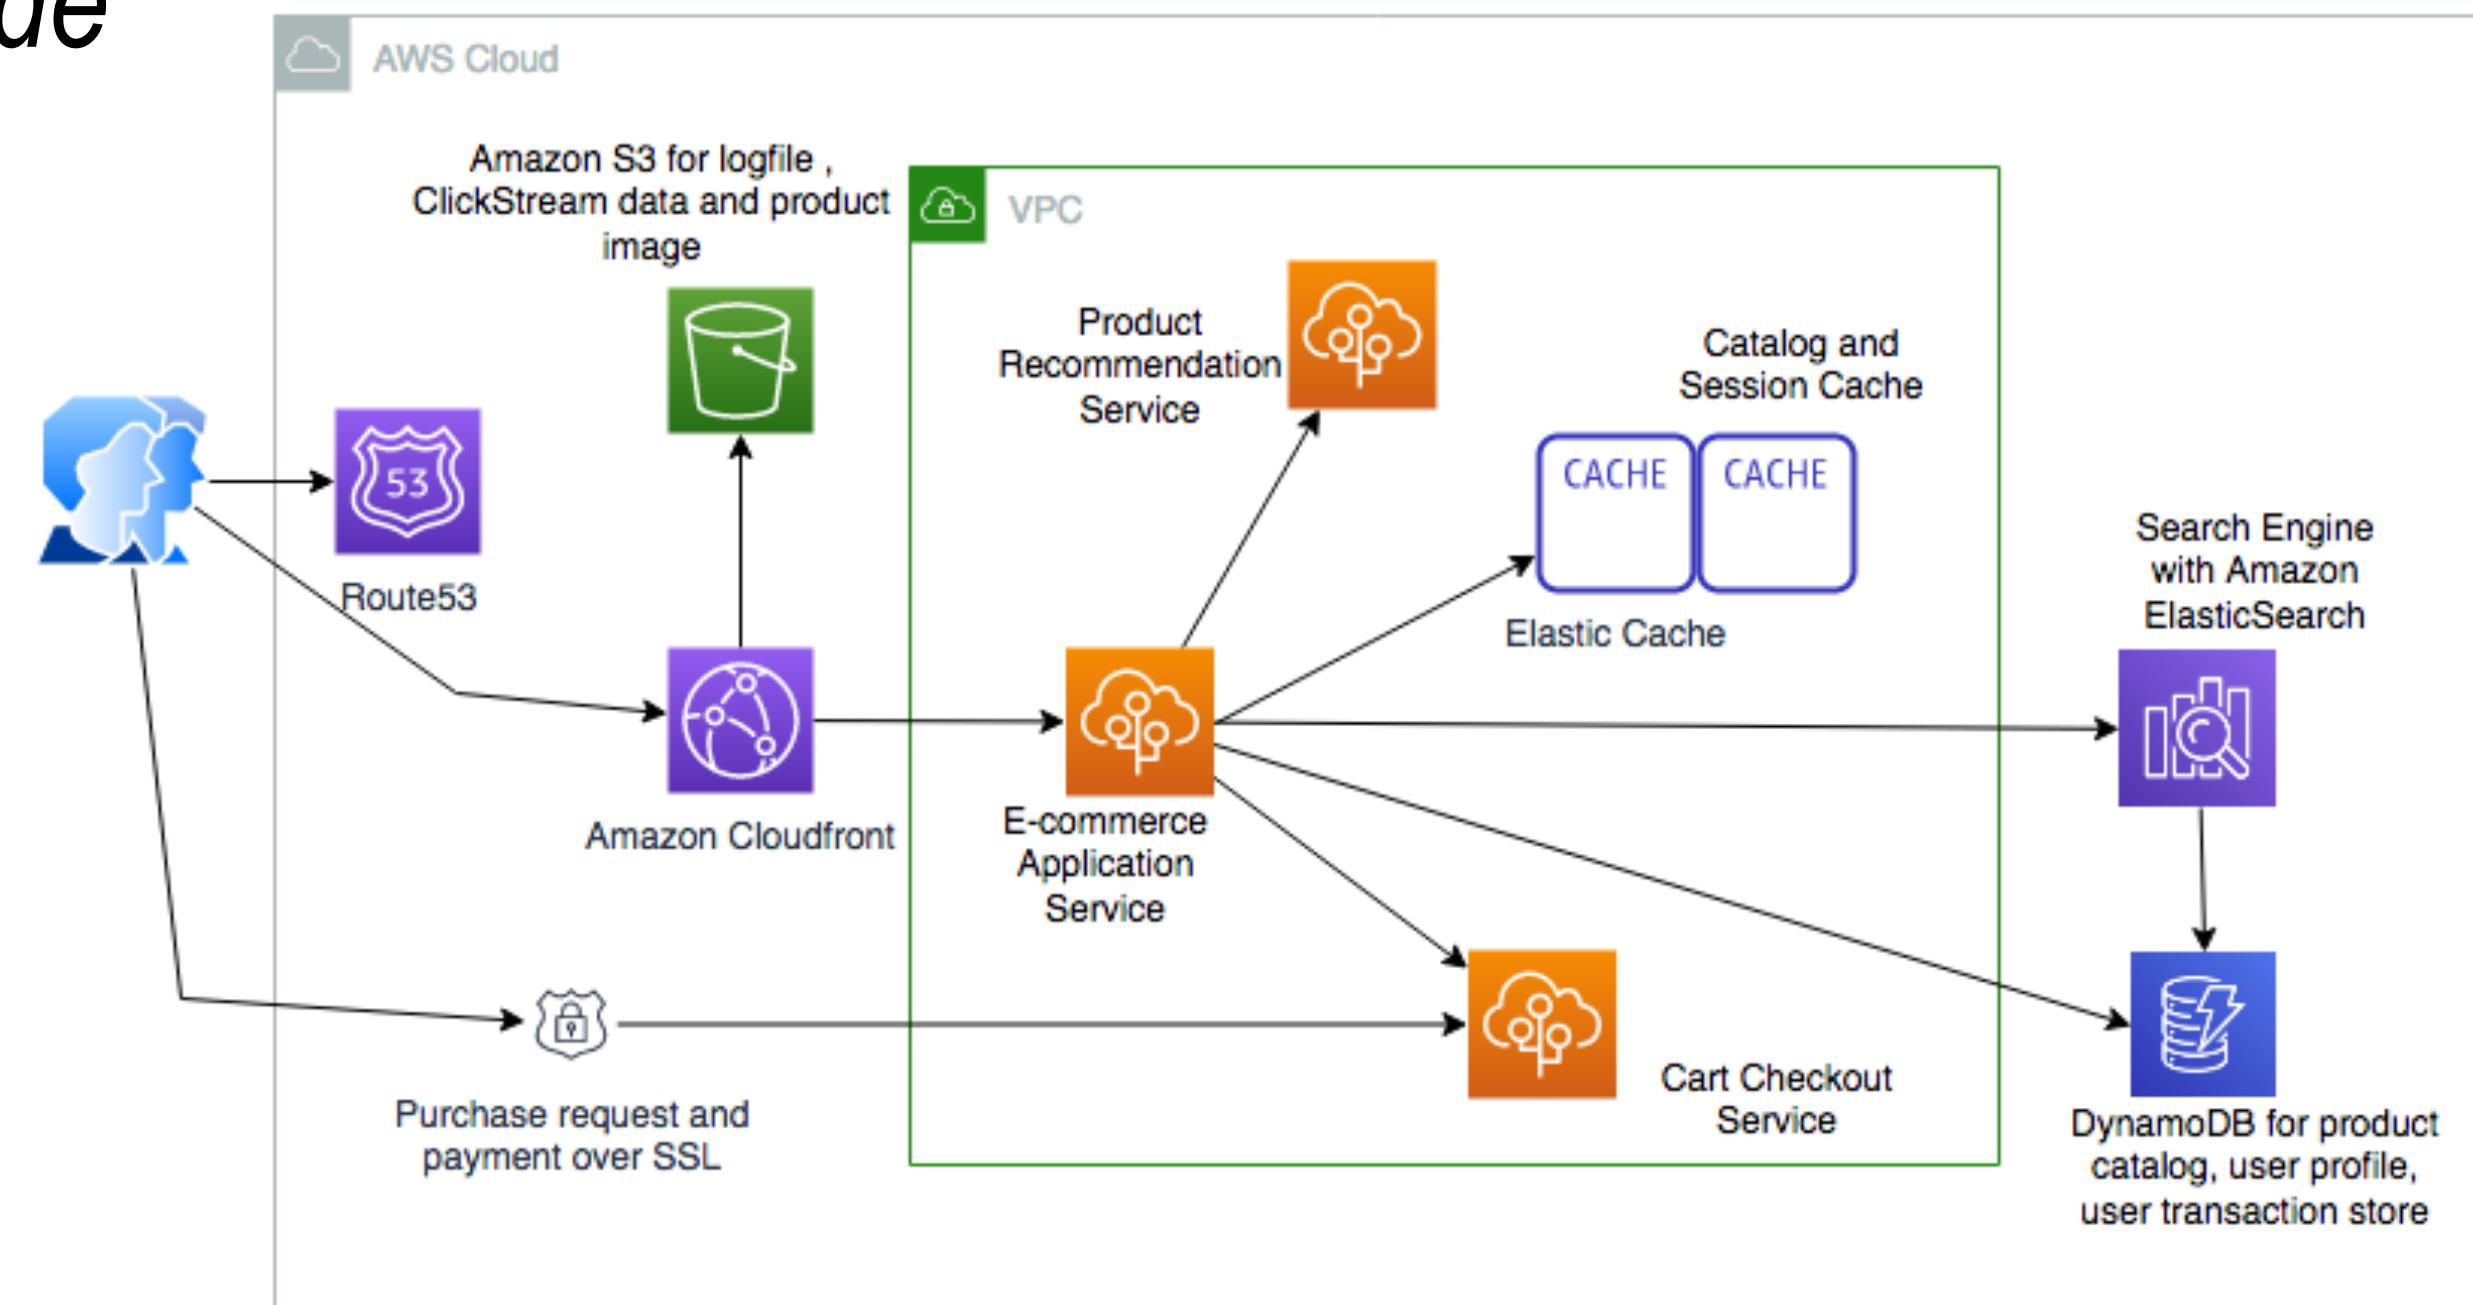
Functional decoupling

- Goal
 - *reduced disruption of system during restart*



Functional decoupling

- Goal
 - *reduced disruption of system during restart*
- No direct references (e.g., no pointers) across component boundaries
 - *Store cross-component references outside component*
 - Naming indirection through runtime
 - Marshall names into state store



<https://subscription.packtpub.com/book/web-development/9781838645649/6/ch06/l1sec28/building-an-soa-based-e-commerce-website-architecture>

Problems with microrebooting

1. Component reboot can induce state corruption/inconsistency that persists across microrebooting
2. A component I depend on (i.e., need to call) is microrebooting when I need it
3. How to avoid resource leakage after arbitrary microrebooting?
4. How does a component re-integrate after microrebooting?

...

Leased resources

- Goal: avoid resource leakage without fancy resource tracking

Leased resources

- Goal: avoid resource leakage without fancy resource tracking
- Lease = timed ownership
 - *File descriptors, memory, ...*
 - *Persistent long-term state*
 - *CPU execution time*

Leased resources

- Goal: avoid resource leakage without fancy resource tracking
- Lease = timed ownership
 - *File descriptors, memory, ...*
 - *Persistent long-term state*
 - *CPU execution time*
- Requests carry TTL => automatically purged when TTL runs out

Problems with microrebooting

1. Component reboot can induce state corruption/inconsistency that persists across microrebooting
2. A component I depend on (i.e., need to call) is microrebooting when I need it
3. How to avoid resource leakage after arbitrary microrebooting?
4. How does a component re-integrate after microrebooting?

...

Retryable interactions

- Goal
 - *seamless reintegration of microrebooted component by recovering in-flight requests transparently*

Retryable interactions

- Goal
 - *seamless reintegration of microrebooted component by recovering in-flight requests transparently*
- Interact via timed RPCs or equivalent
 - *if no response, caller can gracefully recover*
 - *timeouts help turn non-Byzantine failures into fail-stop events*
 - *RPC to a microrebooting module throws RetryAfter(t) exception*

Retryable interactions

- Goal
 - *seamless reintegration of microrebooted component by recovering in-flight requests transparently*
- Interact via timed RPCs or equivalent
 - *if no response, caller can gracefully recover*
 - *timeouts help turn non-Byzantine failures into fail-stop events*
 - *RPC to a microrebooting module throws RetryAfter(t) exception*
- Action depends on whether RPC is idempotent or not

Problems with microrebooting

1. Component reboot can induce state corruption/inconsistency that persists across microrebooting
2. A component I depend on (i.e., need to call) is microrebooting when I need it
3. How does a component reintegrate after microrebooting?
4. How to avoid resource leakage after arbitrary microrebooting?

...